

1



**SNIA**®

2

**Cloud Data Management Interface (CDMI)**

3

*Version 2.0.0a*

4

5

**SNIA Working Draft**

6

**Oct 21, 2019**

### USAGE

Copyright © 2019 SNIA. All rights reserved. All other trademarks or registered trademarks are the property of their respective owners.

The SNIA hereby grants permission for individuals to use this document for personal use only, and for corporations and other business entities to use this document for internal use only (including internal copying, distribution, and display) provided that:

1. Any text, diagram, chart, table or definition reproduced shall be reproduced in its entirety with no alteration, and,

2. Any document, printed or electronic, in which material from this document (or any portion hereof) is reproduced shall acknowledge the SNIA copyright on that material, and shall credit the SNIA for granting permission for its reuse.

Other than as explicitly provided above, you may not make any commercial use of this document, sell any excerpt or this entire document, or distribute this document to third parties. All rights not explicitly granted are expressly reserved to SNIA.

Permission to use this document for purposes other than those enumerated above may be requested by emailing [tcmd@snia.org](mailto:tcmd@snia.org). Please include the identity of the requesting individual or company and a brief description of the purpose, nature, and scope of the requested use.

All code fragments, scripts, data tables, and sample code in this SNIA document are made available under the following license:

BSD 3-Clause Software License

Copyright (c) 2019, The Storage Networking Industry Association.

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

- \* Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer.

- \* Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.

- \* Neither the name of The Storage Networking Industry Association (SNIA) nor the names of its contributors may be used to endorse or promote products derived from this software without specific prior written permission.

THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT OWNER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

### 42 DISCLAIMER

43 The information contained in this publication is subject to change without notice. The SNIA makes no warranty of any  
44 kind with regard to this specification, including, but not limited to, the implied warranties of merchantability and fitness for  
45 a particular purpose. The SNIA shall not be liable for errors contained herein or for incidental or consequential damages  
46 in connection with the furnishing, performance, or use of this specification.

47 Suggestions for revisions should be directed to <https://www.snia.org/feedback/>.

48 Copyright © 2018 SNIA. All rights reserved. All other trademarks or registered trademarks are the property of their  
49 respective owners.

# Table of Contents:

50

51	<b>I CDMI Preamble</b>	<b>1</b>
52	<b>Clause 1: Scope</b>	<b>3</b>
53	<b>Clause 2: Normative references</b>	<b>4</b>
54	<b>Clause 3: Terms, acronyms, and definitions</b>	<b>6</b>
55	<b>Clause 4: Conventions</b>	<b>9</b>
56	4.1 Interface format . . . . .	9
57	4.2 Typographical conventions . . . . .	10
58	4.3 Request and response body requirements . . . . .	11
59	4.4 Key Word requirements . . . . .	12
60	<b>Clause 5: Overview of Cloud Storage</b>	<b>13</b>
61	5.1 Overview . . . . .	13
62	5.2 Reference model for cloud storage interfaces . . . . .	16
63	5.3 Cloud data management interface . . . . .	17
64	5.4 Security . . . . .	21
65	5.5 Required HTTP support . . . . .	23
66	5.6 Time representations . . . . .	25
67	5.7 Backwards compatibility . . . . .	26
68	5.8 Object references . . . . .	27
69	<b>II Basic Cloud Storage</b>	<b>29</b>
70	<b>Clause 6: Data Object Resource Operations using HTTP</b>	<b>30</b>
71	6.1 Overview . . . . .	30
72	6.2 Create a data object using HTTP . . . . .	31
73	6.3 Read a data object using HTTP . . . . .	33
74	6.4 Update a data object using HTTP . . . . .	37
75	6.5 Delete a data object using HTTP . . . . .	40
76	<b>Clause 7: Container Object Resource Operations using HTTP</b>	<b>42</b>
77	7.1 Overview . . . . .	42
78	7.2 Create a container object using HTTP . . . . .	43
79	7.3 Read a container object using HTTP . . . . .	45
80	7.4 Update a container object using HTTP . . . . .	46
81	7.5 Delete a container object using HTTP . . . . .	47
82	7.6 Create (POST) a new data object using HTTP . . . . .	49
83	<b>III CDMI Core</b>	<b>52</b>
84	<b>Clause 8: Data Object Resource Operations using CDMI</b>	<b>53</b>
85	8.1 Overview . . . . .	53
86	8.2 Create a data object using CDMI . . . . .	56
87	8.3 Read a data object using CDMI . . . . .	67
88	8.4 Update a data object using CDMI . . . . .	77
89	8.5 Delete a data object using CDMI . . . . .	87
90	<b>Clause 9: Container Object Resource Operations using CDMI</b>	<b>89</b>
91	9.1 Overview . . . . .	89

92	9.2 Create a container object using CDMI . . . . .	92
93	9.3 Read a container object using CDMI . . . . .	99
94	9.4 Update a container object using CDMI . . . . .	104
95	9.5 Delete a container object using CDMI . . . . .	110
96	9.6 Create (POST) a new data object using CDMI . . . . .	112
97	9.7 Create (POST) a new queue object using CDMI . . . . .	122

## 98 IV CDMI Advanced 128

### 99 Clause 10: Domain object resource operations using CDMI 129

100	10.1 Overview . . . . .	129
101	10.2 Domain object summaries . . . . .	133
102	10.3 Domain object membership . . . . .	136
103	10.4 Create a domain object using CDMI . . . . .	139
104	10.5 Read a domain object using CDMI . . . . .	143
105	10.6 Update a domain object using CDMI . . . . .	147
106	10.7 Delete a domain object using CDMI . . . . .	151

### 107 Clause 11: Queue object resource operations using CDMI 153

108	11.1 Overview . . . . .	153
109	11.2 Create a queue object using CDMI . . . . .	156
110	11.3 Read a queue object using CDMI . . . . .	162
111	11.4 Update a queue object using CDMI . . . . .	169
112	11.5 Delete a queue object using CDMI . . . . .	173
113	11.6 Enqueue a new queue object value using CDMI . . . . .	175
114	11.7 Delete a queue object value using CDMI . . . . .	181

### 115 Clause 12: Capability object resource operations using CDMI 183

116	12.1 Overview . . . . .	183
117	12.2 Read a capabilities object using CDMI . . . . .	201

### 118 Clause 13: Exported protocols 205

119	13.1 Overview . . . . .	205
120	13.2 NFS exported protocol . . . . .	209
121	13.3 CIFS exported protocol . . . . .	211
122	13.4 OCCL exported protocol . . . . .	212
123	13.5 iSCSI exported protocol . . . . .	214
124	13.6 WebDAV exported protocol . . . . .	216

### 125 Clause 14: CDMI snapshots 217

126	14.1 Overview . . . . .	217
127	14.2 Creating a snapshot . . . . .	218
128	14.3 Deleting a snapshot . . . . .	219

### 129 Clause 15: Serialization/deserialization 220

130	15.1 Overview . . . . .	220
131	15.2 Canonical format . . . . .	221
132	15.3 Exporting serialized data . . . . .	223
133	15.4 Importing serialized data . . . . .	224

### 134 Clause 16: Metadata 225

135	16.1 Overview . . . . .	225
136	16.2 Support for storage system metadata . . . . .	226
137	16.3 Support for data system metadata . . . . .	228
138	16.4 Support for provided data system metadata . . . . .	236
139	16.5 Support for user metadata . . . . .	238
140	16.6 Metadata update operations . . . . .	239

### 141 Clause 17: Access control 240

142	17.1 Overview . . . . .	240
143	17.2 Access control lists . . . . .	243

### 144 Clause 18: Retention and hold management 253

145	18.1 Overview . . . . .	253
-----	-------------------------	-----

146	18.2 Retention management disciplines . . . . .	254
147	18.3 CDMI retention . . . . .	255
148	18.4 CDMI hold . . . . .	257
149	18.5 CDMI auto-deletion . . . . .	260
150	18.6 Retention security considerations . . . . .	261
151	<b>Clause 19: Scope specification</b>	<b>262</b>
152	19.1 Overview . . . . .	262
153	19.2 Examples . . . . .	263
154	19.3 Query matching expressions . . . . .	265
155	<b>Clause 20: Results specification</b>	<b>268</b>
156	20.1 Overview . . . . .	268
157	20.2 Examples . . . . .	269
158	<b>Clause 21: Notification queues</b>	<b>270</b>
159	21.1 Overview . . . . .	270
160	<b>Clause 22: Query queues</b>	<b>274</b>
161	22.1 Overview . . . . .	274
162	22.2 Extending CDMI query . . . . .	276
163	<b>Clause 23: Encrypted objects</b>	<b>277</b>
164	23.1 Overview . . . . .	277
165	23.2 Encryption operations . . . . .	278
166	23.3 Example uses of encrypted objects . . . . .	281
167	23.4 KMS integration . . . . .	282
168	23.5 CMS format . . . . .	283
169	23.6 JOSE format . . . . .	284
170	23.7 Signature/digest verification . . . . .	285
171	23.8 Error handling . . . . .	286
172	<b>Clause 24: Delegated access control</b>	<b>287</b>
173	24.1 Overview . . . . .	287
174	24.2 Delegated access control (DAC) . . . . .	289
175	24.3 Delegated access control message exchange . . . . .	291
176	24.4 Client header passthrough . . . . .	293
177	24.5 DAC request . . . . .	294
178	24.6 Packaged DAC request . . . . .	296
179	24.7 DAC response . . . . .	298
180	24.8 Packaged DAC response . . . . .	299
181	24.9 Error handling . . . . .	301
182	24.10 Examples . . . . .	302
183	<b>Clause 25: Data object versions</b>	<b>314</b>
184	25.1 Overview . . . . .	314
185	25.2 Traversing version-enabled data objects . . . . .	316
186	25.3 Concurrent updates and version-enabled data objects . . . . .	317
187	25.4 Capabilities for version-enabled data objects . . . . .	319
188	25.5 Updates triggering version creation . . . . .	320
189	25.6 Operations on version-enabled data objects . . . . .	321
190	25.7 Operations on data object versions . . . . .	322
191	25.8 Query of data object versions . . . . .	323
192	25.9 Version-enabled data object serialization . . . . .	324
193	<b>V CDMI Annexes</b>	<b>326</b>
194	<b>Clause 26: Extensions</b>	<b>327</b>
195	26.1 Overview . . . . .	327
196	26.2 Summary metadata for bandwidth . . . . .	328
197	26.3 Expiring access control entries (ACEs) . . . . .	330
198	26.4 Group storage system metadata . . . . .	331
199	26.5 Header-based metadata . . . . .	332
200	26.6 Immediate query . . . . .	340

201	<b>VI References</b>	<b>343</b>
202	<b>Bibliography</b>	<b>344</b>

# List of Figures

204	Figure 1: Existing data storage interface standards . . . . .	14
205	Figure 2: Storage interfaces for object storage client data . . . . .	15
206	Figure 3: Cloud storage reference model . . . . .	16
207	Figure 4: CDMI object model . . . . .	18
208	Figure 5: Object transitions between named and ID-only . . . . .	19
209	Figure 6: Hierarchy of domains . . . . .	129
210	Figure 7: Hierarchy of capabilities . . . . .	185
211	Figure 8: CDMI and OCCI in an integrated cloud computing environment . . . . .	212
212	Figure 9: Snapshot container structure . . . . .	217
213	Figure 10: Access control flow . . . . .	241
214	Figure 11: Object retention . . . . .	255
215	Figure 12: Object hold . . . . .	257
216	Figure 13: Object hold on object with retention . . . . .	257
217	Figure 14: Object with multiple holds . . . . .	258
218	Figure 15: Encrypted object state transistions . . . . .	278
219	Figure 16: Non-delegated (ACL-based) access control data flow . . . . .	287
220	Figure 17: Delegated access control data flow example for non-encrypted object . . . . .	291
221	Figure 18: Delegated access control data flow example for encrypted object . . . . .	292
222	Figure 19: Updates to a non-version-enabled data object . . . . .	314
223	Figure 20: Updates to a version-enabled data object . . . . .	315
224	Figure 21: Linkages between a version-enabled data object and data object versions . . . . .	316
225	Figure 22: Overlapping concurrent updates . . . . .	317
226	Figure 23: Linkages for overlapping updates . . . . .	317
227	Figure 24: Nested concurrent updates . . . . .	318
228	Figure 25: Linkages for nested updates . . . . .	318
229	Figure 26: Version to <code>capabilityURI</code> relationships . . . . .	319



# List of Tables

230

231	Table 1: Overview of this document . . . . .	2
232	Table 2: Interface format . . . . .	9
233	Table 3: Key word requirements . . . . .	12
234	Table 4: Types of resources in the CDMI object model . . . . .	18
235	Table 5: Creation/consumption of storage system metadata . . . . .	19
236	Table 6: Object ID format . . . . .	20
237	Table 7: Relative URIs resolved against root URIs . . . . .	24
238	Table 8: Capabilities - Create a CDMI data object using HTTP . . . . .	31
239	Table 9: Request headers - Create a CDMI data object using HTTP . . . . .	31
240	Table 10: HTTP status codes - Create a data object using HTTP . . . . .	32
241	Table 11: Capabilities - Read a CDMI data object using HTTP . . . . .	33
242	Table 12: Request header - Read a CDMI data object using HTTP . . . . .	34
243	Table 13: Response headers - Read a CDMI Data Object using HTTP . . . . .	34
244	Table 14: HTTP status codes - Read a CDMI data object using HTTP . . . . .	35
245	Table 15: Capabilities - Update a CDMI data object using HTTP . . . . .	37
246	Table 16: Request headers - Update a CDMI data object using HTTP . . . . .	37
247	Table 17: Response header - Update a CDMI data object using HTTP . . . . .	38
248	Table 18: HTTP status codes - Update a CDMI data object using HTTP . . . . .	38
249	Table 19: Capabilities - Delete a CDMI data object using HTTP . . . . .	40
250	Table 20: HTTP status codes - Delete a CDMI data object using HTTP . . . . .	41
251	Table 21: Capabilities - Create a CDMI container object using HTTP . . . . .	43
252	Table 22: HTTP status codes - Create a container object using HTTP . . . . .	44
253	Table 23: Capabilities - Delete a CDMI container object using HTTP . . . . .	47
254	Table 24: HTTP status codes - Delete a CDMI container object using HTTP . . . . .	48
255	Table 25: Capabilities - Create a CDMI data object using HTTP POST . . . . .	49
256	Table 26: Request header - Create a new data object using HTTP . . . . .	50
257	Table 27: Response header - Create a new data object using HTTP . . . . .	50
258	Table 28: HTTP status codes - Create a new data object using HTTP . . . . .	50
259	Table 29: Capabilities - Create a CDMI data object using CDMI . . . . .	57
260	Table 30: Request headers - Create a CDMI data object using CDMI . . . . .	57
261	Table 31: Request message body - Create a data object using CDMI . . . . .	58
262	Table 32: Response headers - Create a data object using CDMI . . . . .	61
263	Table 33: Response message body - Create a data object using CDMI . . . . .	61
264	Table 34: HTTP status codes - Create a data object using CDMI . . . . .	62
265	Table 35: Capabilities - Read a CDMI data object using CDMI . . . . .	67
266	Table 36: Request headers - Read a CDMI data object using CDMI . . . . .	67
267	Table 37: Response headers - Read a CDMI data object using CDMI . . . . .	68
268	Table 38: Response message body - Read a CDMI data object using CDMI . . . . .	68
269	Table 39: HTTP status codes - Read a CDMI data object using CDMI . . . . .	70
270	Table 40: Capabilities - Update a CDMI data object using CDMI . . . . .	77
271	Table 41: Request headers - Update a CDMI data object using CDMI . . . . .	78
272	Table 42: Request message body - Update a CDMI data object using CDMI . . . . .	79
273	Table 43: Response header - Update a CDMI data object using CDMI . . . . .	82
274	Table 44: HTTP status codes - Update a CDMI data object using CDMI . . . . .	83
275	Table 45: Capabilities - Delete a CDMI data object using CDMI . . . . .	87
276	Table 46: HTTP status codes - Delete a CDMI data object using CDMI . . . . .	88
277	Table 47: Container metadata . . . . .	91

278	Table 48: Capabilities - Create a CDMI container object using CDMI . . . . .	93
279	Table 49: Request headers - Create a container object using CDMI . . . . .	93
280	Table 50: Request message body - Create a container object using CDMI . . . . .	93
281	Table 51: Response headers - Create a container object using CDMI . . . . .	95
282	Table 52: Response message body - Create a container object using CDMI . . . . .	95
283	Table 53: HTTP status codes - Create a CDMI container object using CDMI . . . . .	96
284	Table 54: Capabilities - Read a CDMI Container Object using CDMI . . . . .	99
285	Table 55: Request headers - Read a container object using CDMI . . . . .	99
286	Table 56: Response headers - Read a container object using CDMI . . . . .	100
287	Table 57: Response message body - Read a container object using CDMI . . . . .	100
288	Table 58: HTTP status codes - Read a container object using CDMI . . . . .	102
289	Table 59: Capabilities - Update a CDMI container object using CDMI . . . . .	105
290	Table 60: Request headers - Update a container object using CDMI . . . . .	105
291	Table 61: Request message body - Update a container object using CDMI . . . . .	105
292	Table 62: Response header - Update a container object using CDMI . . . . .	108
293	Table 63: HTTP status codes - Update a container object using CDMI . . . . .	108
294	Table 64: Capabilities - Delete a CDMI container object using CDMI . . . . .	110
295	Table 65: HTTP status codes - Delete a container object using CDMI . . . . .	111
296	Table 66: Capabilities - Create a CDMI data object using CDMI . . . . .	113
297	Table 67: Request headers - Create a new data object Using CDMI . . . . .	113
298	Table 68: Request message body - Create a new data object Using CDMI . . . . .	114
299	Table 69: Response headers - Create a new data object using CDMI . . . . .	118
300	Table 70: Response message body - Create a new data object using CDMI . . . . .	118
301	Table 71: HTTP status codes - Create a new data object using CDMI . . . . .	119
302	Table 72: Capabilities - Create a CDMI Queue object using CDMI . . . . .	123
303	Table 73: Request headers - Create a new queue object using CDMI . . . . .	123
304	Table 74: Request message body - Create a new queue object using CDMI . . . . .	124
305	Table 75: Response headers - Create a new queue object using CDMI . . . . .	125
306	Table 76: Response message body - Create a new queue object using CDMI . . . . .	125
307	Table 77: HTTP status codes - Create a new queue object using CDMI . . . . .	126
308	Table 78: Required metadata for a domain object . . . . .	131
309	Table 79: Contents of domain summary objects . . . . .	133
310	Table 80: Required settings for domain member user objects . . . . .	136
311	Table 81: Required settings for domain member delegation objects . . . . .	137
312	Table 82: Capabilities - Create a CDMI domain object using CDMI . . . . .	139
313	Table 83: Request headers - Create a domain object using CDMI . . . . .	139
314	Table 84: Request message body - Create a domain object using CDMI . . . . .	140
315	Table 85: Response headers - Create a domain object using CDMI . . . . .	140
316	Table 86: Response message body - Create a domain object using CDMI . . . . .	141
317	Table 87: HTTP status codes - Create a domain object using CDMI . . . . .	141
318	Table 88: Capabilities - Read a CDMI domain object using CDMI . . . . .	143
319	Table 89: Request headers - Read a domain object using CDMI . . . . .	143
320	Table 90: Response headers - Read a domain object using CDMI . . . . .	144
321	Table 91: Response message body - Read a domain object using CDMI . . . . .	144
322	Table 92: HTTP status codes - Read a domain object using CDMI . . . . .	145
323	Table 93: Capabilities - Update a CDMI domain object using CDMI . . . . .	147
324	Table 94: Request headers - Update a domain object using CDMI . . . . .	147
325	Table 95: Request message body - Update a domain object using CDMI . . . . .	148
326	Table 96: Response header - Update a domain object using CDMI . . . . .	149
327	Table 97: HTTP status codes - Update a domain object using CDMI . . . . .	149
328	Table 98: Capabilities - Delete a CDMI domain object using CDMI . . . . .	151
329	Table 99: HTTP status codes - Delete a domain object using CDMI . . . . .	152
330	Table 100: Capabilities - Create a CDMI queue object using CDMI . . . . .	156
331	Table 101: Request headers - Create a queue object Using CDMI . . . . .	157
332	Table 102: Request message body - Create a queue object using CDMI . . . . .	157
333	Table 103: Response headers - Create a queue object Using CDMI . . . . .	159
334	Table 104: Response message body - Create a queue object using CDMI . . . . .	159
335	Table 105: HTTP status codes - Create a queue object using CDMI . . . . .	160
336	Table 106: Capabilities - Read a CDMI queue object using CDMI . . . . .	162
337	Table 107: Request headers - Read a queue object using CDMI . . . . .	163
338	Table 108: Response headers - Read a queue object using CDMI . . . . .	163
339	Table 109: Response message body - Read a queue object using CDMI . . . . .	163
340	Table 110: HTTP status codes - Read a queue object using CDMI . . . . .	166

341	Table 111: Capabilities - Update a queue object using CDMI . . . . .	169
342	Table 112: Request headers - Update a queue object Using CDMI . . . . .	169
343	Table 113: Request message body - Update a queue object Using CDMI . . . . .	169
344	Table 114: Response header - Update a queue object Using CDMI . . . . .	172
345	Table 115: HTTP status codes - Update a queue object using CDMI . . . . .	172
346	Table 116: Capabilities - Delete a queue object using CDMI . . . . .	173
347	Table 117: HTTP status codes - Delete a queue object Using CDMI . . . . .	174
348	Table 118: Capabilities - Enqueue a new queue object value using CDMI . . . . .	175
349	Table 119: Request headers - Enqueue a new queue object value using CDMI . . . . .	175
350	Table 120: Request message body - Enqueue a new queue object value using CDMI . . . . .	176
351	Table 121: HTTP status codes - Enqueue a new queue object value Using CDMI . . . . .	178
352	Table 122: Capabilities - Delete a queue object value using CDMI . . . . .	181
353	Table 123: HTTP status codes - Delete a queue object value using CDMI . . . . .	182
354	Table 124: System-wide capabilities . . . . .	186
355	Table 125: Capabilities for storage system metadata . . . . .	190
356	Table 126: Capabilities for data system metadata . . . . .	192
357	Table 127: Capabilities for data objects . . . . .	195
358	Table 128: Capabilities for container objects . . . . .	196
359	Table 129: Capabilities for domain objects . . . . .	198
360	Table 130: Capabilities for queue objects . . . . .	200
361	Table 131: Capabilities - Read a capabilities object using CDMI . . . . .	201
362	Table 132: Request headers - Read a capabilities object using CDMI . . . . .	201
363	Table 133: Response headers - Read a capabilities object Using CDMI . . . . .	202
364	Table 134: Response message body - Read a capabilities object using CDMI . . . . .	202
365	Table 135: HTTP status codes - Read a capabilities object using CDMI . . . . .	202
366	Table 136: Required members of the NFS protocol structure . . . . .	209
367	Table 137: Optional members of the NFS protocol structure . . . . .	209
368	Table 138: Required members of the CIFS protocol structure . . . . .	211
369	Table 139: Serialization import behaviour . . . . .	224
370	Table 140: Storage system metadata . . . . .	226
371	Table 141: Data system metadata . . . . .	228
372	Table 142: Provided values of data system metadata . . . . .	236
373	Table 143: ACE types . . . . .	243
374	Table 144: Who identifiers . . . . .	243
375	Table 145: ACE flags . . . . .	244
376	Table 146: ACE masks bits . . . . .	245
377	Table 147: ACE bit mask/string . . . . .	251
378	Table 148: Query matching expressions . . . . .	265
379	Table 149: Required metadata for a notification queue . . . . .	270
380	Table 150: Notification status metadata . . . . .	273
381	Table 151: Required metadata for a query queue . . . . .	274
382	Table 152: Query status metadata . . . . .	275
383	Table 153: Access modes for DAC . . . . .	289
384	Table 154: DAC request . . . . .	294
385	Table 155: Packaged DAC request . . . . .	296
386	Table 156: DAC response . . . . .	298
387	Table 157: Packaged DAC response . . . . .	299
388	Table 159: Response headers - Inspect a data object using HTTP . . . . .	333
389	Table 160: HTTP status codes - Inspect a data object using HTTP . . . . .	334
390	Table 161: Request headers - Create a container object using HTTP . . . . .	336
391	Table 162: Response Headers - Inspect a container object using HTTP . . . . .	337
392	Table 163: HTTP status codes - Inspect a container object using HTTP . . . . .	337
393	Table 165: Required metadata for a query queue . . . . .	341

394

## Part I

395

# CDMI Preamble

This Cloud Data Management Interface (CDMI™) International Standard is intended for application developers who are implementing or using cloud storage. It documents how to access cloud storage and to manage the data stored there.

This document is organized as follows:

Table 1: Overview of this document

Clause 1	Scope	Defines the scope of this document
Clause 2	Normative references	Lists the normative references for this document
Clause 3	Terms	Provides terminology used in this document
Clause 4	Conventions	Describes the conventions used in presenting the interfaces and the typographical conventions used in this document
Clause 5	Overview of Cloud Storage	Provides a brief overview of cloud storage and details the philosophy behind this International Standard as a model for the operations
Clause 6	Data Object Resource Operations using HTTP	Provides the normative standard of data object resource operations using HTTP
Clause 7	Container Object Resource Operations using HTTP	Provides the normative standard of container object resource operations using HTTP
Clause 8	Data Object Resource Operations using CDMI	Provides the normative standard of data object resource operations using CDMI
Clause 9	Container Object Resource Operations using CDMI	Provides the normative standard of container object resource operations using CDMI
Clause 10	Domain Object Resource Operations using CDMI	Provides the normative standard of domain object resource operations using CDMI
Clause 11	Queue Object Resource Operations using CDMI	Provides the normative standard of queue object resource operations using CDMI
Clause 12	Capability Object Resource Operations using CDMI	Provides the normative standard of capability object resource operations using CDMI
Clause 13	Exported Protocols	Discusses how virtual machines in the cloud computing environment can use the exported protocols from CDMI containers
Clause 14	Snapshots	Discusses how snapshots are accessed under CDMI containers
Clause 15	Serialization/Deserialization	Discusses serialization and deserialization, including import and export of serialized data under CDMI
Clause 16	Metadata	Provides the normative standard of the metadata used in the interface
Clause 18	Retention and Hold Management	Describes the optional retention management disciplines to be implemented into the system management functions
Clause 19	Scope Specification	Describes the structure of the scope specification for JSON objects
Clause 20	Results Specification	Provides a standardized mechanism to define subsets of CDMI object contents
Clause 21	Notification Queues	Describes how CDMI clients can efficiently discover what changes have occurred to the system
Clause 22	Query Queues	Describes how CDMI clients can efficiently discover what content matches a given set of metadata query criteria or full-content search criteria
Clause 23	Encrypted Objects	Describes how to work with transparently encrypted objects
Clause 24	Delegated Access Control	Describes how to delegate access control to external systems
Clause 25	Data Object Versions	Describes how to work with versioned data objects
Clause 26	Extensions	Provides informative vendor extensions. Each extension is added to the standard when at least two vendors implement the extension.

# 401 **Clause 1**

## 402 **Scope**

403 This CDMI™ International Standard specifies the interface to access cloud storage and to manage the data stored  
404 therein. This International Standard applies to developers who are implementing or using cloud storage.

## Clause 2

## Normative references

The following documents, in whole or in part, are normatively referenced in this document and are indispensable for its application. For dated references, only the edition cited applies. For undated references, the latest edition of the referenced document (including any amendments) applies.

The provisions of the referenced specifications other than ISO/IEC, IEC, ISO, and ITU documents, as identified in this clause, are valid within the context of this International Standard. The reference to such a specification within this International Standard does not give it any further status within ISO/IEC. In particular, it does not give the referenced specifications the status of an International Standard.

ISO 3166:2013, *Codes for the representation of names of countries and their subdivisions (Parts 1, 2 and 3)* - see [36][37][38]

ISO 4217:2015, *Codes for the representation of currencies and funds* - see [39]

ISO 8601:2019, *Data elements and interchange formats – Information interchange – Representation of dates and times* - see [33][34]

ISO/IEC 9594-8:2017, *Information technology – Open Systems Interconnection – The Directory: Public-key and attribute certificate frameworks*

ISO 14721:2012, *Space data and information transfer systems – Open archival information system (OAIS) – Reference model* - see [35]

ISO/IEC 14776-414:2009, *SCSI Architecture Model - 4 (SAM-4)* - see [30]

ISO/IEC 17788:2014, *Information technology – Cloud computing – Overview and vocabulary* - see [32]

ISO/IEC 27040:2015, *Information technology – Security techniques – Storage security* - see [31]

ISO/IEC Directives, Part 2, *Rules for the structure and drafting of International Standards*, 6th edition, 2011\* - see [29]

IEEE 1003.1-2017, *IEEE Standard for Information Technology–Portable Operating System Interface (POSIX(R)) Base Specifications, Issue 7* - see [41]

RFC 1867, *Form-based File Upload in HTML* - see [20]

RFC 2045, *Multipurpose Internet Mail Extensions (MIME) Part One: Format of Internet Message Bodies* - see [9]

RFC 2046, *Multipurpose Internet Mail Extensions (MIME) Part Two: Media Types* - see [10]

RFC 2119, *Key Words for Use in RFCs to Indicate Requirement Levels* - see [3]

RFC 2578, *Structure of Management Information Version 2 (SMIv2)* - see [21]

RFC 2616, *Hypertext Transfer Protocol – HTTP/1.1* - see [23]

RFC 2617, *HTTP Authentication: Basic and Digest Access Authentication* - see [8]

RFC 3280, *Internet X.509 Public Key Infrastructure Certificate and Certificate Revocation List (CRL) Profile* - see [13]

RFC 3530, *Network File System (NFS) Version 4 Protocol* - see [1]

RFC 7143, *Internet Small Computer System Interface (iSCSI) Protocol (Consolidated)* - see [4]

RFC 3986, *Uniform Resource Identifier (URI): Generic Syntax* - see [2]

RFC 4627, *The Application/JSON Media Type for JavaScript Object Notation (JSON)* - see [5]

## Cloud Data Management Interface 2.0.0a

- 441 RFC 4648, *The Base16, Base32, and Base64 Data Encodings* - see [19]
- 442 RFC 4918, *HTTP Extensions for Web Distributed Authoring and Versioning (WebDAV)* - see [6]
- 443 RFC 5246, *The Transport Layer Security (TLS) Protocol Version 1.2* - see [24]
- 444 RFC 6068, *The 'mailto' URI Scheme* - see [27]
- 445 RFC 5652, *Cryptographic Message Syntax (CMS)* - see [12]
- 446 RFC 6208, *Cloud Data Management Interface (CDMI) Media Types* - see [25]
- 447 RFC 6839, *Additional Media Type Structured Syntax Suffixes* - see [11]
- 448 RFC 7515, *JSON Web Signatures* - see [17]
- 449 RFC 7516, *JSON Web Encryption* - see [18]
- 450 RFC 7518, *JSON Web Algorithms* - see [15]
- 451 SNIA TLS, *TLS Specification for Storage Systems, version 1.0.1* - see [42]



## Clause 3

# Terms, acronyms, and definitions

For the purposes of this document, the terms and definitions given in Rec. ITU-T Y.3500 | ISO/IEC 17788:2014 and the following apply.

**Access Control List (ACL)** a persistent list, commonly composed of Access Control Entries (ACEs), that enumerates the rights of principals (users and groups) to access resources

**API** Application Programming Interface

**CDMI™** Cloud Data Management Interface

**CDMI capabilities** an object that describes what operations are supported for a given cloud or cloud object

The mimetype for this object is `application/cdmi-capability`.

**CDMI container** an object that stores zero or more children objects and associated metadata

The mimetype for this object is `application/cdmi-container`.

**CDMI data object** an object that stores an array of bytes (value) and associated metadata

The mimetype for this object is `application/cdmi-object`.

**CDMI domain** an object that stores zero or more children domains and associated metadata describing object administrative ownership

The mimetype for this object is `application/cdmi-domain`.

**CDMI object** one of CDMI capabilities, CDMI container, CDMI data object, CDMI domain, or CDMI queue

**CDMI queue** an object that stores a first-in, first-out set of values and associated metadata

The mimetype for this object is `application/cdmi-queue`.

**CIFS** Common Internet File System

**cloud storage** See Data storage as a Service

**CRC** cyclic redundancy check

**CRUD** create, retrieve, update, delete

**current data object version** the most recent version of a version-enabled data object

**data object version** either the current data object version or an historical data object version

**Data Storage as a Service (DSaaS)** delivery of appropriately configured virtual storage and related data services over a network, based on a request for a given service level

**delegated access control (DAC)** the process of delegating an access control decision to a third party

**delegated access control provider (DAC provider)** a third-party system that is capable of making access control decisions

**delegated access control request (DAC request)** a request made to a DAC provider for an access control decision

**delegated access control response (DAC response)** a response from a DAC provider indicating the result of a request for an access control decision

**domain** a shared user authorization database that contains users, groups, and their security policies and associated accounting information

Each CDMI object belongs to a single domain, and each domain provides user mapping and accounting information.

**eventual consistency** a behavior of transactional systems that does not provide immediate consistency guarantees to provide enhanced system availability and tolerance to network partitioning

**FC** Fibre Channel

**FCoE** Fibre Channel over Ethernet

**historical data object version** a non-current state of a version-enabled data object

**HTTP** HyperText Transfer Protocol

**Infrastructure as a Service (IaaS)** delivery over a network of an appropriately configured virtual computing environment, based on a request for a given service level

Typically, IaaS is either self-provisioned or provisionless and is billed based on consumption.

**intermediary CDMI server** a CDMI server that is capable of forwarding DAC requests and responses

**iSCSI** Internet Small Computer Systems Interface (see RFC 7143 [4])

**JOSE** JavaScript Object Signing and Encryption

**JWA** JSON Web Algorithm

**JWE** JSON Web Encryption

**JWS** JSON Web Signing

**JSON** JavaScript Object Notation

**LDAP** Lightweight Directory Access Protocol

**LUN** Logical Unit Number (see *ISO/IEC 14776-414*)

**metadata** data about other data (see [35])

**MIME** Multipurpose Internet Mail Extensions (see RFC 2045 [9])

**NFS** Network File System (see RFC 3530 [1])

**object** an entity that has an object ID, has a unique URI, and contains state

Types of CDMI objects include data objects, container objects, capability objects, domain objects, and queue objects.

**object identifier** a globally-unique value assigned at creation time to identify an object

**OCCT** Open Cloud Computing Interface (see [40])

**Platform as a Service (PaaS)** delivery over a network of a virtualized programming environment, consisting of an application deployment stack based on a virtual computing environment

Typically, PaaS is based on IaaS, is either self-provisioned or provisionless, and is billed based on consumption.

**POSIX** Portable Operating System Interface (see *IEEE Std 1003.1*)

**private cloud** delivery of SaaS, PaaS, IaaS, and/or DaaS to a restricted set of customers, usually within a single organization

Private clouds are created due to issues of trust.

**public cloud** delivery of SaaS, PaaS, IaaS, and/or DaaS to, in principle, a relatively unrestricted set of customers

**Representational State Transfer (REST)** a specific set of principles for defining, addressing, and interacting with resources addressable by URIs (see [7])

**RPO** recovery point objective

**RTO** recovery time objective

**service level** performance targets for a service

**SNMP** Simple Network Management Protocol

530 **Software as a Service (SaaS)** delivery over a network, on demand, of the use of an application  
531       technology that allocates the physical capacity of a volume or file system as applications write data, rather than  
532       pre-allocating all the physical capacity at the time of provisioning

533 **Uniform Resource Identifier (URI)** compact sequence of characters that identifies an abstract or physical resource  
534       (see RFC 3986 [2])

535 **version-enabled data object** a CDMI data object with versioning enabled

536 **VIM** Vendor Interface Module

537 **virtualization** presentation of resources as if they are physical, when in fact, they are decoupled from the underlying  
538       physical resources

539 **WebDAV** Web Distributed Authoring and Versioning (see RFC 4918 [6])

540 **XAM** eXtensible Access Method (see [INCITS 464-2010]\_)

541 **Clause 4**

542 **Conventions**

543 **4.1 Interface format**

544 Each interface description has nine components, as described in [Table 2](#).

545 Table 2: Interface format

546

Component	Description
Synopsis	The GET, PUT, POST, and DELETE semantics
Delayed completion	For long-running operations, a description of behavior when the operation does not immediately complete
Capabilities	A description of the supported operations
Request headers	The request headers, such as Accept, Authorization, Content-Length, Content-Type
Request message body	A description of the message body contents
Response headers	The response headers, such as Content-Length, Content-Type
Response message body	A description of the message body contents
Response Status	A list of HTTP status codes
Example	An example of the operation

## 547 4.2 Typographical conventions

548 All code text and HTTP status codes are shown in a fixed-width font, as follows:

```
PUT /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-object
Content-Type: application/cdmi-object

{
  "mimetype" : "text/plain",
  "metadata" : {

  },
  "value" : "This is the Value of this Data Object"
}
```

549 Requesting an optional field that is not present shall result in an HTTP status code of 404 Not Found.

## 4.3 Request and response body requirements

In request and response body tables, the Requirement column contains one of the following three values:

- **Mandatory.** The value specified in this row shall be provided.
- **Conditional.** If the conditions specified in the Description cell is met, the value specified in this row shall be provided. Otherwise, it can be provided unless the Description specifically prohibits it, in which case it shall not be provided.
- **Optional.** The value specified in this row can be provided.

## 4.4 Key Word requirements

In this International Standard, the key words in Table 3 shall be interpreted as described in ISO/IEC Directives, Part 2.

Table 2 — Key word requirements

Table 3: Key word requirements

Key words	Denotes	Description	Equivalent expressions (for exception cases only)
shall	requirement	An action that is unconditionally required. <ul style="list-style-type: none"> <li>Do not use must as an alternative to shall.</li> <li>To express a direct instruction, for example, when referring to steps to be taken in a test method, use the imperative mood in English.</li> <li>EXAMPLE: Switch on the recorder.</li> </ul>	<ul style="list-style-type: none"> <li>is to</li> <li>is required to</li> <li>it is required that</li> <li>has to</li> <li>only ... is permitted</li> <li>it is necessary</li> </ul>
shall not	requirement	An action that is unconditionally prohibited. Do not use may not instead of shall not to express a prohibition.	<ul style="list-style-type: none"> <li>is not allowed [permitted] [acceptable] [permissible]</li> <li>is required to be not</li> <li>is required that ... be not</li> <li>is not to be</li> </ul>
should	recommendation	An action that is recommended when choosing among several possibilities, or an action that is preferred but not necessarily required.	<ul style="list-style-type: none"> <li>it is recommended that</li> <li>ought to</li> </ul>
should not	recommendation	An action or certain possibility or course of action that is deprecated but not prohibited.	<ul style="list-style-type: none"> <li>it is not recommended that</li> <li>ought not to</li> </ul>
may	permission	An action that indicates what is allowed within the limits of the document. Do not use possible or can in this context. May signifies permission expressed by the document, whereas can refers to the ability of a user of the document or to a possibility open to him or her.	<ul style="list-style-type: none"> <li>is permitted</li> <li>is allowed</li> <li>is permissible</li> </ul>
need not	permission	An action that indicates what is not required within the limits of the document. Do not use impossible in this context.	<ul style="list-style-type: none"> <li>it is not required that</li> <li>no ... is required</li> </ul>

## Clause 5

# Overview of Cloud Storage

### 5.1 Overview

When discussing cloud storage and standards, it is important to distinguish the various resources that are being offered as services. These resources are exposed to clients as functional interfaces (i.e., data paths) and are managed by management interfaces (i.e., control paths). This International Standard explores the various types of interfaces that are part of cloud services today and shows how they are related. This International Standard defines a model for the interfaces that can be mapped to the various cloud services and a model that forms the basis for cloud storage interfaces into the future.

Another important concept in this International Standard is that of metadata. When managing large amounts of data with differing requirements, metadata is a convenient mechanism to express those requirements in such a way that underlying data services can differentiate their treatment of the data to meet those requirements.

The appeal of cloud storage is due to some of the same attributes that define other cloud services: pay as you go, the illusion of infinite capacity (elasticity), and the simplicity of use/management. It is therefore important that any interface for cloud storage support these attributes, while allowing for a multitude of business use cases.

#### 5.1.1 What is Cloud Storage?

The use of the term cloud in describing these new models arose from architecture drawings that typically used a cloud as the icon for a network. The cloud represents any-to-any network connectivity in an abstract way. In this abstraction, the network connectivity in the cloud is represented without concern for how it is made to happen.

The cloud abstraction of complexity produces a simple base on which other features can be built. The general cloud model extends this base by adding a pool of resources. An important part of the cloud model is the concept of a pool of resources that is drawn from, on demand, in small increments. A relatively recent innovation that has made this possible is virtualization.

Thus, cloud storage is simply the delivery of virtualized storage on demand. The formal term that is used for this is Data storage as a Service (DaaS).

#### 5.1.2 Data Storage as a Service

By abstracting data storage behind a set of service interfaces and delivering it on demand, a wide range of actual cloud services and implementations are possible. The only type of storage that is excluded from this definition is that which is delivered in fixed-capacity increments instead of that which is based on demand.

An important part of any DaaS system is the support of legacy clients. Support is accommodated with existing standard protocols such as iSCSI (and others) for block network storage and CIFS/NFS or WebDAV for file network storage, as shown in Fig. 1.

The difference between purchasing a dedicated appliance or purchasing cloud storage is not the functional interface, but the fact that the storage is delivered on demand. Customers pay for either what they actually use or what they have allocated for use. For block storage, a Logical Unit Number (LUN), or virtual volume, is the granularity of allocation. For file protocols, a file system is the unit of granularity. In either case, the actual storage space may be thin-provisioned






Figure 1: Existing data storage interface standards

and billed for based on actual usage. Data services, such as compression and deduplication, can be used to further reduce the actual space consumed.

Managing this storage is typically done out of band for these standard data storage interfaces, either through an API, or more commonly, through an administrative browser-based user interface. This out-of-band interface can be used to invoke other data services as well (e.g., snapshots or cloning).

In this model, the underlying storage space that has been exposed by the out-of-band interfaces is abstracted and exposed using the notion of a container. A container is not only a useful abstraction for storage space, but also serves as a grouping of the data stored in it and a point of control for applying data services in the aggregate.

Each data object is created, retrieved, updated, and deleted as a separate resource. In this type of interface, a container, if used, is a simple grouping of data objects for convenience. Nothing prevents the concept of containers from being hierarchical, although any given implementation might support only a single level (see [Fig. 2](#)).

### 5.1.3 Data management for cloud storage

Many of the initial implementations of cloud storage focused on a kind of best effort quality of storage service and ignored most other types of data services. To address the needs of enterprise applications with cloud storage, however, there is an increasing need to offer better quality of service and to deploy additional data services.

Cloud storage can lose its abstraction and simplicity benefits if new data services that require complex management are added. Cloud storage customers are likely to resist new demands on their time (e.g., setting up backup schedules through dedicated interfaces, deploying data services individually for stored objects).

By supporting metadata in a cloud storage interface and prescribing how the storage system and data system metadata is interpreted to meet the requirements of the data, the simplicity required by the cloud storage model can be maintained while still addressing the requirements of enterprise applications and their data.

User metadata is retained by the cloud and can be used to find the data objects and containers by performing a query for specific metadata values. The schema for this metadata may be determined by each application, domain, or user. For more information on support for user metadata, see [16.5](#).

Storage system metadata is produced/interpreted by the cloud service provider and basic storage functions (e.g., modification and access statistics, access control). For more information on support for storage system metadata, see [16.2](#).

Data system metadata is interpreted by the cloud service provider as data requirements that control the operation of underlying data services for that data. Depending on the level of granularity supported by the cloud, data system




Figure 2: Storage interfaces for object storage client data

625 metadata may apply to an aggregation of data objects in a container or to individual data objects, if the cloud service  
626 provider supports this level of granularity. For more information on support for data system metadata, see [16.3](#).

#### 627 5.1.4 Data and container management

628 There is no reason that managing data and managing containers should involve different interfaces. Therefore, the use  
629 of metadata is extended from applying to individual objects to applying to containers of objects as well. Thus, any data  
630 placed into a container inherits the data system metadata of the container into which it was placed. When creating a new  
631 container within an existing container, the new container would similarly inherit the metadata settings of its parent's data  
632 system metadata. After an object is created, the data system metadata can be overridden at the container or individual  
633 object level, as desired.

634 Even if the provided interface does not support setting metadata on individual objects, metadata can still be applied  
635 to the containers. In such a case, the interface does not provide a mechanism to override metadata that an individual  
636 object inherits from its parent container. For file-based interfaces that support extended attributes (e.g., CIFS, NFSv4),  
637 these extended attributes may be used to specify the data system metadata to override that specified for the container.

## 5.2 Reference model for cloud storage interfaces

The cloud storage reference model is shown in Fig. 3.

Clients acting in the role of using a data storage interface Clients can be inside the storage cloud (i.e., providing storage resources to the cloud as well as consuming them) or outside the storage cloud (i.e., only consuming resources).

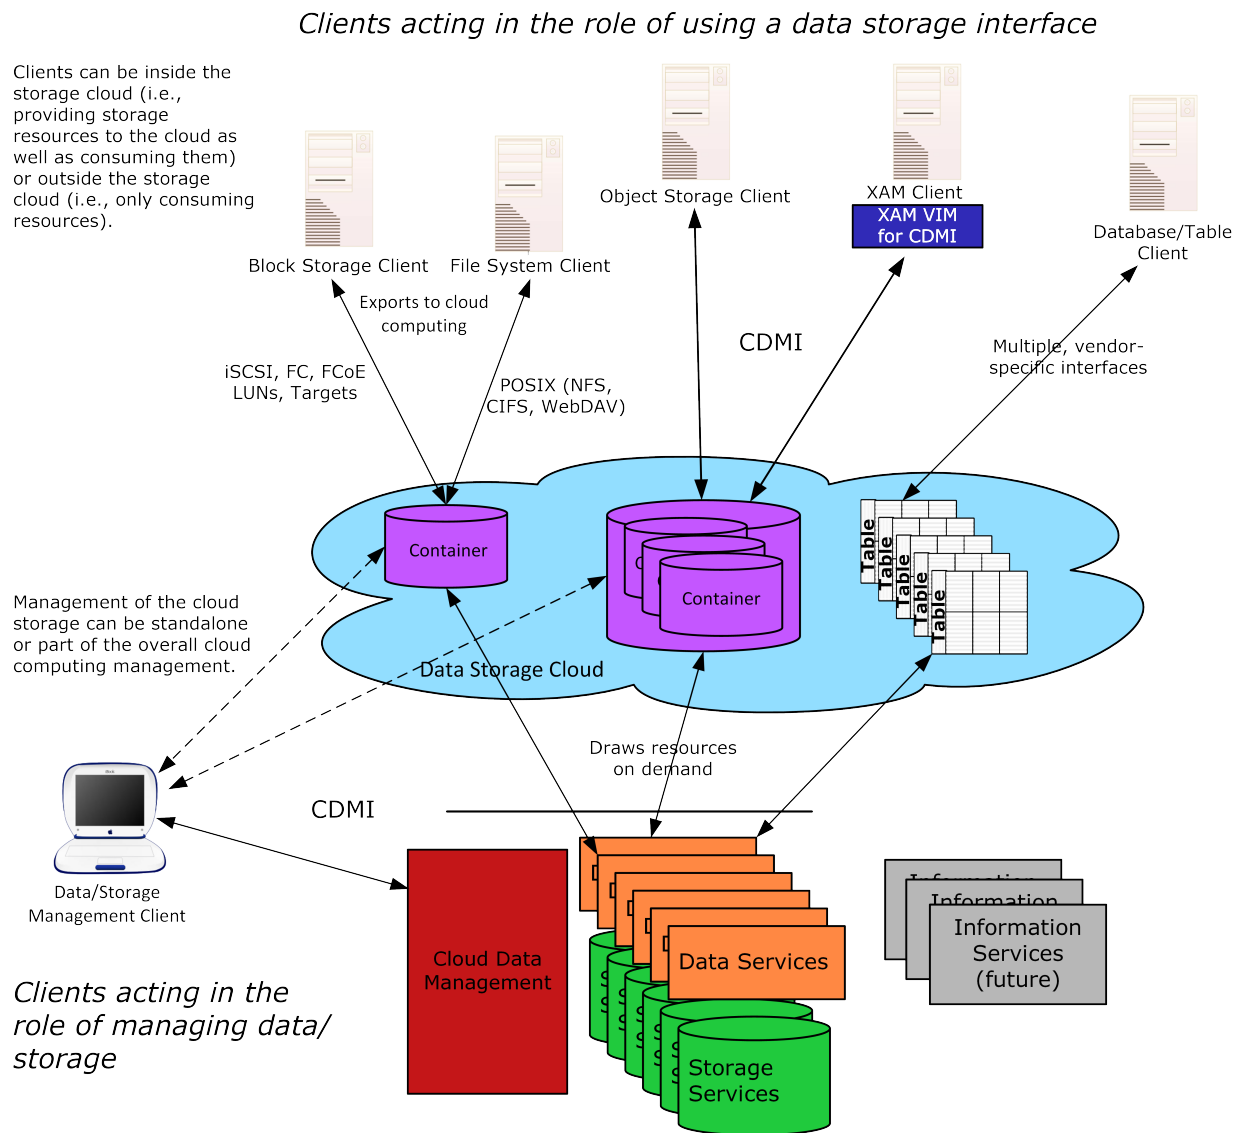


Figure 3: Cloud storage reference model

This model shows multiple types of cloud data storage interfaces that are able to support both legacy and new applications. All of the interfaces allow storage to be provided on demand, drawn from a pool of resources. The storage capacity is drawn from a pool of storage capacity provided by storage services. The data services are applied to individual objects, as determined by the data system metadata. Metadata specifies the data requirements on the basis of individual objects or for groups of objects (containers).

## 5.3 Cloud data management interface

The Cloud Data Management Interface (CDMI™) shown in Fig. 3 may be used to create, retrieve, update, and delete objects in a cloud. The features of the CDMI include functions that:

- allow clients to discover the capabilities available by the cloud service provider,
- manage containers and the data that is placed in them, and
- allow metadata to be associated with containers and the objects they contain.

This International Standard divides operations into two types: those that use a CDMI content type in the HTTP body and those that do not. While much of the same data is available via both types, providing both allows for CDMI-aware clients and non-CDMI-aware clients to interact with a CDMI provider.

CDMI can also be used by administrative and management applications to manage containers, domains, security access, and monitoring/billing information, even for storage that is functionally accessible by legacy or proprietary protocols. The capabilities of the underlying storage and data services are exposed so that clients can understand what services the cloud service provider provides.

Conformant cloud service providers can support a subset of the CDMI, as long as they expose the limitations in the capabilities reported via the interface.

This International Standard uses RESTful principles in the interface design where possible (see [7]).

CDMI defines both a means to manage the data as well as a means to store and retrieve the data. The means by which the storage and retrieval of data is achieved is termed a data path. The means by which the data is managed is termed a control path. CDMI specifies both a data path and control path interface.

CDMI does not need to be used as the only data path and is able to manage cloud storage properties for any data path interface (e.g., standardized or vendor specific).

Container metadata is used to configure the data requirements of the storage provided through the exported protocol (e.g., block protocol or file protocol) that the container exposes. When an implementation is based on an underlying file system to store data for a block protocol (e.g., iSCSI), the CDMI container provides a useful abstraction for representing the data system metadata for the data and the structures that govern the exported protocols.

A cloud service may also support domains that allow administrative ownership to be associated with stored objects. Domains allow this International Standard to (among other things):

- determine how user credentials are mapped to principals used in an Access Control List (ACL),
- allow granting of special cloud-related privileges, and
- allow delegation to external user authorization systems (e.g., LDAP or Active Directory).

Domains may also be hierarchical, allowing for corporate domains with multiple children domains for departments or individuals. The domain concept is also used to aggregate usage data that is used to bill, meter, and monitor cloud use.

Finally, capabilities allow a client to discover the capabilities of a CDMI implementation. Requirements throughout this International Standard shall be understood in the context of CDMI capabilities. Mandatory requirements on functionality that is conditioned on a CDMI capability shall not be interpreted to require implementation of that capability, but rather shall be interpreted to apply only to implementations that support the functionality required by that capability.

For example, in 5.3.3, this International Standard states, “Every cloud storage system shall allow object ID-based access to stored objects.” This requirement shall be understood in the context that access by object ID is predicated on the presence of the `cdmi_object_access_by_ID` capability.

### 5.3.1 Object model for CDMI

The model for CDMI is shown in Fig. 4.

The five types of resources defined are shown in Table 4. The content type in any given operation is specific to each type of resource.

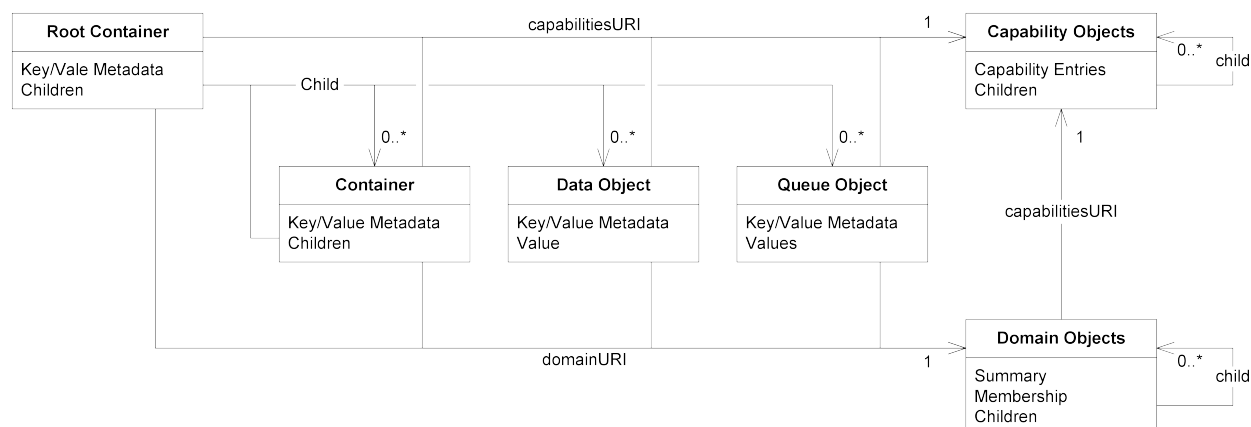


Figure 4: CDMI object model

Table 4: Types of resources in the CDMI object model

Resource type	Description	Reference
Data objects	Data objects are used to store values and provide functionality similar to files in a file system.	See <a href="#">clause 8</a> .
Container objects	Container objects have zero or more children, but do not store values. They provide functionality similar to directories in a file system.	See <a href="#">clause 9</a> .
Domain objects	Domain objects represent administrative groupings for user authentication and accounting purposes.	See <a href="#">clause 10</a> .
Queue objects	Queue objects store zero or more values and are accessed in a first-in-first-out manner.	See <a href="#">clause 11</a> .
Capability objects	Capability objects describe the functionality implemented by a CDMI server and are used by a client to discover supported functionality.	See <a href="#">clause 12</a> .

For data storage operations, the client of the interface only needs to know about container objects and data objects. All data path implementations are required to support at least one level of containers (see 5.1.4). Using the CDMI object model (see Fig. 4), the client can send a PUT via CDMI (see Fig. 3) to the new container URI and create a new container with the specified name. Container metadata are optional and are expressed as a series of name-value pairs. After a container is created, a client can send a PUT to create a data object within the newly created container. A subsequent GET will fetch the data object, including the value field.

Queue objects are also defined (see Fig. 4) and provide in-order-first in-first-out access to enqueued objects. More information on queues can be found in clause 11.

CDMI defines two namespaces that can be used to access stored objects, a flat object ID namespace and a hierarchical path-based namespace. Support for objects accessed by object ID is indicated by the system-wide capability `cdmi_object_access_by_ID`, and support for objects accessed by hierarchical path is indicated by the container capability `cdmi_create_dataobject` found on the root container (and any subcontainers).

Objects are created by ID by performing an HTTP POST against a special URI, designated as `/cdmi_objectid/` (see 9.6). Subsequent to creation, objects are modified by performing PUTs using the object ID assigned by the CDMI server, using the `/cdmi_objectid/` URI (see 8.4). The same URI is used to retrieve and delete objects by ID.

Objects are created by name by performing an HTTP PUT to the desired path URI (see 8.2). Subsequent to creation, objects are modified by performing PUTs using the object path specified by the client (see 8.4). The same URI is used to retrieve and delete objects by path.

CDMI defines mechanisms so that objects having only an object ID can be assigned a path location within the hierarchical namespace, and so that objects having both an object ID and path can have their path dropped, such that the object only has an object ID. This function is accomplished by using a “move” modifier to a PUT or POST operation, as shown in Fig. 5.




Figure 5: Object transitions between named and ID-only

### 5.3.2 CDMI metadata

CDMI uses many different types of metadata, including HTTP metadata, data system metadata, user metadata, and storage system metadata.

HTTP metadata is metadata that is related to the use of the HTTP protocol (e.g., Content-Length, Content-Type, etc.). HTTP metadata is not specifically related to this International Standard but needs to be discussed to explain how CDMI uses the HTTP standard.

CDMI data system metadata, user metadata, and storage system metadata is defined in the form of name-value pairs. Vendor-defined data system metadata and storage system metadata names shall begin with the reverse domain name of the vendor.

Data system metadata is metadata that is specified by a CDMI client and is a component of objects. Data system metadata abstractly specifies the data requirements associated with data services that are deployed in the cloud storage system.

User metadata consists of client-defined JSON strings, arrays, and objects that are stored in the metadata field. The namespace used for user metadata names is self-administered (e.g., using the reverse domain name), and user metadata names shall not begin with the prefix “cdmi\_”.

Storage system metadata is metadata that is generated by the storage services in the system (e.g., creation time, size) to provide useful information to a CDMI client.

The matrix of the creation and consumption of storage system metadata is shown in Table 5.

Table 5: Creation/consumption of storage system metadata

	Created by user	Created By system
Consumed by user	User metadata	Storage system metadata
Consumed by system	Data system metadata	N/A

### 5.3.3 CDMI object IDs

Every object stored within a CDMI-compliant system shall have a globally unique object identifier (ID) assigned at creation time. The CDMI object ID is a string with requirements for how it is generated and how it obtains its uniqueness. Each cloud service that implements CDMI shall generate these identifiers such that the probability of conflicting with identifiers generated by other CDMI Servers and the probability of generating an identifier that has already been used is effectively zero.

Every cloud storage system shall allow object ID-based access to stored objects by allowing the object’s ID to be appended to the root container URI. If the data object “MyDataObject.txt”, located in the root container, has an object ID of “00006FFD001001CCE3B2B4F602032653”, the following pair of URIs access the same data object:

- `https://cloud.example.com/root/MyDataObject.txt`
- `https://cloud.example.com/root/cdmi_objectid/00006FFD001001CCE3B2B4F602032653`

If containers are supported, they shall also be accessible by object ID. If the container “MyContainer”, located in the root container, has an object ID of “00006FFD0010AA33D8CEF9711E0835CA”, the following pairs of URIs access the same object:

- `https://cloud.example.com/root/MyContainer/`
- `https://cloud.example.com/root/cdmi_objectid/00006FFD0010AA33D8CEF9711E0835CA/`
- `https://cloud.example.com/root/MyContainer/MyDataObject.txt`
- `https://cloud.example.com/root/cdmi_objectid/00006FFD0010AA33D8CEF9711E0835CA/MyDataObject.txt`

## 5.3.4 CDMI object ID format

The CDMI Server shall create the object ID, which identifies an object. The object ID shall be globally unique and shall conform to the format defined in Table 6. The native format of an object ID is a variable-length byte sequence and shall be a maximum length of 40 bytes. A client should treat object IDs as opaque byte strings. However, the object ID format is defined such that its integrity may be validated, and independent CDMI Servers may assign unique object ID values independently.

Table 6: Object ID format

0	1	2	3	4	5	6	7	8	9	10	...	38	39
Reserved (zero)	Enterprise Number			Reserved (zero)	Length	CRC		Opaque Data					

The fields shown in Table 6 are defined as follows:

- The reserved bytes shall be set to zero.
- The Enterprise Number field shall be the SNMP enterprise number of the offering organization that developed the system that created the object ID, in network byte order. See RFC 2578 [21] and <https://www.iana.org/assignments/enterprise-numbers>. 0 is a reserved value.
- The byte at offset 5 shall contain the full length of the object ID, in bytes.
- The CRC field shall contain a 2-byte (16-bit) CRC in network byte order. The CRC field enables the object ID to be validated for integrity. The CRC field shall be generated by running the algorithm (see ISO 14701:2012) across all bytes of the object ID, as defined by the Length field, with the CRC field set to zero. The CRC function shall have the following fields:
  - Name : “CRC-16”,
  - Width : 16,
  - Poly : 0x8005,
  - Init : 0x0000,
  - RefIn : True,
  - RefOut : True,
  - XorOut : 0x0000, and
  - Check : 0xBB3D.

This function defines a 16-bit CRC with polynomial 0x8005, reflected input, and reflected output. This CRC-16 is specified in ISO 14701:2012.

- Opaque data in each object ID shall be unique for a given Enterprise Number.

The native format for an object ID is binary. When necessary, such as when included in URIs and JSON strings, the object ID textual representation shall be encoded using Base16 encoding rules described in RFC 4648 [19] and shall be case insensitive.

## 5.4 Security

### 5.4.1 Security objectives

Security, in the context of CDMI, refers to the protective measures employed in managing and accessing data and storage. The specific objectives to be addressed by security include providing a mechanism that:

- assures that the communications between a CDMI client and server cannot be read or modified by a third party;
- allows CDMI clients and servers to assure their identity;
- allows control of the actions a CDMI client is permitted to perform on a CDMI server;
- allows records to be generated for actions performed by a CDMI client on a CDMI server;
- protects data at rest;
- eliminates data in a controlled manner; and
- discovers the security capabilities of a particular implementation.

Security measures within CDMI are summarized as:

- transport security,
- user and entity authentication,
- authorization and access controls,
- data integrity,
- data and media sanitization,
- data retention,
- protections against malware,
- data at-rest encryption, and
- security capabilities.

With the exception of both the transport security and the security capabilities, which are mandatory to implement, the security measures can vary significantly from implementation to implementation.

When security is a concern, the CDMI client should begin with a series of security capability lookups (see 12.1.7 to determine the exact nature of the security features that are available. Based on the values of these capabilities, a risk-based decision should be made as to whether the CDMI server should be used. This is particularly true when the data to be stored in the cloud storage is sensitive or regulated in a way that requires stored data to be protected (e.g., encrypted) or handled in a particular manner (e.g., full accountability and traceability of management and access).

### 5.4.2 HTTP security

HTTP is the mandatory transport mechanism for this version of CDMI. It is important to note that HTTP, by itself, offers no confidentiality or integrity protections. As CDMI is built on top of HTTP, HTTP over Transport Layer Security (TLS) (i.e., HTTPS) is the mechanism that is used to secure the communications between CDMI clients and servers.

To ensure both security and interoperability, all CDMI implementations:

- shall implement the TLS protocol as described in the latest version of the “SNIA TLS Specification for Storage Systems” [42]; with a six-month transition period for implementations. The TLS specification is updated when new vulnerabilities are found, and CDMI implementations shall support the latest specification within six months of its publication announcement;
- shall support both HTTP over TLS and HTTP without TLS; and
- shall allow HTTP without TLS to be disabled.

When TLS is used to secure HTTP, the client and server typically perform some form of entity authentication. However, the specific nature of this entity authentication depends on the cipher suite negotiated; a cipher suite specifies the encryption algorithm and digest algorithm to use on a TLS connection. A very common scenario involves using server-side certificates, which the client trusts, as the basis for unidirectional entity authentication. It is possible that mutual authentication involving both client-side and server-side certificates are required.



### 5.4.3 Client Authentication

A CDMI client shall comply with all security requirements for HTTP that apply to clients.

CDMI clients can be responsible for initiating user authentication for each CDMI operation that is performed. The CDMI server functions as the authenticator and receives and validates authentication credentials from the client.

RFC 2616 [23] and RFC 2617 [8] define requirements for HTTP authentication, which generally starts with an HTTP client request. If the client request does not include an "Authorization" header and authentication is required, the server responds with an HTTP status code of 401 `Unauthorized` and a WWW-Authenticate response header. The HTTP client shall then respond with the appropriate Authorization header in a subsequent request. The format of the WWW-Authenticate and Authorization headers varies depending on the type of authentication required.

- HTTP basic authentication involves sending the user name and password in the clear, and it should only be used on a secure network or in conjunction with TLS.
- HTTP digest authentication sends a secure digest of the user name and password (and other information such as a nonce value), and can be used on an insecure network without TLS.
- HTTP status codes of 401 `Unauthorized` should not include a choice of authentication.
- HTTP basic authentication and/or HTTP digest authentication should be implemented.
- Authentication credentials used with one type of HTTP authentication (i.e., basic or digest) should never be subsequently used with the other type of HTTP authentication.

Once a user is authenticated, the provided principal name shall be mapped by the CDMI domain to a domain user (or used directly as the ACE "who" if domains are not supported). This mapping is then used to determine authorization.

A CDMI server typically relies on an authentication service (local and/or external) to validate client credentials. Differing authentication schemes may be supported, including host-based authentication, Kerberos, PKI, or other; the authentication service is beyond the scope of this International Standard.

### 5.4.4 Use of TLS and HTTP

Recommendations for using HTTP and TLS are as follows:

- A client connecting to a CDMI server using TLS should use TCP port 443, and a client connecting without TLS should use TCP port 80.
- A client that fails to connect to a CDMI server on port 443 should retry without TLS on TCP port 80 if their security policy allows it.
- Servers may respond to HTTP requests on port 80 with an HTTP REDIRECT to the appropriate TLS URI (using port 443). Clients should honor such redirects in this situation.

### 5.4.5 Further information

For further information pertaining to storage security techniques, see ISO 14701:2012.

## 5.5 Required HTTP support

### 5.5.1 RFC 2616 support requirements

A conformant implementation of CDMI shall also be a conformant implementation of RFC 2616 [23] (i.e., HTTP 1.1). The subclauses below list the sections of RFC 2616 [23] that shall be supported; however, this list is not comprehensive.

### 5.5.2 Content-Type negotiation

For CDMI operations, media types for CDMI objects are used as defined in RFC 6208 [25]. All CDMI representations follow the rules established for “application/json” as defined in RFC 4627 [5]. The use of the CDMI media types with the “+json” suffix shall be supported as defined in RFC 6839 [11].

A client can optionally supply an HTTP Accept header, as per section 14.1 of RFC 2616 [23]. If a client is restricting the response to a specific CDMI media type, the corresponding media type shall be specified in the Accept header. Otherwise, the Accept header can contain “\*/\*” or a list of media types, or it may be omitted.

If a request body is present, the client shall include a Content-Type header, as per section 14.17 of RFC 2616 [23]. If the client does not provide a Content-Type header when required or provides a media type in the Content-Type header that does not match with the existing resource media type, the server shall return an HTTP status code of 400 Bad Request.

If a response body is present, the server shall provide a Content-Type header.

This International Standard may further qualify content negotiation (e.g., in 9.3, the absence of a Content-Type header has a specific meaning).

### 5.5.3 Range support

The server shall support HTTP Range headers and partial content responses (see Section 14.16 of RFC 2616 [23]).

The values of the chunkrange, valuerange and queuerange fields are formatted based on the HTTP byte-range-respec, as defined in clause 14.16 of RFC 2616 [23].

### 5.5.4 URI escaping

Percent escaping of reserved characters specified in RFC 3986 [2] shall be applied to all text strings used in HTTP request URIs and HTTP header URIs. This includes user-supplied field names, metadata names, data object names, container object names, queue object names, and domain object names when used in HTTP request URIs and HTTP header URIs.

Field names and values shall not be escaped when stored and when sent in request body and response bodies.

A client retrieving a metadata item named “@user” from a container object with the name of “@MyContainer” would perform the following request:

```
GET /%40MyContainer/?objectName;metadata:%40user HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-container
```

The response shall be:

```
HTTP/1.1 200 OK
Content-Type: application/cdmi-container

{
  "objectName": "@MyContainer/",
  "metadata": {
    "@user": "test",
    ...
  }
}
```

## 5.5.5 Use of URIs

The format and syntax of URIs are defined by RFC 3986 [2].

Every CDMI client shall maintain one or more root URIs that each correspond to a root container on the CDMI server. Since all URIs to CDMI containers end in a trailing slash, all root URIs will end in a trailing slash.

All URIs in this International Standard are relative to the root URI unless otherwise noted. As a consequence, the algorithm used for calculating the resolved URI is as described in Section 5.2 of RFC 3986 [2].

Table 7 shows how relative URIs are resolved against root URIs

Table 7: Relative URIs resolved against root URIs

Root URI	+ Relative URI	=> Resolved URI
https://cloud.example.com/	container/testObject	https://cloud.example.com/container/testObject
https://cloud.example.com/	/container/testObject	https://cloud.example.com/container/testObject
https://cloud.example.com/p1/	container/testObject	https://cloud.example.com/p1/ccontainer/testObject
https://cloud.example.com/p1/	/container/testObject	https://cloud.example.com/container/testObject
https://cloud.example.com/p1/p2/	container/testObject	https://cloud.example.com/p1/p2/container/testObject
https://cloud.example.com/p1/p2/	/container/testObject	https://cloud.example.com/container/testObject

This International Standard places no restrictions on root and relative URIs. All of the examples in this specification use a root URI of https://cloud.example.com/ and return absolute path references as shown in the second line of Table 7.

- If the root URI is "/", the container located at the root URI shall omit the `parentID` field and shall return an empty string ("") for the value of the `parentURI` field.
- If the root URI is not "/" and the parent is a CDMI container, the container located at the root URI shall populate `parentID` field with the CDMI object ID of the CDMI container corresponding to the parent path component, and populate the `parentURI` field with the URI of the parent path component.
- If the root URI is not "/" and the parent is not a CDMI container, the container located at the root URI shall omit the `parentID` field, and populate the `parentURI` field with the URI of the parent path component.
- If the root URI is not "/" and the parent is not accessible, the server may omit the `parentID` field and return an empty string ("") for the value of the `parentURI` field.

## 5.5.6 Reserved characters

The name of CDMI data objects, container objects, queue objects, domain objects and capability objects shall not contain the "/" or "?" characters, as these characters are reserved for delimiters.

## 5.6 Time representations

Unless otherwise specified, all date/time values are in the ISO 8601:2004 extended representation ("YYYY-MM-DDThh:mm:ss.ssssssZ"). The full precision shall be specified, the sub-second separator shall be a ".", the "Z" UTC zone indicator shall be included, and all timestamps shall be in UTC time zone. The "YYYY-MM-DDT24:00:00.000000Z" hour shall not be used, and instead, it shall be represented as "YYYY-MM-DDT00:00:00.000000Z".

Unless otherwise specified, all date/time intervals are in the ISO 8601:2004 start date/end date representation ("YYYY-MM-DDThh:mm:ss.ssssssZ/YYYY-MM-DDThh:mm:ss.ssssssZ"). The end date shall be equal to or later than the start date. The full precision shall be specified, the sub-second separator shall be a ".", the "Z" UTC zone indicator shall be included, and all timestamps shall be in UTC time zone. The "YYYY-MM-DDT24:00:00.000000Z" hour shall not be used, and instead, it shall be represented as "YYYY-MM-DDT00:00:00.000000Z".

## 5.7 Backwards compatibility

CDMI client and server implementations shall implement the following measures to ensure backwards compability with earlier versions of this Interational Standard.

See the CDMI 1.1.1 Specification for details on backwards compatiblity specific to the 1.x versions of CDMI.

### 5.7.1 Specification version detection

CDMI 2.x clients shall not include the `X-CDMI-Specification-Version` custom header. When a CDMI 2.x client performs an operation against a CDMI 1.x Server, the absence of this header shall result in an error response from the CDMI 1.x server. The client may use the presence of the `X-CDMI-Specification-Version` header in an error response as an indication to use CDMI 1.x (which mandates the use of this custom header), if supported.

CDMI 2.x servers may use the presence of the `X-CDMI-Specification-Version` custom header from a CDMI 1.x client as an indication to use CDMI 1.x, if supported.

### 5.7.2 JSON value transfer encoding

CDMI 2.x servers may support the “json” value transfer encoding. When a CDMI server supports both CDMI 2.x and CDMI 1.x, data objects with a value transfer encoding of json shall be made accessible to CDMI 1.x clients using a value transfer encoding of UTF-8, with the server adding in the required escaping.

## 5.8 Object references

Object references are URIs within the cloud storage namespace that redirect to another URI within the same or another cloud storage namespace. References are similar to soft links in a file system. The cloud does not guarantee that the referenced URI will be valid after the time of creation.

References are visible as children in a container and are distinguished from non-references in container children listings by the presence of a trailing “?” character added to the reference name. Performing an operation (with the exception of create or delete) to a reference URI will result in an HTTP status code of 302 Found, with the HTTP Location header containing the absolute redirect destination URI that was specified at the time the reference was created. The reference’s destination URI shall not be changed after a reference has been created.

To continue, when CDMI clients receive an HTTP status code of 302 Found, they should retry the operation using the URI contained within the Location header.

A delete operation on a reference URI shall delete the reference. References cannot be updated. To update the destination of a redirect, the client shall first delete the reference and then create a new reference to the desired destination.

EXAMPLE 1: GET to a URI, where the URI is a reference:

```
GET /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Accept: application/cdm-object
```

The following shows the response.

```
HTTP/1.1 302 Found
Location: https://cloud.example.com/MyContainer/MyOtherDataObject.txt
```

References by object ID shall always redirect to a URI that ends with the same object ID as the request URI.

EXAMPLE 2: GET to an object ID URI, where the URI is a reference:

```
GET /cdmi_objectid/00006FFD0010AA33D8CEF9711E0835CA HTTP/1.1
Host: cloud.example.com
Accept: application/cdm-object
```

The following shows the response.

```
HTTP/1.1 302 Found
Location: https://archive.example.com/cdm-objectid/00006FFD0010AA33D8CEF9711E0835CA
```

EXAMPLE 3: PUT to create a reference:

```
PUT /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com Accept: application/cdm-object
Content-Type: application/cdm-object

{
  "reference": "https://cloud.example.com/MyContainer/MyOtherDataObject.txt"
}
```

The following shows the response.

```
HTTP/1.1 201 Created
```

963 **EXAMPLE 4: POST to create a reference:**

```
POST /cdmi_objectid/ HTTP/1.1
Host: cloud.example.com Accept: application/cdmi-object
Content-Type: application/cdmi-object

{
  "reference": "https://cloud.example.com/MyContainer/MyOtherDataObject.txt"</P>
}
```

964 The following shows the response.

```
HTTP/1.1 201 Created
Location: https://cloud.example.com/cdmi_objectid/00007ED90010DF417BAD70A0C7F5CDDA
```

965 **EXAMPLE 5: DELETE to delete a reference:**

```
DELETE /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
```

966 The following shows the response.

```
HTTP/1.1 204 No Content
```

967

## Part II

968

# Basic Cloud Storage



## Clause 6

# Data Object Resource Operations using HTTP

### 6.1 Overview

Data objects is the fundamental storage components within CDMI™, and is analogous to files in a file system.

As CDMI builds on top of, and is compatible with, the HTTP standard (RFC 2616 [23]), this allows unmodified HTTP clients to communicate with a CDMI server. This also allows CDMI operations to coexist with other HTTP-based storage protocols, such as WebDAV, S3, and OpenStack Swift.

A CDMI server differentiates between HTTP and CDMI operations using the standard Content-Type and Accept headers. When CDMI MIME types defined in RFC 6208 [25] are used in these headers, this indicates that CDMI behaviors, as described in [clause 8](#), are used in addition to the standard HTTP behaviors.

In CDMI 1.0.2, basic HTTP operations were described as “Non-CDMI” operations to distinguish them from operations using CDMI MIME types.

A CDMI implementation that supports data objects shall include support for basic data object HTTP operations corresponding with the CDMI capabilities that are published by the implementation. Capabilities allow a client to discover which operations (such as create, update, delete, etc.) are supported and are described in [clause 9](#).

Ciphertext representation of encrypted objects are created, accessed, and updated by explicitly specifying a MIME type “application/cms” or “application/jose+json”. Otherwise, a plaintext representation is created, accessed, and updated. For more details on encrypted updates, see [clause 23](#).

## 6.2 Create a data object using HTTP

### 6.2.1 Synopsis

The following HTTP PUT operation creates a new data object in the specified container:

- PUT <root URI>/<ContainerName>/<DataObjectName>

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate containers that already exist, with one slash (i.e., "/" ) between each pair of container names.
- <DataObjectName> is the name specified for the data object to be created.

After it is created, the data object shall also be accessible at <root URI>/cdmi\_objectid/<objectID>.

### 6.2.2 Capabilities

Capabilities that indicate which operations are supported are shown in Table 8.

Table 8: Capabilities - Create a CDMI data object using HTTP

Capability	Location	Description
cdmi_create_dataobject	Parent Container	Ability to create a new data object
cdmi_create_value_range	System Wide Capability	Ability to create a data object using a specified byte range

### 6.2.3 Request headers

The HTTP request headers for creating a CDMI data object using HTTP are shown in Table 9.

Table 9: Request headers - Create a CDMI data object using HTTP

Header	Type	Description	Requirement
Content-Type	Header string	The content type of the data to be stored as a data object. The value specified in this header shall be converted to lower case and stored in the <code>mimetype</code> field of the CDMI data object. <ul style="list-style-type: none"> <li>• If the <code>Content-Type</code> header includes the charset parameter as defined in RFC 2616 [23] of "utf-8 (e.g., "; charset=utf-8)", the <code>valuetransferencoding</code> field of the CDMI data object shall be set to "utf-8". Otherwise, the <code>valuetransferencoding</code> field of the CDMI data object shall be set to "base64".</li> <li>• If not specified, the <code>mimetype</code> field shall be set to "application/octet-stream".</li> </ul>	Optional
X-CDMI-Partial	Header String	Indicates that the newly created object is part of a series of writes and has not yet been fully created. When set to "true", the <code>completionStatus</code> field shall be set to "Processing". X-CDMI-Partial works across CDMI and non-CDMI operations.	Optional
Content-Range	Header String	A valid ranges-specifier (see RFC 2616 [23] Section 14.35.1)	Optional

### 6.2.4 Request message body

The request message body contains the data to be stored in the value of the data object.

## 6.2.5 Response headers

No response headers are specified.

## 6.2.6 Response message body

No response message body fields are specified.

## 6.2.7 Response status

The HTTP status codes that occur when creating a data object using HTTP are described in [Table 10](#).

Table 10: HTTP status codes - Create a data object using HTTP

HTTP Status	Description
201 Created	The new data object was created.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server.

## 6.2.8 Example

EXAMPLE 1: PUT to the container URI the data object name and contents.

```
PUT /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Content-Type: text/plain;charset=utf-8
Content-Length: 37

This is the Value of this Data Object
```

The following shows the response:

```
HTTP/1.1 201 Created
```

EXAMPLE 2: Put to the container URI to create an encrypted object:

```
PUT /MyContainer/MyEncryptedObject.txt HTTP/1.1
Host: cloud.example.com
Content-Type: application/cms
Content-Length: 1425

<CMS Encrypted Object>
```

The following shows the response:

```
HTTP/1.1 201 Created
```

EXAMPLE 3: PUT to the container URI to create an encrypted object:

The following shows the response:

```
HTTP/1.1 201 Created
```

## 6.3 Read a data object using HTTP

### 6.3.1 Synopsis

The following HTTP GET operations read from an existing data object at the specified URI:

- GET <root URI>/<ContainerName>/<DataObjectName>
- GET <root URI>/cdmi\_objectid/<DataObjectID>

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate containers.
- <DataObjectName> is the name of the data object to be read from.
- <DataObjectID> is the ID of the data object to be read from.

### 6.3.2 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 11](#).

Table 11: Capabilities - Read a CDMI data object using HTTP

Capability	Location	Description
cdmi_read_value	Data Object	Ability to read the value of an existing data object
cdmi_read_value_range	Data Object	Ability to read a sub-range of the value of an existing data object
cdmi_object_access_by_ID	System Wide Capability	Ability to access the object by ID

### 6.3.3 Request header

The HTTP request header for reading a CDMI data object using HTTP is shown in [Table 12](#).

Table 12: Request header - Read a CDMI data object using HTTP

Header	Type	Description	Requirement
Range	Header string	A valid ranges-specifier (see RFC 2616 [23] Section 14.35.1)	Optional
Accept	Header string	<p>“*/*” or a value as described in: 5.5.2.</p> <ul style="list-style-type: none"> <li>If the object has a mimetype of “application/cms” or “application/jose+json”, and the mimetype “application/cms” or “application/jose+json” is included in the Accept header mimetype, the CDMI server shall return the CMS or JOSE value in the response message body.</li> <li>Otherwise, the decrypted plaintext shall be returned in the response message body, along with the encapsulated mimetype in the Content-Type response header. If decryption is not possible, an error result code shall be returned. (See clause 23 – Encrypted Objects)</li> <li>If the Accept header mimetype list includes “*/*” before “application/cms” and/or “application/jose+json”, the server will first try to return the decrypted plaintext, and shall return the CMS or JOSE value when decryption fails.</li> <li>If the Accept header mimetype list excludes “*/*”, decrypted plaintext shall only be returned if the encapsulated mimetype is included in the Accept header mimetype list.</li> </ul>	Optional

### 6.3.4 Request message body

A request body shall not be provided.

### 6.3.5 Response headers

The HTTP response headers for reading a data object using HTTP are shown in Table 13.

Table 13: Response headers - Read a CDMI Data Object using HTTP

Header	Type	Description	Requirement
Content-Type	Header string	The content type returned shall be the mimetype field in the data object.	Mandatory
Location	Header string	The server shall respond with the URI that the reference redirects to if the object is a reference.	Conditional

### 6.3.6 Response message body

When reading a data object using HTTP, the following applies:

- The response message body shall be the contents of the data object’s value field.
- When reading a value, zeros shall be returned for any gaps resulting from non-contiguous writes.

### 6.3.7 Response status

The HTTP status codes that occur when reading a data object using HTTP are described in Table 14.

Table 14: HTTP status codes - Read a CDMI data object using HTTP

HTTP Status	Description
200 OK	The data object content was returned in the response.
206 Partial Content	A requested range of the data object content was returned in the response.
302 Found	The resource is a reference to another resource.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI, or a requested field within the resource was not found.

### 6.3.8 Examples

**EXAMPLE 1: GET to the data object URI to read the value of the data object:**

```
GET /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
```

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: text/plain
Content-Length: 37

This is the Value of this Data Object
```

**EXAMPLE 2: GET to the data object URI to read the first 11 bytes of the value of the data object:**

```
GET /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Range: bytes=0-10
```

The following shows the response.

```
HTTP/1.1 206 Partial Content
Content-Type: text/plain
Content-Range: bytes 0-10/37
Content-Length: 11

This is the
```

**EXAMPLE 3: GET to the data object URI to always return the ciphertext of an encrypted object:**

```
GET /MyContainer/MyEncryptedObject.txt HTTP/1.1
Host: cloud.example.com
Accept: application/cms, application/jose+json
```

The following shows the response:

```
HTTP/1.1 200 OK
Content-Type: application/cms
Content-Length: 1425

<CMS Encrypted Object>
```

**EXAMPLE 4: GET to the data object URI to read the plaintext of an encrypted object, if possible; otherwise, get the ciphertext:**

```
GET /MyContainer/MyEncryptedObject.txt HTTP/1.1
Host: cloud.example.com
Accept: */*, application/cms, application/jose+json
<Header credentials used to authenticate and access the decryptionkey>
```

The following shows the response:

```
HTTP/1.1 200 OK
Content-Type: text/plain
Content-Length: 252

<Decrypted contents of Encrypted Value>
```

1061 **EXAMPLE 5: GET to the data object URI to read the plaintext of an encrypted object:**

```
GET /MyContainer/MyEncryptedObject.txt HTTP/1.1
Host: cloud.example.com
<Header credentials used to authenticate and access the decryption key>
```

1062 **The following shows the response:**

```
HTTP/1.1 200 OK
Content-Type: text/plain
Content-Length: 252

<Decrypted contents of Encrypted Value>
```

## 6.4 Update a data object using HTTP

### 6.4.1 Synopsis

The following HTTP PUT operation updates an existing data object at the specified URI:

- PUT <root URI>/<ContainerName>/<DataObjectName>
- PUT <root URI>/cdmi\_objectid/<DataObjectID>

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate containers.
- <DataObjectName> is the name of the data object to be updated.
- <DataObjectID> is the ID of the data object to be read from.

### 6.4.2 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 15](#).

Table 15: Capabilities - Update a CDMI data object using HTTP

Capability	Location	Description
cdmi_modify_value	Data Object	Ability to modify the value of an existing data object
cdmi_modify_value_range	Data Object	Ability to modify a sub-range of the value of an existing data object
cdmi_object_access_by_ID	System Wide Capability	Ability to access the object by ID

### 6.4.3 Request headers

The HTTP request headers for updating a CDMI data object using HTTP are shown in [Table 16](#).

Table 16: Request headers - Update a CDMI data object using HTTP

Header	Type	Description	Requirement
Content-Type	Header string	<p>The content type of the data to be stored as a data object. The value specified in this header shall be converted to lower case and stored in the <code>mimetype</code> field of the CDMI data object.</p> <ul style="list-style-type: none"> <li>• If the <code>Content-Type</code> header includes the <code>charset</code> parameter as defined in RFC 2616 [23] of “utf-8 (e.g., “; charset=utf-8”), the <code>valuetransferencoding</code> field of the CDMI data object shall be set to “utf-8”. Otherwise, the <code>valuetransferencoding</code> field of the CDMI data object shall be set to “base64”.</li> <li>• If not specified, the existing <code>mimetype</code> field value shall be preserved.</li> </ul>	Optional
Content-Range	Header string	A valid ranges-specifier (see RFC 2616 [23] Section 14.35.1)	Optional

Continued on next page



Table 16 – continued from previous page

Header	Type	Description	Requirement
X-CDMI-Partial	Header string	Indicates that the operation is part of a series of updates and has not yet been fully created. When set to “true”, the <code>completionStatus</code> field shall be set to “Processing”. X-CDMI-Partial works across CDMI and non-CDMI operations.  If the <code>completionStatus</code> field had previously been set to “Processing” by including this header in a create or update, the next update without this field shall change the <code>completionStatus</code> field back to “Complete”.	Optional

## 6.4.4 Request message body

The request message body contains the data to be stored in the value of the data object.

## 6.4.5 Response header

The HTTP response header for updating a data object using HTTP is shown in Table 17.

Table 17: Response header - Update a CDMI data object using HTTP

Header	Type	Description	Requirement
Location	Header string	The server shall respond with the URI to which the reference redirects if the object is a reference.	Conditional

## 6.4.6 Response message body

A response body may be provided as per RFC 2616 [23].

## 6.4.7 Response status

The HTTP status codes that occur when updating a data object using HTTP are described in Table 18.

Table 18: HTTP status codes - Update a CDMI data object using HTTP

HTTP Status	Description
204 No Content	The data object content was returned in the response.
302 Found	The resource is a reference to another resource.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server.

## 6.4.8 Examples

EXAMPLE 1: PUT to the data object URI to update the value of the data object:

```
PUT /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Content-Type: text/plain
Content-Length: 37
```

This is the value of this data object

1091

The following shows the response.

```
HTTP/1.1 204 No Content
```

1092

EXAMPLE 2: PUT to the data object URI to update four bytes within the value of the data object:

```
PUT /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Content-Range: bytes 21-24/37
Content-Type: text/plain
Content-Length: 4
```

that

1093

The following shows the response.

```
HTTP/1.1 204 No Content
```

## 6.5 Delete a data object using HTTP

### 6.5.1 Synopsis

The following HTTP DELETE operations delete an existing data object at the specified URI:

- DELETE <root URI>/<ContainerName>/<DataObjectName>
- DELETE <root URI>/cdmi\_objectid/<DataObjectID>

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate containers.
- <DataObjectName> is the name of the data object to be deleted.
- <DataObjectID> is the ID of the data object to be deleted.

### 6.5.2 Capability

Capabilities that indicate which operations are supported are shown in [Table 19](#).

Table 19: Capabilities - Delete a CDMI data object using HTTP

Capability	Location	Description
cdmi_delete_dataobject	Data Object	Ability to delete an existing data object
cdmi_object_access_by_ID	System Wide Capability	Ability to access the object by ID

### 6.5.3 Request headers

Request headers may be provided as per RFC 2616 [23].

### 6.5.4 Request message body

A request body may be provided as per RFC 2616 [23].

### 6.5.5 Response headers

Response headers may be provided as per RFC 2616 [23].

### 6.5.6 Response message body

A response body may be provided as per RFC 2616 [23].

### 6.5.7 Response status

[Table 20](#) describes the HTTP status codes that occur when deleting a data object using HTTP.

Table 20: HTTP status codes - Delete a CDMI data object using HTTP

HTTP Status	Description
204 No Content	The data object was successfully deleted.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock, has caused a state transition error on the server, or the data object cannot be deleted.

### 6.5.8 Example

EXAMPLE 1: DELETE to the data object URI:

```
DELETE /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

## Clause 7

# Container Object Resource Operations using HTTP

### 7.1 Overview

Container objects is the fundamental grouping mechanism for stored data within CDMI, and is analogous to directories in a file system. Each container object has zero or more child objects.

Following the URI conventions for hierarchical paths, container URIs shall consist of one or more container names that are separated by forward slashes ("/") and that end with a forward slash ("/").

As basic HTTP operations do not use the CDMI MIME types that distinguish data object operations from container object operations, a CDMI implementation shall use the presence or absence of a forward slash at the end of a URI to distinguish between a container object create or a data object create, respectively.

If a basic HTTP read, update, or delete operation is performed against an existing container resource and the trailing slash at the end of the URI is omitted, the server shall respond with an HTTP status code of 301 Moved Permanently. In addition, a Location header containing the URI with the trailing slash added shall be returned.

A CDMI server differentiates between HTTP and CDMI operations using the standard Content-Type and Accept headers. When CDMI MIME types defined in RFC 6208 [25] are used in these headers, this indicates that CDMI behaviors, as described in [Clause 9](#) are used in addition to the standard HTTP behaviors.

A CDMI implementation that supports container objects shall include support for basic container object HTTP operations corresponding with the CDMI capabilities that are published by the implementation. Capabilities allow a client to discover which operations (such as create, update, delete, etc.) are supported and are described in [Clause 12](#).

## 7.2 Create a container object using HTTP

### 7.2.1 Synopsis

To create a new container object, the following request shall be performed:

- PUT <root URI>/<ContainerName>/<ContainerObjectName>/

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate container objects that already exist, with one slash (i.e., "/") between each pair of container object names.
- <ContainerObjectName> is the name specified for the container object to be created.

After it is created, the container object shall also be accessible at <root URI>/cdmi\_objectid/<objectID>/.

The presence of a trailing slash at the end of the HTTP PUT URI indicates that a container object is being created and distinguishes it from a request to create a data object.

### 7.2.2 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 21](#).

Table 21: Capabilities - Create a CDMI container object using HTTP

Capability	Location	Description
cdmi_create_container	Parent Container	Ability to create a new data object

### 7.2.3 Request headers

Request headers can be provided as per RFC 2616 [23].

### 7.2.4 Request message body

A request body shall not be provided.

### 7.2.5 Response headers

Response headers can be provided as per RFC 2616 [23].

### 7.2.6 Response message body

A response body can be provided as per RFC 2616 [23].

### 7.2.7 Response status

[Table 22](#) describes the HTTP status codes that occur when creating a container object using HTTP.

Table 22: HTTP status codes - Create a container object using HTTP

HTTP Status	Description
201 Created	The new container object was created.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server.

## 7.2.8 Example

EXAMPLE 1: PUT to the URI the container object name:

```
PUT /MyContainer/ HTTP/1.1
Host: cloud.example.com
```

The following shows the response.

```
HTTP/1.1 201 Created
```

### 7.3 Read a container object using HTTP

Reading a container object using HTTP is not defined by this version of this International Standard. A server is allowed to implement responses such as an Apache directory listing or an S3-style bucket listing.

To read a container object using CDMI, see [9.3](#).



## 7.4 Update a container object using HTTP

1174

1175

Updating a container object using HTTP is not defined by this version of this International Standard.

1176

To update a container object using CDMI, see [9.4](#).

## 7.5 Delete a container object using HTTP

### 7.5.1 Synopsis

The following HTTP DELETE operations delete an existing container object at the specified URI, including all contained children and snapshots:

- DELETE <root URI>/<ContainerName>/<ContainerObjectName>/
- DELETE <root URI>/cdmi\_objectid/<ContainerObjectID>

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate container objects.
- <ContainerObjectName> is the name of the container object to be deleted.
- <ContainerObjectID> is the ID of the container object to be deleted.

### 7.5.2 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 23](#).

Table 23: Capabilities - Delete a CDMI container object using HTTP

Capability	Location	Description
cdmi_delete_container	Parent Container	Ability to delete an existing container object
cdmi_object_access_by_ID	System Wide Capability	Ability to access the object by ID

### 7.5.3 Request headers

Request headers can be provided as per RFC 2616 [\[23\]](#).

### 7.5.4 Request message body

A request body can be provided as per RFC 2616 [\[23\]](#).

### 7.5.5 Response headers

Response headers can be provided as per RFC 2616 [\[23\]](#).

### 7.5.6 Response message body

A response body can be provided as per RFC 2616 [\[23\]](#).

### 7.5.7 Response status

[Table 24](#) describes the HTTP status codes that occur when deleting a container object using HTTP.

Table 24: HTTP status codes - Delete a CDMI container object using HTTP

HTTP Status	Description
204 No Content	The container object was successfully deleted.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server.

## 7.5.8 Example

EXAMPLE 1: DELETE to the container object URI:

```
DELETE /MyContainer/ HTTP/1.1
Host: cloud.example.com
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

## 7.6 Create (POST) a new data object using HTTP

### 7.6.1 Synopsis

To create a new data object in a specified container where the name of the data object is a server-assigned object identifier, the following request shall be performed:

```
POST <root URI>/<ContainerName>/
```

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate container objects that already exist, with one slash (i.e., "/") between each pair of container object names.

The data object shall be accessible as a child of the container with a server-assigned name and shall also be accessible at <root URI>/cdmi\_objectid/<objectID>.

HTTP POST to a container is used to enable CDMI servers to support RFC 1867 [20] form-based file uploading. When implementing RFC 1867 [20], the CDMI server-assigned name can be set to, or derived from, the user-provided file name.

### 7.6.2 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 25](#).

Table 25: Capabilities - Create a CDMI data object using HTTP POST

Capability	Location	Description
cdmi_create_dataobject cdmi_post_dataobject	Parent Container	Ability to create a new data object
cdmi_post_dataobject_by_ID	System Wide Capability	Ability to create a data object in "/cdmi_objectid/"
cdmi_create_value_range	System Wide Capability	Ability to create a data object using a specified byte range
cdmi_create_value_range_by_ID	System Wide Capability	Ability to create a data object in "/cdmi_objectid/" using a specified byte range
cdmi_multipart_mime	System Wide Capability	Ability to create a data object using multi-part MIME

### 7.6.3 Request headers

The HTTP request header for creating a new CDMI data object using HTTP is shown in [Table 26](#).

Table 26: Request header - Create a new data object using HTTP

Header	Type	Description	Requirement
Content-Type	Header String	<p>The content type of the data to be stored as a data object. The value specified here shall be converted to lower case and stored in the <code>mimetype</code> field of the CDMI data object.</p> <ul style="list-style-type: none"> <li>If the content type includes the charset parameter as defined in RFC 2616 [23] of “utf-8 (e.g., “; charset=utf-8”), the <code>valuetransferencoding</code> field of the CDMI data object shall be set to “utf-8”. Otherwise, the <code>valuetransferencoding</code> field of the CDMI data object shall be set to “base64”.</li> <li>If not specified, the <code>mimetype</code> field shall be set to “application/octet-stream”.</li> </ul>	Optional
X-CDMI-Partial	Header String	<p>Indicates that the newly created object is part of a series of writes and has not yet been fully created. When set to “true”, the <code>completionStatus</code> field shall be set to “Processing”. X-CDMI-Partial works across CDMI and non-CDMI operations.</p>	Optional

## 7.6.4 Request message body

The message body shall contain the contents (value) of the data object to be created.

## 7.6.5 Response headers

The HTTP response header for creating a new CDMI data object using HTTP is shown in Table 27.

Table 27: Response header - Create a new data object using HTTP

Header	Type	Description	Requirement
Location	Header string	<p>The unique absolute URI for the new data object as assigned by the system.</p> <p>In the absence of file name information from the client, the system shall assign the URI in the form:</p> <p>http://host:port/&lt;root URI&gt;/&lt;ContainerName&gt;/&lt;ObjectID&gt; or https://host:port/&lt;root URI&gt;/&lt;ContainerName&gt;/&lt;ObjectID&gt;.</p>	Mandatory

## 7.6.6 Response message body

A response body can be provided as per RFC 2616 [23].

## 7.6.7 Response status

Table 28 describes the HTTP status codes that occur when creating a new data object using HTTP.

Table 28: HTTP status codes - Create a new data object using HTTP

HTTP Status	Description
201 Created	The new data object was created.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.

## 7.6.8 Examples

EXAMPLE 1: POST to the container object URI the data object contents:

```
POST /MyContainer/ HTTP/1.1
Host: cloud.example.com
Content-Type: text/plain;charset=utf-8

<object contents>
```

The following shows the response.

```
HTTP/1.1 201 Created
Location: https://cloud.example.com/MyContainer/00007ED900104E1D14771DC67C27BF8B
```

1240

## **Part III**

1241

# **CDMI Core**

## Clause 8

# Data Object Resource Operations using CDMI

## 8.1 Overview

Data objects are the fundamental storage component within CDMI™ and are analogous to files within a file system. Each data object has a set of well-defined fields that include:

- a mandatory value,
- mandatory fields generated by the cloud storage system,
- mandatory metadata items generated by the cloud storage system,
- optional metadata generated by the cloud storage system; and
- optional metadata specified by the cloud user.

All cloud storage systems shall support data objects, but the ability to create a data object is determined by the presence or absence of the `cdmi_create_dataobject` and `cdmi_post_dataobject` capabilities in the parent container, and by the `cdmi_post_dataobject_by_ID` system-wide capability for creation by ID.

Each CDMI data object is represented as a JSON object, containing one or more “fields”. For example, the “metadata” field contains metadata items.

EXAMPLE 1: CDMI Data Object

```
{
  "objectType" : "application/cdmi-object",
  "objectID" : "00007ED90010D891022876A8DE0BC0FD",
  "objectName" : "MyDataObject.txt",
  "parentURI" : "/MyContainer/",
  "parentID" : "00007E7F00102E230ED82694DAA975D2",
  "domainURI" : "/cdmi_domains/MyDomain/",
  "capabilitiesURI" : "/cdmi_capabilities/dataobject/",
  "completionStatus" : "Complete",
  "mimetype" : "text/plain",
  "metadata" : {
    "cdmi_size" : "37"
  },
  "valuetransferencoding" : "utf-8",
  "valuerange" : "0-36",
  "value" : "This is the Value of this Data Object"
}
```

The meaning, use, and permitted values of each field is described in each operation that creates, modifies or retrieves CDMI data objects.



## 8.1.1 Data object addressing

Data objects are addressed in CDMI in two ways:

- by name (e.g. `https://cloud.example.com/dataobject`); and
- by ID (e.g. `https://cloud.example.com/cdmi_objectid/00007ED90010D891022876A8DE0BC0FD`).

Every data object has a single, globally-unique object identifier (ID) that remains constant for the life of the object. Each data object shall have one or more URI addresses that allow the object to be accessed.

## 8.1.2 Data object fields

Individual fields within a data object can be accessed by specifying the field name after a question mark “?” that is appended to the end of the data object URI.

EXAMPLE 2: The following URI returns the value field in the response body:

```
https://cloud.example.com/dataobject?value
```

A list of unique fields, separated by a semicolon “;” can be specified, allowing multiple fields to be accessed in a single request.

EXAMPLE 3: The following URI returns the value and metadata fields in the response body:

```
https://cloud.example.com/dataobject?value;metadata
```

When a client provides fields that are not defined in this International Standard or deserializes an object containing fields that are not defined in this International Standard, these fields shall be persisted, but shall not be interpreted.

## 8.1.3 Data Object Value

The encoding of the data transported in the data object value field depends on the data object `valuetransferencoding` field.

- If the value transfer encoding of the object is set to “`utf-8`”, the data stored in the value of the data object shall be a valid UTF-8 string and shall be transported as a UTF-8 string in the value field.
- If the value transfer encoding of the object is set to “`base64`”, the data stored in the value of the data object can contain arbitrary binary sequences, and it shall be transported as a base 64-encoded string in the value field.
- If the value transfer encoding of the object is set to “`json`”, the data stored in the value of the data object shall contain a valid JSON object, and the value field shall contain a valid JSON object. The JSON stored and returned shall be semantically equivalent but may not be syntactically identical. For example, whitespace outside of JSON-quoted strings may be removed or added by either client libraries or by the server. This means that the number of bytes sent may not be the same as the number of bytes stored.

Specific ranges of the value of a data object can be accessed by specifying a byte range after the value field name.

EXAMPLE 4: The following URI returns the first thousand bytes in the value field:

```
https://cloud.example.com/dataobject?value:0-999
```

Because a byte range of a UTF-8 string is often not a valid UTF-8 string, the response to a range request shall always be transported in the value field as a base 64-encoded string. Likewise, when updating a range of bytes within the value of a data object, the contents of the value field shall be transported as a base 64-encoded string.

Byte ranges are specified as single inclusive byte ranges as per Section 14.35.1 of RFC 2616 [23].

The value of a data object can also be specified and retrieved using multipart MIME, where the CDMI JSON is transferred in the first MIME part, and the raw object value is transferred in the second MIME part. Each MIME part, including any header fields, shall conform to RFC 2045 [9], RFC 2046 [10], and RFC 2616 [23]. The length of each part can optionally be specified by a Content-Length header in addition to the MIME boundary delimiter.

Multiple non-overlapping ranges of the value of a data object can also be accessed or updated in a multipart MIME operation by transferring one MIME part for each range of the value. The byte ranges for these operations shall be specified as per Section 14.35.1 of RFC 2616 [23].

Multipart MIME enables the efficient transfer of binary data alongside CDMI object metadata without incurring the overhead of the UTF-8 or Base64 encoding and validation required to represent binary data in JSON.

### 8.1.4 Data object metadata

Data object metadata can also include arbitrary user-supplied metadata, storage system metadata, and data system metadata, as specified in [clause 16](#). Metadata shall be stored as a valid UTF-8 string. Binary data stored in user metadata shall be first encoded such that it can be contained in a UTF-8 string, with the use of base 64 encoding recommended.

Every data object has a parent object from which the data object inherits data system metadata that is not explicitly specified in the data object itself.

EXAMPLE 5: The “budget.xls” data object stored at the following URI would inherit data system metadata from its parent container, “finance”:

```
https://cloud.example.com/finance/budget.xls
```

### 8.1.5 Data object access control

If read access to any of the requested fields is not permitted by the object ACL, only the permitted fields shall be returned. If no requested fields are permitted to be read, an HTTP status code of 403 *Forbidden* shall be returned to the client.

If write access to any of the requested fields is not permitted by the object ACL, no updates shall be performed, and an HTTP status code of 403 *Forbidden* shall be returned to the client.

### 8.1.6 Data object consistency

Writing to a data object is an atomic operation.

- If a client reads a data object simultaneously with a write to that same data object, the reading client shall get either the old version or the new version, but not a mixture of both.
- If a write is terminated due to errors, the contents of the data object shall be as if the write never occurred (i.e., writes are atomic in the face of errors).

Create and update timestamps that are returned in response to multiple client writes to a given object can indicate that a specific write is the newest (i.e., the write whose data is expected to be returned to subsequent reads until another write is processed). However, there is no guarantee that the write with the latest timestamp is the one whose data is returned on subsequent reads.

Range writes can result in a gap in an object value that have had no data written to them. Reading from a gap in a data object value shall return zero for each byte read.

Implementations of this International Standard shall provide the atomicity features described in this subclause for data objects that are accessed via CDMI. The atomicity properties of data objects that are accessed by protocols other than CDMI are outside the scope of this International Standard.

### 8.1.7 Data object representations

The representations in this clause are shown using JSON notation. Both clients and servers shall support UTF-8 JSON representation. The request and response body JSON fields can be specified or returned in any order, with the exception that, if present, for data objects, the “valuerange” and “value” fields shall appear last and in that order.

### 8.1.8 Encrypted objects

CDMI data object operations only permit management operations and access to the ciphertext of encrypted objects. For more details on encrypted objects, see [clause 23](#).

## 8.2 Create a data object using CDMI

### 8.2.1 Synopsis

To create a new data object, the following request shall be performed:

- PUT <root URI>/<ContainerName>/<DataObjectName>

To create a new data object by ID, see 9.6.

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate containers that already exist, with one slash (i.e., "/") between each pair of container names.
- <DataObjectName> is the name specified for the data object to be created.

After it is created, the data object shall also be accessible at <root URI>/cdmi\_objectid/<objectID>.

### 8.2.2 Delayed completion of create

In response to a create operation for a data object, the server can return an HTTP status code of 202 *Accepted* to indicate that the object is in the process of being created. This response is useful for long-running operations (e.g., copying a large data object from a source URI). Such a response has the following implications.

- The server shall return a Location header with an absolute URI to the object to be created along with an HTTP status code of 202 *Accepted*.
- With an HTTP status code of 202 *Accepted*, the server implies that the following checks have passed:
  - user authorization for creating the object;
  - user authorization for read access to any source object for move, copy, serialize, or deserialize; and
  - availability of space to create the object or at least enough space to create a URI to report an error.
- A client might not be able to immediately access the created object, e.g., due to delays resulting from the implementation's use of eventual consistency.

The client performs GET operations to the URI to track the progress of the operation. In response, the server returns two fields in its response body to indicate progress.

- A mandatory `completionStatus` text field contains either "Processing", "Complete", or an error string starting with the value "Error".
- An optional `percentComplete` field contains the percentage of the operation that has completed (0 to 100).

GET shall not return any value for the data object when `completionStatus` is not "Complete". If the final result of the create operation is an error, the URI is created with the `completionStatus` field set to the error message. It is the client's responsibility to delete the URI after the error has been noted.

## 8.2.3 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 29](#).

Table 29: Capabilities - Create a CDMI data object using CDMI

Capability	Location	Description
cdmi_create_dataobject	Parent Container	Ability to create a new data object
cdmi_create_reference	Parent Container	Ability to create a new reference
cdmi_copy_dataobject	Parent Container	Ability to create a data object that is a copy of another data object
cdmi_move_dataobject	Parent Container	Ability to move a data object from another container
cdmi_deserialize_dataobject	Parent Container	Ability to create a data object that is deserialized from the contents of the PUT or the contents of another data object
cdmi_serialize_dataobject cdmi_serialize_container cdmi_serialize_domain cdmi_serialize_queue	Parent Container	Ability to create a data object that contains a serialized representation of an existing data object, container, domain or queue
cdmi_create_value_range	Parent Container	Ability to create a data object using a specified byte range
cdmi_multipart_mime	System Wide Capability	Ability to create a data object using multi-part MIME

## 8.2.4 Request headers

The HTTP request headers for creating a CDMI data object using CDMI are shown in [Table 30](#).

Table 30: Request headers - Create a CDMI data object using CDMI

Header	Type	Description	Requirement
Accept	Header string	"application/cdmi-object" or a consistent value defined in <a href="#">5.5.2</a>	Optional
Content-Type	Header string	"application/cdmi-object" or "multipart/mixed" <ul style="list-style-type: none"> <li>If "multipart/mixed" is specified, the body shall consist of at least two MIME parts, where the first part shall contain a body of content-type "application/cdmi-object", and the second and subsequent parts shall contain one or more byte ranges of the value.</li> <li>If multiple byte ranges are included and the Content-Range header is omitted for a part, the data in the part shall be appended to the data in the preceding part, with the first part having a byte offset of zero.</li> </ul>	Mandatory
X-CDMI-Partial	Header string	Indicates that the newly created object is part of a series of writes and has not yet been fully created. When set to "true", the completionStatus field shall be set to "Processing". X-CDMI-Partial works across CDMI and non-CDMI operations.	Optional

## 8.2.5 Request message body

The request message body fields for creating a data object using CDMI are shown in [Table 31](#).

Table 31: Request message body - Create a data object using CDMI

Field Name	Type	Description	Requirement
mimetype	JSON string	MIME type of the data contained within the value field of the data object <ul style="list-style-type: none"> <li>This field may be included when creating by value or when deserializing, serializing, copying, and moving a data object.</li> <li>If this field is not included and multi-part MIME is not being used, the value of "text/plain" shall be assigned as the field value.</li> <li>If this field is not included and multi-part MIME is being used, the value of the <code>Content-Type</code> header of the second MIME part shall be assigned as the field value.</li> <li>This field value shall be converted to lower case before being stored.</li> </ul>	Optional
metadata	JSON object	Metadata for the data object <ul style="list-style-type: none"> <li>If this field is included, the contents of the JSON object provided in this field shall be used as data object metadata.</li> <li>If this field is included when deserializing, serializing, copying, or moving a data object, the contents of the JSON object provided in this field shall be used as object metadata instead of the metadata from the source URI.</li> <li>If this field is not included, no user-specified metadata shall be added to the object.</li> <li>If this field is not included when deserializing, serializing, copying, or moving a data object, metadata from the source URI shall be used.</li> <li>This field shall not be included when creating a reference to a data object.</li> </ul>	Optional
domainURI	JSON string	URI of the owning domain <ul style="list-style-type: none"> <li>If different from the parent domain, the user shall have the "cross-domain" privilege (see <code>cdmi_member_privileges</code> in <a href="#">Table 80</a>).</li> <li>If not specified, the domain of the parent container shall be used.</li> </ul>	Optional
deserialize	JSON string	URI of a CDMI data object with a value that contains a data object serialized as specified in <a href="#">clause 15</a> . The serialized data object shall be deserialized to create the new data object.	Optional <sup>1</sup>
serialize	JSON String	URI of a CDMI object that shall be serialized into the new data object	Optional <sup>1</sup>

Continued on next page

Table 31 – continued from previous page

Field Name	Type	Description	Requirement
copy	JSON string	<p>URI of a source CDMI data object or queue object that shall be copied into the new destination data object.</p> <ul style="list-style-type: none"> <li>If the destination data object URI and the copy source object URI both do not specify individual fields, the destination data object shall be a complete copy of the source data object.</li> <li>If the destination data object URI or the copy source object URI specifies individual fields, only the fields specified shall be used to create the destination data object. If specified fields are not present in the source, default field values shall be used.</li> <li>If the destination data object URI and the copy source object URI both specify fields, an HTTP status code of <code>400 Bad Request</code> shall be returned to the client.</li> <li>If the copy source object URI points to a queue object, as part of the copy operation, multiple queue values shall be concatenated into a single data object value.</li> <li>If the copy source object URI points to one or more queue object values, as part of the copy operation, the specified queue values shall be concatenated into a single data object value.</li> <li>If there are insufficient permissions to read the data object at the source URI or create the data object at the destination URI, or if the read operation fails, the copy shall return an HTTP status code of <code>400 Bad Request</code>, and the destination object shall not be created.</li> </ul>	Optional <sup>1</sup>
move	JSON string	<p>URI of an existing local or remote CDMI data object (source URI) that shall be relocated to the URI specified in the PUT. The contents of the object, including the object ID, shall be preserved by a move, and the data object at the source URI shall be removed after the data object at the destination has been successfully created.</p> <p>If there are insufficient permissions to read the data object at the source URI, write the data object at the destination URI, or delete the data object at the source URI, or if any of these operations fail, the move shall return an HTTP status code of <code>400 Bad Request</code>, and the source and destination are left unchanged.</p>	Optional <sup>1</sup>
reference	JSON string	URI of a CDMI data object that shall be redirected to by a reference. If any other fields are supplied when creating a reference, the server shall respond with an HTTP status code of <code>400 Bad Request</code> .	Optional <sup>1</sup>
deserializevalue	JSON string	<p>A data object serialized as specified in <a href="#">clause 15</a> and encoded using base 64 encoding rules described in RFC 4648 [19], that shall be deserialized to create the new data object.</p> <ul style="list-style-type: none"> <li>If multi-part MIME is being used and this field contains the value of the MIME boundary parameter, the contents of the second MIME part shall be assigned as the field value.</li> <li>If the serialized data object in the second MIME part does not include a value field, the contents of the third MIME part shall be assigned as the field value of the value field.</li> </ul>	Optional <sup>1</sup>

Continued on next page

Table 31 – continued from previous page

Field Name	Type	Description	Requirement
valuetransferencoding	JSON string	<p>The value transfer encoding used for the data object value. Three value transfer encodings are defined.</p> <ul style="list-style-type: none"> <li>“utf-8” indicates that the data object contains a valid UTF-8 string, and it shall be transported as a UTF-8 string in the value field.</li> <li>“base64” indicates that the data object may contain arbitrary binary sequences, and it shall be transported as a base 64-encoded string in the value field. Setting the contents of the data object value field to any value other than a valid base 64 string shall result in an HTTP status code of 400 <i>Bad Request</i> being returned to the client.</li> <li>“json” indicates that the data object contains a valid JSON object, and the <i>value</i> field shall be a JSON object containing valid JSON data. If the contents of the <i>value</i> field are set to any value other than a valid JSON object, an HTTP status code of 400 <i>Bad Request</i> shall be returned to the client.</li> <li>This field shall only be included when creating a data object by value.</li> <li>If this field is not included and multi-part MIME is not being used, the value of “utf-8” shall be assigned as the field value.</li> <li>If this field is not included and multi-part MIME is being used, the value of “utf-8” shall be assigned as the field value if the Content-Type header of the second and all MIME parts includes the charset parameter as defined in RFC 2046 of “utf-8” (e.g., “; charset=utf-8”). Otherwise, the value of “base64” shall be assigned as the field value. This field applies only to the encoding of the value when represented in CDMI; the <i>Content-Transfer-Encoding</i> header of the part specifies the encoding of the value within a multi-part MIME request, as defined in RFC 2045 [9].</li> </ul>	Optional <sup>1</sup>
value	JSON string	<p>The data object value</p> <ul style="list-style-type: none"> <li>If this field is not included and multi-part MIME is not being used, an empty JSON String (i.e., “”) shall be assigned as the field value.</li> <li>If this field is not included and multi-part MIME is being used, the contents of the second MIME part shall be assigned as the field value.</li> <li>If the <i>valuetransferencoding</i> field indicates UTF-8 encoding, the value shall be a UTF-8 string escaped using the JSON escaping rules described in RFC 4627 [5].</li> <li>If the <i>valuetransferencoding</i> field indicates base 64 encoding, the value shall be first encoded using the base 64 encoding rules described in RFC 4648 [19].</li> <li>If the <i>valuetransferencoding</i> field indicates JSON encoding, the value shall contain a valid JSON object.</li> </ul>	Optional <sup>1</sup>

<sup>1</sup> Only one of these fields shall be specified in any given operation. Except for value, these fields shall not be stored. If more than one of these fields is supplied, the server shall respond with an HTTP status code of 400 *Bad Request*.

## 8.2.6 Response headers

The HTTP response headers for creating a data object using CDMI are shown in Table 32.

Table 32: Response headers - Create a data object using CDMI

Header	Type	Description	Requirement
Content-Type	Header string	"application/cdm-object"	Mandatory
Location	Header string	When an HTTP status code of 202 Accepted is returned, the server shall respond with the absolute URL of the object that is in the process of being created.	Conditional

## 8.2.7 Response message body

The response message body fields for creating a data object using CDMI are shown in Table 33.

Table 33: Response message body - Create a data object using CDMI

Field Name	Type	Description	Requirement
objectType	JSON string	"application/cdm-object"	Mandatory
objectID	JSON string	Object ID of the object	Mandatory
objectName	JSON string	Name of the object	Mandatory
parentURI	JSON string	URI for the parent object. Appending the objectName to the parentURI shall always produce a valid URI for the object.	Mandatory
parentID	JSON string	Object ID of the parent container object	Mandatory
domainURI	JSON string	URI of the owning domain	Mandatory
capabilitiesURI	JSON string	URI to the capabilities for the object	Mandatory
completionStatus	JSON string	A string indicating if the object is still in the process of being created or updated by another operation, and after that operation is complete, indicates if it was successfully created or updated or if an error occurred.  The value shall be the string "Processing", the string "Complete", or an error string starting with the value "Error".	Mandatory
percentComplete	JSON string	A string indicating the percentage of completion if the object is still in the process of being created or updated by another operation. <ul style="list-style-type: none"> <li>When the value of completionStatus is "Processing", this field, if provided, shall indicate the percentage of completion as a numeric integer value from "0" through "100".</li> <li>When the value of completionStatus is "Complete", this field, if provided, shall contain the value "100".</li> <li>When the value of completionStatus is "Error", this field, if provided, can contain any integer value from "0" through "100".</li> </ul>	Optional
mimetype	JSON string	MIME type of the value of the data object	Mandatory

Continued on next page



Table 33 – continued from previous page

Field Name	Type	Description	Requirement
metadata	JSON object	Metadata for the data object. This field includes any user and data system metadata specified in the request body metadata field, along with storage system metadata generated by the cloud storage system. See <a href="#">clause 16</a> for a further description of metadata.	Mandatory

## 8.2.8 Response status

The HTTP status codes that occur when creating a data object using CDMI are described in [Table 34](#).

Table 34: HTTP status codes - Create a data object using CDMI

HTTP Status	Description
201 Created	The new data object was created.
202 Accepted	The data object is in the process of being created. The CDMI client should monitor the <code>completionStatus</code> and <code>percentComplete</code> fields to determine the current status of the operation.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server.

## 8.2.9 Examples

EXAMPLE 1: PUT to the container URI the data object name and contents:

```
PUT /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Accept: application/cdm-object
Content-Type: application/cdm-object

{
  "mimetype" : "text/plain",
  "metadata" : {

  },
  "value" : "This is the Value of this Data Object"
}
```

The following shows the response.

```
HTTP/1.1 201 Created
Content-Type: application/cdm-object

{
  "objectType" : "application/cdm-object",
  "objectID" : "00007ED90010D891022876A8DE0BC0FD",
  "objectName" : "MyDataObject.txt",
  "parentURI" : "/MyContainer/",
  "parentID" : "00007E7F00102E230ED82694DAA975D2",
  "domainURI" : "/cdmi_domains/MyDomain/",
  "capabilitiesURI" : "/cdmi_capabilities/dataobject/",
  "completionStatus" : "Complete",
  "mimetype" : "text/plain",
  "metadata" : {
    "cdmi_size" : "37"
  }
}
```

1393 **EXAMPLE 2: PUT to the container URI the data object name and binary contents:**

```
PUT /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-object
Content-Type: application/cdmi-object

{
  "mimetype" : "text/plain",
  "metadata" : { },
  "valuetransferencoding" : "base64"
  "value" : "VGhpcyBpcyB0aGUgVmFsdWUgb2YgdGhpcyBEYXRhIE9iamVjdA=="
}
```

1394 **The following shows the response.**

```
HTTP/1.1 201 Created
Content-Type: application/cdmi-object

{
  "objectType": "application/cdmi-object",
  "objectID": "00007ED9001008C174ABCE6AC3287E5F",
  "objectName": "MyDataObject.txt",
  "parentURI": "/MyContainer/",
  "parentID" : "00007E7F00102E230ED82694DAA975D2",
  "domainURI": "/cdmi_domains/MyDomain/",
  "capabilitiesURI": "/cdmi_capabilities/dataobject/",
  "completionStatus": "Complete",
  "mimetype": "text/plain",
  "metadata": {
    "cdmi_size": "37"
  }
}
```

1395 **EXAMPLE 3: PUT to the container URI the data object name and binary contents using multi-part MIME:**

```
PUT /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-object
Content-Type: multipart/mixed; boundary=gc0p4Jq0M2Yt08j34c0p

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/cdmi-object

{
  "domainURI": "/cdmi_domains/MyDomain/",
  "metadata": {
    "colour": "blue"
  }
}

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/octet-stream
Content-Transfer-Encoding: binary

<37 bytes of binary data>

--gc0p4Jq0M2Yt08j34c0p--
```

1396 **The following shows the response.**

```
HTTP/1.1 201 Created
Content-Type: application/cdmi-object

{
  "objectType": "application/cdmi-object",
  "objectID": "00007ED900103ADE9DE3A8D1CF5436A3",
  "objectName": "MyDataObject.txt",
  "parentURI": "/MyContainer/",
  "parentID" : "00007E7F00102E230ED82694DAA975D2",
```

(continues on next page)

(continued from previous page)

```

    "domainURI": "/cdmi_domains/MyDomain/",
    "capabilitiesURI": "/cdmi_capabilities/dataobject/",
    "completionStatus": "Complete",
    "mimetype": "application/octet-stream",
    "metadata": {
        "cdmi_size": "37",
        "colour": "blue",
        ...
    }
}

```

EXAMPLE 4: PUT to the container URI the data object name and binary contents using multi-part MIME with optional content-lengths for the parts:

```

PUT /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-object
Content-Type: multipart/mixed; boundary=gc0p4Jq0M2Yt08j34c0p

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/cdmi-object
Content-Length: 82

{
    "domainURI": "/cdmi_domains/MyDomain/",
    "metadata": {
        "colour": "blue"
    }
}

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/octet-stream
Content-Transfer-Encoding: binary
Content-Length: 37

<37 bytes of binary data>

--gc0p4Jq0M2Yt08j34c0p--

```

The following shows the response.

```

HTTP/1.1 201 Created
Content-Type: application/cdmi-object

{
    "objectType": "application/cdmi-object",
    "objectID": "00007ED900103ADE9DE3A8D1CF5436A3",
    "objectName": "MyDataObject.txt",
    "parentURI": "/MyContainer/",
    "parentID": "00007E7F00102E230ED82694DAA975D2",
    "domainURI": "/cdmi_domains/MyDomain/",
    "capabilitiesURI": "/cdmi_capabilities/dataobject/",
    "completionStatus": "Complete",
    "mimetype": "application/octet-stream",
    "metadata": {
        "cdmi_size": "37",
        "colour": "blue",
    }
}

```

EXAMPLE 5: PUT to the container URI the data object name and JSON contents:

```

PUT /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-object
Content-Type: application/cdmi-object

{

```

(continues on next page)

(continued from previous page)

```

"mimetype" : "text/plain",
"metadata" : { },
"valuetransferencoding" : "json"
"value" : {
  "test" : "value"
}
}

```

The following shows the response.

```

HTTP/1.1 201 Created
Content-Type: application/cdm-object

{
  "objectType": "application/cdm-object",
  "objectID": "0000706D0010374085EF1A5C7018D774",
  "objectName": "MyDataObject.txt",
  "parentURI": "/MyContainer/",
  "parentID" : "00007ED90010067404EDED32860C086A",
  "domainURI": "/cdmi_domains/MyDomain/",
  "capabilitiesURI": "/cdmi_capabilities/dataobject/",
  "completionStatus": "Complete",
  "mimetype": "text/plain",
  "metadata": {
    "cdmi_size": "21"
  }
}

```

EXAMPLE 6: PUT to the container URI to create an encrypted object:

```

PUT /MyContainer/MyEncryptedObject.txt HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdm-object

{
  "mimetype" : "application/cms",
  "metadata" : {
    "cdmi_enc_key_id" : "testkey"
  },
  "valuetransferencoding" : "base64"
  "value" : "<CMS Encrypted Object in Base64>"
}

```

The following shows the response:

```

HTTP/1.1 201 Created

```

EXAMPLE 7: PUT to the container URI to create an encrypted object:

```

PUT /MyContainer/MyEncryptedObject2.txt HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdm-object

{
  "mimetype" : "application/jose+json",
  "metadata" : {
    "cdmi_enc_key_id" : "77c7e2b8-6e13-45cf-8672-617b5b45243a"
  },
  "valuetransferencoding" : "json",
  "value" : {
    "protected": "eyJhbGciOiJkaXIiLCJraWQiOiI3N2M3ZTJi
      OC02ZTEzLTQ1Y2YtODY3Mi02MTdiNWlONTI0
      M2EiLCJlbmMiOiJBMTI4R0NNIn0",
    "iv": "refa467QzzKx6QAB",
    "ciphertext": "JW_i_f52hww_ELQPGaYyeAB6HYGcr55919T
      YnSovc23XJoBcW29rHP8yZOZG7YhLpT1bjF
      uvZPjQS-m0IFtVcXkZXdh_lr_FrdYt9HRUY
      kshtRmMIUayGmUnd9zMDB2n0cRDIHAzFVeJ

```

(continues on next page)

(continued from previous page)

```
UDxkUwVAE7_YGRPdcqMyiBoCO-FBdE-Nceb
4h3-FtBP-c_BIwCPTjb9o0SbdcdREEMJMyZ
BH8ySWMvilgPD9yxi-aQpGbSv_F9N4IZAxs
cj5g-NJsUPbjk29-s7LJAGb15wEBtXphVCg
yy53CoIKLHHeJHXex45Uz9aKZSRSInZI-wj
sY0yu3cT4_aQ3i1o-tiE-F8Ios61EKgyIQ4
CWao8PFMj8TTnp",
"tag": "vbb32Xv1lea2OtmHAdccRQ",
"cty": "text/plain"
}
}
```

1405

The following shows the response:

```
HTTP/1.1 201 Created
```

## 8.3 Read a data object using CDMI

### 8.3.1 Synopsis

To read an existing data object, the following requests shall be performed:

- GET <root URI>/<ContainerName>/<DataObjectName>
- GET <root URI>/<ContainerName>/<DataObjectName>?<fieldname>;<fieldname>;...
- GET <root URI>/<ContainerName>/<DataObjectName>?value:<range>;...
- GET <root URI>/<ContainerName>/<DataObjectName>?metadata:<prefix>;...
- GET <root URI>/cdmi\_objectid/<DataObjectID>
- GET <root URI>/cdmi\_objectid/<DataObjectID>?<fieldname>;<fieldname>;...
- GET <root URI>/cdmi\_objectid/<DataObjectID>?value:<range>;...
- GET <root URI>/cdmi\_objectid/<DataObjectID>?metadata:<prefix>;...

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate containers.
- <DataObjectName> is the name of the data object to be read from.
- <fieldname> is the name of a field.
- <range> is a byte range of the data object value to be returned in the value field.
- <prefix> is a matching prefix that returns all metadata items that start with the prefix value.
- <DataObjectID> is the ID of the data object to be read from.

### 8.3.2 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 35](#).

Table 35: Capabilities - Read a CDMI data object using CDMI

Capability	Location	Description
cdmi_read_metadata	Data Object	Ability to read the metadata of an existing data object
cdmi_read_value	Data Object	Ability to read the value of an existing data object
cdmi_read_value_range	Data Object	Ability to read a sub-range of the value of an existing data object
cdmi_multipart_mime	System Wide Capability	Ability to read a data object using multi-part MIME
cdmi_object_access_by_ID	System Wide Capability	Ability to access the object by ID

### 8.3.3 Request headers

The HTTP request headers for reading a CDMI data object using CDMI are shown in [Table 36](#).

Table 36: Request headers - Read a CDMI data object using CDMI

Header	Type	Description	Requirement
Accept	Header string	"application/cdmi-object", "multipart/mixed", or a consistent value defined in <a href="#">5.5.2</a>	Optional

### 8.3.4 Request message body

A request body shall not be provided.

### 8.3.5 Response headers

The HTTP response headers for reading a data object using CDMI are shown in [Table 37](#).

Table 37: Response headers - Read a CDMI data object using CDMI

Header	Type	Description	Requirement
Content-Type	Header string	<p>“application/cdm-object” or “multipart/mixed”</p> <ul style="list-style-type: none"> <li>If “multipart/mixed”, the body shall consist of at least two MIME parts, where the first part shall contain a body of content-type “application/cdm-object” and the second and subsequent parts shall contain the requested byte ranges of the value.</li> <li>If multiple byte ranges are included and the Content-Range header is omitted for a part, the data in the part shall be appended to the data in the preceding part, with the first part having a byte offset of zero.</li> </ul>	Mandatory
Location	Header string	The server shall respond with the URI that the reference redirects to if the object is a reference.	Conditional

### 8.3.6 Response message body

The response message body fields for reading a CDMI data object using CDMI are shown in [Table 38](#).

Table 38: Response message body - Read a CDMI data object using CDMI

Field Name	Type	Description	Requirement
objectType	JSON string	“application/cdm-object”	Mandatory
objectID	JSON string	Object ID of the object	Mandatory
objectName	JSON string	<p>Name of the object</p> <ul style="list-style-type: none"> <li>For objects in a container, the objectName field shall be returned.</li> <li>For objects not in a container (objects that are only accessible by ID), the “objectName” field does not exist and shall not be returned.</li> </ul>	Conditional
parentURI	JSON string	<p>URI for the parent object</p> <ul style="list-style-type: none"> <li>For objects in a container, the parentURI field shall be returned.</li> <li>For objects not in a container (objects that are only accessible by ID), the “parentURI” field does not exist and shall not be returned.</li> </ul> <p>Appending the “objectName” to the “parentURI” shall always produce a valid URI for the object.</p>	Conditional

Continued on next page

Table 38 – continued from previous page

Field Name	Type	Description	Requirement
parentID	JSON string	Object ID of the parent container object <ul style="list-style-type: none"> <li>For objects in a container, the “parentID” field shall be returned.</li> <li>For objects not in a container (objects that are only accessible by ID), the “parentID” field does not exist and shall not be returned.</li> </ul>	Conditional
domainURI	JSON string	URI of the owning domain	Mandatory
capabilitiesURI	JSON string	URI to the capabilities for the object	Mandatory
completionStatus	JSON string	A string indicating if the object is still in the process of being created or updated by another operation, and after that operation is complete, indicates if it was successfully created or updated or if an error occurred.  The value shall be the string “Processing”, the string “Complete”, or an error string starting with the value “Error”.	Mandatory
percentComplete	JSON string	A string indicating the percentage of completion if the object is still in the process of being created or updated by another operation. <ul style="list-style-type: none"> <li>When the value of completionStatus is “Processing”, this field, if provided, shall indicate the percentage of completion as a numeric integer value from 0 through 100.</li> <li>When the value of completionStatus is “Complete”, this field, if provided, shall contain the value “100”.</li> <li>When the value of completionStatus is “Error”, this field, if provided, may contain any integer value from “0” through “100”.</li> </ul>	Optional
mimetype	JSON string	MIME type of the value of the data object	Mandatory
metadata	JSON object	Metadata for the data object. This field includes any user and data system metadata specified in the request body metadata field, along with storage system metadata generated by the cloud storage system.  See <a href="#">clause 16</a> for a further description of metadata.	Mandatory
valuerange	JSON string	The range of bytes of the data object to be returned in the value field <ul style="list-style-type: none"> <li>If a specific value range has been requested, the valuerange field shall correspond to the bytes requested. If the request extends beyond the end of the value, the valuerange field shall indicate the smaller byte range returned.</li> <li>If the object value has gaps (due to PUTs with non-contiguous value ranges), the value range will indicate the range to the first gap in the object value.</li> <li>The cdmi_size storage system metadata of the data object shall always indicate the complete size of the object, including zero-filled gaps.</li> </ul>	Mandatory

Continued on next page



Table 38 – continued from previous page

Field Name	Type	Description	Requirement
valuetransferencoding	JSON string	The value transfer encoding used for the data object value. Three value transfer encodings are defined: <ul style="list-style-type: none"> <li>“utf-8” indicates that the data object contains a valid UTF-8 string, and it shall be transported as a UTF-8 string in the value field.</li> <li>“base64” indicates that the data object may contain arbitrary binary sequences, and it shall be transported as a base 64-encoded string in the value field.</li> <li>“json” indicates that the data object contains a valid JSON object, and the value field shall contain a valid JSON object.</li> </ul>	Mandatory
value	JSON string	The data object value <ul style="list-style-type: none"> <li>If the valuetransferencoding field indicates UTF-8 encoding, the value field shall contain a UTF-8 string using JSON escaping rules described in RFC 4627 [5].</li> <li>If the valuetransferencoding field indicates base 64 encoding, the value field shall contain a base 64-encoded string as described in RFC 4648 [19].</li> <li>If the valuetransferencoding field indicates JSON encoding, the value field shall contain a valid JSON object.</li> <li>The value field shall not be provided when using multi-part MIME.</li> <li>The value field shall only be provided when the completionStatus field contains “Complete”.</li> <li>When reading a value, zeros shall be returned for any gaps resulting from non-contiguous writes.</li> </ul>	Conditional

If individual fields are specified in the GET request, only these fields are returned in the result body. Optional fields that are requested but do not exist are omitted from the result body.

### 8.3.7 Response status

The HTTP status codes that occur when reading a data object using CDMI are described in Table 39.

Table 39: HTTP status codes - Read a CDMI data object using CDMI

HTTP Status	Description
200 OK	The data object content was returned in the response.
202 Accepted	The data object is in the process of being created. The CDMI client should monitor the completionStatus and percentComplete fields to determine the current status of the operation.
302 Found	The resource is a reference to another resource.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
406 Not Acceptable	The server is unable to provide the object in the specified in the Accept header.

### 8.3.8 Examples

EXAMPLE 1: GET to the data object URI to read all fields of the data object:

```
GET /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Accept: application/cdm-object
```

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: application/cdm-object

{
  "objectType" : "application/cdm-object",
  "objectID" : "00007ED90010D891022876A8DE0BC0FD",
  "objectName" : "MyDataObject.txt",
  "parentURI" : "/MyContainer/",
  "parentID" : "00007E7F00102E230ED82694DAA975D2",
  "domainURI" : "/cdm_domains/MyDomain/",
  "capabilitiesURI" : "/cdm_capabilities/dataobject/",
  "completionStatus" : "Complete",
  "mimetype" : "text/plain",
  "metadata" : {
    "cdm_size" : "37"
  },
  "valuerange" : "0-36",
  "valuetransferencoding" : "utf-8",
  "value" : "This is the Value of this Data Object"
}
```

EXAMPLE 2: GET to the data object URI by ID to read all fields of the data object:

```
GET /cdm_objectid/00007ED90010D891022876A8DE0BC0FD HTTP/1.1
Host: cloud.example.com
Accept: application/cdm-object
```

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: application/cdm-object

{
  "objectType" : "application/cdm-object",
  "objectID" : "00007ED90010D891022876A8DE0BC0FD",
  "objectName" : "MyDataObject.txt",
  "parentURI" : "/MyContainer/",
  "parentID" : "00007E7F00102E230ED82694DAA975D2",
  "domainURI" : "/cdm_domains/MyDomain/",
  "capabilitiesURI" : "/cdm_capabilities/dataobject/",
  "completionStatus" : "Complete",
  "mimetype" : "text/plain",
  "metadata" : {
    "cdm_size" : "37"
  },
  "valuetransferencoding" : "utf-8",
  "valuerange" : "0-36",
  "value" : "This is the Value of this Data Object"
}
```

EXAMPLE 3: GET to the data object URI to read the value and mimetype fields of the data object:

```
GET /MyContainer/MyDataObject.txt?value;mimetype HTTP/1.1
Host: cloud.example.com
Accept: application/cdm-object
```

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: application/cdmi-object

{
  "value" : "This is the Value of this Data Object",
  "mimetype" : "text/plain"
}
```

EXAMPLE 4: GET to the data object URI to read the first 11 bytes of the value of the data object:

```
GET /MyContainer/MyDataObject.txt?valuerange;value:0-10 HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-object
```

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: application/cdmi-object

{
  "valuerange" : "0-10",
  "value" : "VGhpcyBpcyB0aGU="
}
```

EXAMPLE 5: GET to the data object URI to read the data object using multi-part MIME:

```
GET /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Accept: multipart/mixed
```

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: multipart/mixed; boundary=gc0p4Jq0M2Yt08j34c0p

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/cdmi-object

{
  "objectType": "application/cdmi-object",
  "objectID": "00007ED90010C2414303B5C6D4F83170",
  "objectName": "MyDataObject.txt",
  "parentURI": "/MyContainer/",
  "parentID": "00007E7F00102E230ED82694DAA975D2",
  "domainURI": "/cdmi_domains/MyDomain/",
  "capabilitiesURI": "/cdmi_capabilities/dataobject/",
  "completionStatus": "Complete",
  "mimetype": "application/octet-stream",
  "metadata": {
    "cdmi_size": "37",
    "colour": "blue",
    ...
  },
  "valuerange": "0-36",
  "valuetransferencoding": "base64"
}

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/octet-stream
Content-Transfer-Encoding: binary

<37 bytes of binary data>

--gc0p4Jq0M2Yt08j34c0p--
```

EXAMPLE 6: GET to the data object URI to read the data object using multi-part MIME, with optional content-lengths for the parts:

```
GET /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Accept: multipart/mixed
```

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: multipart/mixed; boundary=gc0p4Jq0M2Yt08j34c0p

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/cdm-object
Content-Length: 505

{
  "objectType": "application/cdm-object",
  "objectID": "00007ED90010C2414303B5C6D4F83170",
  "objectName": "MyDataObject.txt",
  "parentURI": "/MyContainer/",
  "parentID": "00007E7F00102E230ED82694DAA975D2",
  "domainURI": "/cdmi_domains/MyDomain/",
  "capabilitiesURI": "/cdmi_capabilities/dataobject/",
  "completionStatus": "Complete",
  "mimetype": "application/octet-stream",
  "metadata": {
    "cdmi_size": "37",
    "colour": "blue",
    ...
  },
  "valuerange": "0-36",
  "valuetransferencoding": "base64"
}

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/octet-stream
Content-Transfer-Encoding: binary
Content-Length: 37

<37 bytes of binary data>

--gc0p4Jq0M2Yt08j34c0p--
```

**EXAMPLE 7: GET to the data object URI to read the metadata and multiple byte ranges of the binary contents using multi-part MIME:**

```
GET /MyContainer/MyDataObject.txt?metadata;value:0-10;value:21-24 HTTP/1.1
Host: cloud.example.com
Accept: multipart/mixed
```

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: multipart/mixed; boundary=gc0p4Jq0M2Yt08j34c0p

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/cdm-object

{
  "metadata": {
    "cdmi_size": "37",
    "colour": "blue",
    ...
  }
}

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/octet-stream
Content-Transfer-Encoding: binary
Content-Range: bytes 0-10/37
```

(continues on next page)

(continued from previous page)

```

<11 bytes of binary data>

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/octet-stream
Content-Transfer-Encoding: binary
Content-Range: bytes 21-24/37

<4 bytes of binary data>

--gc0p4Jq0M2Yt08j34c0p--

```

EXAMPLE 8: GET to the data object URI to read the value and valuetransferencoding fields of a data object storing JSON data:

```

GET /cdmi_objectid/0000706D0010374085EF1A5C7018D774?valuetransferencoding;value HTTP/
↪1.1
Host: cloud.example.com
Accept: application/cdmi-object

```

The following shows the response.

```

Content-Type: application/cdmi-object

{
  "valuetransferencoding" : "json"
  "value" : {
    "test" : "value"
  }
}

```

EXAMPLE 9: GET to the data object URI to read a newly-created data object with a current version:

```

GET /MyContainer/MyVersionedDataObject.txt HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-object

```

The following shows the response.

```

Content-Type: application/cdmi-object

{
  "objectType" : "application/cdmi-object",
  "objectID" : "00007ED900100DA32EC94351F8970400",
  "objectName" : "MyVersionedDataObject.txt",
  "parentURI" : "/MyContainer/",
  "parentID" : "00007E7F00102E230ED82694DAA975D2",
  "domainURI" : "/cdmi_domains/MyDomain/",
  "capabilitiesURI" : "/cdmi_capabilities/dataobject/",
  "completionStatus" : "Complete",
  "mimetype" : "text/plain",
  "metadata" : {
    "cdmi_size" : "33",
    "cdmi_versioning" : "user",
    "cdmi_version_object" : "/cdmi_objectid/00007ED900100DA32EC94351F8970400",
    "cdmi_version_current" : "/cdmi_objectid/00007ED90010512EB55A9304EAC5D4AA",
    "cdmi_version_oldest" : [
      "/cdmi_objectid/00007ED90010512EB55A9304EAC5D4AA"
    ],
    ...
  },
  "valuerange" : "0-32",
  "valuetransferencoding" : "utf-8",
  "value" : "First version of this Data Object"
}

```

EXAMPLE 10: GET to the data object URI to read a data object with two historical versions:

```
GET /MyContainer/MyVersionedDataObject.txt HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-object
```

The following shows the response.

```
Content-Type: application/cdmi-object

{
  "objectType" : "application/cdmi-object",
  "objectID" : "00007ED900100DA32EC94351F8970400",
  "objectName" : "MyDataObject.txt",
  "parentURI" : "/MyContainer/",
  "parentID" : "00007E7F00102E230ED82694DAA975D2",
  "domainURI" : "/cdmi_domains/MyDomain/",
  "capabilitiesURI" : "/cdmi_capabilities/dataobject/",
  "completionStatus" : "Complete",
  "mimetype" : "text/plain",
  "metadata" : {
    "cdmi_size" : "33",
    "cdmi_versioning" : "user",
    "cdmi_version_object" : "/cdmi_objectid/00007ED900100DA32EC94351F8970400",
    "cdmi_version_current" : "/cdmi_objectid/00007ED90010F077F4EB1C99C87524CC",
    "cdmi_version_oldest" : [
      "/cdmi_objectid/00007ED90010512EB55A9304EAC5D4AA"
    ],
    ...
  },
  "valuerange" : "0-32",
  "valuetransferencoding" : "utf-8",
  "value" : "Third version of this Data Object"
}
```

**EXAMPLE 11: GET to the URI of a data object version:**

```
GET /cdmi_objectid/00007ED9001005192891EEBE599D94BB HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-object
```

The following shows the response.

```
Content-Type: application/cdmi-object

{
  "objectType" : "application/cdmi-object",
  "objectID" : "00007ED9001005192891EEBE599D94BB",
  "objectName" : "MyVersionedDataObject.txt",
  "parentURI" : "/MyContainer/",
  "parentID" : "00007E7F00102E230ED82694DAA975D2",
  "domainURI" : "/cdmi_domains/MyDomain/",
  "capabilitiesURI" : "/cdmi_capabilities/dataobject/dataobject_version/",
  "completionStatus" : "Complete",
  "mimetype" : "text/plain",
  "metadata" : {
    "cdmi_size" : "34",
    "cdmi_version_object" : "/cdmi_objectid/00007ED900100DA32EC94351F8970400",
    "cdmi_version_current" : "/cdmi_objectid/00007ED90010F077F4EB1C99C87524CC",
    "cdmi_version_oldest" : [
      "/cdmi_objectid/00007ED90010512EB55A9304EAC5D4AA"
    ],
    "cdmi_version_parent" : "/cdmi_objectid/00007ED90010512EB55A9304EAC5D4AA",
    "cdmi_version_children" : [
      "/cdmi_objectid/00007ED90010F077F4EB1C99C87524CC"
    ],
    ...
  },
  "valuerange" : "0-33",
  "valuetransferencoding" : "utf-8",

```

(continues on next page)

(continued from previous page)

```
"value" : "Second version of this Data Object"  
}
```

## 8.4 Update a data object using CDMI

### 8.4.1 Synopsis

To update an existing data object, the following requests shall be performed:

- PUT <root URI>/<ContainerName>/<DataObjectName>
- PUT <root URI>/<ContainerName>/<DataObjectName>?value:<range>
- PUT <root URI>/<ContainerName>/<DataObjectName>?metadata:<metadataname>;....
- PUT <root URI>/cdmi\_objectid/<DataObjectID>
- PUT <root URI>/cdmi\_objectid/<DataObjectID>?value:<range>
- PUT <root URI>/cdmi\_objectid/<DataObjectID>?metadata:<metadataname>;....

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate containers.
- <DataObjectName> is the name of the data object to be updated.
- <range> is a byte range for the data object value to be updated.
- <DataObjectID> is the ID of the data object to be updated.

### 8.4.2 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 40](#).

Table 40: Capabilities - Update a CDMI data object using CDMI

Capability	Location	Description
cdmi_modify_metadata	Data Object	Ability to modify the metadata of an existing data object
cdmi_modify_value	Data Object	Ability to modify the value of an existing data object
cdmi_modify_value_range	Data Object	Ability to modify a sub-range of the value of an existing data object
cdmi_multipart_mime	System Wide Capability	Ability to modify a data object using multi-part MIME
cdmi_object_access_by_ID	System Wide Capability	Ability to access the object by ID



### 8.4.3 Request headers

The HTTP request headers for updating a CDMI data object using CDMI are shown in [Table 41](#).

Table 41: Request headers - Update a CDMI data object using CDMI

Header	Type	Description	Requirement
Content-Type	Header string	<p>“application/cdm-object” or “multipart/mixed”</p> <ul style="list-style-type: none"> <li>If “multipart/mixed” is specified, the body shall consist of at least two MIME parts, where the first part shall contain a body of content-type “application/cdm-object”, and the second and subsequent parts shall contain one or more byte ranges of the value.</li> <li>If multiple byte ranges are included and the Content-Range header is omitted for a part, the data in the part shall be appended to the data in the preceding part, with the first part having a byte offset of zero.</li> </ul>	Mandatory
X-CDMI-Partial	Header string	<p>Indicates that the newly created object is part of a series of writes and has not yet been fully created. When set to “true”, the completionStatus field shall be set to “Processing”. X-CDMI-Partial works across CDMI and non-CDMI operations.</p> <p>If the completionStatus field had previously been set to “Processing” by including this header in a create or update, the next update without this field shall change the completionStatus field back to “Complete”.</p>	Optional

### 8.4.4 Request message body

The request message body fields for updating a data object using CDMI are shown in [Table 42](#).

Table 42: Request message body - Update a CDMI data object using CDMI

Field Name	Type	Description	Requirement
mimetype	JSON string	<p>MIME type of the data contained within the value field of the data object. If present, this value replaces the existing mimetype field value.</p> <ul style="list-style-type: none"> <li>This field may be included when updating by value, deserializing, and copying a data object.</li> <li>If this field is not included, the existing value of the mimetype field shall be left unchanged.</li> <li>This field value shall be converted to lower case before being stored.</li> </ul> <p>If this field is set to “application/cms” or “application/jose+json”, the CDMI server shall encrypt or reencrypt the value of the object in place, using the key specified by the “cdmi_enc_key_id” metadata item. If the “cdmi_enc_key_id” metadata item is not present, the object ID shall be used as the key identifier. The mimetype of the plaintext shall be stored in the CMS or JWE JSON representation.</p> <p>If a “cdmi_enc_value_sign_id” metadata item is present, the encrypted object shall also be signed.</p> <p>If this field is changed from “application/cms” or “application/jose+json” to any other mimetype, the CDMI server shall decrypt the value of the object in place, replacing the specified mimetype with the mimetype of the encrypted object, if stored as part of the encrypted object.</p> <p>For more details on encrypted objects, see <a href="#">clause 23</a>.</p>	Optional
metadata	JSON object	Metadata for the data object. If present, the new metadata specified replaces the existing object metadata. If individual metadata items are specified in the URI, only those items are replaced; other items are preserved. See <a href="#">clause 16</a> for a further description of metadata.	Optional
domainURI	JSON string	<p>URI of the owning domain</p> <ul style="list-style-type: none"> <li>If different from the parent domain, the user shall have the “cross-domain” privilege (see <a href="#">cdmi_member_privileges</a> in <a href="#">Table 80</a>).</li> <li>If not specified, the existing domain shall be preserved.</li> </ul>	Optional
deserialize	JSON string	<p>URI of a CDMI data object with a value that contains a data object serialized as specified in <a href="#">clause 15</a>. The serialized data object shall be deserialized to update the existing data object.</p> <p>The object ID of the serialized data object shall match the object ID of the destination data object. Otherwise, the server shall return an HTTP status code of 400 Bad Request.</p>	Optional <sup>1</sup>

Continued on next page

Table 42 – continued from previous page

Field Name	Type	Description	Requirement
copy	JSON string	<p>URI of a source CDMI data object or queue object that shall be copied into an existing destination data object.</p> <ul style="list-style-type: none"> <li>If the destination data object URI and the copy source object URI both do not specify individual fields, the destination data object shall be replaced with the source data object.</li> <li>If the destination data object URI or the copy source object URI specifies individual fields, only the fields specified shall be used to update the destination data object. If specified fields are not present in the source, these fields shall be ignored.</li> <li>If the destination data object URI and the copy source object URI both specify fields, an HTTP status code of <code>400 Bad Request</code> shall be returned to the client.</li> </ul> <p>If the copy source object URI points to a queue object, as part of the copy operation, multiple queue values shall be concatenated into a single data object value.</p> <p>If there are insufficient permissions to read the data object at the source URI, update the data object at the destination URI, or if the read operation fails, the copy shall return an HTTP status code of <code>400 Bad Request</code>, and the destination shall be left unchanged.</p>	Optional <sup>1</sup>
deserializevalue	JSON string	<p>A data object serialized as specified in <a href="#">clause 15</a> and encoded using base 64 encoding rules described in RFC 4648 [19], that shall be deserialized to update the existing data object.</p> <p>The object ID of the serialized data object shall match the object ID of the destination data object. Otherwise, the server shall return an HTTP status code of <code>400 Bad Request</code>.</p>	Optional <sup>1</sup>

Continued on next page

Table 42 – continued from previous page

Field Name	Type	Description	Requirement
valuetransferencoding	JSON string	<p>The value transfer encoding used for the data object value. Three value transfer encodings are defined:</p> <ul style="list-style-type: none"> <li>• “utf-8” indicates that the data object contains a valid UTF-8 string and shall be transported as a UTF-8 string in the value field. If the contents of the data object value field are set or updated to any value other than a valid UTF-8 string, an HTTP status code of 400 Bad Request shall be returned to the client.</li> <li>• “base64” indicates that the data object may contain arbitrary binary sequence and shall be transported as a base 64 encoded string in the value field. Setting the contents of the data object value field to any value other than a valid base 64 string shall result in an HTTP status code of 400 Bad Request being returned to the client.</li> <li>• “json” indicates that the data object contains a valid JSON object and shall be transported as a JSON object in the value field. If the contents of the data object value field are set or updated to any value other than a valid JSON object, an HTTP status code of 400 Bad Request shall be returned to the client.</li> </ul> <p>This field shall only be included when updating a data object by value.</p> <ul style="list-style-type: none"> <li>• If this field is not included and multi-part MIME is not being used, the existing value of “valuetransferencoding” shall be left unchanged.</li> <li>• If this field is not included and multi-part MIME is being used, the value of “utf-8” shall be assigned as the field value if the “Content-Type” header of the second and all subsequent MIME parts includes the charset parameter as defined in RFC 2046 of “utf-8” (e.g., “; charset=utf-8”). Otherwise, the value of “base64” shall be assigned as the field value. This field applies only to the encoding of the value when represented in JSON; the “Content-Transfer-Encoding” header of the part specifies the encoding of the value within a multi-part MIME request, as defined in RFC 2045.</li> </ul>	Optional

Continued on next page

Table 42 – continued from previous page

Field Name	Type	Description	Requirement
value	JSON string	<p>This field contains the new data for the object. If present, this value replaces the existing value.</p> <ul style="list-style-type: none"> <li>If this field is not included and multi-part MIME is being used, the contents of the second and subsequent MIME parts shall be assigned to the corresponding byte ranges of the field value.</li> <li>If the <code>valuetransferencoding</code> field indicates UTF-8 encoding, the value shall be a UTF-8 string escaped using the JSON escaping rules described in RFC 4627 [5].</li> <li>If the <code>valuetransferencoding</code> field indicates base 64 encoding, the value shall be first encoded using the base 64 encoding rules described in RFC 4648 [19].</li> <li>If the <code>valuetransferencoding</code> field indicates JSON encoding, the value field shall contain a valid JSON object.</li> <li>If a value range was specified in the request, the new data shall be inserted at the location specified by the range. Any resulting gaps between ranges shall be treated as if zeros had been written and shall be included when calculating the size of the value. When storing a range, the value shall be encoded using base 64, and the <code>valuetransferencoding</code> field shall be set to “base64”.</li> </ul>	Optional <sup>1</sup>

#### 8.4.5 Response header

The HTTP response header for updating a data object using CDMI is shown in Table 43.

Table 43: Response header - Update a CDMI data object using CDMI

Header	Type	Description	Requirement
Location	Header string	The server shall respond with the URI to which the reference redirects if the object is a reference.	Conditional

#### 8.4.6 Response message body

A response body can be provided as per RFC 2616 [23].

#### 8.4.7 Response status

The HTTP status codes that occur when updating a data object using CDMI are described in Table 44.

<sup>1</sup> Only one of these fields shall be specified in any given operation. Except for value, these fields shall not be stored. If more than one of these fields is supplied, the server shall respond with an HTTP status code of 400 `Bad Request`.

Table 44: HTTP status codes - Update a CDMI data object using CDMI

HTTP Status	Description
204 No Content	The data object content was returned in the response.
302 Found	The resource is a reference to another resource.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server.

## 8.4.8 Examples

EXAMPLE 1: PUT to the data object URI to set new field values:

```
PUT /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdm-object

{
  "mimetype" : "text/plain",
  "metadata" : {
    "colour" : "blue",
    "length" : "10"
  },
  "value" : "This is the Value of this Data Object"
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

EXAMPLE 2: PUT to the data object URI to set a new MIME type:

```
PUT /MyContainer/MyDataObject.txt?mimetype HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdm-object

{
  "mimetype" : "text/plain"
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

EXAMPLE 3: PUT to the data object URI to update a range of the value:

```
PUT /MyContainer/MyDataObject.txt?value:21-24 HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdm-object

{
  "value" : "dGhhZA=="
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

When updating a value without specifying a value transfer encoding, the client must be aware of the current value transfer encoding of the object.

- If a client sends a value containing a UTF-8 string that is not a valid base 64 string to update an existing object with a value transfer encoding of “base64”, the server shall return an error.
- If a client sends a value containing a base 64 string to update an existing object with a value transfer encoding of “utf-8”, the server shall not return an error. Instead, the server shall store the literal base 64 character sequence in the data object instead of the data encoded in the base 64 string.

EXAMPLE 4: PUT to the data object URI to replace all metadata with new metadata:

```
PUT /MyContainer/MyDataObject.txt?metadata HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-object

{
  "metadata" : {
    "colour" : "red",
    "number" : "7"
  }
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

EXAMPLE 5: PUT to the data object URI to add a new metadata item while preserving existing metadata:

```
PUT /MyContainer/MyDataObject.txt?metadata:shape HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-object

{
  "metadata" : {
    "shape" : "round"
  }
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

EXAMPLE 6: PUT to the data object URI to replace just one metadata item with a new value:

```
PUT /MyContainer/MyDataObject.txt?metadata:colour HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-object

{
  "metadata" : {
    "colour" : "green"
  }
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

EXAMPLE 7: Delete a single metadata item:

```
PUT /MyContainer/MyDataObject.txt?metadata:colour HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-object

{
  "metadata": {}
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

EXAMPLE 8: Add, update, and delete metadata items. Assume a starting condition where the object has a metadata item “colour” with value “green” and a metadata item “shape” with value “round” and does not have a metadata item “size”. After the update, “colour” has value “red”, “shape” is deleted, and “size” has been added with value “10”.

```
PUT /MyContainer/MyDataObject.txt?metadata:colour;metadata:shape;metadata:size HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-object

{
  "metadata": {
    "colour": "red",
    "size": "10"
  }
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

EXAMPLE 9: PUT to the data object URI to set new field values and the binary contents using multi-part MIME:

```
PUT /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Content-Type: multipart/mixed; boundary=gc0p4Jq0M2Yt08j34c0p

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/cdmi-object

{
  "metadata": {
    "colour": "red",
    "number": "7"
  }
}

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/octet-stream
Content-Transfer-Encoding: binary

<37 bytes of binary data>

--gc0p4Jq0M2Yt08j34c0p--
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

EXAMPLE 10: PUT to the data object URI to replace just one metadata item and update multiple byte ranges within the binary contents of the data object using multi-part MIME:

```
PUT /MyContainer/BinaryObject.txt?metadata:colour HTTP/1.1
Host: cloud.example.com
Content-Type: multipart/mixed; boundary=gc0p4Jq0M2Yt08j34c0p

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/cdmi-object

{
  "metadata": {
    "colour": "green"
  }
}

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/octet-stream
Content-Range: bytes 0-10/37
```

(continues on next page)



(continued from previous page)

```
<11 bytes of binary data>

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/octet-stream
Content-Range: bytes 21-24/37

<4 bytes of binary data>

--gc0p4Jq0M2Yt08j34c0p--
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

**EXAMPLE 11: PUT to the data object URI to encrypt an existing object:**

```
PUT /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdm-object

{
  "mimetype" : "application/cms",
  "metadata" : {
    "cdmi_enc_key_id" : "testkey"
  }
}
```

The following shows the response:

```
HTTP/1.1 204 No Content
```

**EXAMPLE 12: PUT to the data object URI to decrypt an existing encrypted object:**

The following shows the response:

```
HTTP/1.1 204 No Content
```

## 8.5 Delete a data object using CDMI

### 8.5.1 Synopsis

To delete an existing data object, the following requests shall be performed:

- DELETE <root URI>/<ContainerName>/<DataObjectName>
- DELETE <root URI>/cdmi\_objectid/<DataObjectID>

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate containers.
- <DataObjectName> is the name of the data object to be deleted.
- <DataObjectID> is the ID of the data object to be deleted.

### 8.5.2 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 45](#).

Table 45: Capabilities - Delete a CDMI data object using CDMI

Capability	Location	Description
cdmi_delete_dataobject	Data Object	Ability to delete an existing data object
cdmi_object_access_by_ID	System Wide Capability	Ability to access the object by ID

### 8.5.3 Request headers

Request headers can be provided as per RFC 2616 [\[23\]](#).

### 8.5.4 Request message body

A request body can be provided as per RFC 2616 [\[23\]](#).

### 8.5.5 Response headers

Response headers can be provided as per RFC 2616 [\[23\]](#).

### 8.5.6 Response message body

A response body can be provided as per RFC 2616 [\[23\]](#).

### 8.5.7 Response status

[Table 46](#) describes the HTTP status codes that occur when deleting a data object using CDMI.

Table 46: HTTP status codes - Delete a CDMI data object using CDMI

HTTP Status	Description
204 No Content	The data object was successfully deleted.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server or the data object cannot be deleted.

### 8.5.8 Example

#### EXAMPLE 1: DELETE by data object URI:

```
DELETE /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

#### EXAMPLE 2: DELETE by data object ID:

```
DELETE /cdmi_objectid/00007ED90010D891022876A8DE0BC0FD HTTP/1.1
Host: cloud.example.com
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

## Clause 9

# Container Object Resource Operations using CDMI

## 9.1 Overview

Container objects are the fundamental grouping of stored data within CDMI™ and are analogous to directories within a file system. Each container object has a set of well-defined fields that include:

- zero or more child objects,
- mandatory fields generated by the cloud storage system,
- mandatory metadata items generated by the cloud storage system,
- optional metadata generated by the cloud storage system; and
- optional metadata specified by the cloud user.

All cloud storage systems shall support containers, but the ability to create a containers is determiend by the presence or absence of the `cdmi_create_container` capability in the parent container.

Each CDMI container object is represented as a JSON object, containing one or more “fields”. For example, the “metadata” field contains metadata items.

EXAMPLE 1: CDMI Container Object

```
{
  "objectType" : "application/cdmi-container",
  "objectID" : "00007ED900104E1D14771DC67C27BF8B",
  "objectName" : "MyContainer/",
  "parentURI" : "/",
  "parentID" : "00007E7F0010128E42D87EE34F5A6560",
  "domainURI" : "/cdmi_domains/MyDomain/",
  "capabilitiesURI" : "/cdmi_capabilities/container/",
  "completionStatus" : "Complete",
  "metadata" : {
    "cdmi_ctime" : "2018-05-16T08:01:02.353Z"
  },
  "childrenrange" : "0-4",
  "children" : [
    "red",
    "green",
    "yellow",
    "orange/",
    "purple/"
  ]
}
```

The meaning, use, and permitted values of each field is described in each operation that creates, modifies or retrieves CDMI container objects.

## 9.1.1 Container object addressing

Container objects are addressed in CDMI in two ways:

- by name (e.g. `https://cloud.example.com/container/`); and
- by ID (e.g. `https://cloud.example.com/cdmi_objectid/00007ED900104E1D14771DC67C27BF8B/`).

Every container object has a single, globally-unique object ID that remains constant for the life of the object. Each container object may also have one or more URI addresses that allow the container object to be accessed.

When a container object is addressed via more than one unique URIs, all operations may be performed through any of these URIs. For example, a container object may be accessible via multiple virtual hosting paths, where `https://cloud.example.com/users/snua/cdmi/` is also accessible through `https://snua.example.com/cdmi/`. Conflicting writes via different paths shall be managed the same way that conflicting writes via one path are managed, via the principle of eventual consistency (see 9.2).

Following the URI conventions for hierarchical paths, container URIs shall consist of one or more container names that are separated by forward slashes ("/") and that end with a forward slash ("/").

If a request is performed against an existing container resource and the trailing slash at the end of the URI is omitted, the server shall respond with an HTTP status code of 301 Moved Permanently. In addition, a Location header containing the URI with the trailing slash added shall be returned.

If a CDMI request is performed to create a new container resource and the trailing slash at the end of the URI is omitted, the server shall respond with an HTTP status code of 400 Bad Request.

Non-CDMI requests to create a container resource shall include the trailing slash at the end of the URI; otherwise, the request shall be considered a request to create a data object.

Containers may also be nested.

EXAMPLE 2: The following URI represents a nested container:

```
https://cloud.example.com/container/subcontainer/
```

A nested container has a parent container object, shall be included in the children field of the parent container object, and shall inherit data system metadata and ACLs from its parent container.

This model allows direct mapping between CDMI-managed cloud storage and file systems (e.g., NFSv4 or WebDAV). If a CDMI container object is exported as a file system, then the file system may make the CDMI metadata accessible via file system-specific mechanisms. As files and directories are created by the file system, they become visible through the CDMI interface acting as a data path. The mapping between file system constructs and CDMI data objects, container objects, and metadata is outside the scope of this International Standard.

## 9.1.2 Container object fields

Individual fields within a container object may be accessed by specifying the field name after a question mark "?" appended to the end of the container object URI.

EXAMPLE 3: The following URI returns just the children field in the response body:

```
https://cloud.example.com/container/?children
```

EXAMPLE 4: By specifying a range after the children field name, specific ranges of the children field may be accessed.

```
https://cloud.example.com/container/?children:0-2
```

Children ranges are specified in a way that is similar to byte ranges as per Section 14.35.1 of RFC 2616 [23]. A client can determine the number of children present by requesting the childrenrange field without requesting a range of children.

A list of fields, separated by a semicolon ";", may be specified, allowing multiple fields to be accessed in a single request.

EXAMPLE 5: The following URI would return the children and metadata fields in the response body:

```
https://cloud.example.com/container/?children;metadata
```

When a client provides fields that are not defined in this International Standard or deserializes an object containing fields that are not defined in this International Standard, these fields shall be persisted, but shall not be interpreted.

### 9.1.3 Container object metadata

The following optional container-specific data system metadata may be provided (see Table 47).

Table 47: Container metadata

Metadata Name	Type	Description	Requirement
cdmi_assignedsize	JSON string	The number of bytes that is reported via exported protocols (e.g., the device may be thin provisioned). This number may limit <code>cdmi_size</code> .	Optional

Container metadata may also include arbitrary user-supplied metadata, storage system metadata, and data system metadata as described in clause 16.

### 9.1.4 Container object access control

If read access to any of the requested fields is not permitted by the object ACL, only the permitted fields shall be returned. If no requested fields are permitted to be read, an HTTP status code of 403 `Forbidden` shall be returned to the client.

If write access to any of the requested fields is not permitted by the object ACL, no updates shall be performed, and an HTTP status code of 403 `Forbidden` shall be returned to the client.

### 9.1.5 Reserved container object names

This International Standard defines reserved container names that should not be used by clients when creating new containers. These container names are reserved for use by this International Standard, and if an attempt is made to create or delete them, an HTTP status code of 400 `Bad Request` shall be returned to the client.

Reserved container names defined in this specification include:

- “cdmi\_objectid”
- “cdmi\_domains”
- “cdmi\_capabilities”
- “cdmi\_snapshots”
- “cdmi\_versions”

As additional names may be added in future versions of this International Standard, server implementations shall prevent the creation of user-defined containers if the container name starts with “cdmi\_”.

### 9.1.6 Container object representations

The representations in this clause are shown using JSON notation. Both clients and servers shall support UTF-8 JSON representation. The request and response body JSON fields may be specified or returned in any order, with the exception that, if present, for container objects, the “childrenrange” and “children” fields shall appear last and in that order.

## 9.2 Create a container object using CDMI

### 9.2.1 Synopsis

To create a new container object, the following request shall be performed:

- PUT <root URI>/<ContainerName>/<NewContainerName>/

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate container objects that already exist, with one slash (i.e., "/") between each pair of container object names.
- <NewContainerName> is the name specified for the container object to be created.

After it is created, the container object shall also be accessible at <root URI>/cdmi\_objectid/<objectID>/.

### 9.2.2 Delayed completion of create

In response to a create operation for a container object, the server may return an HTTP status code of 202 *Accepted* to indicate that the object is in the process of being created. This response is useful for long-running operations (e.g., deserializing a source data object to create a large container object hierarchy). Such a response has the following implications.

- The server shall return a Location header with an absolute URI to the object to be created along with an HTTP status code of 202 *Accepted*.
- With an HTTP status code of 202 *Accepted*, the server implies that the following checks have passed:
  - user authorization for creating the container object;
  - user authorization for read access to any source object for move, copy, serialize, or deserialize; and
  - availability of space to create the container object or at least enough space to create a URI to report an error.
- A client might not be able to immediately access the created object, e.g., due to delays resulting from the implementation's use of eventual consistency.

The client performs GET operations to the URI to track the progress of the operation. In response, the server returns two fields in its response body to indicate progress.

- A mandatory `completionStatus` text field contains either "Processing", "Complete", or an error string starting with the value "Error".
- An optional `percentComplete` field contains the percentage that the accepted PUT has completed (0 to 100). GET does not return any children for the container object when `completionStatus` is not "Complete".

When the final result of the create operation is an error, the URI is created with the `completionStatus` field set to the error message. It is the client's responsibility to delete the URI after the error has been noted.

## 9.2.3 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 48](#).

Table 48: Capabilities - Create a CDMI container object using CDMI

Capability	Location	Description
cdmi_create_container	Parent container	Ability to create a new container object
cdmi_create_reference	Parent container	Ability to create a new reference
cdmi_copy_container	Parent container	Ability to create a container object that is a copy of another container object
cdmi_move_container	Parent container	Ability to move a container object from another location
cdmi_deserialize_container	Parent container	Ability to create a container object that is deserialized from the contents of the PUT or the contents of a data object

## 9.2.4 Request headers

The HTTP request headers for creating a CDMI container object using CDMI are shown in [Table 49](#).

Table 49: Request headers - Create a container object using CDMI

Header	Type	Description	Requirement
Accept	Header string	“application/cdm-container” or a consistent value described in <a href="#">5.5.2</a>	Optional
Content-Type	Header string	“application/cdm-container”	Mandatory

## 9.2.5 Request message body

The request message body fields for creating a container object using CDMI are shown in [Table 50](#).

Table 50: Request message body - Create a container object using CDMI

Field Name	Type	Description	Requirement
metadata	JSON object	Metadata for the container object <ul style="list-style-type: none"> <li>If this field is included, the contents of the JSON object provided in this field shall be used as container object metadata.</li> <li>If this field is included when deserializing, serializing, copying, or moving a container object, the contents of the JSON object provided in this field shall be used as object metadata instead of the metadata from the source URI.</li> <li>If this field is not included, no user-specified metadata shall be added to the object.</li> <li>If this field is not included when deserializing, serializing, copying, or moving a container object, metadata from the source URI shall be used.</li> <li>This field shall not be included when creating a reference to a container object.</li> </ul>	Optional

Continued on next page



Table 50 – continued from previous page

Field Name	Type	Description	Requirement
domainURI	JSON string	URI of the owning domain <ul style="list-style-type: none"> <li>If different from the parent domain, the user shall have the “cross-domain” privilege (see <code>cdmi_member_privileges</code> in Table 80).</li> <li>If not specified, the existing domain shall be preserved.</li> </ul>	Optional
exports	JSON object	A structure for each protocol enabled for this container object (see clause 13). This field shall not be included when referencing a container object.	Optional
deserialize	JSON string	URI of a CDMI data object with a value that contains a container object serialized as specified in clause 15. The serialized container object shall be deserialized to create the new container object, including all child objects.  When deserializing a container object, any exported protocols from the original serialized container object are not applied to the newly created container object(s).	Optional <sup>1</sup>
copy	JSON string	URI of a source CDMI container object that shall be copied into the new destination container object. <ul style="list-style-type: none"> <li>If the destination container object URI and the copy source object URI both do not specify individual fields, the destination container object shall be a complete copy of the source container object, including all child objects under the source container object.</li> <li>If the destination container object URI or the copy source object URI specifies individual fields, only the fields specified shall be used to create the destination container object. If specified fields are not present in the source, default field values shall be used.</li> <li>If the destination container object URI and the copy source object URI both specify fields, an HTTP status code of 400 <code>Bad Request</code> shall be returned to the client.</li> </ul> When copying a container object, exported protocols are not preserved across the copy.  If there are insufficient permissions to read the container object at the source URI or create the container object at the destination URI, or if the read operation fails, the copy shall return an HTTP status code of 400 <code>Bad Request</code> , and the destination container object shall not be created.	Optional <sup>1</sup>
move	JSON string	URI of an existing local or remote CDMI container object (source URI) that shall be relocated, along with all child objects, to the URI specified in the PUT. The contents of the container object and all children, including the object ID, shall be preserved by a move, and the container object and all children of the source URI shall be removed after the objects at the destination have been successfully created.  If there are insufficient permissions to read the objects at the source URI, write the objects at the destination URI, or delete the objects at the source URI, or if any of these operations fail, the move shall return an HTTP status code of 400 <code>Bad Request</code> , and the source and destination are left unchanged.	Optional <sup>1</sup>
reference	JSON string	URI of a CDMI container object that shall be redirected to by a reference. If other fields are supplied when creating a reference, the server shall respond with an HTTP status code of 400 <code>Bad Request</code> .	Optional <sup>1</sup>

Continued on next page

Table 50 – continued from previous page

Field Name	Type	Description	Requirement
deserializevalue	JSON string	A container object serialized as specified in <a href="#">clause 15</a> and encoded using base 64 encoding rules described in RFC 4648 [19], that shall be deserialized to create the new container object, including all child objects.	Optional <sup>1</sup>

## 9.2.6 Response headers

The HTTP response headers for creating a CDMI container object using CDMI are shown in [Table 51](#).

Table 51: Response headers - Create a container object using CDMI

Header	Type	Description	Requirement
Content-Type	Header string	"application/cdm-container"	Mandatory
Location	Header string	When an HTTP status code of 202 <i>Accepted</i> is returned, the server shall respond with the absolute URL of the object that is in the process of being created.	Conditional

## 9.2.7 Response message body

The response message body fields for creating a CDMI container object using CDMI are shown in [Table 52](#).

Table 52: Response message body - Create a container object using CDMI

Field Name	Type	Description	Requirement
objectType	JSON string	"application/cdm-container"	Mandatory
objectID	JSON string	Object ID of the object	Mandatory
objectName	JSON string	Name of the object	Mandatory
parentURI	JSON string	URI for the parent object Appending the <code>objectName</code> to the <code>parentURI</code> shall always produce a valid URI for the object.	Mandatory
parentID	JSON string	Object ID of the parent container object	Mandatory
domainURI	JSON string	URI of the owning domain	Mandatory
capabilitiesURI	JSON string	URI to the capabilities for the object	Mandatory
completionStatus	JSON string	A string indicating if the object is still in the process of being created or updated by another operation, and after that operation is complete, indicates if it was successfully created or updated or if an error occurred.  The value shall be the string "Processing", the string "Complete", or an error string starting with the value "Error".	Mandatory

Continued on next page

<sup>1</sup> Only one of these fields shall be specified in any given operation. Except for value, these fields shall not be stored. If more than one of these fields is supplied, the server shall respond with an HTTP status code of 400 *Bad Request*.

Table 52 – continued from previous page

Field Name	Type	Description	Requirement
percentComplete	JSON string	A string indicating the percentage of completion if the object is still in the process of being created or updated by another operation. <ul style="list-style-type: none"> <li>When the value of <code>completionStatus</code> is "Processing", this field, if provided, shall indicate the percentage of completion as a numeric integer value from "0" through "100".</li> <li>When the value of <code>completionStatus</code> is "Complete", this field, if provided, shall contain the value "100".</li> <li>When the value of <code>completionStatus</code> is "Error", this field, if provided, may contain any integer value from "0" through "100".</li> </ul>	Optional
metadata	JSON object	Metadata for the container object. This field includes any user and data system metadata specified in the request body metadata field, along with storage system metadata generated by the cloud storage system. See <a href="#">clause 16</a> for a further description of metadata.	Mandatory
exports	JSON object	A structure for each protocol that is enabled for this container object. See <a href="#">clause 13</a> .	Optional <sup>2</sup>
snapshots	JSON array of JSON strings	URI(s) of the snapshot container objects. See <a href="#">clause 14</a> .	Optional <sup>2</sup>
childrenrange	JSON string	The children of the container expressed as a range. If a range of children is requested, this field indicates the children returned as a range.  This field should not be returned in the response message body that is associated with a copy, move, deserialize, or deserialize value operation.	Optional
children	JSON array of JSON strings	Names of the children objects in the container object. Child container objects end with "/".  This field should not be returned in the response message body that is associated with a copy, move, deserialize, or deserialize value operation.	Optional

## 9.2.8 Response status

Table 53 describes the HTTP status codes that occur when creating a container object using CDMI.

Table 53: HTTP status codes - Create a CDMI container object using CDMI

HTTP status	Description
201 Created	The new container object was created.
202 Accepted	The container is in the process of being created. The CDMI client should monitor the <code>completionStatus</code> and <code>percentComplete</code> fields to determine the current status of the operation.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or had caused a state transition error on the server.

<sup>2</sup> Returned only if present.

## 9.2.9 Example

### EXAMPLE 1: Create a new container with no metadata:

```
PUT /MyContainer/ HTTP/1.1
Host: cloud.example.com
Accept: application/cdm-container
Content-Type: application/cdm-container

{
}
```

The following shows the response.

```
HTTP/1.1 201 Created
Content-Type: application/cdm-container

{
  "objectType" : "application/cdm-container",
  "objectID" : "00007ED900104E1D14771DC67C27BF8B",
  "objectName" : "MyContainer/",
  "parentURI" : "/",
  "parentID" : "00007E7F0010128E42D87EE34F5A6560",
  "domainURI" : "/cdm_domains/MyDomain/",
  "capabilitiesURI" : "/cdm_capabilities/container/",
  "completionStatus" : "Complete",
  "metadata" : {
    ...
  },
  "childrenrange": "",
  "children": []
}
```

### EXAMPLE 2: Create a container with metadata:

```
PUT /MyContainer/ HTTP/1.1
Host: cloud.example.com
Accept: application/cdm-container
Content-Type: application/cdm-container

{
  "metadata": {
    "Colour": "Yellow"
  }
}
```

The following shows the response.

```
HTTP/1.1 201 Created
Content-Type: application/cdm-container

{
  "objectType" : "application/cdm-container",
  "objectID" : "00007ED900104E1D14771DC67C27BF8B",
  "objectName" : "MyContainer/",
  "parentURI" : "/",
  "parentID" : "00007E7F0010128E42D87EE34F5A6560",
  "domainURI" : "/cdm_domains/MyDomain/",
  "capabilitiesURI" : "/cdm_capabilities/container/",
  "completionStatus" : "Complete",
  "metadata" : {
    "Colour": "Yellow",
    ...
  },
  "childrenrange": "",
  "children": []
}
```

### EXAMPLE 3: Create a container that is a copy of a container:

```
PUT /MyContainerCopy/ HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-container
Content-Type: application/cdmi-container

{
  "copy": "/MyContainer/"
}
```

1706

The following shows the response.

```
HTTP/1.1 201 Created
Content-Type: application/cdmi-container

{
  "objectType" : "application/cdmi-container",
  "objectID" : "00007ED900104E1D14771DC67C27BF8B",
  "objectName" : "MyContainerCopy/",
  "parentURI" : "/",
  "parentID" : "00007E7F0010128E42D87EE34F5A6560",
  "domainURI" : "/cdmi_domains/MyDomain/",
  "capabilitiesURI" : "/cdmi_capabilities/container/",
  "completionStatus" : "Complete",
  "metadata" : {
    "Colour": "Yellow",
    ...
  }
}
```

1707

**EXAMPLE 4: Rename a container:**

```
PUT /MyContainerRenamed/ HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-container
Content-Type: application/cdmi-container

{
  "move": "/MyContainer/"
}
```

1708

The following shows the response.

```
HTTP/1.1 201 Created
Content-Type: application/cdmi-container

{
  "objectType" : "application/cdmi-container",
  "objectID" : "00007ED900104E1D14771DC67C27BF8B",
  "objectName" : "MyContainerRenamed/",
  "parentURI" : "/",
  "parentID" : "00007E7F0010128E42D87EE34F5A6560",
  "domainURI" : "/cdmi_domains/MyDomain/",
  "capabilitiesURI" : "/cdmi_capabilities/container/",
  "completionStatus" : "Complete",
  "metadata" : {
    "Colour": "Yellow",
    ...
  }
}
```

## 9.3 Read a container object using CDMI

### 9.3.1 Synopsis

To read an existing container object, the following requests shall be performed:

- GET <root URI>/<ContainerName>/<TheContainerName>/
- GET <root URI>/<ContainerName>/<TheContainerName>/?<fieldname>;<fieldname>;...
- GET <root URI>/<ContainerName>/<TheContainerName>/?children:<range>;...
- GET <root URI>/<ContainerName>/<TheContainerName>/?metadata:<prefix>;...
- GET <root URI>/cdmi\_objectid/<ContainerObjectID>/
- GET <root URI>/cdmi\_objectid/<ContainerObjectID>/?<fieldname>;<fieldname>;...
- GET <root URI>/cdmi\_objectid/<ContainerObjectID>/?children:<range>;...
- GET <root URI>/cdmi\_objectid/<ContainerObjectID>/?metadata:<prefix>;...

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate container objects.
- <TheContainerName> is the name specified for the container object to be read from.
- <fieldname> is the name of a field.
- <range> is a numeric range within the list of children.
- <prefix> is a matching prefix that returns all metadata items that start with the prefix value.
- <ContainerObjectID> is the ID of the data object to be read from.

### 9.3.2 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 54](#).

Table 54: Capabilities - Read a CDMI Container Object using CDMI

Capability	Location	Description
cdmi_read_metadata	Container object	Ability to read the metadata of an existing container object
cdmi_list_children	Container object	Ability to list the children of an existing container object
cdmi_list_children_range	Container object	Ability to list a specific range of children of an existing container object
cdmi_object_access_by_ID	System wide capability	Ability to access the object by ID

### 9.3.3 Request headers

The HTTP request headers for reading a CDMI container object using CDMI are shown in [Table 55](#).

Table 55: Request headers - Read a container object using CDMI

Header	Type	Description	Requirement
Accept	Header string	“application/cdm-container” or a consistent value as described in <a href="#">5.5.2</a>	Optional

### 9.3.4 Request message body

A request body shall not be provided.

### 9.3.5 Response headers

The HTTP response headers for reading a CDMI container object using CDMI are shown in [Response headers - Read a container object using CDMI](#).

Table 56: Response headers - Read a container object using CDMI

Header	Type	Description	Requirement
Content-Type	Header string	"application/cdm-container"	Mandatory
Location	Header string	The server shall respond with an absolute URI to which the reference redirects if the object is a reference.	Conditional

### 9.3.6 Response message body

The response message body fields for reading a CDMI container object using CDMI are shown in [Table 57](#)

Table 57: Response message body - Read a container object using CDMI

Field Name	Type	Description	Requirement
objectType	JSON string	"application/cdm-container"	Mandatory
objectID	JSON string	Object ID of the object	Mandatory
objectName	JSON string	Name of the object <ul style="list-style-type: none"> <li>For objects in a container, the objectName field shall be returned.</li> <li>For objects not in a container (objects that are only accessible by ID), the "objectName" field does not exist and shall not be returned.</li> </ul>	Conditional
parentURI	JSON string	URI for the parent object <ul style="list-style-type: none"> <li>For objects in a container, the parentURI field shall be returned.</li> <li>For objects not in a container (objects that are only accessible by ID), the "parentURI" field does not exist and shall not be returned.</li> </ul> Appending the "objectName" to the "parentURI" shall always produce a valid URI for the object.	Conditional
parentID	JSON string	Object ID of the parent container object <ul style="list-style-type: none"> <li>For objects in a container, the "parentID" field shall be returned.</li> <li>For objects not in a container (objects that are only accessible by ID), the "parentID" field does not exist and shall not be returned.</li> </ul>	Conditional
domainURI	JSON string	URI of the owning domain	Mandatory
capabilitiesURI	JSON string	URI to the capabilities for the object	Mandatory

Continued on next page

Table 57 – continued from previous page

Field Name	Type	Description	Requirement
completionStatus	JSON string	A string indicating if the object is still in the process of being created or updated by another operation, and after that operation is complete, indicates if it was successfully created or updated or if an error occurred.  The value shall be the string “Processing”, the string “Complete”, or an error string starting with the value “Error”.	Mandatory
percentComplete	JSON string	A string indicating the percentage of completion if the object is still in the process of being created or updated by another operation. <ul style="list-style-type: none"> <li>When the value of completionStatus is “Processing”, this field, if provided, shall indicate the percentage of completion as a numeric integer value from 0 through 100.</li> <li>When the value of completionStatus is “Complete”, this field, if provided, shall contain the value “100”.</li> <li>When the value of completionStatus is “Error”, this field, if provided, may contain any integer value from “0” through “100”.</li> </ul>	Optional
metadata	JSON object	Metadata for the container object. This field includes any user and data system metadata specified in the request body metadata field, along with storage system metadata generated by the cloud storage system. See <a href="#">clause 16</a> for a further description of metadata.	Mandatory
exports	JSON object	A structure for each protocol that is enabled for this container object (see <a href="#">clause 13</a> )	Optional <sup>1</sup>
snapshots	JSON array of JSON strings	URIs of the snapshot container objects	Optional <sup>1</sup>
childrenrange	JSON string	The children of the container expressed as a range. If a range of children is requested, this field indicates the children returned as a range.	Mandatory
children	JSON array of JSON strings	Names of the children objects in the container object. When a client uses a child name in a request URI or a header URI, the client shall escape reserved characters according to RFC 3986 [2], e.g., a “%” character in a child name shall be replaced with “%25”. <ul style="list-style-type: none"> <li>Children that are container objects shall have “/” appended to the child name.</li> <li>Children that are references shall have “?” appended to the child name.</li> </ul>	Mandatory

If individual fields are specified in the GET request, only these fields are returned in the result body. Optional fields that are requested but do not exist are omitted from the result body.

### 9.3.7 Response status

Table 58 describes the HTTP status codes that occur when reading a container object using CDML.

<sup>1</sup> Returned only if present.



Table 58: HTTP status codes - Read a container object using CDMI

HTTP status	Description
200 OK	The metadata for the container object is provided in the message body.
302 Found	The resource is a reference to another resource.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
406 Not Acceptable	The server is unable to provide the object in the content type specified in the Accept header.

### 9.3.8 Examples

EXAMPLE 1: GET to the container object URI to read all the fields of the container object:

```
GET /MyContainer/ HTTP/1.1
Host: cloud.example.com
Accept: application/cdm-container
```

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: application/cdm-container

{
  "objectType" : "application/cdm-container",
  "objectID" : "00007ED900104E1D14771DC67C27BF8B",
  "objectName" : "MyContainer/",
  "parentURI" : "/",
  "parentID" : "00007E7F0010128E42D87EE34F5A6560",
  "domainURI" : "/cdmi_domains/MyDomain/",
  "capabilitiesURI" : "/cdmi_capabilities/container/",
  "completionStatus" : "Complete",
  "metadata" : {
    ...
  },
  "exports" : {
    "OCCE/iSCSI" : {
      "identifier": "00007E7F00104BE66AB53A9572F9F51E",
      "permissions": [
        "https://example.com/compute/0/",
        "https://example.com/compute/1/"
      ]
    },
    "Network/NFSv4" : {
      "identifier" : "/users",
      "permissions" : "domain"
    },
    "childrenrange" : "0-4",
    "children" : [
      "red",
      "green",
      "yellow",
      "orange/",
      "purple/"
    ]
  }
}
```

EXAMPLE 2: GET to the container object URI to read parentURI and children of the container object:

```
GET /MyContainer/?parentURI;children HTTP/1.1
Host: cloud.example.com
Accept: application/cdm-container
```

1749

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: application/cdmi-container

{
  "parentURI" : "/",
  "children" : [
    "red",
    "green",
    "yellow",
    "orange/",
    "purple/"
  ]
}
```

1750

**EXAMPLE 3:** GET to the container object URI to read children 0..2 and childrenrange of the container object:

```
GET /MyContainer/?childrenrange;children:0-2 HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-container
```

1751

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: application/cdmi-container

{
  "childrenrange" : "0-2",
  "children" : [
    "red",
    "green",
    "yellow"
  ]
}
```

1752

**EXAMPLE 4:** GET to the container object by ID to read children 0..2 and childrenrange of the container object:

```
GET /cdmi_objectid/0000706D0010B84FAD185C425D8B537E/?childrenrange;children:0-2 HTTP/
↪1.1
Host: cloud.example.com
Accept: application/cdmi-container
```

1753

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: application/cdmi-container

{
  "childrenrange": "0-2",
  "children": [
    "red",
    "green",
    "yellow"
  ]
}
```

## 9.4 Update a container object using CDMI

### 9.4.1 Synopsis

To update an existing container object, the following requests shall be performed:

- PUT <root URI>/<ContainerName>/<TheContainerName>
- PUT <root URI>/<ContainerName>/<TheContainerName>?metadata:<metadataname>;...
- PUT <root URI>/cdmi\_objectid/<ContainerObjectID>
- PUT <root URI>/cdmi\_objectid/<ContainerObjectID>?metadata:<metadataname>;...

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate container objects.
- <TheContainerName> is the name of the container object to be updated.
- <ContainerObjectID> is the ID of the data object to be updated.

### 9.4.2 Delayed completion of snapshot

If the creation of a snapshot (see [clause 14](#)) is requested by including a snapshot field in the request message body, the server may return an HTTP status code of 202 Accepted. Such a response has the following implications:

- With an HTTP status code of 202 Accepted, the server implies that the following checks have passed:
  - user authorization for creating the snapshot,
  - user authorization for read access to the container object, and
  - availability of space to create the snapshot or at least enough space to create a URI to report an error.
- A client might not be able to immediately access the snapshot, e.g., due to delays resulting from the implementation's use of eventual consistency.

The client performs GET operations to the snapshot URI to track the progress of the operation. In particular, the server returns two fields in its response body to indicate progress:

- A `completionStatus` field contains either "Processing", "Complete", or an error string starting with the value "Error".
- An optional `percentComplete` field contains the percentage that the accepted PUT has completed ("0" to "100"). GET does not return any value for the object when `completionStatus` is not "Complete".

When the final result of the snapshot operation is an error, the snapshot URI is created with the `completionStatus` field set to the error message. It is the client's responsibility to delete the URI after the error has been noted.

### 9.4.3 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 59](#).

Table 59: Capabilities - Update a CDMI container object using CDMI

Capability	Location	Description
cdmi_modify_metadata	Container object	Ability to modify the metadata of an existing container object
cdmi_snapshot	Container object	Ability to create a new snapshot of an existing container object
cdmi_export_<protocol>	Container object	Ability to add and modify exports for an existing container object
cdmi_object_access_by_ID	System wide capability	Ability to access the object by ID

### 9.4.4 Request headers

The HTTP request headers for updating a CDMI container object using CDMI are shown in [Table 60](#).

Table 60: Request headers - Update a container object using CDMI

Header	Type	Description	Requirement
Content-Type	Header string	"application/cdmi-container"	Mandatory

### 9.4.5 Request message body

The request message body fields for updating a container object using CDMI are shown in [Table 61](#).

Table 61: Request message body - Update a container object using CDMI

Field Name	Type	Description	Requirement
metadata	JSON object	Metadata for the container object. If present, the new metadata specified replaces the existing object metadata. If individual metadata items are specified in the URI, only those items are replaced; other items are preserved. See <a href="#">clause 16</a> for a further description of metadata.	Optional
domainURI	JSON string	URI of the owning domain <ul style="list-style-type: none"> <li>If different from the parent domain, the user shall have the "cross-domain" privilege (see <a href="#">cdmi_member_privileges</a> in <a href="#">Table 80</a>).</li> <li>If not specified, the parent domain shall be used.</li> </ul>	Optional

Continued on next page

Table 61 – continued from previous page

Field Name	Type	Description	Requirement
snapshot	JSON string	<p>Name of the snapshot to be taken. This is not a URL, but rather, the final component of the absolute URL where the snapshot will exist when the snapshot operation successfully completes.</p> <ul style="list-style-type: none"> <li>• If a snapshot is added or changed, the PUT operation only returns after the snapshot is added to the snapshot list.</li> <li>• After they are created, snapshots may be accessed as children container objects under the <code>cdmi_snapshots</code> child container object of the container object receiving a snapshot.</li> <li>• When creating a snapshot with the same name as an existing snapshot, the new snapshot will replace the existing snapshot.</li> </ul>	Optional
deserialize	JSON string	<p>URI of a CDMI data object with a value that contains a container object serialized as specified in <a href="#">clause 15</a>. The serialized container object shall be deserialized to update the existing container object.</p> <p>The object ID of the serialized container object shall match the object ID of the destination container object. Otherwise, the server shall return an HTTP status code of <code>400 Bad Request</code>.</p> <ul style="list-style-type: none"> <li>• If the serialized container object does not contain children, the update is applied only to the container object, and any existing children are left as is.</li> <li>• If the serialized container object does contain children, then creates, updates, and deletes are recursively applied for each child, depending on the differences between the provided serialized state and the current state of the child.</li> </ul>	Optional <sup>1</sup>
copy	JSON string	<p>URI of a CDMI container object that shall be copied into the existing container object. Only the contents of the container object itself shall be copied, not any children of the container object.</p> <ul style="list-style-type: none"> <li>• If the destination container object URI and the copy source object URI both do not specify individual fields, the destination container object shall be replaced with the source container object, including all child objects under the source container object.</li> <li>• If the destination container object URI or the copy source object URI specifies individual fields, only the fields specified shall be used to update the destination container object. If specified fields are not present in the source, these fields shall be ignored.</li> <li>• If the destination container object URI and the copy source object URI both specify fields, an HTTP status code of <code>400 Bad Request</code> shall be returned to the client.</li> </ul> <p>When copying a container object, exported protocols are not preserved across the copy.</p> <p>If there are insufficient permissions to read the container object at the source URI or create the container object at the destination URI, or if the read operation fails, the copy shall return an HTTP status code of <code>400 Bad Request</code>, and the destination container object shall not be updated.</p>	Optional <sup>1</sup>

Continued on next page

Table 61 – continued from previous page

Field Name	Type	Description	Requirement
deserializevalue	JSON sting	<p>A container object serialized as specified in <a href="#">clause 15</a> and encoded using base 64 encoding rules described in RFC 4648 [19], that shall be deserialized to update the existing container object.</p> <p>The object ID of the serialized container object shall match the object ID of the destination container object. Otherwise, the server shall return an HTTP status code of 400 <i>Bad Request</i>.</p> <ul style="list-style-type: none"> <li>• If the serialized container object does not contain children, the update is applied only to the container object, and any existing children are left as is.</li> <li>• If the serialized container object does contain children, then creates, updates, and deletes are recursively applied for each child, depending on the differences between the provided serialized state and the current state of the children.</li> </ul>	Optional <sup>1</sup>
exports	JSON object	A structure for each protocol that is enabled for this container object (see <a href="#">clause 13</a> ). If an exported protocol is added or changed, the PUT operation only returns after the export operation has completed.	Optional

<sup>1</sup> Only one of these fields shall be specified in any given operation. Except for value, these fields shall not be stored.

## 9.4.6 Response headers

The HTTP response header for updating a CDMI container object using CDMI is shown in Table 62.

Table 62: Response header - Update a container object using CDMI

Header	Type	Description	Requirement
Location	Header string	The server shall respond with an absolute URI to which the reference redirects if the object is a reference.	Conditional

## 9.4.7 Response message body

A response body can be provided as per RFC 2616 [23].

## 9.4.8 Response status

Table 63 describes the HTTP status codes that occur when updating a container object using CDMI.

Table 63: HTTP status codes - Update a container object using CDMI

HTTP status	Description
204 No Content	The data object content was returned in the response.
202 Accepted	The container or snapshot (subcontainer object) is in the process of being created. The CDMI client should monitor the <code>completionStatus</code> and <code>percentComplete</code> fields to determine the current status of the operation.
302 Found	The resource is a reference to another resource.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server.

## 9.4.9 Examples

EXAMPLE 1: PUT to the container object URI to replace all metadata with new metadata:

```
PUT /MyContainer/ HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdm-container

{
  "metadata": {
    "colour": "red",
    "number": "7"
  }
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

EXAMPLE 2: PUT to the container object URI to set a new exported protocol value:

```
PUT /MyContainer/?exports HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdm-container
```

(continues on next page)

(continued from previous page)

```
{
  "exports" : {
    "OCCT/iSCSI" : {
      "identifier" : "00007ED900104E1D14771DC67C27BF8B",
      "permissions" : "00007E7F00104EB781F900791C70106C"
    },
    "Network/NFSv4" : {
      "identifier" : "/users",
      "permissions" : "domain"
    }
  }
}
```

1801

The following shows the response.

```
HTTP/1.1 204 No Content
```



## 9.5 Delete a container object using CDMI

### 9.5.1 Synopsis

To delete an existing container object, including all contained children and snapshots, the following requests shall be performed:

- DELETE <root URI>/<ContainerName>/<TheContainerName>
- DELETE <root URI>/cdmi\_objectid/<ContainerObjectID>

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate container objects.
- <TheContainerName> is the name of the container object to be deleted.
- <ContainerObjectID> is the ID of the container object to be deleted.

### 9.5.2 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 64](#).

Table 64: Capabilities - Delete a CDMI container object using CDMI

Capability	Location	Description
cdmi_delete_container	Container object	Ability to delete an existing container object
cdmi_object_access_by_ID	System wide capability	Ability to access the object by ID

### 9.5.3 Request headers

Request headers can be provided as per RFC 2616 [23].

### 9.5.4 Request message body

A request body can be provided as per RFC 2616 [23].

### 9.5.5 Response headers

Response headers can be provided as per RFC 2616 [23].

### 9.5.6 Response message body

A response body can be provided as per RFC 2616 [23].

### 9.5.7 Response status

[Table 65](#) describes the HTTP status codes that occur when deleting a container object using CDMI.

Table 65: HTTP status codes - Delete a container object using CDMI

HTTP status	Description
204 No Content	The container object was successfully deleted.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server.

### 9.5.8 Example

EXAMPLE 1: DELETE to the container object URI:

```
DELETE /MyContainer/ HTTP/1.1
Host: cloud.example.com
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

EXAMPLE 2: DELETE by container object ID:

```
DELETE /cdmi_objectid/00007ED900104E1D14771DC67C27BF8B/ HTTP/1.1
Host: cloud.example.com
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

## 9.6 Create (POST) a new data object using CDMI

### 9.6.1 Synopsis

To create a new data object in a specified container, the following request shall be performed:

- POST <root URI>/<ContainerName>/

To create a new data object where the data object does not belong to a container and is only accessible by ID (see 5.3.1), the following request shall be performed:

- POST <root URI>/cdmi\_objectid/

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate container objects that already exist, with one slash (i.e., "/") between each pair of container object names.
- <DataObjectName> is the name specified for the data object to be created.

If created in a container, the data object shall be accessible as a child of the container with a server-assigned name, and shall also be accessible at <root URI>/cdmi\_objectid/<objectID>.

If created in "/cdmi\_objectid/", the data object shall only be accessible at <root URI>/cdmi\_objectid/<objectID>.

### 9.6.2 Delayed completion of create

In response to a create operation for a data object, the server may return an HTTP status code of 202 Accepted to indicate that the object is in the process of being created. This response is useful for long-running operations (e.g., copying a large data object from a source URI). Such a response has the following implications.

- The server shall return a Location header with an absolute URI to the object to be created along with an HTTP status code of 202 Accepted.
- With an HTTP status code of 202 Accepted, the server implies that the following checks have passed:
  - user authorization for creating the object;
  - user authorization for read access to any source object for move, copy, serialize, or deserialize; and
  - availability of space to create the object or at least enough space to create a URI to report an error.
- A client might not be able to immediately access the created object, e.g., due to delays resulting from the implementation's use of eventual consistency.

The client performs GET operations to the URI to track the progress of the operation. In response, the server returns two fields in its response body to indicate progress.

- A mandatory completionStatus text field contains either "Processing", "Complete", or an error string starting with the value "Error".
- An optional percentComplete field contains the percentage of the operation that has completed (0 to 100).

GET shall not return any value for the data object when completionStatus is not "Complete". If the final result of the create operation is an error, the URI is created with the completionStatus field set to the error message. It is the client's responsibility to delete the URI after the error has been noted.

### 9.6.3 Capabilities

Capabilities that indicate which operations are supported are shown in Table 66.

Table 66: Capabilities - Create a CDMI data object using CDMI

Capability	Location	Description
cdmi_post_dataobject cdmi_create_dataobject	Parent container	Ability to create a new data object
cdmi_create_reference	Parent container	Ability to create a new reference
cdmi_copy_dataobject	Parent container	Ability to create a data object that is a copy of another data object
cdmi_move_dataobject	Parent container	Ability to move a data object from another container
cdmi_deserialize_dataobject	Parent container	Ability to create a data object that is deserialized from the contents of the PUT or the contents of another data object
cdmi_serialize_dataobject cdmi_serialize_container cdmi_serialize_domain cdmi_serialize_queue	Parent container	Ability to create a data object that contains a serialized representation of an existing data object, container, domain or queue
cdmi_create_value_range	Parent container	Ability to create a data object using a specified byte range
cdmi_post_dataobject_by_ID	System wide capability	Ability to create a new data object in “/cdmi_objectid/”
cdmi_create_reference_by_ID	System wide capability	Ability to create a new reference in “/cdmi_objectid/”
cdmi_copy_dataobject_by_ID	System wide capability	Ability to create a data object in “/cdmi_objectid/” that is a copy of another data object
cdmi_object_move_to_ID	System wide capability	Ability to move a data object to “/cdmi_objectid/” from another container
cdmi_deserialize_dataobject_by_ID	System wide capability	Ability to create a data object in “/cdmi_objectid/” that is deserialized from the contents of the PUT or the contents of another data object
cdmi_serialize_dataobject_to_ID cdmi_serialize_container_to_ID cdmi_serialize_domain_to_ID cdmi_serialize_queue_to_ID	System wide capability	Ability to create a data object in “/cdmi_objectid/” that contains a serialized representation of an existing data object, container, domain or queue
cdmi_create_value_range_by_ID	System wide capability	Ability to create a data object in “/cdmi_objectid/” using a specified byte range
cdmi_multipart_mime	System wide capability	Ability to create a data object using multi-part MIME

### 9.6.4 Request headers

The HTTP request headers for creating a new CDMI data object using CDMI are shown in Table 67.

Table 67: Request headers - Create a new data object Using CDMI

Header	Type	Description	Requirement
Accept	Header string	“application/cdmi-object” or a consistent value as described in 5.5.2	Optional

Continued on next page

Table 67 – continued from previous page

Header	Type	Description	Requirement
Content-Type	Header string	<p>“application/cdm-object” or “multipart/mixed”</p> <ul style="list-style-type: none"> <li>If “multipart/mixed” is specified, the body shall consist of at least two MIME parts, where the first part shall contain a body of content-type “application/cdm-object”, and the second and subsequent parts shall contain one or more byte ranges of the value.</li> <li>If multiple byte ranges are included and the Content-Range header is omitted for a part, the data in the part shall be appended to the data in the preceding part, with the first part having a byte offset of zero.</li> </ul>	Mandatory
X-CDMI-Partial	Header string	Indicates that the newly created object is part of a series of writes and has not yet been fully created. When set to “true”, the completionStatus field shall be set to “Processing”. X-CDMI-Partial works across CDMI and non-CDMI operations.	Optional

## 9.6.5 Request message body

The request message body fields for creating a new data object using CDMI are shown in Table 68.

Table 68: Request message body - Create a new data object Using CDMI

Field Name	Type	Description	Requirement
mimetype	JSON string	<p>MIME type of the data contained within the value field of the data object</p> <ul style="list-style-type: none"> <li>This field may be included when creating by value or when deserializing, serializing, copying, and moving a data object.</li> <li>If this field is not included and multi-part MIME is not being used, the value of “text/plain” shall be assigned as the field value.</li> <li>If this field is not included and multi-part MIME is being used, the value of the Content-Type header of the second MIME part shall be assigned as the field value.</li> <li>This field value shall be converted to lower case before being stored.</li> </ul>	Optional
metadata	JSON object	<p>Metadata for the data object</p> <ul style="list-style-type: none"> <li>If this field is included, the contents of the JSON object provided in this field shall be used as data object metadata.</li> <li>If this field is included when deserializing, serializing, copying, or moving a data object, the contents of the JSON object provided in this field shall be used as object metadata instead of the metadata from the source URI.</li> <li>If this field is not included, no user-specified metadata shall be added to the object.</li> <li>If this field is not included when deserializing, serializing, copying, or moving a data object, metadata from the source URI shall be used.</li> <li>This field shall not be included when creating a reference to a data object.</li> </ul>	Optional

Continued on next page

Table 68 – continued from previous page

Field Name	Type	Description	Requirement
domainURI	JSON string	URI of the owning domain <ul style="list-style-type: none"> <li>If different from the parent domain, the user shall have the “cross-domain” privilege (see <code>cdmi_member_privileges</code> in Table 80).</li> <li>If not specified, the domain of the parent container shall be used.</li> </ul>	Optional
deserialize	JSON string	URI of a CDMI data object with a value that contains a data object serialized as specified in clause 15. The serialized data object shall be deserialized to create the new data object.	Optional <sup>1</sup>
serialize	JSON string	URI of a CDMI object that shall be serialized into the new data object	Optional <sup>1</sup>
copy	JSON string	URI of a source CDMI data object or queue object that shall be copied into the new destination data object. <ul style="list-style-type: none"> <li>If the destination data object URI and the copy source object URI both do not specify individual fields, the destination data object shall be a complete copy of the source data object.</li> <li>If the destination data object URI or the copy source object URI specifies individual fields, only the fields specified shall be used to create the destination data object. If specified fields are not present in the source, default field values shall be used.</li> <li>If the destination data object URI and the copy source object URI both specify fields, an HTTP status code of 400 <i>Bad Request</i> shall be returned to the client.</li> <li>If the copy source object URI points to a queue object, as part of the copy operation, multiple queue values shall be concatenated into a single data object value.</li> <li>If the copy source object URI points to one or more queue object values, as part of the copy operation, the specified queue values shall be concatenated into a single data object value.</li> <li>If there are insufficient permissions to read the data object at the source URI or create the data object at the destination URI, or if the read operation fails, the copy shall return an HTTP status code of 400 <i>Bad Request</i>, and the destination object shall not be created.</li> </ul>	Optional <sup>1</sup>
move	JSON string	URI of an existing local or remote CDMI data object (source URI) that shall be relocated to the URI specified in the PUT. The contents of the object, including the object ID, shall be preserved by a move, and the data object at the source URI shall be removed after the data object at the destination has been successfully created. <p>If there are insufficient permissions to read the data object at the source URI, write the data object at the destination URI, or delete the data object at the source URI, or if any of these operations fail, the move shall return an HTTP status code of 400 <i>Bad Request</i>, and the source and destination are left unchanged.</p>	Optional <sup>1</sup>
reference	JSON string	URI of a CDMI data object that shall be redirected to by a reference. If any other fields are supplied when creating a reference, the server shall respond with an HTTP status code of 400 <i>Bad Request</i> .	Optional <sup>1</sup>

Continued on next page

Table 68 – continued from previous page

Field Name	Type	Description	Requirement
deserializevalue	JSON string	<p>A data object serialized as specified in <a href="#">clause 15</a> and encoded using base 64 encoding rules described in RFC 4648 [19], that shall be deserialized to create the new data object.</p> <ul style="list-style-type: none"> <li>• If multi-part MIME is being used and this field contains the value of the MIME boundary parameter, the contents of the second MIME part shall be assigned as the field value.</li> <li>• If the serialized data object in the second MIME part does not include a value field, the contents of the third MIME part shall be assigned as the field value of the value field.</li> </ul>	Optional <sup>1</sup>
valuetransferencoding	JSON string	<p>The value transfer encoding used for the data object value. Three value transfer encodings are defined:</p> <ul style="list-style-type: none"> <li>• “utf-8” indicates that the data object contains a valid UTF-8 string, and it shall be transported as a UTF-8 string in the value field.</li> <li>• “base64” indicates that the data object may contain arbitrary binary sequences, and it shall be transported as a base 64-encoded string in the value field. Setting the contents of the data object value field to any value other than a valid base 64 string shall result in an HTTP status code of 400 <i>Bad Request</i> being returned to the client.</li> <li>• “json” indicates that the data object contains a valid JSON object, and the <i>value</i> field shall be a JSON object containing valid JSON data. If the contents of the <i>value</i> field are set to any value other than a valid JSON object, an HTTP status code of 400 <i>Bad Request</i> shall be returned to the client.</li> <li>• This field shall only be included when creating a data object by value.</li> <li>• If this field is not included and multi-part MIME is not being used, the value of “utf-8” shall be assigned as the field value.</li> <li>• If this field is not included and multi-part MIME is being used, the value of “utf-8” shall be assigned as the field value if the Content-Type header of the second and all MIME parts includes the charset parameter as defined in RFC 2046 of “utf-8” (e.g., “; charset=utf-8”). Otherwise, the value of “base64” shall be assigned as the field value. This field applies only to the encoding of the value when represented in JSON; the <i>Content-Transfer-Encoding</i> header of the part specifies the encoding of the value within a multi-part MIME request, as defined in RFC 2045 [9].</li> </ul>	Optional <sup>1</sup>

Continued on next page

Table 68 – continued from previous page

Field Name	Type	Description	Requirement
value	JSON string	<p>The data object value</p> <ul style="list-style-type: none"> <li>• If this field is not included and multi-part MIME is not being used, an empty JSON String (i.e., "") shall be assigned as the field value.</li> <li>• If this field is not included and multi-part MIME is being used, the contents of the second MIME part shall be assigned as the field value.</li> <li>• If the <code>valuetransferencoding</code> field indicates UTF-8 encoding, the value shall be a UTF-8 string escaped using the JSON escaping rules described in RFC 4627 [5].</li> <li>• If the <code>valuetransferencoding</code> field indicates base 64 encoding, the value shall be first encoded using the base 64 encoding rules described in RFC 4648 [19].</li> <li>• If the <code>valuetransferencoding</code> field indicates JSON encoding, the value shall contain a valid JSON object.</li> </ul>	Optional <sup>1</sup>

<sup>1</sup> Only one of these fields shall be specified in any given operation. Except for value, these fields shall not be stored. If more than one of these fields is supplied, the server shall respond with an HTTP status code of 400 `Bad Request`.



## 9.6.6 Response headers

The HTTP response headers for creating a new CDMI data object using CDMI are shown in Table 69.

Table 69: Response headers - Create a new data object using CDMI

Header	Type	Description	Requirement
Content-Type	Header string	"application/cdm-object"	Mandatory
Location	Header string	The unique absolute URI for the new data object as assigned by the system. In the absence of file name information from the client, the system shall assign the URI in the form: http://host:port/<root URI>/<ContainerName>/<ObjectID> or https://host:port/<root URI>/<ContainerName>/<ObjectID>.	Mandatory

## 9.6.7 Response message body

The response message body fields for creating a new CDMI data object using CDMI are shown in Table 70.

Table 70: Response message body - Create a new data object using CDMI

Field Name	Type	Description	Requirement
objectType	JSON string	"application/cdm-object"	Mandatory
objectID	JSON string	Object ID of the object	Mandatory
objectName	JSON string	Name of the object <ul style="list-style-type: none"> <li>For objects in a container, the objectName field shall be returned.</li> <li>For objects not in a container (objects that are only accessible by ID), the objectName field does not exist and shall not be returned.</li> </ul>	Conditional
parentURI	JSON string	URI for the parent object <ul style="list-style-type: none"> <li>For objects in a container, the parentURI field shall be returned.</li> <li>For objects not in a container (objects that are only accessible by ID), the parentURI field does not exist and shall not be returned.</li> </ul> Appending the objectName to the parentURI shall always produce a valid URI for the object.	Conditional
parentID	JSON string	Object ID of the parent container object <ul style="list-style-type: none"> <li>For objects in a container, the parentID field shall be returned.</li> <li>For objects not in a container (objects that are only accessible by ID), the parentID field does not exist and shall not be returned.</li> </ul>	Conditional
domainURI	JSON string	URI of the owning domain	Mandatory
capabilitiesURI	JSON string	URI to the capabilities for the object	Mandatory

Continued on next page

Table 70 – continued from previous page

Field Name	Type	Description	Requirement
completionStatus	JSON string	A string indicating if the object is still in the process of being created or updated by another operation, and after that operation is complete, indicates if it was successfully created or updated or if an error occurred.  The value shall be the string “Processing”, the string “Complete”, or an error string starting with the value “Error”.	Mandatory
percentComplete	JSON string	A string indicating the percentage of completion if the object is still in the process of being created or updated by another operation. <ul style="list-style-type: none"> <li>When the value of completionStatus is “Processing”, this field, if provided, shall indicate the percentage of completion as a numeric integer value from “0” through “100”.</li> <li>When the value of completionStatus is “Complete”, this field, if provided, shall contain the value “100”.</li> <li>When the value of completionStatus is “Error”, this field, if provided, may contain any integer value from “0” through “100”.</li> </ul>	Optional
mimetype	JSON string	MIME type of the value of the data object	Mandatory
metadata	JSON object	Metadata for the data object. This field includes any user and data system metadata specified in the request body metadata field, along with storage system metadata generated by the cloud storage system. See <a href="#">clause 16</a> for a further description of metadata.	Mandatory

## 9.6.8 Response status

Table 71 describes the HTTP status codes that occur when creating a new data object using CDMI.

Table 71: HTTP status codes - Create a new data object using CDMI

HTTP status	Description
201 Created	The new data object was created.
202 Accepted	The data object is in the process of being created. The CDMI client should monitor the completionStatus and percentComplete fields to determine the current status of the operation.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server.

## 9.6.9 Examples

EXAMPLE 1: POST to the container object URI the data object contents:

```
POST /MyContainer/ HTTP/1.1
Host: cloud.example.com
Accept: application/cdm-object
Content-Type: application/cdm-object

{
```

(continues on next page)

(continued from previous page)

```

    "mimetype" : "text/plain",
    "metadata" : {

    },
    "value" : "This is the Value of this Data Object"
}

```

1886 The following shows the response.

```

HTTP/1.1 201 Created
Content-Type: application/cdm-object
Location: https://cloud.example.com/MyContainer/00007ED900104E1D14771DC67C27BF8B

{
  "objectType" : "application/cdm-object",
  "objectID" : "00007ED900104E1D14771DC67C27BF8B",
  "objectName" : "00007ED900104E1D14771DC67C27BF8B",
  "parentURI" : "/MyContainer/",
  "parentID" : "00007ED900104E1D14771DC67C27BF8B",
  "domainURI" : "/cdmi_domains/MyDomain/",
  "capabilitiesURI" : "/cdmi_capabilities/dataobject/",
  "completionStatus" : "Complete",
  "mimetype" : "text/plain",
  "metadata" : {
    ...
  }
}

```

1887 EXAMPLE 2: POST to the object ID URI the data object contents:

```

POST /cdmi_objectid/ HTTP/1.1
Host: cloud.example.com
Accept: application/cdm-object
Content-Type: application/cdm-object

{
  "mimetype": "text/plain",
  "domainURI": "/cdmi_domains/MyDomain/",
  "value": "This is the Value of this Data Object"
}

```

1888 The following shows the response.

```

HTTP/1.1 201 Created
Location: https://cloud.example.com/cdm_objectid/00007ED900104E1D14771DC67C27BF8B
Content-Type: application/cdm-object

{
  "objectType": "application/cdm-object",
  "objectID": "00007ED900104E1D14771DC67C27BF8B",
  "domainURI": "/cdmi_domains/MyDomain/",
  "capabilitiesURI": "/cdmi_capabilities/dataobject/",
  "completionStatus": "Complete",
  "mimetype": "text/plain",
  "metadata": {
    "cdmi_acl": [
      {
        "acetype": "ALLOW",
        "identifier": "OWNER@",
        "aceflags": "NO_FLAGS",
        "acemask": "ALL_PERMS"
      },
      ...
    ],
    ...
  }
}

```

1889 EXAMPLE 3: POST to the object ID URI the data object fields and binary contents using multi-part MIME:

```
POST /cdmi_objectid/ HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-object
Content-Type: multipart/mixed; boundary=gc0p4Jq0M2Yt08j34c0p

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/cdmi-object

{
  "domainURI": "/cdmi_domains/MyDomain/",
  "metadata": {
    "colour": "blue"
  }
}

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/octet-stream
Content-Transfer-Encoding: binary

<37 bytes of binary data>

--gc0p4Jq0M2Yt08j34c0p--
```

1890 The following shows the response.

```
HTTP/1.1 201 Created
Location: https://cloud.example.com/cdmi_objectid/00007ED90010C2414303B5C6D4F83170
Content-Type: application/cdmi-object

{
  "objectType": "application/cdmi-object",
  "objectID": "00007ED90010C2414303B5C6D4F83170",
  "domainURI": "/cdmi_domains/MyDomain/",
  "capabilitiesURI": "/cdmi_capabilities/dataobject/",
  "completionStatus": "Complete",
  "mimetype": "application/octet-stream",
  "metadata": {
    "cdmi_size": "37",
    "colour": "blue",
    ...
  }
}
```

## 9.7 Create (POST) a new queue object using CDMI

### 9.7.1 Synopsis

To create a new queue object (see [clause 11](#)) in a specified container where the name of the queue object is a server-assigned object identifier, the following request shall be performed:

- POST <root URI>/<ContainerName>/

To create a new queue object where the queue object does not belong to a container and is only accessible by ID (see [5.3.1](#)), the following request shall be performed:

- POST <root URI>/cdmi\_objectid/

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate container objects that already exist, with one slash (i.e., "/") between each pair of container object names.

If created in a container, the queue object shall be accessible as a child of the container with a server-assigned name, and shall also be accessible at <root URI>/cdmi\_objectid/<objectID>.

If created in "/cdmi\_objectid/", the queue object shall only be accessible at <root URI>/cdmi\_objectid/<objectID>.

### 9.7.2 Delayed completion of create

In response to a create operation for a queue object, the server may return an HTTP status code of 202 Accepted to indicate that the object is in the process of being created. This response is useful for long-running operations (e.g., copying a large number of queue values from a source URI). Such a response has the following implications.

- The server shall return a Location header with an absolute URI to the object to be created along with an HTTP status code of 202 Accepted.
- With an HTTP status code of 202 Accepted, the server implies that the following checks have passed:
  - user authorization for creating the object;
  - user authorization for read access to any source object for move, copy, serialize, or deserialize; and
  - availability of space to create the object or at least enough space to create a URI to report an error.
- A client might not be able to immediately access the created object, e.g., due to delays resulting from the implementation's use of eventual consistency.

The client performs GET operations to the URI to track the progress of the operation. In response, the server returns two fields in its response body to indicate progress.

- A mandatory `completionStatus` text field contains either "Processing", "Complete", or an error string starting with the value "Error".
- An optional `percentComplete` field contains the percentage of the operation that has completed (0 to 100).

GET shall not return any value for the queue object when `completionStatus` is not "Complete". If the final result of the create operation is an error, the URI is created with the `completionStatus` field set to the error message. It is the client's responsibility to delete the URI after the error has been noted.

### 9.7.3 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 72](#).

Table 72: Capabilities - Create a CDMI Queue object using CDMI

Capability	Location	Description
cdmi_post_queue cdmi_create_queue	Parent container	Ability to create a new queue object
cdmi_create_reference	Parent container	Ability to create a new reference
cdmi_copy_queue	Parent container	Ability to create a queue object that is a copy of another queue object
cdmi_move_queue	Parent container	Ability to move a queue object from another container
cdmi_deserialize_queue	Parent container	Ability to create a queue object that is deserialized from the contents of the PUT or the contents of another queue object
cdmi_post_queue_by_ID	System wide capability	Ability to create a new queue object in <code>/cdmi_objectid/</code>
cdmi_create_reference_by_ID	System wide capability	Ability to create a new reference in <code>/cdmi_objectid/</code>
cdmi_copy_queue_by_ID	System wide capability	Ability to create a queue object in <code>/cdmi_objectid/</code> that is a copy of another queue object
cdmi_object_move_to_ID	System wide capability	Ability to move a queue object to <code>/cdmi_objectid/</code> from another container
cdmi_deserialize_queue_by_ID	System wide capability	Ability to create a queue object in <code>/cdmi_objectid/</code> that is deserialized from the contents of the PUT or the contents of another data object
cdmi_serialize_dataobject_to_ID cdmi_serialize_container_to_ID cdmi_serialize_domain_to_ID cdmi_serialize_queue_to_ID	System wide capability	Ability to create a data object in <code>/cdmi_objectid/</code> that contains a serialized representation of an existing data object, container, domain or queue

### 9.7.4 Request headers

The HTTP request headers for creating a new CDMI queue object using CDMI are shown in [Table 73](#).

Table 73: Request headers - Create a new queue object using CDMI

Header	Type	Description	Requirement
Accept	Header string	<code>application/cdmi-object</code> or a consistent value as described in <a href="#">5.5.2</a>	Optional
Content-Type	Header string	<code>application/cdmi-queue</code>	Mandatory
Content-Range	Header string	A valid ranges-specifier (see RFC 2616 [23] Section 14.35.1)	Optional

### 9.7.5 Request message body

The request message body fields for creating a new queue object using CDMI are shown in [Table 74](#).

Table 74: Request message body - Create a new queue object using CDMI

Field Name	Type	Description	Requirement
metadata	JSON object	Metadata for the queue object <ul style="list-style-type: none"> <li>If this field is included, the contents of the JSON object provided in this field shall be used as queue object metadata.</li> <li>If this field is included when deserializing, serializing, copying, or moving a queue object, the contents of the JSON object provided in this field shall be used as object metadata instead of the metadata from the source URI.</li> <li>If this field is not included, no user-specified metadata shall be added to the object.</li> <li>If this field is not included when deserializing, serializing, copying, or moving a queue object, metadata from the source URI shall be used.</li> <li>This field shall not be included when creating a reference to a queue object.</li> </ul>	Optional
domainURI	JSON string	URI of the owning domain <ul style="list-style-type: none"> <li>If different from the parent domain, the user shall have the “cross-domain” privilege (see <code>cdmi_member_privileges</code> in Table 80).</li> <li>If not specified, the domain of the parent container shall be used.</li> </ul>	Optional
deserialize	JSON string	URI of a CDMI data object with a value that contains a queue object serialized as specified in clause 15. The serialized queue object shall be deserialized to create the new queue object.	Optional <sup>1</sup>
copy	JSON string	URI of a CDMI queue object that will be copied into the new queue object	Optional <sup>1</sup>
move	JSON string	URI of a CDMI queue object that will be copied into the new queue object. When the copy is successfully completed, the queue object at the source URI is removed.	Optional <sup>1</sup>
reference	JSON string	URI of a CDMI queue object that shall be redirected to by a reference. If other fields are supplied when creating a reference, the server shall respond with an HTTP status code of 400 <i>Bad Request</i> .	Optional <sup>1</sup>
deserializevalue	JSON string	A queue object serialized as specified in clause 15 and encoded using base 64 encoding rules described in RFC 4648 [19], that shall be deserialized to create the new queue object.	Optional <sup>1</sup>

<sup>1</sup> Only one of these fields shall be specified in any given operation. Except for value, these fields shall not be stored. If more than one of these fields is supplied, the server shall respond with an HTTP status code of 400 *Bad Request*.

## 9.7.6 Response headers

The response headers for creating a new CDMI queue object using CDMI are shown in [Table 75](#).

Table 75: Response headers - Create a new queue object using CDMI

Header	Type	Description	Requirement
Content-Type	Header string	"application/cdm-queue"	Mandatory
Location	Header string	The unique absolute URI for the new queue object as assigned by the system.	Mandatory

## 9.7.7 Response message body

The response message body fields for creating a new CDMI queue object using CDMI are shown in [Table 76](#).

Table 76: Response message body - Create a new queue object using CDMI

Field Name	Type	Description	Requirement
objectType	JSON string	"application/cdm-queue"	Mandatory
objectID	JSON string	Object ID of the object	Mandatory
objectName	JSON string	Name of the object <ul style="list-style-type: none"> <li>For objects in a container, the objectName field shall be returned.</li> <li>For objects not in a container (objects that are only accessible by ID), the objectName field does not exist and shall not be returned.</li> </ul>	Conditional
parentURI	JSON string	URI for the parent object <ul style="list-style-type: none"> <li>For objects in a container, the parentURI field shall be returned.</li> <li>For objects not in a container (objects that are only accessible by ID), the parentURI field does not exist and shall not be returned.</li> </ul> Append the objectName to the parentURI shall always produce a valid URI for the object.	Conditional
parentID	JSON string	Object ID of the parent container object <ul style="list-style-type: none"> <li>For objects in a container, the parentID field shall be returned.</li> <li>For objects not in a container (objects that are only accessible by ID), the parentID field does not exist and shall not be returned.</li> </ul>	Conditional
domainURI	JSON string	URI of the owning domain	Mandatory
capabilitiesURI	JSON string	URI to the capabilities for the object	Mandatory
completionStatus	JSON string	A string indicating if the object is still in the process of being created or updated by another operation, and after that operation is complete, indicates if it was successfully created or updated or if an error occurred.  The value shall be the string "Processing", the string "Complete", or an error string starting with the value "Error".	Mandatory

Continued on next page



Table 76 – continued from previous page

Field Name	Type	Description	Requirement
percentComplete	JSON string	<p>A string indicating the percentage of completion if the object is still in the process of being created or updated by another operation.</p> <ul style="list-style-type: none"> <li>When the value of <code>completionStatus</code> is "Processing", this field, if provided, shall indicate the percentage of completion as a numeric integer value from "0" through "100".</li> <li>When the value of <code>completionStatus</code> is "Complete", this field, if provided, shall contain the value "100".</li> <li>When the value of <code>completionStatus</code> is "Error", this field, if provided, may contain any integer value from "0" through "100".</li> </ul>	Optional
metadata	JSON object	Metadata for the queue object. This field includes any user and data system metadata specified in the request body metadata field, along with storage system metadata generated by the cloud storage system. See <a href="#">clause 16</a> for a further description of metadata.	Mandatory
queueValues	JSON string	The range of designators for enqueued values. Every enqueued value shall be assigned a unique, monotonically-incrementing positive integer designator, starting from 0. If no values are enqueued, an empty string shall be returned. If values are enqueued, the lowest designator, followed by a hyphen ("-"), followed by the highest designator shall be returned.	Mandatory

### 9.7.8 Response status

Table 77 describes the HTTP status codes that occur when creating a new queue object using CDMI.

Table 77: HTTP status codes - Create a new queue object using CDMI

HTTP status	Description
201 Created	The new queue object was created.
202 Accepted	The queue object is in the process of being created. The CDMI client should monitor the <code>completionStatus</code> and <code>percentComplete</code> fields to determine the current status of the operation.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server.

### 9.7.9 Example

EXAMPLE 1: POST to the container object URI the queue object contents:

```
POST /MyContainer/ HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-queue
Accept: application/cdmi-queue

{
}
```

The following shows the response.

```
HTTP/1.1 201 Created
Content-Type: application/cdmi-queue
Location: https://cloud.example.com/MyContainer/00007ED900104E1D14771DC67C27BF8B

{
  "objectType" : "application/cdmi-queue",
  "objectID" : "00007ED900104E1D14771DC67C27BF8B",
  "objectName" : "00007ED900104E1D14771DC67C27BF8B",
  "parentURI" : "/MyContainer/",
  "parentID" : "00007ED900104E1D14771DC67C27BF8B",
  "domainURI" : "/cdmi_domains/MyDomain/",
  "capabilitiesURI" : "/cdmi_capabilities/queue/",
  "completionStatus" : "Complete",
  "metadata" : {
    ...
  },
  "queueValues" : ""
}
```

1948

## **Part IV**

1949

# **CDMI Advanced**

## Clause 10

# Domain object resource operations using CDMI

## 10.1 Overview

Domain objects represent the concept of administrative ownership of stored data within a CDMI™ storage system. Each object may be owned and managed by a different administrative entity, which is expressed as a domain.

If a cloud storage system supports domains, the `cdmi_domains` system-wide capability shall be present, and the `cdmi_domains` container shall be present in the CDMI root container.

A cloud storage system may include a hierarchy of domains that provide access to domain-related information within a CDMI context. This domain hierarchy is a series of CDMI objects that correspond to parent and child domains, with each domain corresponding to logical groupings of objects that are to be managed together. Domain measurement information about objects that are associated with each domain flow up to parent domains, facilitating billing and management operations that are typical for a cloud storage environment.

Fig. 6 shows the hierarchy of domains and shows how the `domainURI` links data objects, container objects and queue objects into the domain hierarchy.

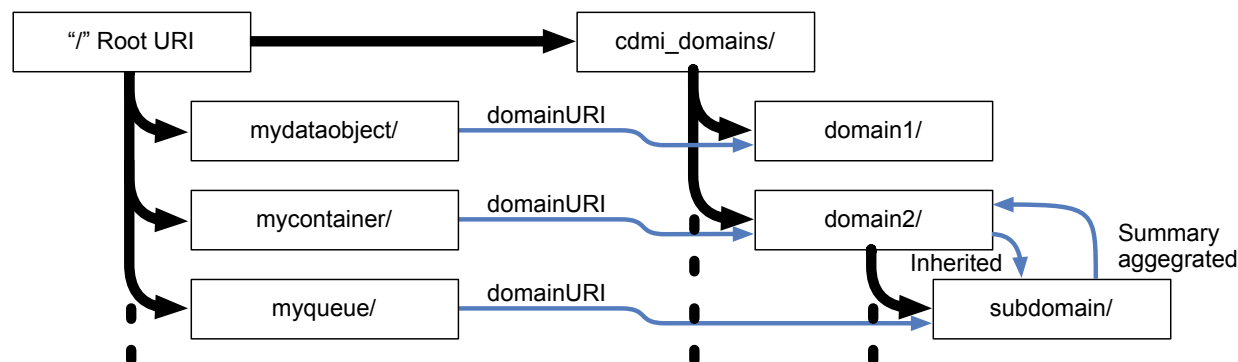


Figure 6: Hierarchy of domains

Each CDMI domain object is represented as a JSON object, containing one or more “fields”. For example, the “metadata” field contains metadata items.

EXAMPLE 1: CDMI domain object

```
{
  "objectType" : "application/cdmi-domain",
  "objectID" : "00007E7F00104BE66AB53A9572F9F51E",
  "objectName" : "MyDomain/",
  "parentURI" : "/cdmi_domains/",
  "parentID" : "00007E7F0010C058374D08B0AC7B3550",
```

(continues on next page)

(continued from previous page)

```

"domainURI" : "/cdmi_domains/MyDomain/",
"capabilitiesURI" : "/cdmi_capabilities/domain/",
"metadata" : {
    "cdmi_domain_enabled": "true",
    "cdmi_authentication_methods": "anonymous, basic",
    ...
},
"childrenrange" : "0-1",
"children" : [
    "cdmi_domain_summary/",
    "cdmi_domain_members/"
]
}

```

The meaning, use, and permitted values of each field is described in each operation that creates, modifies or retrieves CDMI domain objects.

### 10.1.1 Domain object addressing

Domain objects are created as children of a special `cdmi_domains` container object, which is present in the root URI for the cloud storage system when domains are supported. The `cdmi_domains` container object is system-generated, read-only, cannot be deleted, and only permits the creation of children domain objects, as indicated by the presence of the `cdmi_create_domain` capability. The ability to create a sub-domain under an existing domain object is indicated by the presence of the `cdmi_create_domain` capability for a given domain object.

Domain objects are addressed in CDMI in two ways:

- by name (e.g., `https://cloud.example.com/cdmi_domains/myDomain/`); and
- by ID (e.g., `https://cloud.example.com/cdmi_objectid/00007ED90010329E642EBFBC8B57E9AD/`).

Every domain object has a single, globally-unique object ID that remains constant for the life of the object. Each domain object shall also have at least one URI address that allows the domain object to be accessed. Following the URI conventions for hierarchical paths, domain URIs shall start with “<root URI>/cdmi\_domains/” and consist of one or more domain names that are separated by forward slashes (“/”) and that end with a forward slash (“/”).

If a request is performed against an existing domain resource and the trailing slash at the end of the URI is omitted, the server shall respond with an HTTP status code of 301 Moved Permanently, and a Location header containing the URI with the trailing slash will be added.

If a CDMI request is performed to create a new domain resource and the trailing slash at the end of the URI is omitted, the server shall respond with an HTTP status code of 400 Bad Request.

Domain objects may also be nested.

EXAMPLE 2: The following URI represents a nested domains:

```
https://cloud.example.com/cdmi_domains/myDomain/subDomain/
```

A sub-domain has a parent domain object, shall be included in the children field of the parent domain object, and shall inherit Domain Membership from its parent domain (if not specified in the sub-domain).

### 10.1.2 Domain object fields

Individual fields within a domain object may be accessed by specifying the field name after a question mark “?” appended to the end of the domain object URI.

EXAMPLE 3: The following URI returns just the children field in the response message body:

```
https://cloud.example.com/cdmi_domains/myDomain/?children
```

EXAMPLE 4: By specifying a range after the children field name, specific ranges of the children field may be accessed.

```
https://cloud.example.com/cdmi_domains/myDomain/?children:0-2
```

Children ranges are specified in a way that is similar to byte ranges as per Section 14.35.1 of RFC 2616 [23]. A client can determine the number of children present by requesting the `childrenrange` field without requesting a range of children.

A list of fields separated by a semicolon “;” may be specified, allowing multiple fields to be accessed in a single request.

EXAMPLE 5: The following URI would return the children and metadata fields in the response body:

```
https://cloud.example.com/cdmi_domains/myDomain/?children;metadata
```

When a client provides fields that are not defined in this International Standard or deserializes an object containing fields that are not defined in this International Standard, these fields shall be persisted, but shall not be interpreted.

### 10.1.3 Domain object metadata

The following domain-specific field shall be present for each domain (see Table 78).

Table 78: Required metadata for a domain object

Metadata name	Type	Description	Requirement
<code>cdmi_domain_enabled</code>	JSON string	Indicates if the domain is enabled and specified at the time of creation. Values shall be “true” or “false”. <ul style="list-style-type: none"> <li>If this metadata item is not present at the time of domain creation, the value is set to “false”.</li> <li>If a domain is disabled, the cloud storage system shall not permit any operations to be performed against any URI managed by that domain.</li> <li>When a domain is disabled, all operations that are performed against URIs that are managed by a disabled domain shall return an HTTP status code of 403 Forbidden.</li> </ul>	Mandatory
<code>cdmi_domain_delete_reassign</code>	JSON string	If the domain is deleted, indicates to which domain the objects that belong to the domain shall be reassigned. <ul style="list-style-type: none"> <li>To delete a domain that contains objects, this metadata item shall be present.</li> <li>If this metadata item is not present or does not contain the URI of a valid domain that is different from the URI of the domain being deleted, an attempt to delete a domain that has objects shall result in an HTTP status code of 400 Bad Request.</li> </ul>	Conditional
<code>cdmi_authentication_methods</code>	JSON array of JSON strings	Indicates a list of which authentication methods are enabled for the domain.  Supported authentication method values are indicated by the <code>cdmi_authentication_methods</code> capability.	Optional

Domains may also contain domain-specific data system metadata items as defined in 16.3 and 16.4. Domain data system metadata shall be inherited to child domain objects.

### 10.1.4 Domain object access control

If read access to any of the requested fields is not permitted by the object ACL, only the permitted fields shall be returned. If no requested fields are permitted to be read, an HTTP status code of 403 Forbidden shall be returned to the client.

If write access to any of the requested fields is not permitted by the object ACL, no updates shall be performed, and an HTTP status code of 403 Forbidden shall be returned to the client.

### 10.1.5 Domain usage in access control

When a transaction is performed against a CDMI object, the associated domain object (i.e., the domain object indicated by the `domainURI`) specifies the authentication context. The user identity and credentials presented as part of the transaction are compared to the domain membership list to determine if the user is authorized within the domain and to resolve the user's principal. If resolved, the user's principal is evaluated against the object's ACL to determine if the transaction is permitted.

When evaluating members within a domain, delegations are evaluated first, in any order, followed by user records, in any order. If there is at least one matching record and none of the matching records indicate that the user is disabled, the user is considered to be a member of the domain.

When a sub-domain is initially created, the membership container contains one member record that is a delegation in which the delegation URI is set to the URI of the parent domain.

### 10.1.6 Domain object representations

The representations in this clause are shown using JSON notation. Both clients and servers shall support UTF-8 JSON representation. The request and response body JSON fields may be specified or returned in any order, with the exception that, if present, for domain objects, the `childrenrange` and `children` fields shall appear last and in that order.

## 10.2 Domain object summaries

Domain object summaries provide summary measurement information about domain usage and billing. If supported, a domain summary container named “cdmi\_domain\_summary” shall be present under each domain container. Like any container, the domain summary subcontainer may have an Access Control List (ACL) (see 17.1) that restricts access to this information.

Within each domain summary container are a series of domain summary data objects that are generated by the cloud storage system. The “yearly”, “monthly”, and “daily” containers of these data objects contain domain summary data objects corresponding to each year, month, and day, respectively. These containers are organized into the following structures:

```
https://example.com/cdmi_domains/domain/
https://example.com/cdmi_domains/domain/cdmi_domain_summary/
https://example.com/cdmi_domains/domain/cdmi_domain_summary/cumulative
https://example.com/cdmi_domains/domain/cdmi_domain_summary/daily/
https://example.com/cdmi_domains/domain/cdmi_domain_summary/daily/2009-07-01
https://example.com/cdmi_domains/domain/cdmi_domain_summary/daily/2009-07-02
https://example.com/cdmi_domains/domain/cdmi_domain_summary/daily/2009-07-03
https://example.com/cdmi_domains/domain/cdmi_domain_summary/monthly/
https://example.com/cdmi_domains/domain/cdmi_domain_summary/monthly/2009-07
https://example.com/cdmi_domains/domain/cdmi_domain_summary/monthly/2009-08
https://example.com/cdmi_domains/domain/cdmi_domain_summary/monthly/2009-10
https://example.com/cdmi_domains/domain/cdmi_domain_summary/yearly/
https://example.com/cdmi_domains/domain/cdmi_domain_summary/yearly/2009
https://example.com/cdmi_domains/domain/cdmi_domain_summary/yearly/2010
```

The “cumulative” summary data object covers the entire time period, from the time the domain is created to the time it is accessed. Each data object at the daily, monthly, and yearly level contains domain summary information for the time period specified, bounded by domain creation time and access time.

If a time period extends earlier than the domain creation time, the summary information includes the time from when the domain was created until the end of the time period.

EXAMPLE 1: If a domain were created on July 4, 2009, at noon, the daily summary “2009-07-04” would contain information from noon until midnight, the monthly summary “2009-07” would contain information from noon on July 4 until midnight on July 31, and the yearly summary “2009” would contain information from noon on July 4 until midnight on December 31.

If a time period starts after the time when the domain was created and ends earlier than the time of access, the summary data object contains complete information for that time period.

EXAMPLE 2: If a domain were created on July 4, 2009, and on July 10, the “2009-07-06” daily summary data object was accessed, it would contain information for the complete day.

If a time period ends after the current access time, the domain summary data object contains partial information from the start of the time period (or the time the domain was created) until the time of access.

EXAMPLE 3: If a domain were created on July 4, 2009, and at noon on July 10, the “2009-07-10” daily summary data object was accessed, it would contain information from the beginning of the day until noon.

The information in Table 79 shall be present within the contents of each domain summary object, which are in JSON representation.

Table 79: Contents of domain summary objects

Metadata name	Type	Description	Requirement
cdmi_domainURI	JSON string	Domain name corresponding to the domain that is summarized	Mandatory

Continued on next page



Table 79 – continued from previous page

Metadata name	Type	Description	Requirement
cdmi_summary_start	JSON string	An ISO-8601 time indicating the start of the time range that the summary information is presenting	Mandatory
cdmi_summary_end	JSON string	An ISO-8601 time indicating the end of the time range that the summary information is presenting	Mandatory
cdmi_summary_objecthours	JSON string	The sum of the time each object belonging to the domain existed during the summary time period	Optional
cdmi_summary_objectsmin	JSON string	The minimum number of objects belonging to the domain during the summary time period	Optional
cdmi_summary_objectsmax	JSON string	The maximum number of objects belonging to the domain during the summary time period	Optional
cdmi_summary_objectsaverage	JSON string	The average number of objects belonging to the domain during the summary time period	Optional
cdmi_summary_puts	JSON string	The number of objects written to the domain	Optional
cdmi_summary_gets	JSON string	The number of objects read from the domain	Optional
cdmi_summary_bytehours	JSON string	The sum of the time each byte belonging to the domain existed during the summary time period	Optional
cdmi_summary_bytesmin	JSON string	The minimum number of bytes belonging to the domain during the summary time period	Optional
cdmi_summary_bytesmax	JSON string	The maximum number of bytes belonging to the domain during the summary time period	Optional
cdmi_summary_bytesaverage	JSON string	The average number of bytes belonging to the domain during the summary time period	Optional
cdmi_summary_writes	JSON string	The number of bytes written to the domain	Optional
cdmi_summary_reads	JSON string	The number of bytes read from the domain	Optional
cdmi_summary_charge	JSON string	An ISO 4217 currency code (see [39]) that is followed or preceded by a numeric value and separated by a space, where the numeric value represents the closing charge in the indicated currency for the use of the service associated with the domain over the summary time period	Optional
cdmi_summary_kwhours	JSON string	The sum of energy consumed (in kilowatt hours) by the domain during the summary time period	Optional
cdmi_summary_kwmin	JSON string	The minimum rate at which energy is consumed (in kilowatt hours per hour) by the domain during the summary time period	Optional
cdmi_summary_kwmax	JSON string	The maximum rate at which energy is consumed (in kilowatt hours per hour) by the domain during the summary time period	Optional

Continued on next page

Table 79 – continued from previous page

Metadata name	Type	Description	Requirement
cdmi_summary_kwaverage	JSON string	The average rate at which energy is consumed (in kilowatt hours per hour) by the domain during the summary time period	Optional

2074 **EXAMPLE 4:** An example of a daily domain summary object is as follows:

```

{
  "cdmi_domainURI" : "/cdmi_domains/MyDomain/",
  "cdmi_summary_start" : "2009-12-10T00:00:00",
  "cdmi_summary_end" : "2009-12-10T23:59:59",
  "cdmi_summary_objecthours" : "382239734",
  "cdmi_summary_puts" : "234234",
  "cdmi_summary_gets" : "489432",
  "cdmi_summary_bytehours" : "334895798347",
  "cdmi_summary_writes" : "7218368343",
  "cdmi_summary_reads" : "11283974933",
  "cdmi_summary_charge" : "4289.23 USD"
}

```

2075 If the charge value is provided, the value is for the operational cost (excluding fixed fees) of service already performed  
 2076 and storage and bandwidth already consumed. Pricing of services is handled separately.

2077 Domain summary information may be extended by vendors to include additional metadata or domain reports beyond  
 2078 the metadata items specified by this International Standard, as long as the field names for those metadata items do not  
 2079 begin with "cdmi\_".

## 10.3 Domain object membership

In cloud storage environments, in the same way that domains are often created programmatically, domain user membership and credential mapping also shall be populated using such interfaces. By providing access to user membership, this capability enables self-enrollment, automatic provisioning, and other advanced self-service capabilities, either directly using CDMI or through software systems that interface with CDMI.

The domain membership capability provides information about, and allows the specification of, end users and groups of users that are allowed to access the domain via CDMI and other access protocols. The concept of domain membership is not intended to replace or supplant ACLs (see 17.1), but rather to provide a single, unified place to map identities and credentials to principals used by ACLs within the context of a domain (see model described in 10.1.5). It also provides a place for authentication mappings to external authentication providers, such as LDAP and Active Directory, to be specified.

If supported, a domain membership container named `cdmi_domain_members` shall be present under each domain. Like any container, the domain membership container has an Access Control List (see 17.1) that restricts access to this information.

Within each domain membership container are a series of user objects that are specified through CDMI to define each user known to the domain. These objects are formatted into the following structure:

```
https://example.com/cdmi_domains/domain/
https://example.com/cdmi_domains/domain/cdmi_domain_members/
https://example.com/cdmi_domains/domain/cdmi_domain_members/john_doe
https://example.com/cdmi_domains/domain/cdmi_domain_members/john_smith
```

The domain membership container may also contain subcontainers with data objects. Data objects in these subcontainers are treated the same as data objects in the domain membership container, and no meaning is inferred from the subcontainer name. This organization is used to create different access security relationships for groups of user objects and to allow delegation to a common set of members.

Table 80 lists the domain settings that shall be present within each domain member user object.

Table 80: Required settings for domain member user objects

Metadata name	Type	Description	Requirement
<code>cdmi_member_enabled</code>	JSON string	If true, this field indicates that requests associated with this domain member are allowed. If false, all requests performed by this domain member shall result in an HTTP status code of 403 <code>Forbidden</code> .	Mandatory
<code>cdmi_member_type</code>	JSON string	This field indicates the type of member record. Values include “user”, “group”, and “delegation”.	Mandatory
<code>cdmi_member_name</code>	JSON string	This field contains the user or group name as presented by the client. This will normally be the standard full name of the principal.	Mandatory
<code>cdmi_member_credentials</code>	JSON string	This field contains credentials to be matched against the credentials as presented by the client. If this field is not present, one or more delegations shall be present and shall be used to resolve user credentials. As one cannot log in as a group but only as a member of a group, the “group” type member records shall not have credentials.	Optional
<code>cdmi_member_principal</code>	JSON string	This field indicates to which principal name (used in ACLs) the user or group is mapped. If this field is not present, one or more delegations shall be present and shall be used to resolve the principal.	Optional

Continued on next page

Table 80 – continued from previous page

Metadata name	Type	Description	Requirement
cdmi_member_privileges	JSON array of JSON strings	This field explicitly confers zero or more special privileges to a user or group. When delegated, privileges are conferred based on the information returned from the external system to which the delegation points. The following privileges are defined: <ul style="list-style-type: none"> <li>• “administrator”. Allows the principal to take ownership of any object/container.</li> <li>• “backup_operator”. Bypass regular ACL checks to allow backup and restore of objects and containers, including all associated attributes, metadata, ACLs and ownership.</li> <li>• “cross_domain”. Operations specifying a domain other than the domain of the parent object are permitted. Unless this privilege is conferred by the user record or a group (possibly nested) to which the user or group belongs, all attempts to change the domain of objects to a domain other than the parent domain shall fail.</li> </ul>	Mandatory
cdmi_member_groups	JSON array of JSON strings	This field contains a JSON array of group names to which the user or group belongs.	Optional

2105 Table 81 lists the domain settings that shall be present within each domain member delegation object.

Table 81: Required settings for domain member delegation objects

Metadata name	Type	Description	Requirement
cdmi_member_enabled	JSON string	If true, this field indicates that requests associated with this domain member are allowed. If false, all requests performed by this domain member shall result in an HTTP status code of 403 Forbidden.	Mandatory
cdmi_member_type	JSON string	This field indicates the type of member record. Values include “user” and “delegation”.	Mandatory
cdmi_delegation_URI	JSON string	This field contains the URI of an external identity resolution provider (such as LDAP or Active Directory) or the URI of a domain membership container object.  External delegations are expressed in the form of <code>ldap://&lt;uri&gt;</code> or <code>ad://&lt;uri&gt;</code> .	Mandatory

2106 EXAMPLE 1: An example of a domain membership object for a user is as follows:

```
{
  "cdmi_member_enabled" : "true",
  "cdmi_member_type" : "user",
  "cdmi_member_name" : "John Doe",
  "cdmi_member_credentials" : "p+5/oX1cmExfOIrUxhX1lw==",
```

(continues on next page)

(continued from previous page)

```
"cdmi_member_groups" : [
  "users"
],
"cdmi_member_principal" : "jdoe",
"cdmi_privileges" : [
  "administrator",
  "cross_domain"
]
}
```

2107 **EXAMPLE 2:** An example of a domain membership object for a delegation is as follows:

```
{
  "cdmi_member_enabled" : "true",
  "cdmi_member_type" : "delegation",
  "cdmi_delegation_URI" : "/cdmi_domains/MyDomain/"
}
```

## 10.4 Create a domain object using CDMI

### 10.4.1 Synopsis

To create a new domain object, the following request shall be performed:

- PUT <root URI>/cdmi\_domains/<DomainName>/<NewDomainName>/

Where:

- <root URI> is the path to the CDMI cloud.
- <DomainName> is zero or more intermediate domains that already exist, with one slash (i.e., “/”) between each pair of domain names.
- <NewDomainName> is the name specified for the domain to be created.

After it is created, the domain shall also be accessible at <root URI>/cdmi\_objectid/<objectID>/.

### 10.4.2 Delayed completion of create

Delayed completion shall not be supported for creating domain objects.

### 10.4.3 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 82](#).

Table 82: Capabilities - Create a CDMI domain object using CDMI

Capability	Location	Description
cdmi_create_domain	Parent container	Ability to create a new domain object
cdmi_copy_domain	Parent container	Ability to create a domain object that is a copy of another domain object
cdmi_deserialize_domain	Parent container	Ability to create a domain object that is deserialized from the contents of the PUT or the contents of another data object

### 10.4.4 Request headers

The HTTP request headers for creating a CDMI domain object using CDMI are shown in [Table 83](#)

Table 83: Request headers - Create a domain object using CDMI

Header	Type	Description	Requirement
Accept	Header string	“application/cdmi-domain” or a consistent value as described in <a href="#">5.5.2</a>	Optional
Content-Type	Header string	“application/cdmi-domain”	Mandatory

### 10.4.5 Request message body

The request message body fields for creating a domain object using CDMI are shown in [Table 84](#).

Table 84: Request message body - Create a domain object using CDMI

Field Name	Type	Description	Requirement
metadata	JSON object	Metadata for the domain object <ul style="list-style-type: none"> <li>If this field is included, the contents of the JSON object provided in this field shall be used as domain object metadata.</li> <li>If this field is included when deserializing, serializing, copying, or moving a domain object, the contents of the JSON object provided in this field shall be used as object metadata instead of the metadata from the source URI.</li> <li>If this field is not included, no user-specified metadata shall be added to the object.</li> <li>If this field is not included when deserializing, serializing, copying, or moving a domain object, metadata from the source URI shall be used.</li> </ul>	Optional
copy	JSON string	URI of a CDMI domain that shall be copied into the new domain, including all child domains and membership from the source domain	Optional <sup>1</sup>
move	JSON string	URI of an existing local CDMI domain object (source URI) that shall be relocated, along with all child domains, to the URI specified in the PUT. The contents of the domain and all sub-domains, including the object ID, shall be preserved by a move, and the domain and sub-domains of the source URI shall be removed after the objects at the destination have been successfully created.  If there are insufficient permissions to read the objects at the source URI, write the objects at the destination URI, or delete the objects at the source URI, or if any of these operations fail, the move shall return an HTTP status code of 400 <i>Bad Request</i> , and the source and destination are left unchanged.	Optional <sup>1</sup>
deserialize	JSON string	URI of a CDMI data object with a value that contains a domain object serialized as specified in <a href="#">clause 15</a> . The serialized domain object shall be deserialized to create the new domain object, including all child objects.	Optional <sup>1</sup>
deserializevalue	JSON string	A domain object serialized as specified in <a href="#">clause 15</a> and encoded using base 64 encoding rules described in RFC 4648 [19], that shall be deserialized to create the new domain object, including all child objects.	Optional <sup>1</sup>

## 10.4.6 Response headers

The HTTP response headers for creating a domain object using CDMI are shown in [Table 85](#)

Table 85: Response headers - Create a domain object using CDMI

Header	Type	Description	Requirement
Content-Type	Header string	"application/cdm-domain"	Mandatory

<sup>1</sup> Only one of these fields shall be specified in any given operation. Except for value, these fields shall not be stored. If more than one of these fields is supplied, the server shall respond with an HTTP status code of 400 *Bad Request*.

## 10.4.7 Response message body

The response message body fields for creating a domain object using CDMI are shown in [Table 86](#)

Table 86: Response message body - Create a domain object using CDMI

Field Name	Type	Description	Requirement
objectType	JSON string	"application/cdm-domain"	Mandatory
objectID	JSON string	Object ID of the domain	Mandatory
objectName	JSON string	Name of the object	Mandatory
parentURI	JSON string	URI for the parent object Appending the objectName to the parentURI shall always produce a valid URI for the object.	Mandatory
parentID	JSON string	Object ID of the parent container object	Mandatory
domainURI	JSON string	URI of the owning domain. A domain object is always owned by itself.	Mandatory
capabilitiesURI	JSON string	URI to the capabilities for the object	Mandatory
metadata	JSON object	Metadata for the domain object. This field includes any user and data system metadata specified in the request body metadata field, along with storage system metadata generated by the cloud storage system. See <a href="#">clause 16</a> for a further description of metadata.	Mandatory
childrenrange	JSON string	The sub-domains of the domain expressed as a range. If a range of sub-domains is requested, this field indicates the children returned as a range.	Mandatory
children	JSON array of JSON strings	Names of the children domains in the domain. Child containers end with "/".	Mandatory

## 10.4.8 Response status

[Table 87](#) describes the HTTP status codes that occur when creating a domain object using CDMI.

Table 87: HTTP status codes - Create a domain object using CDMI

HTTP Status	Description
201 Created	The new domain object was created.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server.

## 10.4.9 Example

EXAMPLE 1: PUT to the domain URI the domain name and metadata:



```
PUT /cdmi_domains/MyDomain/ HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-domain
Content-Type: application/cdmi-domain

"metadata":
{
  "cdmi_domain_enabled": "true"
}
```

2134

The following shows the response.

```
HTTP/1.1 201 Created
Content-Type: application/cdmi-domain

{
  "objectType" : "application/cdmi-domain",
  "objectID" : "00007E7F00104BE66AB53A9572F9F51E",
  "objectName" : "MyDomain/",
  "parentURI" : "/cdmi_domains/",
  "parentID" : "00007E7F0010C058374D08B0AC7B3550",
  "domainURI" : "/cdmi_domains/MyDomain/",
  "capabilitiesURI" : "/cdmi_capabilities/domain/",
  "metadata" : {
    "cdmi_domain_enabled": "true",
    "cdmi_authentication_methods": "anonymous, basic",
    ...
  },
  "childrenrange" : "0-1",
  "children" : [
    "cdmi_domain_summary/",
    "cdmi_domain_members/"
  ]
}
```

## 10.5 Read a domain object using CDMI

### 10.5.1 Synopsis

To read an existing domain object, the following requests shall be performed:

- GET <root URI>/cdmi\_domains/<DomainName>/<TheDomainName>/
- GET <root URI>/cdmi\_domains/<DomainName>/<TheDomainName>/?<fieldname>;<fieldname>;...
- GET <root URI>/cdmi\_domains/<DomainName>/<TheDomainName>/?children:<range>;...
- GET <root URI>/cdmi\_domains/<DomainName>/<TheDomainName>/?metadata:<prefix>;...
- GET <root URI>/cdmi\_objectid/<DomainObjectID>/
- GET <root URI>/cdmi\_objectid/<DomainObjectID>/?<fieldname>;<fieldname>;...
- GET <root URI>/cdmi\_objectid/<DomainObjectID>/?children:<range>;...
- GET <root URI>/cdmi\_objectid/<DomainObjectID>/?metadata:<prefix>;...

Where:

- <root URI> is the path to the CDMI cloud.
- <DomainName> is zero or more parent domains.
- <TheDomainName> is the name specified for the domain to be read from.
- <fieldname> is the name of a field.
- <range> is a numeric range within the list of children.
- <prefix> is a matching prefix that returns all metadata items that start with the prefix value.
- <DomainObjectID> is the ID of the domain object to be read from.

### 10.5.2 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 88](#).

Table 88: Capabilities - Read a CDMI domain object using CDMI

Capability	Location	Description
cdmi_read_metadata	Domain object	Ability to read the metadata of an existing domain object
cdmi_list_children	Domain object	Ability to list the children of an existing domain object
cdmi_object_access_by_ID	System wide capability	Ability to access the object by ID

### 10.5.3 Request headers

The HTTP request headers for reading a CDMI domain object using CDMI are shown in [Table 89](#).

Table 89: Request headers - Read a domain object using CDMI

Header	Type	Description	Requirement
Accept	Header string	“application/cdmi-domain” or a consistent value as described in <a href="#">5.5.2</a>	Optional

### 10.5.4 Request message body

A request body shall not be provided.

## 10.5.5 Response headers

The HTTP response headers for reading a CDMI domain object using CDMI are shown in [Table 90](#).

Table 90: Response headers - Read a domain object using CDMI

Header	Type	Description	Requirement
Content-Type	Header string	"application/cdm-domain"	Mandatory
Location	Header string	The server shall respond with an absolute URI to which the reference redirects if the object is a reference.	Conditional

## 10.5.6 Response message body

The response message body fields for reading a CDMI domain object using CDMI are shown in [Table 91](#)

Table 91: Response message body - Read a domain object using CDMI

Field Name	Type	Description	Requirement
objectType	JSON string	"application/cdm-domain"	Mandatory
objectID	JSON string	Object ID of the domain	Mandatory
objectName	JSON string	Name of the object	Mandatory
parentURI	JSON string	URI for the parent object Appending the "objectName" to the "parentURI" shall always produce a valid URI for the object.	Mandatory
parentID	JSON string	Object ID of the parent domain object <ul style="list-style-type: none"> <li>For domain objects directly under "cdmi_domains", the object ID of "cdmi_domains" container shall be returned.</li> <li>For domain objects under another domain, the object ID of the parent domain shall be returned.</li> </ul>	Mandatory
domainURI	JSON string	URI of the owning domain. A domain object shall always be owned by itself.	Mandatory
capabilitiesURI	JSON string	URI to the capabilities for the object	Mandatory
metadata	JSON object	Metadata for the domain object. This field includes any user and data system metadata specified in the request body metadata field, along with storage system metadata generated by the cloud storage system. See <a href="#">clause 16</a> for a further description of metadata.	Mandatory
childrenrange	JSON string	The sub-domains of the domain expressed as a range. If a range of sub-domains is requested, this field indicates the children returned as a range.	Mandatory
children	JSON array of JSON strings	The children of the domain. Sub-domains end with "/".	Mandatory

If individual fields are specified in the GET request, only these fields are returned in the result body. Optional fields that are requested but do not exist are omitted from the result body.

## 10.5.7 Response status

[Table 92](#) describes the HTTP status codes that occur when reading a domain object using CDMI.

Table 92: HTTP status codes - Read a domain object using CDMI

HTTP Status	Description
200 OK	The domain object content was returned in the response.
302 Found	The resource is a reference to another resource.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
406 Not Acceptable	The server is unable to provide the object in the content type specified in the Accept header.

## 10.5.8 Examples

EXAMPLE 1: GET to the domain URI to read all the fields of the domain:

```
GET /cdmi_domains/MyDomain/ HTTP/1.1
Host: cloud.example.com
Accept: application/cdm-domain
```

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: application/cdm-domain

{
  "objectType": "application/cdm-domain",
  "objectID": "00007E7F00104BE66AB53A9572F9F51E",
  "objectName": "MyDomain/",
  "parentURI": "/cdmi_domains/",
  "parentID": "00007E7F0010C058374D08B0AC7B3550",
  "domainURI": "/cdmi_domains/MyDomain/",
  "capabilitiesURI": "/cdmi_capabilities/domain/",
  "metadata": {
    "cdmi_domain_enabled": "true",
    "cdmi_authentication_methods": "anonymous, basic",
    ...
  },
  "childrenrange": "0-1",
  "children": [
    "cdmi_domain_summary/",
    "cdmi_domain_members/"
  ]
}
```

EXAMPLE 2: GET to the domain URI to read the parentURI and children of the domain:

```
GET /MyDomain/?parentURI;children HTTP/1.1
Host: cloud.example.com
Accept: application/cdm-domain
```

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: application/cdm-domain

{
  "parentURI" : "/cdmi_domains/",
  "children" : [
    "cdmi_domain_summary/",
    "cdmi_domain_members/"
  ]
}
```

EXAMPLE 3: GET to the domain URI to read the first two children of the domain:

```
GET /MyDomain/?childrenrange;children:0-1 HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-domain
```

2177

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: application/cdmi-domain

{
  "childrenrange" : "0-1",
  "children" : [
    "cdmi_domain_summary/",
    "cdmi_domain_members/"
  ]
}
```

## 10.6 Update a domain object using CDMI

### 10.6.1 Synopsis

To update an existing domain object, the following requests shall be performed:

- PUT <root URI>/cdmi\_domains/<DomainName>/<TheDomainName>/
- PUT <root URI>/cdmi\_domains/<DomainName>/<TheDomainName>/?metadata:<metadataname>;
- ...
- PUT <root URI>/cdmi\_objectid/<DomainObjectID>
- PUT <root URI>/cdmi\_objectid/<DomainObjectID>?metadata:<metadataname>;....

Where:

- <root URI> is the path to the CDMI cloud.
- <DomainName> is zero or more parent domains.
- <TheDomainName> is the name specified for the domain to be read from.
- <DomainObjectID> is the ID of the data object to be read from.

### 10.6.2 Delayed completion of update

Delayed completion shall not be supported for creating domain objects.

### 10.6.3 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 93](#).

Table 93: Capabilities - Update a CDMI domain object using CDMI

Capability	Location	Description
cdmi_modify_metadata	Domain object	Ability to modify the metadata of an existing domain object
cdmi_object_access_by_ID	System wide capability	Ability to access the object by ID

### 10.6.4 Request headers

The HTTP request headers for updating a CDMI domain object using CDMI are shown in [Table 94](#).

Table 94: Request headers - Update a domain object using CDMI

Header	Type	Description	Requirement
Content-Type	Header string	"application/cdmi-domain"	Mandatory

### 10.6.5 Request message body

The request message body fields for updating a domain object using CDMI are shown in [Table 95](#).

Table 95: Request message body - Update a domain object using CDMI

Field Name	Type	Description	Requirement
metadata	JSON object	Metadata for the domain object. If present, the new metadata specified replaces the existing object metadata. If individual metadata items are specified in the URI, only those items are replaced; other items are preserved. See <a href="#">clause 16</a> for a further description of metadata.	Optional
copy	JSON string	URI of a CDMI domain object that shall be copied into the existing domain object. Only the metadata and membership of the domain object itself shall be copied, not any sub-domains of the domain object. <ul style="list-style-type: none"> <li>If the destination domain object URI and the copy source domain object URI both do not specify individual fields, the destination domain object metadata and membership shall be replaced with the source domain object metadata and membership.</li> <li>If the destination domain object URI or the copy source domain object URI specifies individual fields, only the fields specified shall be used to update the destination domain object. If specified fields are not present in the source, these fields shall be ignored.</li> <li>If the destination domain object URI and the copy source domain object URI both specify fields, an HTTP status code of 400 <i>Bad Request</i> shall be returned to the client.</li> </ul> <p>If there are insufficient permissions to read the domain object at the source URI or create the domain object at the destination URI, or if the read operation fails, the copy shall return an HTTP status code of 400 <i>Bad Request</i>, and the destination domain object shall not be updated.</p>	Optional <sup>2</sup>
deserialize	JSON string	URI of a CDMI data object with a value that contains a domain object serialized as specified in <a href="#">clause 15</a> . The serialized domain object shall be deserialized to update the existing domain object.  The object ID of the serialized domain object shall match the object ID of the destination domain object. Otherwise, the server shall return an HTTP status code of 400 <i>Bad Request</i> . <ul style="list-style-type: none"> <li>If the serialized domain object does not contain sub-domains, the update is applied only to the domain object, and any existing sub-domains are left as is.</li> <li>If the serialized domain object does contain sub-domains, then creates, updates, and deletes are recursively applied for each sub-domain, depending on the differences between the provided serialized state and the current state of the sub-domains.</li> </ul>	Optional <sup>2</sup>

Continued on next page

Table 95 – continued from previous page

Field Name	Type	Description	Requirement
deserializevalue	JSON string	<p>A domain object serialized as specified in <a href="#">clause 15</a> and encoded using base 64 encoding rules described in RFC 4648 [19], that shall be deserialized to update the existing domain object.</p> <p>The object ID of the serialized domain object shall match the object ID of the destination domain object. Otherwise, the server shall return an HTTP status code of 400 Bad Request.</p> <ul style="list-style-type: none"> <li>If the serialized domain object does not contain sub-domains, the update is applied only to the domain object, and any existing sub-domains are left as is.</li> <li>If the serialized domain object does contain sub-domains, then creates, updates, and deletes are recursively applied for each sub-domain, depending on the differences between the provided serialized state and the current state of the sub-domains.</li> </ul>	Optional <sup>2</sup>

## 10.6.6 Response header

The HTTP response header for updating a CDMI domain object using CDMI is shown in [Table 96](#)

Table 96: Response header - Update a domain object using CDMI

Header	Type	Description	Requirement
Location	Header string	The server shall respond with an absolute URI to which the reference redirects if the object is a reference.	Conditional

## 10.6.7 Response message body

A response body may be provided as per RFC 2616 [23].

## 10.6.8 Response status

[Table 97](#) describes the HTTP status codes that occur when updating a domain object using CDMI.

Table 97: HTTP status codes - Update a domain object using CDMI

HTTP Status	Description
204 No Content	The data object content was returned in the response.
302 Found	The resource is a reference to another resource.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server.

## 10.6.9 Example

EXAMPLE 1: PUT to the domain URI to set new field values:

<sup>2</sup> Only one of these fields shall be specified in any given operation. Except for value, these fields shall not be stored.



```
PUT /cdmi_domains/MyDomain/ HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-domain

{
  "metadata" : {
    "test" : "value"
  }
}
```

2213

The following shows the response.

```
HTTP/1.1 204 No Content
```

## 10.7 Delete a domain object using CDMI

### 10.7.1 Synopsis

To delete an existing domain object, and transfer all objects associated with that domain to another domain (to preserve access), the following request shall be performed:

- DELETE <root URI>/cdmi\_domains/<DomainName>/<TheDomainName>/
- DELETE <root URI>/cdmi\_objectid/<DomainObjectID>

Where:

- <root URI> is the path to the CDMI cloud.
- <DomainName> is zero or more parent domains.
- <TheDomainName> is the name specified for the domain to be deleted.
- <DomainObjectID> is the ID of the domain object to be deleted.

### 10.7.2 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 98](#).

Table 98: Capabilities - Delete a CDMI domain object using CDMI

Capability	Location	Description
cdmi_delete_domain	Domain object	Ability to delete an existing domain object
cdmi_object_access_by_ID	System wide capability	Ability to access the object by ID

### 10.7.3 Request headers

Request headers may be provided as per RFC 2616 [\[23\]](#).

### 10.7.4 Request message body

A request body may be provided as per RFC 2616 [\[23\]](#).

### 10.7.5 Response headers

Response headers may be provided as per RFC 2616 [\[23\]](#).

### 10.7.6 Response message body

A response body may be provided as per RFC 2616 [\[23\]](#).

### 10.7.7 Response status

[Table 99](#) describes the HTTP status codes that occur when deleting a domain object using CDMI.

2237

Table 99: HTTP status codes - Delete a domain object using CDMI

2238

HTTP Status	Description
204 No Content	The domain object was successfully deleted.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server.

2239

## 10.7.8 Example

2240

EXAMPLE 1: DELETE to the domain object URI:

```
DELETE /cdmi_domains/MyDomain/ HTTP/1.1
Host: cloud.example.com
```

2241

The following shows the response.

```
HTTP/1.1 204 No Content
```

## Clause 11

# Queue object resource operations using CDMI

## 11.1 Overview

Queue objects are similar to data objects, only with first-in, first-out access “queue”-style access semantics when storing and retrieving value data.

If a cloud storage system supports queues, the `cdmi_queues` system-wide capability shall be present. The ability to create a queue object is determined by the presence or absence of the `cdmi_create_queue` and `cdmi_post_queue` capabilities in the parent container, and by the `cdmi_post_queue_by_ID` system-wide capability for creation by ID.

A queue object writer POSTs data into a queue object, and a queue object reader GETs value(s) from the queue object and subsequently deletes the value(s) to acknowledge receipt of the value(s) that it received. Queues provides a simple mechanism for one or more writers to send data to a single reader in a reliable way. If supported by the cloud storage system, cloud clients create the queue objects by using the mechanism described in 9.7 and this clause.

FIXME - Add diagram (Issue #120)

Each CDMI queue object is represented as a JSON object, containing one or more “fields”. For example, the “metadata” field contains metadata items.

EXAMPLE 1: CDMI queue object

```
{
  "objectType": "application/cdmi-queue",
  "objectID": "00007E7F00104BE66AB53A9572F9F51E",
  "objectName": "MyQueue",
  "parentURI": "/MyContainer/",
  "parentID": "00007ED900104F67307652BAC9A37C93",
  "domainURI": "/cdmi_domains/MyDomain/",
  "capabilitiesURI": "/cdmi_capabilities/queue/",
  "completionStatus": "Complete",
  "metadata": {},
  "queueValues": "1-1",
  "mimetype": [
    "text/plain"
  ],
  "valueRange": [
    "0-19"
  ],
  "valueTransferEncoding": [
    "utf-8"
  ],
  "value": [
    "First Enqueued Value"
  ]
}
```

The meaning, use, and permitted values of each field is described in each operation that creates, modifies or retrieves CDMI queue objects.

### 11.1.1 Queue object addressing

Queue objects are addressed in CDMI in two ways:

- by name (e.g., `https://cloud.example.com/queueobject`); and
- by ID (e.g., `https://cloud.example.com/cdmi_objectid/00007ED900104F67307652BAC9A37C93/`).

Every queue object has a single, globally-unique object ID that remains constant for the life of the object. Each queue object may also have one or more URI addresses that allow the queue object to be accessed.

### 11.1.2 Queue object fields

Individual fields within a queue object can be accessed by specifying the field name after a question mark “?” appended to the end of the queue object URI.

EXAMPLE 2: The following URI returns just the number of values stored in the queue object in the response body:

```
https://cloud.example.com/queueobject?queueValues
```

A list of unique fields, separated by a semicolon “;” can be specified, allowing multiple fields to be accessed in a single request.

EXAMPLE 3: The following URI returns the number of values stored and metadata fields in the response body:

```
https://cloud.example.com/queueobject?queueValues;metadata
```

When a client provides fields that are not defined in this International Standard or deserializes an object containing fields that are not defined in this International Standard, these fields shall be persisted, but shall not be interpreted.

### 11.1.3 Queue object value

The encoding of the data stored in the queue object value field depends on the queue object value transfer encoding field:

- If the value transfer encoding of the object is set to “utf-8”, the data stored in the value of the queue object shall be a valid UTF-8 string, and it shall be transported as a UTF-8 string in the value field.
- If the value transfer encoding of the object is set to “base64”, the data stored in the value of the queue object can contain arbitrary binary sequences, and it shall be transported as a base 64-encoded string in the value field.
- If the value transfer encoding of the object is set to “json”, the data stored in the value of the queue object shall be a valid JSON object, and the value field shall contain a valid JSON object.

Specific ranges of the value of a queue object can be accessed by specifying a byte range after the value field name.

EXAMPLE 4: The following URI returns the first thousand bytes of the oldest value enqueued:

```
https://cloud.example.com/queueobject?value:0-999
```

Because a byte range of a UTF-8 string is often not a valid UTF-8 string, the response to a range request shall always be transported in the value field as a base 64-encoded string.

Byte ranges are specified as single, inclusive byte ranges as per Section 14.35.1 of RFC 2616 [23].

If read access to any of the requested fields is not permitted by the object ACL, only the permitted fields shall be returned. If no requested fields are permitted to be read, an HTTP status code of 403 *Forbidden* shall be returned to the client.

If write access to any of the requested fields is not permitted by the object ACL, no updates shall be performed, and an HTTP status code of 403 *Forbidden* shall be returned to the client.

When a client provides or includes deserialization fields that are not defined in this International Standard, these fields shall be stored as part of the object.

The value of a queue object may also be specified and retrieved using multi-part MIME, where the CDMI JSON is transferred in the first MIME part and the raw queue values are transferred in the subsequent MIME parts. Each MIME part, including any header fields, shall conform to RFC 2045 [9], RFC 2046 [10], and RFC 2616 [23], and the length of each part may optionally be specified by a Content-Length header in addition to the MIME boundary delimiter.

Multiple non-overlapping ranges of the value of a queue object may also be accessed or updated in a multi-part MIME operation by transferring one MIME part for each range of the value. The byte ranges for these operations shall be specified as per Section 14.35.1 of RFC 2616 [23].

Multi-part MIME enables the efficient transfer of binary data alongside CDMI object metadata without incurring the overhead of the UTF-8 or Base64 encoding and validation required to represent binary data in JSON.

### 11.1.4 Queue object metadata

Queue object metadata may also include arbitrary user-supplied metadata, storage system metadata, and data system metadata, as specified in [clause 16](#). Metadata shall be stored as a valid UTF-8 string. Binary data stored in user metadata shall be first encoded such that it can be contained in a UTF-8 string, with the use of base 64 encoding recommended.

Every queue object has a parent object from which the queue object inherits data system metadata that is not explicitly specified in the data object itself.

EXAMPLE 5: The “pages” queue object stored at the following URI would inherit data system metadata from its parent container, “OCR”:

`https://cloud.example.com/OCR/pages`

### 11.1.5 Queue object access control

If read access to any of the requested fields is not permitted by the object ACL, only the permitted fields shall be returned. If no requested fields are permitted to be read, an HTTP status code of 403 `Forbidden` shall be returned to the client.

If write access to any of the requested fields is not permitted by the object ACL, no updates shall be performed, and an HTTP status code of 403 `Forbidden` shall be returned to the client.

### 11.1.6 Queue object consistency

Writing to a queue object is an atomic operation.

For non-value-related fields:

- If a client reads a queue object simultaneously with a write to that same queue object, the reading client shall get either the old version or the new version, but not a mixture of both.
- If a write is terminated due to errors, the contents of the queue object shall be as if the write never occurred (i.e., writes are atomic in the face of errors).

For value-related fields:

- If a client dequeues or deletes one or more queue values simultaneously with one or more queue values being enqueued to that same queue object, the order of operations shall be as if the dequeue/delete operation happens before the enqueue operation.
- If a dequeue, delete or enqueue is terminated due to errors, the contents of the queue object shall be as if the dequeue/delete/enqueue never occurred (i.e., writes are atomic in the face of errors).

Create and update timestamps that are returned in response to multiple client writes to a given object may indicate that a specific write is the newest (i.e., the write whose data is expected to be returned to subsequent reads until another write is processed). However, there is no guarantee that the write with the latest timestamp is the one whose data is returned on subsequent reads.

Implementations of this International Standard shall provide the atomicity features described in this subclause for queue objects that are accessed via CDMI. The atomicity properties of queue objects that are accessed by protocols other than CDMI are outside the scope of this International Standard.

### 11.1.7 Queue object representations

The representations in this clause are shown using JSON notation. Both clients and servers shall support UTF-8 JSON representation. The request and response body JSON fields can be specified or returned in any order, with the exception that, if present, for queue objects, the “`valuerange`” and “`value`” fields shall appear last and in that order.

## 11.2 Create a queue object using CDMI

### 11.2.1 Synopsis

To create a new queue object, the following request shall be performed:

- PUT <root URI>/<ContainerName>/<QueueName>

To create a new queue object by ID, see 9.7.

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate containers that already exist, with one slash (i.e., "/") between each pair of container names.
- <QueueName> is the name specified for the queue object to be created.

After it is created, the object shall also be accessible at <root URI>/cdmi\_objectid/<objectID>.

The newly created queue shall have no values unless the queue is created as a result of copying or moving a source queue that has values or as a result of deserializing a serialized queue that has values.

### 11.2.2 Delayed completion of create

In response to a create operation for a queue object, the server may return an HTTP status code of 202 *Accepted* to indicate that the object is in the process of being created. This response is useful for long-running operations (e.g., copying a large queue object from a source URI). Such a response has the following implications.

- The server shall return a Location header with an absolute URI to the object to be created along with an HTTP status code of 202 *Accepted*.
- With an HTTP status code of 202 *Accepted*, the server implies that the following checks have passed:
  - user authorization for creating the object;
  - user authorization for read access to any source object for move, copy, serialize, or deserialize; and
  - availability of space to create the object or at least enough space to create a URI to report an error.
- A client might not be able to immediately access the created object, e.g., due to delays resulting from the implementation's use of eventual consistency.

The client performs GET operations to the URI to track the progress of the operation. In response, the server returns two fields in its response body to indicate progress.

- A mandatory `completionStatus` text field contains either "Processing", "Complete", or an error string starting with the value "Error".
- An optional `percentComplete` field contains the percentage of the operation that has completed (0 to 100).

GET shall not return any value for the queue object when `completionStatus` is not "Complete". If the final result of the create operation is an error, the URI is created with the `completionStatus` field set to the error message. It is the client's responsibility to delete the URI after the error has been noted.

### 11.2.3 Capabilities

Capabilities that indicate which operations are supported are shown in Table 100.

Table 100: Capabilities - Create a CDMI queue object using CDMI

Capability	Location	Description
cdmi_create_queue	Parent container	Ability to create a new queue object
cdmi_create_reference	Parent container	Ability to create a new reference
cdmi_copy_queue	Parent container	Ability to create a queue object that is a copy of another queue object

Continued on next page

Table 100 – continued from previous page

Capability	Location	Description
cdmi_move_queue	Parent container	Ability to move a queue object from another queue object
cdmi_deserialize_queue	Parent container	Ability to create a queue object that is deserialized from the contents of the PUT or the contents of another data object

## 11.2.4 Request headers

The HTTP request headers for creating a CDMI queue object using CDMI are shown in [Table 101](#)

Table 101: Request headers - Create a queue object Using CDMI

Header	Type	Description	Requirement
Accept	Header string	"application/cdmi-queue"	Mandatory
Content-Type	Header string	"application/cdmi-queue"	Mandatory

## 11.2.5 Request message body

The request message body fields for creating a queue object using CDMI are shown in [Table 102](#).

Table 102: Request message body - Create a queue object using CDMI

Field Name	Type	Description	Requirement
metadata	JSON object	Metadata for the queue object <ul style="list-style-type: none"> <li>If this field is included, the contents of the JSON object provided in this field shall be used as queue object metadata.</li> <li>If this field is included when deserializing, serializing, copying, or moving a queue object, the contents of the JSON object provided in this field shall be used as object metadata instead of the metadata from the source URI.</li> <li>If this field is not included, no user-specified metadata shall be added to the object.</li> <li>If this field is not included when deserializing, serializing, copying, or moving a queue object, metadata from the source URI shall be used.</li> <li>This field shall not be included when creating a reference to a queue object.</li> </ul>	Optional
domainURI	JSON string	URI of the owning domain <ul style="list-style-type: none"> <li>If different from the parent domain, the user shall have the "cross_domain" privilege (see <a href="#">cdmi_member_privileges</a> in <a href="#">Table 80</a>).</li> <li>If not specified, the domain of the parent container shall be used.</li> </ul>	Optional
deserialize	JSON string	URI of a CDMI data object with a value that contains a queue object serialized as specified in <a href="#">clause 15</a> . The serialized queue object shall be deserialized to create the new queue object.	Optional <sup>1</sup>

Continued on next page



Table 102 – continued from previous page

Field Name	Type	Description	Requirement
copy	JSON string	<p>URI of a source CDMI queue object that shall be copied into the new destination queue object.</p> <ul style="list-style-type: none"> <li>If the destination queue object URI and the copy source queue object URI both do not specify individual fields, the destination queue object shall be a complete copy of the source queue object, including all enqueued values.</li> <li>If the destination queue object URI or the copy source queue object URI specifies individual fields, only the fields specified shall be used to create the destination queue object. If specified fields are not present in the source, default field values shall be used.</li> <li>If the destination queue object URI and the copy source queue object URI both specify fields, an HTTP status code of 400 <i>Bad Request</i> shall be returned to the client.</li> </ul> <p>If there are insufficient permissions to read the queue object at the source URI or create the queue object at the destination URI, or if the read operation fails, the copy shall return an HTTP status code of 400 <i>Bad Request</i>, and the destination queue object shall not be created.</p>	Optional <sup>1</sup>
move	JSON string	<p>URI of an existing local or remote CDMI queue object (source URI) that shall be relocated to the URI specified in the PUT. The contents of the queue object, including the object ID, shall be preserved by a move, and the queue object at the source URI shall be removed after the queue object at the destination has been successfully created.</p> <p>If there are insufficient permissions to read the queue object at the source URI, write the queue object at the destination URI, or delete the queue object at the source URI, or if any of these operations fail, the move shall return an HTTP status code of 400 <i>Bad Request</i>, and the source and destination are left unchanged.</p>	Optional <sup>1</sup>
reference	JSON string	URI of a CDMI queue object that shall be redirected to by a reference. If other fields are supplied when creating a reference, the server shall respond with an HTTP status code of 400 <i>Bad Request</i> .	Optional <sup>1</sup>
deserializevalue	JSON string	A queue object serialized as specified in <a href="#">clause 15</a> and encoded using base 64 encoding rules described in RFC 4648 [19], that shall be deserialized to create the new queue object.	Optional <sup>1</sup>

### 11.2.6 Response status

The HTTP response headers for creating a CDMI queue object using CDMI are shown in [Table 103](#)

<sup>1</sup> Only one of these fields shall be specified in any given operation. Except for value, these fields shall not be stored. If more than one of these fields is supplied, the server shall respond with an HTTP status code of 400 *Bad Request*.

Table 103: Response headers - Create a queue object Using CDMI

Header	Type	Description	Requirement
Content-Type	Header string	"application/cdmf-queue"	Mandatory
Location	Header string	When an HTTP status code of 202 Accepted is returned, the server shall respond with the absolute URL of the object that is in the process of being created.	Conditional

## 11.2.7 Response message body

The response message body fields for creating a CDMI queue object using CDMI are shown in Table 104

Table 104: Response message body - Create a queue object using CDMI

Field Name	Type	Description	Requirement
objectType	JSON string	"application/cdmf-queue"	Mandatory
objectID	JSON string	Object ID of the object	Mandatory
objectName	JSON string	Name of the object	Mandatory
parentURI	JSON string	URI for the parent object Appending the objectName to the parentURI shall always produce a valid URI for the object.	Mandatory
parentID	JSON string	Object ID of the parent container object	Mandatory
domainURI	JSON string	URI of the owning domain.	Mandatory
capabilitiesURI	JSON string	URI to the capabilities for the object	Mandatory
completionStatus	JSON string	A string indicating if the object is still in the process of being created or updated by another operation, and after that operation is complete, indicates if it was successfully created or updated or if an error occurred.  The value shall be the string "Processing", the string "Complete", or an error string starting with the value "Error".	Mandatory
percentComplete	JSON string	A string indicating the percentage of completion if the object is still in the process of being created or updated by another operation. <ul style="list-style-type: none"> <li>When the value of completionStatus is "Processing", this field, if provided, shall indicate the percentage of completion as a numeric integer value from "0" through "100".</li> <li>When the value of completionStatus is "Complete", this field, if provided, shall contain the value "100".</li> <li>When the value of completionStatus is "Error", this field, if provided, may contain any integer value from "0" through "100".</li> </ul>	Optional

Continued on next page

Table 104 – continued from previous page

Field Name	Type	Description	Requirement
metadata	JSON object	Metadata for the queue object. This field includes any user and data system metadata specified in the request body metadata field, along with storage system metadata generated by the cloud storage system. See <a href="#">clause 16</a> for a further description of metadata.	Mandatory
queueValues	JSON string	The range of designators for enqueued values. Every enqueued value shall be assigned a unique, monotonically-incrementing positive integer designator, starting from 0. If no values are enqueued, an empty string shall be returned. If values are enqueued, the lowest designator, followed by a hyphen ("-"), followed by the highest designator shall be returned.	Mandatory

## 11.2.8 Response status

The HTTP status codes that occur when creating a queue object using CDMI are described in [Table 105](#).

Table 105: HTTP status codes - Create a queue object using CDMI

HTTP Status	Description
201 Created	The new queue object was created.
202 Accepted	The queue object is in the process of being created. The CDMI client should monitor the <code>completionStatus</code> and <code>percentComplete</code> fields to determine the current status of the operation.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server.

## 11.2.9 Examples

Example 1: PUT to the queue URI the queue object name and contents:

```
PUT /MyContainer/MyQueue HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-queue
Content-Type: application/cdmi-queue

{
  "metadata" : {
  }
}
```

The following shows the response.

```
HTTP/1.1 201 Created
Content-Type: application/cdmi-queue

{
  "objectType" : "application/cdmi-queue",
  "objectID" : "00007E7F00104BE66AB53A9572F9F51E",
  "objectName" : "MyQueue",
  "parentURI" : "/MyContainer/",
```

(continues on next page)

(continued from previous page)

```

    "parentID" : "00007ED900104F67307652BAC9A37C93",
    "domainURI" : "/cdmi_domains/MyDomain/",
    "capabilitiesURI" : "/cdmi_capabilities/queue/",
    "completionStatus" : "Complete",
    "metadata" : {
        ...
    },
    "queueValues" : ""
}

```

2403 **EXAMPLE 2: PUT to the queue object URI to create a new queue, copying from another queue:**

```

PUT /MyContainer/MyQueue HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-queue

{
    "copy": "/MyContainer/SourceQueue?value:0-9"
}

```

2404 **The following shows the response.**

```

HTTP/1.1 201 Created
Content-Type: application/cdmi-queue

{
    "objectType": "application/cdmi-queue",
    "objectID": "00007E7F00104BE66AB53A9572F9F51E",
    "objectName": "MyQueue",
    "parentURI" : "/MyContainer/",
    "parentID": "00007ED900104F67307652BAC9A37C93",
    "domainURI": "/cdmi_domains/MyDomain/",
    "capabilitiesURI": "/cdmi_capabilities/queue/",
    "completionStatus": "Complete",
    "metadata": {
        ...
    },
    "queueValues": "0-9"
}

```

## 11.3 Read a queue object using CDMI

### 11.3.1 Synopsis

To read all fields from an existing queue object, the following request shall be performed:

- GET <root URI>/<ContainerName>/<QueueName>
- GET <root URI>/<ContainerName>/<QueueName>?<fieldname>;<fieldname>;...
- GET <root URI>/<ContainerName>/<QueueName>?value:<range>;...
- GET <root URI>/<ContainerName>/<QueueName>?metadata:<prefix>;...
- GET <root URI>/<ContainerName>/<QueueName>?values:<count>
- GET <root URI>/cdmi\_objectid/<QueueObjectID>
- GET <root URI>/cdmi\_objectid/<QueueObjectID>?<fieldname>;<fieldname>;...
- GET <root URI>/cdmi\_objectid/<QueueObjectID>?value:<range>;...
- GET <root URI>/cdmi\_objectid/<QueueObjectID>?metadata:<prefix>;...
- GET <root URI>/cdmi\_objectid/<QueueObjectID>?values:<count>

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate containers.
- <QueueName> is the name of the queue object to be read from.
- <fieldname> is the name of a field.
- <range> is a byte range of the queue object value to be returned in the value field. If a byte range is requested, the range returned shall be from the oldest queue object value.
- <prefix> is a matching prefix that returns all metadata items that start with the prefix value.
- <count> is the number of values to be retrieved from the queue object. If more queue object entries are requested to be retrieved than exist in the queue object, the count is processed as if it is equal to the number of entries in the queue object.
- <QueueObjectID> is the ID of the queue object to be read from.

Reading a queue object shall, by default, return the complete value of the oldest item in the queue, unless the queue-Values range is empty.

### 11.3.2 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 106](#).

Table 106: Capabilities - Read a CDMI queue object using CDMI

Capability	Location	Description
cdmi_read_metadata	Queue object	Ability to read the metadata of an existing queue object
cdmi_read_value	Queue object	Ability to read the value of an existing queue object
cdmi_multipart_mime	Queue object	Ability to read a queue object using multi-part MIME
cdmi_object_access_by_ID	System wide capability	Ability to access the object by ID

### 11.3.3 Request headers

The HTTP request headers for reading a CDMI queue object using CDMI are shown in [Table 107](#)

Table 107: Request headers - Read a queue object using CDMI

Header	Type	Description	Requirement
Accept	Header string	<p>“application/cdmqueue”, “multipart/mixed”, or a consistent value as defined in 5.5.2</p> <p>If “multipart/mixed”, the body shall consist of one or more MIME parts, where the first part shall contain a body of content-type “application/cdmqueue”, and the second and subsequent parts shall each contain the corresponding queue value.</p>	Optional

### 11.3.4 Request message body

A request body shall not be provided.

### 11.3.5 Response status

The HTTP response headers for reading a CDMI queue object using CDMI are shown in Table 108.

Table 108: Response headers - Read a queue object using CDMI

Header	Type	Description	Requirement
Content-Type	Header string	“application/cdmqueue” or “multipart/mixed”	Mandatory
Location	Header string	The server shall respond with an absolute URI to which the reference redirects if the object is a reference.	Conditional

### 11.3.6 Response message body

The response message body fields for reading a CDMI queue object using CDMI are shown in Table 109

Table 109: Response message body - Read a queue object using CDMI

Field Name	Type	Description	Requirement
objectType	JSON string	“application/cdmqueue”	Mandatory
objectID	JSON string	Object ID of the object	Mandatory
objectName	JSON string	<p>Name of the object</p> <ul style="list-style-type: none"> <li>For objects in a container, the objectName field shall be returned.</li> <li>For objects not in a container (objects that are only accessible by ID), the “objectName” field does not exist and shall not be returned.</li> </ul>	Conditional
parentURI	JSON string	<p>URI for the parent object</p> <ul style="list-style-type: none"> <li>For objects in a container, the parentURI field shall be returned.</li> <li>For objects not in a container (objects that are only accessible by ID), the “parentURI” field does not exist and shall not be returned.</li> </ul> <p>Appending the “objectName” to the “parentURI” shall always produce a valid URI for the object.</p>	Conditional

Continued on next page

Table 109 – continued from previous page

Field Name	Type	Description	Requirement
parentID	JSON string	Object ID of the parent container object <ul style="list-style-type: none"> <li>For objects in a container, the “parentID” field shall be returned.</li> <li>For objects not in a container (objects that are only accessible by ID), the “parentID” field does not exist and shall not be returned.</li> </ul>	Conditional
domainURI	JSON string	URI of the owning domain	Mandatory
capabilitiesURI	JSON string	URI to the capabilities for the object	Mandatory
completionStatus	JSON string	A string indicating if the object is still in the process of being created or updated by another operation, and after that operation is complete, indicates if it was successfully created or updated or if an error occurred.  The value shall be the string “Processing”, the string “Complete”, or an error string starting with the value “Error”.	Mandatory
percentComplete	JSON string	A string indicating the percentage of completion if the object is still in the process of being created or updated by another operation. <ul style="list-style-type: none"> <li>When the value of completionStatus is “Processing”, this field, if provided, shall indicate the percentage of completion as a numeric integer value from 0 through 100.</li> <li>When the value of completionStatus is “Complete”, this field, if provided, shall contain the value “100”.</li> <li>When the value of completionStatus is “Error”, this field, if provided, may contain any integer value from “0” through “100”.</li> </ul>	Optional
metadata	JSON object	Metadata for the queue object. This field includes any user and data system metadata specified in the request body metadata field, along with storage system metadata generated by the cloud storage system.  See <a href="#">clause 16</a> for a further description of metadata.	Mandatory
queueValues	JSON string	The range of designators for enqueued values. Every enqueued value shall be assigned a unique, monotonically-incrementing positive integer designator, starting from 0. If no values are enqueued, an empty string shall be returned. If values are enqueued, the lowest designator, followed by a hyphen (“-”), followed by the highest designator shall be returned. <ul style="list-style-type: none"> <li>This field shall only be provided when completionStatus is “Complete” and when one or more values are enqueued.</li> </ul>	Mandatory
mimetype	JSON array of JSON strings	MIME types for each queue object value * The MIME types of the values are returned, each corresponding to the value in the same position in the JSON array. * This field shall only be provided when completionStatus is “Complete” and when one or more values are enqueued.	Optional

Continued on next page

Table 109 – continued from previous page

Field Name	Type	Description	Requirement
valuerange	JSON array of JSON strings	<p>The range of bytes of the queue object values to be returned in the value field</p> <ul style="list-style-type: none"> <li>The value ranges of the values are returned, each corresponding to the value in the same position in the JSON array.</li> <li>If a specific value range has been requested, the entry in the valuerange field shall correspond to the bytes requested. If the request extends beyond the end of the value, the valuerange field shall indicate the smaller byte range returned.</li> <li>This field shall only be provided when <code>completionStatus</code> is "Complete" and when one or more values are enqueued.</li> </ul>	Optional
valuetransferencoding	JSON array of JSON strings	<p>The value transfer encoding used for each queue object value. Two value transfer encodings are defined:</p> <ul style="list-style-type: none"> <li>"utf-8" indicates that the queue object value contains a valid UTF-8 string, and it shall be transported as a UTF-8 string in the value field.</li> <li>"base64" indicates that the queue object value may contain arbitrary binary sequences, and it shall be transported as a base 64-encoded string in the value field.</li> <li>"json" indicates that the queue object value contains a valid JSON object, and the value field shall contain a JSON object.</li> </ul> <p>The value transfer encodings are returned, each corresponding to the value in the same position in the JSON array.</p> <ul style="list-style-type: none"> <li>This field shall only be provided when <code>completionStatus</code> is "Complete" and when one or more values are enqueued.</li> </ul>	Optional
value	JSON array of JSON strings	<p>The oldest enqueued queue object values</p> <ul style="list-style-type: none"> <li>The values in the JSON array are returned in order from oldest to newest.</li> <li>If the <code>valuetransferencoding</code> field indicates UTF-8 encoding, the corresponding value field shall contain a UTF-8 string using JSON escaping rules described in RFC 4627 [5].</li> <li>If the <code>valuetransferencoding</code> field indicates base 64 encoding, the corresponding value field shall contain a base 64-encoded string as described in RFC RFC 4648 [19].</li> <li>If the <code>valuetransferencoding</code> field indicates JSON encoding, the corresponding value field shall contain a JSON object.</li> <li>The value field shall not be provided when using multi-part MIME.</li> <li>The value field shall only be provided when the <code>completionStatus</code> field contains "Complete".</li> </ul>	Conditional

2446 If individual fields are specified in the GET request, only these fields are returned in the result body. Optional fields that  
2447 are requested but do not exist are omitted from the result body.



### 11.3.7 Response status

The HTTP status codes that occur when reading a queue object using CDMI are described in [Table 110](#).

Table 110: HTTP status codes - Read a queue object using CDMI

HTTP Status	Description
200 OK	The queue object content was returned in the response.
302 Found	The resource is a reference to another resource.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
406 Not Acceptable	The server is unable to provide the object in the content type specified in the Accept header.

### 11.3.8 Examples

EXAMPLE 1: GET to the queue object URI to read all fields of the queue object:

```
GET /MyContainer/MyQueue HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-queue
```

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: application/cdmi-queue

{
  "objectType": "application/cdmi-queue",
  "objectID": "00007E7F00104BE66AB53A9572F9F51E",
  "objectName": "MyQueue",
  "parentURI": "/MyContainer/",
  "parentID": "00007ED900104F67307652BAC9A37C93",
  "domainURI": "/cdmi_domains/MyDomain/",
  "capabilitiesURI": "/cdmi_capabilities/queue/",
  "completionStatus": "Complete",
  "metadata": {},
  "queueValues": "1-1",
  "mimetype": [
    "text/plain"
  ],
  "valuerange": [
    "0-19"
  ],
  "valuetransferencoding": [
    "utf-8"
  ],
  "value": [
    "First Enqueued Value"
  ]
}
```

EXAMPLE 2: GET to the queue object URI to read the value and queue items of the queue object:

```
GET /MyContainer/MyQueue?value;queueValues HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-queue
```

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: application/cdmi-queue
```

(continues on next page)

(continued from previous page)

```
{
  "queueValues" : "1-1",
  "value" : [
    "First Enqueued Value"
  ]
}
```

2457 **EXAMPLE 3:** GET to the queue object URI to read the first five bytes of the value of the queue object:

```
GET /MyContainer/MyQueue?value:0-4 HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-queue
```

2458 The following shows the response:

```
HTTP/1.1 200 OK
Content-Type: application/cdmi-queue

{
  "value" : [
    "First"
  ]
}
```

2459 **EXAMPLE 4:** GET to the queue object URI to read two values of the queue object:

```
GET /MyContainer/MyQueue?mimetype;valuerange;values:2 HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-queue
```

2460 The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: application/cdmi-queue

{
  "mimetype" : [
    "text/plain",
    "text/plain"
  ],
  "valuerange" : [
    "0-19",
    "0-20"
  ],
  "value" : [
    "First Enqueued Value",
    "Second Enqueued Value"
  ]
}
```

2461 **EXAMPLE 5:** GET to the queue object URI to read the queue object using multi-part MIME:

```
GET /MyContainer/MyQueue HTTP/1.1
Host: cloud.example.com
Accept: multipart/mixed
```

2462 The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: multipart/mixed; boundary=gc0p4Jq0M2Yt08j34c0p

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/cdmi-queue

{
  "objectType": "application/cdmi-queue",
  "objectID": "00007ED9001035E14BD1BA70C2EE98FC",
```

(continues on next page)

(continued from previous page)

```

    "objectName": "MyQueue",
    "parentURI": "/MyContainer/",
    "parentID" : " 00007ED90010C2414303B5C6D4F83170",
    "domainURI": "/cdmi_domains/MyDomain/",
    "capabilitiesURI": "/cdmi_capabilities/queue/",
    "completionStatus": "Complete",
    "metadata": {
        ...
    },
    "queueValues": "1-2",
    "mimetype": [
        "application/octet-stream",
        "application/octet-stream"
    ],
    "valuerange": [
        "0-19",
        "0-36"
    ],
    "valuetransferencoding": [
        "base64",
        "base64"
    ]
}

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/octet-stream
Content-Transfer-Encoding: binary

<20 bytes of binary data>

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/octet-stream
Content-Transfer-Encoding: binary

<37 bytes of binary data>

--gc0p4Jq0M2Yt08j34c0p--

```

## 11.4 Update a queue object using CDMI

### 11.4.1 Synopsis

To update some or all fields in an existing queue object (excluding the enqueueing of values), the following request shall be performed:

- PUT <root URI>/<ContainerName>/<QueueName>
- PUT <root URI>/<ContainerName>/<QueueName>?metadata:<metadataname>;...
- PUT <root URI>/cdmi\_objectid/<QueueObjectID>
- PUT <root URI>/cdmi\_objectid/<QueueObjectID>?metadata:<metadataname>;...

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate containers.
- <QueueName> is the name of the queue object to be updated.
- <QueueObjectID> is the ID of the queue object to be updated.

### 11.4.2 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 111](#).

Table 111: Capabilities - Update a queue object using CDMI

Capability	Location	Description
cdmi_modify_metadata	Queue object	Ability to modify the metadata of an existing queue object
cdmi_object_access_by_ID	System wide capability	Ability to access the object by ID

### 11.4.3 Request headers

The HTTP request headers for updating a CDMI queue object using CDMI are shown in [Table 112](#)

Table 112: Request headers - Update a queue object Using CDMI

Header	Type	Description	Requirement
Content-Type	Header string	"application/cdmi-queue"	Mandatory

### 11.4.4 Request message body

The request message body fields for updating a queue object using CDMI are shown in [Table 113](#).

Table 113: Request message body - Update a queue object Using CDMI

Field Name	Type	Description	Requirement
metadata	JSON object	Metadata for the queue object. If present, the new metadata specified replaces the existing object metadata. If individual metadata items are specified in the URI, only those items are replaced; other items are preserved. See <a href="#">clause 16</a> for a further description of metadata.	Optional

Continued on next page

Table 113 – continued from previous page

Field Name	Type	Description	Requirement
domainURI	JSON string	URI of the owning domain <ul style="list-style-type: none"> <li>If different from the parent domain, the user shall have the “cross-domain” privilege (see <code>cdmi_member_privileges</code> in Table 80).</li> <li>If not specified, the existing domain shall be preserved.</li> </ul>	Optional
deserialize	JSON string	URI of a CDMI data object with a value that contains a queue object serialized as specified in clause 15. The serialized queue object shall be deserialized to update the existing queue object. <ul style="list-style-type: none"> <li>If the destination queue object URI and the source serialized queue object URI both do not specify individual fields, the destination queue object shall be replaced with the contents of the serialized source queue object, with the exception that the destination queue values shall be preserved. See 11.6 to deserialize enqueued items.</li> <li>If the destination queue object URI or the source serialized queue object URI specifies individual fields, only the fields specified shall be used to update the destination queue object. If specified fields are not present in the source, these fields shall be ignored. If the value field is specified, it shall be ignored.</li> <li>If the destination queue object URI and the source serialized queue object URI both specify fields, an HTTP status code of 400 <code>Bad Request</code> shall be returned to the client.</li> </ul> If there are insufficient permissions to read the serialized queue object at the source URI or update the queue object at the destination URI, or if the read operation fails, the update shall return an HTTP status code of 400 <code>Bad Request</code> , and the destination queue object shall not be updated.	Optional <sup>2</sup>

Continued on next page

Table 113 – continued from previous page

Field Name	Type	Description	Requirement
copy	JSON string	<p>URI of a source CDMI queue object that shall be copied into the existing destination queue object.</p> <ul style="list-style-type: none"> <li>If the destination queue object URI and the copy source queue object URI both do not specify individual fields, the destination queue object shall be replaced with the source queue object, with the exception that the destination queue values shall be preserved. See 11.6 to copy enqueued items.</li> <li>If the destination queue object URI or the copy source queue object URI specifies individual fields, only the fields specified shall be used to update the destination queue object. If specified fields are not present in the source, these fields shall be ignored. If the value field is specified, it shall be ignored.</li> <li>If the destination queue object URI and the copy source queue object URI both specify fields, an HTTP status code of 400 <i>Bad Request</i> shall be returned to the client.</li> </ul> <p>If there are insufficient permissions to read the queue object at the source URI or update the queue object at the destination URI, or if the read operation fails, the update shall return an HTTP status code of 400 <i>Bad Request</i>, and the destination queue object shall not be updated.</p>	Optional <sup>2</sup>
deserializevalue	JSON string	<p>A queue object serialized as specified in <a href="#">clause 15</a> and encoded using base 64 encoding rules described in RFC 4648 [19], that shall be deserialized to update the existing queue object.</p> <p>The object ID of the serialized queue object shall match the object ID of the destination queue object. Otherwise, the server shall return an HTTP status code of 400 <i>Bad Request</i>.</p> <ul style="list-style-type: none"> <li>If the destination queue object URI does not specify individual fields, the destination queue object shall be replaced with the contents of the serialized source queue object, with the exception that the destination queue values shall be preserved. See 11.6 to deserialize enqueued items.</li> <li>If the destination queue object URI specifies individual fields, only the fields specified shall be used to update the destination queue object. If specified fields are not present in the source, these fields shall be ignored. If the value field is specified, it shall be ignored.</li> </ul> <p>If there are insufficient permissions update the queue object at the destination URI, the update shall return an HTTP status code of 400 <i>Bad Request</i>, and the destination queue object shall not be updated.</p>	Optional <sup>2</sup>

#### 11.4.5 Response header

The HTTP response header for updating a CDMI queue object using CDMI is shown in [Table 114](#)

<sup>2</sup> Only one of these fields shall be specified in any given operation. Except for value, these fields shall not be stored.

Table 114: Response header - Update a queue object Using CDMI

Header	Type	Description	Requirement
Location	Header string	The server shall respond with an absolute URI to which the reference redirects if the object is a reference.	Conditional

### 11.4.6 Response message body

A response body can be provided as per RFC 2616 [23].

### 11.4.7 Response status

Table 115 describes the HTTP status codes that occur when updating a queue object using CDMI.

Table 115: HTTP status codes - Update a queue object using CDMI

HTTP Status	Description
204 No Content	The data object content was returned in the response.
302 Found	The resource is a reference to another resource.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server.

### 11.4.8 Examples

EXAMPLE 1: PUT to the queue object URI to set new metadata:

```
PUT /MyContainer/MyQueue HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdm-queue

{
  "metadata" : {
  }
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

EXAMPLE 2: PUT to the queue object URI to move six queue values from another queue:

```
PUT /MyContainer/MyQueue HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdm-queue

{
  "move": "/MyContainer/SourceQueue?value:10-15"
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

## 11.5 Delete a queue object using CDMI

### 11.5.1 Synopsis

To delete an existing queue object, along with all enqueued values, the following request shall be performed:

- DELETE <root URI>/<ContainerName>/<QueueName>
- DELETE <root URI>/cdmi\_objectid/<QueueObjectID>

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate containers.
- <QueueName> is the name of the queue object to be deleted.
- <QueueObjectID> is the ID of the queue object to be deleted.

### 11.5.2 Capability

Capabilities that indicate which operations are supported are shown in [Table 116](#).

Table 116: Capabilities - Delete a queue object using CDMI

Capability	Location	Description
cdmi_delete_queue	Queue object	Ability to delete an existing queue object
cdmi_object_access_by_ID	System wide capability	Ability to access the object by ID

### 11.5.3 Request header

Request headers can be provided as per RFC 2616 [\[23\]](#).

### 11.5.4 Request message body

A request body can be provided as per RFC 2616 [\[23\]](#).

### 11.5.5 Response headers

Response headers can be provided as per RFC 2616 [\[23\]](#).

### 11.5.6 Response message body

A response body can be provided as per RFC 2616 [\[23\]](#).

### 11.5.7 Response status

[Table 117](#) describes the HTTP status codes that occur when deleting a queue object using CDMI.



Table 117: HTTP status codes - Delete a queue object Using CDMI

HTTP Status	Description
204 No Content	The queue object was successfully deleted.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server.

### 11.5.8 Example

EXAMPLE 1: DELETE to the queue object URI:

```
DELETE /MyContainer/MyQueue HTTP/1.1
Host: cloud.example.com
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

## 11.6 Enqueue a new queue object value using CDMI

### 11.6.1 Synopsis

To enqueue one or more values into an existing queue object, the following request shall be performed:

- POST <root URI>/<ContainerName>/<QueueName>
- POST <root URI>/cdmi\_objectid/<QueueObjectID>

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate containers that already exist, with one slash (i.e., “/”) between each pair of container names.
- <QueueName> is the name of the queue object to be enqueued into.
- <QueueObjectID> is the ID of the queue object to be enqueued into.

### 11.6.2 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 118](#).

Table 118: Capabilities - Enqueue a new queue object value using CDMI

Capability	Location	Description
cdmi_modify_value	Queue object	Ability to enqueue a value into an existing queue object
cdmi_multipart_mime	System wide capability	Ability to modify a queue object using multi-part MIME
cdmi_object_access_by_ID	System wide capability	Ability to access the object by ID

### 11.6.3 Request headers

The HTTP request headers for enqueueing a new CDMI queue object value using CDMI are shown in [Table 119](#)

Table 119: Request headers - Enqueue a new queue object value using CDMI

Header	Type	Description	Requirement
Content-Type	Header string	“application/cdmi-queue” or “multipart/mixed”  If “multipart/mixed”, the first part shall contain a body of content-type “application/cdmi-queue”, and the subsequent parts shall contain the queue values as described in <a href="#">8.3</a> .	Mandatory

### 11.6.4 Request message body

The request message body fields for enqueueing a new queue object value using CDMI are shown in [Table 120](#).

Table 120: Request message body - Enqueue a new queue object value using CDMI

Field Name	Type	Description	Requirement
mimetype	JSON array of JSON strings	<p>MIME type(s) of the data value(s) to be enqueued into the queue object.</p> <ul style="list-style-type: none"> <li>If this field is not included and multi-part MIME is not being used, the value of "text/plain" shall be assigned as the field value.</li> <li>If this field is not included and multi-part MIME is being used, the value of the "Content-Type" header of the corresponding MIME part shall be assigned as the field value.</li> <li>The same number of array elements shall be present as is present in the value field, and the mimetype field shall be associated with the value in the corresponding position.</li> <li>This mimetype field value shall be converted to lower case before being stored.</li> </ul>	Optional
copy	JSON string	<p>URI of a source CDMI data object or queue object from which the value shall be copied and enqueued.</p> <ul style="list-style-type: none"> <li>If a copy source object URI to a data object is provided, the value, mimetype, and valuetransferencoding field values from the source data object are used to enqueue the new item into the destination queue object.</li> <li>If a copy source object URI to a queue object is provided, the corresponding value, mimetype, and valuetransferencoding field values of the specified number of enqueued items in the source queue object are copied to the destination queue object.</li> </ul>	Optional <sup>3</sup>
move	JSON string	<p>URI of a source CDMI data object or queue object from which the value shall be moved and enqueued.</p> <ul style="list-style-type: none"> <li>If a move source object URI to a data object is provided, the value, mimetype, and valuetransferencoding field values from the source data object are used to enqueue the new item into the destination queue object, and the source data object is atomically deleted.</li> <li>If a move source object URI to a queue object is provided, the corresponding value, mimetype, and valuetransferencoding field values of the specified number of enqueued items in the source queue object are transferred to the destination queue object and atomically removed from the source queue object.</li> </ul>	Optional <sup>3</sup>

Continued on next page

Table 120 – continued from previous page

Field Name	Type	Description	Requirement
valuetransferencoding	JSON array of JSON strings	<p>The value transfer encoding used for the queue object value. Two value transfer encodings are defined:</p> <ul style="list-style-type: none"> <li>• “utf-8” indicates that the queue object value contains a valid UTF-8 string, and shall be transported as a UTF-8 string in the value field.</li> <li>• “base64” indicates that the queue object value may contain arbitrary binary sequences, and shall be transported as a base 64 encoded string in the value field. Setting the contents of the queue object value field to any value other than a valid base 64 string shall result in an HTTP status code of 400 Bad Request being returned to the client.</li> <li>• “json” indicates that the queue object value contains a valid JSON object, and the value field shall contain a JSON object. Setting the contents of the queue object value field to any value other than a valid JSOM object shall result in an HTTP status code of 400 Bad Request being returned to the client.</li> <li>• If this field is not included and multi-part MIME is not being used, the value of “utf-8” shall be assigned as the field value.</li> <li>• If this field is not included and multi-part MIME is being used, the value of “utf-8” shall be assigned as the field value if the “Content-Type” header of the corresponding MIME part includes the charset parameter as defined in RFC 2046 of “utf-8” (e.g., “; charset=utf-8”). Otherwise, the value of “base64” shall be assigned as the field value. This field applies only to the encoding of the value when represented in JSON; the “Content-Transfer-Encoding” header of the part specifies the encoding of the value within a multi-part MIME request, as defined in RFC 2045 [9].</li> </ul>	Optional
value	JSON array of JSON strings	<p>Data to be enqueued into the queue object.</p> <ul style="list-style-type: none"> <li>• If this field is not included and multi-part MIME is being used, the contents of the MIME parts shall be assigned as the field value.</li> <li>• If the corresponding valuetransferencoding field indicates UTF-8 encoding, the value shall be a UTF-8 string escaped using the JSON escaping rules described in RFC 4627 [5].</li> <li>• If the corresponding valuetransferencoding field indicates base 64 encoding, the value shall be first encoded using the base 64 encoding rules as described in RFC 4648 [19].</li> <li>• If the corresponding valuetransferencoding field indicates JSON encoding, the value shall contain a JSON object.</li> </ul>	Optional <sup>3</sup>

<sup>3</sup> Only one of these fields shall be specified in any given operation. Except for value, these fields shall not be stored. If more than one of these fields is supplied, the server shall respond with an HTTP status code of 400 Bad Request.

### 11.6.5 Response headers

Response headers can be provided as per RFC 2616 [23].

### 11.6.6 Response message body

A response body can be provided as per RFC 2616 [23].

### 11.6.7 Response status

Table 121 describes the HTTP status codes that occur when enqueueing a new queue object using CDMI.

Table 121: HTTP status codes - Enqueue a new queue object value Using CDMI

HTTP Status	Description
204 No Content	The new queue object values were enqueued.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server.

### 11.6.8 Examples

EXAMPLE 1: POST to the queue object URI a new value:

```
POST /MyContainer/MyQueue HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-queue

{
  "mimetype" : [
    "text/plain"
  ],
  "value" : [
    "Value to Enqueue"
  ]
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

EXAMPLE 2: POST to the queue object URI to copy an existing value:

```
POST /MyContainer/MyQueue HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-object

{
  "copy" : "/MyContainer/MyDataObject.txt"
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

EXAMPLE 3: POST to the queue object URI to transfer 20 values from another queue object:

```
POST /MyContainer/MyQueue HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-object

{
  "move" : "/MyContainer/FirstQueue?values:20"
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

**EXAMPLE 4: POST to the queue object URI two new values:**

```
POST /MyContainer/MyQueue HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-object

{
  "mimetype" : [
    "text/plain",
    "text/plain"
  ],
  "value" : [
    "First",
    "Second"
  ]
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

**EXAMPLE 5: POST to the queue object URI two new values, one with base 64 transfer encoding and one with utf-8 transfer encoding:**

```
POST /MyContainer/MyQueue HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-object

{
  "mimetype": [
    "text/plain",
    "text/plain",
    "application/json"
  ],
  "valuetransferencoding": [
    "utf-8",
    "base64",
    "json"
  ],
  "value": [
    "First",
    "U2Vjb25k",
    {
      "value" : "test"
    }
  ]
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

**EXAMPLE 6: POST to the queue object URI the binary contents of two new values using multi-part MIME:**

```
POST /MyContainer/MyQueue HTTP/1.1
Host: cloud.example.com
```

(continues on next page)

(continued from previous page)

```
Content-Type: multipart/mixed; boundary=gc0p4Jq0M2Yt08j34c0p

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/cdmi-queue

{}

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/octet-stream
Content-Transfer-Encoding: binary

<20 bytes of binary data>

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/octet-stream
Content-Transfer-Encoding: binary

<37 bytes of binary data>

--gc0p4Jq0M2Yt08j34c0p--
```

The following shows the response.

```
HTTP/1.1 204 No content
```

**EXAMPLE 7: POST to the queue object URI the mime types and binary contents of two new values using multi-part MIME:**

```
POST /MyContainer/MyQueue HTTP/1.1
Host: cloud.example.com
Content-Type: multipart/mixed; boundary=gc0p4Jq0M2Yt08j34c0p

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/cdmi-queue

{
  "mimetype" : [
    "application/pdf",
    "image/jpeg"
  ]
}

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/octet-stream
Content-Transfer-Encoding: binary

<20 bytes of binary data>

--gc0p4Jq0M2Yt08j34c0p
Content-Type: application/octet-stream
Content-Transfer-Encoding: binary

<37 bytes of binary data>

--gc0p4Jq0M2Yt08j34c0p--
```

The following shows the response.

```
HTTP/1.1 204 No content
```

## 11.7 Delete a queue object value using CDMI

### 11.7.1 Synopsis

To delete one or more of the oldest enqueued values in an existing queue, the following request shall be performed:

- DELETE <root URI>/<ContainerName>/<QueueName>?value
- DELETE <root URI>/<ContainerName>/<QueueName>?values:<count>
- DELETE <root URI>/<ContainerName>/<QueueName>?values:<range>
- DELETE <root URI>/cdmi\_objectid/<QueueObjectID>?value
- DELETE <root URI>/cdmi\_objectid/<QueueObjectID>?values:<count>
- DELETE <root URI>/cdmi\_objectid/<QueueObjectID>?values:<range>

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate containers.
- <QueueName> is the name of the queue object to be deleted from.
- <QueueObjectID> is the ID of the queue object to be deleted from.
- <count> is the number of values, starting from the oldest, to be removed from the queue object. If more queue object entries are requested to be deleted than exist in the queue object, the count shall be considered equal to the number of entries in the queue object.
- <range> is the lowest to highest numbers as found in the queueValues field that are to be removed from the queue object. The first range value shall be smaller or equal to the lowest queue value. If the first range value is smaller than the lowest queue value, the lowest existing queue value shall be used. If the first range value is larger than the lowest queue value, an HTTP status code of 400 Bad Request shall be returned to the client. If the second range value is higher than the highest existing queue value, the highest existing queue value shall be used, which allows for idempotent queue value deletion.

The “?value” suffix at the end of the queue resource URI shall be included to distinguish the deletion of the oldest value from the deletion of the queue object itself, as described in 11.5 (which deletes all enqueued values).

### 11.7.2 Capabilities

Capabilities that indicate which operations are supported are shown in Table 122.

Table 122: Capabilities - Delete a queue object value using CDMI

Capability	Location	Description
cdmi_modify_value	Queue object	Ability to delete a value from an existing queue object
cdmi_object_access_by_ID	System wide capability	Ability to access the object by ID

### 11.7.3 Request header

Request headers can be provided as per RFC 2616 [23].

### 11.7.4 Request message body

A request body can be provided as per RFC 2616 [23].



## 11.7.5 Response headers

Response headers can be provided as per RFC 2616 [23].

## 11.7.6 Response message body

A response body can be provided as per RFC 2616 [23].

## 11.7.7 Response status

Table 123 describes the HTTP status codes that occur when deleting a queue object value using CDMI.

Table 123: HTTP status codes - Delete a queue object value using CDMI

HTTP Status	Description
204 No Content	The queue object value was successfully deleted.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
409 Conflict	The operation conflicts with a non-CDMI access protocol lock or has caused a state transition error on the server.

## 11.7.8 Example

EXAMPLE 1: DELETE to the queue object URI value to delete the oldest enqueued value:

```
DELETE /MyContainer/MyQueue?value HTTP/1.1
Host: cloud.example.com
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

EXAMPLE 2: DELETE to the queue object URI value to remove the ten oldest values:

```
DELETE /MyContainer/MyQueue?values:10 HTTP/1.1
Host: cloud.example.com
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

EXAMPLE 3: DELETE to the queue object URI value to remove queue values 10 through 19:

```
DELETE /MyContainer/MyQueue?values:10-19 HTTP/1.1
Host: cloud.example.com
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

## Clause 12

# Capability object resource operations using CDMI

## 12.1 Overview

Capability objects indicate what specific functionality and operations are supported by a given CDMI server, and allow CDMI clients to discover what subset of this International Standard is implemented.

All CDMI servers shall support capabilities and the ability for CDMI clients to read capabilities.

Each CDMI capability object is represented as a JSON object, containing one or more “fields”. For example, the “capabilities” field contains specific capability items.

EXAMPLE 1: CDMI capability object

```
{
  "objectType": "application/cdm-capability",
  "objectID": "00007E7F00104BE66AB53A9572F9F51E",
  "objectName": "cdmi_capabilities/",
  "parentURI": "/",
  "parentID": "00007E7F0010128E42D87EE34F5A6560",
  "capabilities": {
    "cdmi_domains": "true",
    "cdmi_export_nfs": "true",
    "cdmi_export_iscsi": "true",
    "cdmi_queues": "true",
    "cdmi_notification": "true",
    "cdmi_query": "true",
    "cdmi_metadata_maxsize": "4096",
    "cdmi_metadata_maxitems": "1024"
  },
  "childrenrange": "0-3",
  "children": [
    "domain/",
    "container/",
    "dataobject/",
    "queue/"
  ]
}
```

The meaning, use, and permitted values of each field is described in [12.2](#).

## 12.1.1 Capability object addressing

Capability objects are addressed in CDMI in two ways:

- by name (e.g. `https://cloud.example.com/cdmi_capabilities/`); and
- by ID (e.g. `https://cloud.example.com/cdmi_objectid/00007E7F00104BE66AB53A9572F9F51E/`).

Every capability object has a single, globally-unique object ID that remains constant for the life of the object. Each capability object may also have one or more URI addresses that allow the capability object to be accessed.

When a capability object is addressed via more than one unique URIs, all operations may be performed through any of these URIs. For example, a capability object may be accessible via multiple virtual hosting paths, where `https://cloud.example.com/users/snia/cdmi/cdmi_capabilities/` is also accessible through `https://snia.example.com/cdmi/cdmi_capabilities/`.

Following the URI conventions for hierarchical paths, capability URIs shall consist of one or more capability names that are separated by forward slashes ("/") and that end with a forward slash ("/").

If a request is performed against an existing capability resource and the trailing slash at the end of the URI is omitted, the server shall respond with an HTTP status code of 301 Moved Permanently. In addition, a Location header containing the URI with the trailing slash added shall be returned.

Capabilities may also be nested.

EXAMPLE 2: The following URI represents a nested capability:

`https://cloud.example.com/cdmi_capabilities/container/`

A nested capability has a parent capability object, and shall be included in the children field of the parent capability object.

## 12.1.2 Capability object fields

Every CDMI object (excluding capability objects) includes a server-generated "capabilitiesURI" field that contains the URI of the capabilities object that describes which operations are permitted for that CDMI object.

Fig. 7 (shown on the next page) shows the hierarchy of capabilities and shows how the capabilitiesURI links data objects, container objects, queue objects and domain objects into the capabilities tree.

System-wide capabilities are described by the root capabilities object, which is accessible at "<root URI>/cdmi\_capabilities/".

Capabilities cannot be altered by clients, but may be changed by the CDMI server to reflect configuration changes or operational changes. For example, if a CDMI server is upgraded or reconfigured, additional capabilities may become present, or existing capabilities may no longer be present. In practice, capabilities rarely change, and a client can assume that they shall remain constant for the duration of a client-server HTTP/HTTPS session.

Cloud clients should use capabilities to discover what operations are supported. If an operation is attempted on a CDMI object that does not have a corresponding capability, an HTTP status code of 400 Bad Request shall be returned to the client.

The capabilities defined as part of this International Standard are described starting in 12.1.7. Vendor-defined capabilities not specified in this International Standard shall not start with "cdmi\_".

## 12.1.3 Capability object metadata

Capability objects do not have metadata.

## 12.1.4 Capability object access control

Capability objects are not subject to CDMI ACLs. Any authenticated CDMI client shall be capable of reading all Capability objects<sup>1</sup>.

<sup>1</sup> A CDMI Server may filter the visibility of capability objects and/or capability items for security purposes, for example, to prevent the client discovery of the names and characteristics of classification levels above the client's maximum classification level. Such filtering is out of scope of this International Standard.

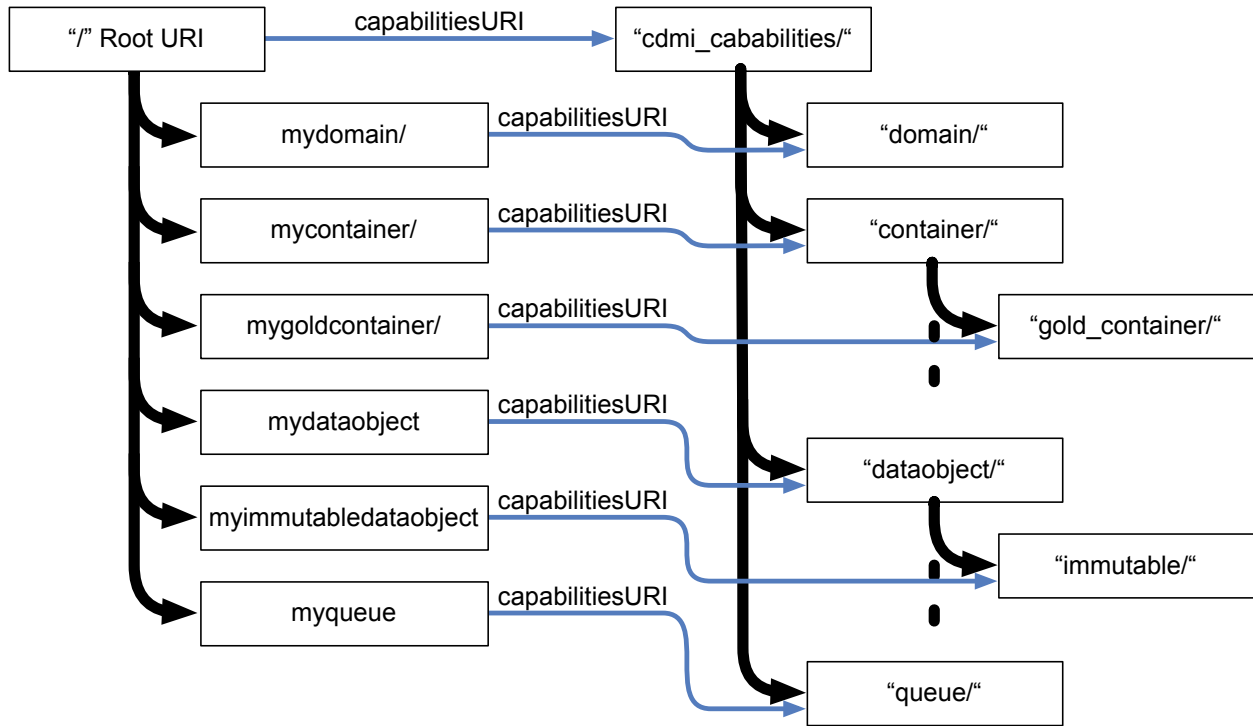


Figure 7: Hierarchy of capabilities

Capabilities may differ from the operations permitted by an Access Control List (ACL) (see 17.1) associated with a given object. For example, a read-only cloud may not permit write access to a container or object, despite the presence of an ACL allowing write access.

### 12.1.5 Queue object consistency

Capability objects are read-only.

### 12.1.6 Capability object representations

The representations in this clause are shown using JSON notation. Both clients and servers shall support UTF-8 JSON representation. The request and response body JSON fields may be specified or returned in any order, with the exception that, if present, for capability objects, the "childrenrange" and "children" fields shall appear last and in that order.

## 12.1.7 Cloud storage system-wide capabilities

Table 124 defines the system-wide capabilities in a cloud storage system. These capabilities, which are found in the capabilities object, are referred to by the root URI (root capabilities).

Table 124: System-wide capabilities

Capability name	Type	Definition
cdmi_domains	JSON string	If present and “true”, the CDMI server supports domains.  If not present, the domainURI field shall not be present in response bodies and the “cdmi_domains” URI shall not be present.
cdmi_export_cifs	JSON string	If present and “true”, the CDMI server supports CIFS exports.
cdmi_dataobjects	JSON string	If present and “true”, the CDMI server supports data objects.
cdmi_export_iscsi	JSON string	If present and “true”, the CDMI server supports iSCSI exports.
cdmi_export_nfs	JSON string	If present and “true”, the CDMI server supports NFS protocol exports.
cdmi_export_occi_iscsi	JSON string	If present and “true”, the CDMI server supports OCCI/iSCSI exports.
cdmi_export_webdav	JSON string	If present and “true”, the CDMI server supports WebDAV exports.
cdmi_metadata_maxitems	JSON string	If present, this capability indicates the maximum number of user-defined metadata items supported per object.  If not present, there is no limit placed on the number of user-defined metadata items.
cdmi_metadata_maxsize	JSON string	If present, this capability indicates the maximum size, in bytes, of each user-defined metadata item supported per object.  If not present, there is no limit placed on the size of user-defined metadata items.
cdmi_metadata_maxtotalsize	JSON string	If present, this capability indicates the maximum size, in bytes, of user-defined metadata supported by the CDMI server.  If not present, there is no limit placed on the size of user-defined metadata.
cdmi_notification	JSON string	If present and “true”, the CDMI server supports notification queues.
cdmi_logging	JSON string	If present and “true”, the CDMI server supports logging queues.
cdmi_query	JSON string	If present and “true”, the CDMI server supports query queues.
cdmi_query_regex	JSON string	If present and “true”, the CDMI server supports query with regular expressions.
cdmi_query_contains	JSON string	If present and “true”, the CDMI server supports query with “contains” expressions.
cdmi_query_tags	JSON string	If present and “true”, the CDMI server supports query with tag-matching expressions.
cdmi_query_value	JSON string	If present and “true”, the CDMI server supports query of value fields.
cdmi_queues	JSON string	If present and “true”, the CDMI server supports queue objects.
cdmi_security_access_control	JSON string	If present and “true”, the CDMI server supports ACLs. See 12.1.9 for additional information.

Continued on next page

Table 124 – continued from previous page

Capability name	Type	Definition
cdmi_security_data_integrity	JSON string	If present and “true”, the CDMI server supports data integrity/authenticity. See 12.1.9 for additional information.
cdmi_security_encryption	JSON string	If present and “true”, the CDMI server supports data at-rest encryption. See 12.1.9 for additional information.
cdmi_security_immutability	JSON string	If present and “true”, the CDMI server supports data immutability/retentions. See 12.1.9 for additional information.
cdmi_security_sanitization	JSON string	If present and “true”, the CDMI server supports data/media sanitization. See 12.1.9 for additional information.
cdmi_serialization_json	JSON string	If present and “true”, the CDMI server supports JSON as a serialization format.
cdmi_snapshots	JSON string	If present and “true”, the CDMI server supports snapshots.
cdmi_references	JSON string	If present and “true”, the CDMI server supports references.
cdmi_object_move_from_local	JSON string	If present and “true”, the CDMI server supports moving CDMI objects from URIs within the same storage system.
cdmi_object_move_from_remote	JSON string	If present and “true”, the CDMI server supports moving CDMI objects from URIs within other CDMI storage systems.
cdmi_object_move_from_ID	JSON string	If present and “true”, the CDMI server supports moving CDMI objects without a path from a /cdmi_objectid/ URI within the same storage system. This effectively adds a path, allowing the object to be accessed by ID and by path.
cdmi_object_move_to_ID	JSON string	If present and “true”, the CDMI server supports moving CDMI objects with a path to a /cdmi_objectid/ URI within the same storage system. This effectively removes the path, leaving the object only accessible by ID.
cdmi_object_copy_from_local	JSON string	If present and “true”, the CDMI server supports copying CDMI objects from URIs within the same storage system.
cdmi_object_copy_from_remote	JSON string	If present and “true”, the CDMI server supports copying CDMI objects from URIs within other CDMI storage systems.
cdmi_object_access_by_ID	JSON string	If present and “true”, the CDMI server supports accessing, updating, and deleting objects through /cdmi_objectid/.
cdmi_post_dataobject_by_ID	JSON string	If present and “true”, the CDMI server supports adding a new data object by ID via POST to “/cdmi_objectid/”.
cdmi_post_queue_by_ID	JSON string	If present and “true”, the CDMI server supports adding a new queue object by ID via POST to “/cdmi_objectid/”.
cdmi_deserialize_dataobject_by_ID	JSON string	If present and “true”, the CDMI server supports deserializing serialized data objects when creating a new data object by ID via POST to “/cdmi_objectid/”.
cdmi_deserialize_queue_by_ID	JSON string	If present and “true”, the CDMI server supports deserializing serialized queue objects when creating a new queue object by ID via POST to “/cdmi_objectid/”.

Continued on next page

Table 124 – continued from previous page

Capability name	Type	Definition
cdmi_serialize_dataobject_to_ID	JSON string	If present and “true”, the CDMI server supports serializing data objects when creating a new data object by ID via POST to “/cdmi_objectid/”.
cdmi_serialize_domain_to_ID	JSON string	If present and “true”, the CDMI server supports serializing domain objects when creating a new data object by ID via POST to “/cdmi_objectid/”.
cdmi_serialize_container_to_ID	JSON string	If present and “true”, the CDMI server supports serializing container objects when creating a new data object by ID via POST to “/cdmi_objectid/”.
cdmi_serialize_queue_to_ID	JSON string	If present and “true”, the CDMI server supports serializing queue objects when creating a new data object by ID via POST to “/cdmi_objectid/”.
cdmi_copy_dataobject_by_ID	JSON string	If present and “true”, the CDMI server supports copying an existing data object when creating a new data object by ID via POST to “/cdmi_objectid/”.
cdmi_copy_queue_by_ID	JSON string	If present and “true”, the CDMI server supports copying an existing queue object when creating a new queue object by ID via POST to “/cdmi_objectid/”.
cdmi_create_reference_by_ID	JSON string	If present and “true”, the CDMI server supports creating a new reference via POST to “/cdmi_objectid/”.
cdmi_copy_dataobject_from_queue	JSON string	If present and “true”, the CDMI server supports the ability to copy to a data object from a queue object.
cdmi_multipart_mime	JSON string	If present and “true”, the CDMI server supports storing and retrieving the value of data and queue objects using multi-part MIME.
cdmi_create_value_range_by_ID	JSON string	If present and “true”, the CDMI server supports a new data object's value to be created with byte ranges through “/cdmi_objectid/”.
cdmi_dac	JSON string	If present and “true”, the CDMI server supports delegated access control.
cdmi_dac_methods	JSON array of JSON strings	If present, this capability contains a list of URI schemes supported for DAC URIs, as specified in the IANA URI Schemes registry. The following schemes shall be supported: <ul style="list-style-type: none"> <li>• “https” – mandatory for all DAC implementations</li> </ul> The following schemes may be supported: <ul style="list-style-type: none"> <li>• “http” – optional for DAC implementations</li> <li>• “mailto” – optional for DAC implementations</li> </ul>
cdmi_enc_cms	JSON string	If present and “true”, the CDMI server supports operations against the contents of CMS encrypted objects.
cdmi_enc_jwe	JSON string	If present and “true”, the CDMI server supports operations against the contents of JWE encrypted objects.
cdmi_enc_inplace	JSON string	If present and “true”, the CDMI server supports operations to encrypt and decrypt objects in place, including updates.
cdmi_enc_access	JSON string	If present and “true”, the CDMI server supports operations to decrypt objects on access.
cdmi_cms_encryption	JSON array of JSON strings	If present, this capability lists which CMS ContentEncryptionAlgorithmIdentifier encryption algorithms are supported for operations against the contents of CMS encrypted objects.

Continued on next page

Table 124 – continued from previous page

Capability name	Type	Definition
cdmi_cms_digest	JSON array of JSON strings	If present, this capability lists which CMS <code>MessageAuthenticationCodeAlgorithm</code> digest algorithms are supported for operations against the contents of CMS encrypted objects.
cdmi_cms_signature	JSON array of JSON strings	If present, this capability lists which CMS <code>SignatureAlgorithmIdentifier</code> signature algorithms are supported for operations against the contents of CMS encrypted objects.
cdmi_jwe_enc	JSON array of JSON strings	If present, this capability lists which JOSE “enc” encryption algorithms are supported for operations against the contents of JWE encrypted objects, as defined in RFC 7518 [15].
cdmi_jwe_alg	JSON array of JSON strings	If present, this capability lists which JOSE “alg” encryption algorithms are supported for operations against the contents of JWE encrypted objects, as defined in RFC 7518 [15].
cdmi_jws_alg	JSON array of JSON strings	If present, this capability lists which JOSE “alg” encryption algorithms are supported for operations against the contents of JWS signatures, as defined in RFC 7518 [15].
cdmi_valuetransferencoding_json	JSON string	If present and “true”, the CDMI server supports JSON value transfer encodings.



## 12.1.8 Storage system metadata capabilities

Table 125 defines the capabilities for storage system metadata in a cloud storage system. These capabilities are found in the capabilities objects for domain objects, data objects, container objects, and queue objects. See 16.2 for a description of these storage system metadata items.

Table 125: Capabilities for storage system metadata

Capability name	Type	Definition
cdmi_acl	JSON string	If present and “true”, the CDMI server supports ACLs. When a CDMI implementation supports ACLs for the purpose of access control, the system-wide capability of <code>cdmi_security_access_control</code> specified in 12.1.7 of 12.1.7 shall also be set to “true”.  If not present, there is no support for ACL-based access control.
cdmi_size	JSON string	If present and “true”, the CDMI server shall generate a <code>cdmi_size</code> storage system metadata for each stored object.
cdmi_ctime	JSON string	If present and “true”, the CDMI server shall generate a <code>cdmi_ctime</code> storage system metadata for each stored object.
cdmi_atime	JSON string	If present and “true”, the CDMI server shall generate a <code>cdmi_atime</code> storage system metadata for each stored object.
cdmi_mtime	JSON string	If present and “true”, the CDMI server shall generate a <code>cdmi_mtime</code> storage system metadata for each stored object.
cdmi_acount	JSON string	If present and “true”, the CDMI server shall generate a <code>cdmi_acount</code> storage system metadata for each stored object.
cdmi_mcount	JSON string	If present and “true”, the CDMI server shall generate a <code>cdmi_mcount</code> storage system metadata for each stored object.
cdmi_dac_uri	JSON string	If present and “true”, the CDMI server supports delegated access control metadata.
cdmi_dac_certificate	JSON string	If present and “true”, the CDMI server supports delegated access control metadata.
cdmi_enc_signature	JSON string	If present and “true”, the CDMI server shall generate a <code>cdmi_signature</code> storage system metadata for each stored object when a corresponding <code>sign_id</code> data system metadata item is present.
cdmi_version_object	JSON string	If present and “true”, the CDMI server shall generate a <code>cdmi_version_object</code> storage system metadata for each version-enabled data object and data object version.
cdmi_version_current	JSON string	If present and “true”, the CDMI server shall generate a <code>cdmi_version_current</code> storage system metadata for each version-enabled data object and data object version.
cdmi_version_oldest	JSON array of JSON strings	If present and “true”, the CDMI server shall generate a <code>cdmi_version_oldest</code> storage system metadata for each version-enabled data object and data object version.
cdmi_version_parent	JSON string	If present and “true”, the CDMI server shall generate a <code>cdmi_version_parent</code> storage system metadata for each data object version that has a previous version.

Continued on next page

Table 125 – continued from previous page

Capability name	Type	Definition
cdmi_version_children	JSON array of JSON strings	If present and “true”, the CDMI server shall generate a <code>cdmi_version_children</code> storage system metadata for each data object version.

## 12.1.9 Data system metadata capabilities

Table 126 defines the capabilities that indicate which data system metadata items are interpreted for objects stored in a cloud storage system. These capabilities are found in the capabilities objects for domains, data objects, containers, and queues. See 16.3 for a description of the meaning of the corresponding data system metadata items.

Table 126: Capabilities for data system metadata

Capability name	Type	Definition
cdmi_assignedsize	JSON string	If present and “true”, the CDMI server supports the <code>cdmi_assignedsize</code> data system metadata as defined in 16.3.
cdmi_data_redundancy	JSON string	If present, the CDMI server supports the <code>cdmi_data_redundancy</code> data system metadata as defined in 16.3. The value of the capability shall be set to a positive numeric string representing the maximum value that the server supports.
cdmi_data_dispersion	JSON string	If present and “true”, the CDMI server supports the <code>cdmi_data_dispersion</code> data system metadata as defined in 16.3.
cdmi_data_retention	JSON string	If present and “true”, the CDMI server supports both the <code>cdmi_retention_id</code> and <code>cdmi_retention_period</code> data system metadata as defined in 16.3.
cdmi_data_autodelete	JSON string	If present and “true”, the CDMI server supports the <code>cdmi_data_autodelete</code> data system metadata as defined in 16.3.
cdmi_data_holds	JSON string	If present and “true”, the CDMI server supports the <code>cdmi_hold_id</code> data system metadata as defined in 16.3.  When a cloud storage system supports holds for the purpose of making data immutable, the system-wide capability of <code>cdmi_security_immutability</code> specified in Table 124 of 12.1.7 shall be present and set to “true”.
cdmi_encryption	JSON array of JSON strings	If present, the CDMI server supports the <code>cdmi_encryption</code> data system metadata as defined in 16.3.  When present, this capability shall contain one or more JSON strings, each string corresponding to an algorithm/mode/length value as described in the <code>cdmi_encryption</code> data system metadata in 16.3.  When a cloud storage system supports at-rest encryption, the system-wide capability of <code>cdmi_security_encryption</code> specified in Table 124 of 12.1.7 shall be present and set to “true”.
cdmi_geographic_placement	JSON string	If present and “true”, the CDMI server supports the <code>cdmi_geographic_placement</code> data system metadata as defined in 16.3.
cdmi_immediate_redundancy	JSON string	If present, the CDMI server supports the <code>cdmi_immediate_redundancy</code> data system metadata as defined in 16.3.  When present, this capability shall contain a string set to a positive numeric string representing the maximum value that the server supports.

Continued on next page

Table 126 – continued from previous page

Capability name	Type	Definition
cdmi_infrastructure_redundancy	JSON string	<p>If present, the CDMI server supports the <code>cdmi_infrastructure_redundancy</code> data system metadata as defined in 16.3.</p> <p>When present, this capability shall contain a string set to a positive numeric string representing the maximum value that the server supports.</p>
cdmi_latency	JSON string	<p>If present and “true”, the CDMI server supports the <code>cdmi_latency</code> data system metadata as defined in 16.3.</p>
cdmi_RPO	JSON string	<p>If present and “true”, the CDMI server supports the <code>cdmi_RPO</code> data system metadata as defined in 16.3.</p>
cdmi_RTO	JSON string	<p>If present and “true”, the CDMI server supports the <code>cdmi_RTO</code> data system metadata as defined in 16.3.</p>
cdmi_sanitization_method	JSON array of JSON strings	<p>If present, the CDMI server supports the <code>cdmi_sanitization_method</code> data system metadata as defined in 16.3.</p> <p>When present, this capability shall contain one or more JSON strings, each string corresponding to a sanitization method as described in the <code>cdmi_sanitization_method</code> data system metadata in 16.3.</p> <p>When a cloud storage system supports sanitization, the system-wide capability of <code>cdmi_security_sanitization</code> specified in Table 124 of 12.1.7 shall be present and set to “true”.</p>
cdmi_throughput	JSON string	<p>If present and “true”, the CDMI server supports the <code>cdmi_throughput</code> data system metadata as defined in 16.3.</p>
cdmi_value_hash	JSON array of JSON strings	<p>If present, the CDMI server supports the <code>cdmi_value_hash</code> data system metadata as defined in 16.3.</p> <p>When present, this capability shall contain one or more JSON strings, each string corresponding to an algorithm/length value as described in the <code>cdmi_value_hash</code> data system metadata in 16.3.</p> <p>When a cloud storage system supports value hashing, the system-wide capability of <code>cdmi_security_data_integrity</code> specified in Table 124 of 12.1.7 shall be present and set to “true”.</p>
cdmi_enc_key_id	JSON string	<p>When the cloud storage system supports the <code>cdmi_enc_key_id</code> data system metadata as defined in clause 16.3, the <code>cdmi_enc_key_id</code> capability shall be present and set to the string value “true”. When this capability is absent, or present and set to the string value “false”, <code>cdmi_enc_key_id</code> data system metadata shall not be used.</p>
cdmi_enc_value_sign_id	JSON string	<p>When the cloud storage system supports the <code>cdmi_enc_value_sign_id</code> data system metadata as defined in clause 16.3, the <code>cdmi_enc_value_sign_id</code> capability shall be present and set to the string value “true”. When this capability is absent, or present and set to the string value “false”, <code>cdmi_enc_value_sign_id</code> data system metadata shall not be used.</p>

Continued on next page

Table 126 – continued from previous page

Capability name	Type	Definition
cdmi_enc_value_verify_id	JSON string	When the cloud storage system supports the <code>cdmi_enc_value_verify_id</code> data system metadata as defined in <a href="#">clause 16.3</a> , the <code>cdmi_enc_value_verify_id</code> capability shall be present and set to the string value “true”. When this capability is absent, or present and set to the string value “false”, <code>cdmi_enc_value_verify_id</code> data system metadata shall not be used.
cdmi_enc_object_sign_id	JSON string	When the cloud storage system supports the <code>cdmi_enc_object_sign_id</code> data system metadata as defined in <a href="#">clause 16.3</a> , the <code>cdmi_enc_object_sign_id</code> capability shall be present and set to the string value “true”. When this capability is absent, or present and set to the string value “false”, <code>cdmi_enc_object_sign_id</code> data system metadata shall not be used.
cdmi_enc_object_verify_id	JSON string	When the cloud storage system supports the <code>cdmi_enc_object_verify_id</code> data system metadata as defined in <a href="#">clause 16.3</a> , the <code>cdmi_enc_object_verify_id</code> capability shall be present and set to the string value “true”. When this capability is absent, or present and set to the string value “false”, <code>cdmi_enc_object_verify_id</code> data system metadata shall not be used.
cdmi_versioning	JSON array of JSON strings	If present, this capability indicates that the cloud storage system shall support versioning of data objects and contains a list of which versioning behaviors are supported. The following values are defined: <ul style="list-style-type: none"> <li>• “value” indicates that the system shall support the versioning of the object value.</li> <li>• “user” indicates that the system shall support the versioning of the object value and user metadata.</li> <li>• “all” indicates that the system shall support the versioning of all updates made to a data object.</li> </ul> When present, the system shall support the following storage system metadata: <code>cdmi_version_object</code> , <code>cdmi_version_current</code> , <code>cdmi_version_oldest</code> , <code>cdmi_version_parent</code> , and <code>cdmi_version_children</code> as indicated by the corresponding storage system metadata capabilities.
cdmi_versions_count	JSON string	If present, this capability specifies the maximum number of historical versions that may be specified. If absent, restrictions on the number of historical versions specified shall be ignored.
cdmi_version_age	JSON string	If present, this capability specifies the maximum age of historical versions that may be specified. If absent, restrictions on the age of historical versions specified shall be ignored.
cdmi_versions_size	JSON string	If present, this capability specifies the maximum total size of historical versions that may be specified. If absent, restrictions on the size of historical versions specified shall be ignored.

## 12.1.10 Data object capabilities

Table 127 defines the capabilities for data objects in a cloud storage system.

Table 127: Capabilities for data objects

Capability name	Type	Definition
cdmi_read_value	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to read the object's value.
cdmi_read_value_range	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to read the object's value with byte ranges.
cdmi_read_metadata	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to read the object's metadata.
cdmi_modify_value	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to modify the object's value.
cdmi_modify_value_range	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to modify the object's value with byte ranges.
cdmi_modify_metadata	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to modify the object's metadata.
cdmi_modify_deserialize_dataobject	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability of the data object to deserialize a serialized data object into the data object as an update.
cdmi_delete_dataobject	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to delete the object.

### 12.1.11 Container object capabilities

Table 128 defines the capabilities for containers in a cloud storage system.

Table 128: Capabilities for container objects

Capability name	Type	Definition
cdmi_list_children	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to list the container’s children.
cdmi_list_children_range	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to list the container’s children with ranges.
cdmi_read_metadata	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to read the container’s metadata.
cdmi_modify_metadata	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to modify the container’s metadata.
cdmi_modify_deserialize_container	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability of the container object to deserialize a serialized container object into the container object as an update.
cdmi_snapshot	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability of the container object to create a new snapshot.
cdmi_serialize_dataobject	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to serialize a data object.
cdmi_serialize_container	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to serialize the container and all children’s contents.
cdmi_serialize_queue	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to serialize a queue object.
cdmi_serialize_domain	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to serialize the domain and all child domains.
cdmi_deserialize_container	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability of the container to deserialize the serialized containers and associated serialized children into the container.
cdmi_deserialize_queue	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability of the container to deserialize the serialized queue objects into the container.
cdmi_deserialize_dataobject	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability of the container to deserialize the serialized data objects into the container.
cdmi_create_dataobject	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability of the container to add a new data object.
cdmi_post_dataobject	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability of the container to add a new data object via POST.
cdmi_post_queue	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability of the container to add a new queue object via POST.

Continued on next page

Table 128 – continued from previous page

Capability name	Type	Definition
cdmi_create_container	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to create a new container object via PUT.
cdmi_create_queue	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to create new queue objects..
cdmi_create_reference	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to create a new child reference via PUT.
cdmi_export_container_cifs	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to export a container as a file system via CIFS.
cdmi_export_container_nfs	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to export a container as a file system via NFS.
cdmi_export_container_iscsi	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to export a container as a file system via iSCSI.
cdmi_export_container_occi	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to export a container as a file system via OCCI.
cdmi_export_container_webdav	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to export a container as a file system via WebDAV.
cdmi_delete_container	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to delete a container.
cdmi_move_container	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to move a container object into a container.
cdmi_copy_container	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to copy a container object into a container.
cdmi_move_dataobject	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to move a data object into a container.
cdmi_copy_dataobject	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to copy a data object into a container.
cdmi_create_value_range	JSON string	If present and “true”, this capability indicates that the container allows a new data object’s value to be created with byte ranges.



## 12.1.12 Domain object capabilities

Table 129 defines the capabilities for domains in a cloud storage system. (All capabilities refer to what may be done via CDMI content-type operations.

Table 129: Capabilities for domain objects

Capability name	Type	Definition
cdmi_create_domain	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to add a new subdomain.
cdmi_delete_domain	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to delete a domain.
cdmi_move_domain	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to move a domain.
cdmi_domain_summary	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to support domain summaries.
cdmi_domain_members	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to support domain user management.
cdmi_list_children	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to list the domain's children.
cdmi_read_metadata	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to read the domain's metadata.
cdmi_modify_metadata	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to modify the domain's metadata.
cdmi_modify_deserialize_domain	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to deserialize a serialized domain object into the domain object as an update.
cdmi_copy_domain	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to copy the domain (via PUT) to another URI.
cdmi_deserialize_domain	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to deserialize serialized domains and associated serialized children into the domain.

Continued on next page

Table 129 – continued from previous page

Capability name	Type	Definition
cdmi_authentication_methods	JSON array of JSON strings	<p>If present, the CDMI server supports authentication methods that are supported by a domain.</p> <p>When present, this capability shall contain one or more of the following JSON strings:</p> <ul style="list-style-type: none"> <li>• “anonymous” - Absence of authentication supported</li> <li>• “basic” - HTTP basic authentication supported (RFC 2617 [8])</li> <li>• “digest” - HTTP digest authentication supported (RFC 2617 [8])</li> <li>• “krb5” - Kerberos authentication supported, using the Kerberos domain specified in the CDMI domain (RFC 4559 [14])</li> <li>• “x509” - certificate-based authentication via TLS (RFC 5246 [24])</li> <li>• “s3” - S3 API signed header authentication supported</li> <li>• “openstack” - OpenStack Identity API header authentication supported</li> </ul> <p>Interoperability with these authentication methods are not defined by this International Standard. Servers may include other authentication methods not included in the above list. In these cases, it is up to the CDMI client and CDMI server to ensure interoperability.</p>

### 12.1.13 Queue object capabilities

Table 130 defines the capabilities for queue objects in a cloud storage system.

Table 130: Capabilities for queue objects

Capability name	Type	Definition
cdmi_read_value	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to read a queue’s value.
cdmi_read_metadata	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to read the queue’s metadata.
cdmi_modify_value	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to modify the queue’s value.
cdmi_modify_metadata	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to modify the queue’s metadata.
cdmi_modify_deserialize_queue	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to deserialize a serialized queue into the queue as an update.
cdmi_delete_queue	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to delete a queue.
cdmi_move_queue	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to move a queue to another URI.
cdmi_copy_queue	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to copy a queue to another URI.
cdmi_reference_queue	JSON string	If present and “true”, this capability indicates that the CDMI server shall support the ability to reference a queue from another queue.

## 12.2 Read a capabilities object using CDMI

### 12.2.1 Synopsis

To read an existing capability object, the following requests shall be performed:

- GET <root URI>/cdmi\_capabilities/<Capability>/<TheCapability>/
- GET <root URI>/cdmi\_capabilities/<Capability>/<TheCapability>/?<fieldname>;<fieldname>;...
- GET <root URI>/cdmi\_capabilities/<Capability>/<TheCapability>/?children:<range>;...
- GET <root URI>/cdmi\_objectid/<CapabilityObjectID>/
- GET <root URI>/cdmi\_objectid/<CapabilityObjectID>/?<fieldname>;<fieldname>;...
- GET <root URI>/cdmi\_objectid/<CapabilityObjectID>/?children:<range>;...

Where:

- <root URI> is the path to the CDMI cloud.
- <Capability> is zero or more parent capabilities.
- <TheCapability> is the name specified for the capability to be read from.
- <fieldname> is the name of a field.
- <range> is a numeric range within the list of children.
- <prefix> is a matching prefix that returns all metadata items that start with the prefix value.
- <CapabilityObjectID> is the ID of the capability object to be read from.

### 12.2.2 Capabilities

Capabilities that indicate which operations are supported are shown in [Table 131](#).

Table 131: Capabilities - Read a capabilities object using CDMI

Capability	Location	Description
cdmi_object_access_by_ID	System wide capability	Ability to access the object by ID

### 12.2.3 Request headers

The HTTP request headers for reading a CDMI capabilities object using CDMI are shown in [Table 132](#).

Table 132: Request headers - Read a capabilities object using CDMI

Header	Type	Description	Requirement
Accept	Header string	"application/cdm-capability" or a consistent value as described in <a href="#">5.5.2</a>	Optional

### 12.2.4 Request message body

A request body shall not be provided.

### 12.2.5 Response headers

The HTTP response headers for reading a CDMI capabilities object using CDMI are shown in [Table 133](#).

Table 133: Response headers - Read a capabilities object Using CDMI

Header	Type	Description	Requirement
Content-Type	Header string	"application/cdmf-capability"	Mandatory

## 12.2.6 Response message body

The response message body fields for reading a CDMI capabilities object using CDMI are shown in Table 134.

Table 134: Response message body - Read a capabilities object using CDMI

Field name	Type	Description	Requirement
objectType	JSON string	"application/cdmf-capability"	Mandatory
objectID	JSON string	Object ID of the object	Mandatory
objectName	JSON string	Name of the object	Mandatory
parentURI	JSON string	URI for the parent object Appending the "objectName" to the "parentURI" shall always produce a valid URI for the object.	Mandatory
parentID	JSON string	Object ID of the parent capability object.	Mandatory
capabilities	JSON object	The capabilities supported by the corresponding object. Capabilities in the "/cdmf_capabilities/" object are system-wide capabilities. Capabilities found in children objects under "/cdmf_capabilities/" correspond to the capabilities of a specific subset of objects.	Mandatory
childrenrange	JSON string	The child capabilities of the capability expressed as a range. If a range of child capabilities is requested, this field indicates the children returned as a range.	Mandatory
children	JSON array of JSON strings	Names of the children capabilities objects. For the root container capabilities, this includes "domain/", "container/", "dataobject/", and "queue/". Within each of these capabilities objects, further more specialized capabilities profiles may be specified by the CDMI server.	Mandatory

If individual fields are specified in the GET request, only these fields are returned in the result body. Optional fields that are requested but do not exist are omitted from the result body.

## 12.2.7 Response status

Table 135 describes the HTTP status codes that occur when reading a capabilities object using CDMI.

Table 135: HTTP status codes - Read a capabilities object using CDMI

HTTP status	Description
200 OK	The capabilities object content was returned in the response.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.
406 Not Acceptable	The server is unable to provide the object in the content type specified in the Accept header.

## 12.2.8 Examples

EXAMPLE 1: GET to the root container capabilities URI to read all fields of the container:

```
GET /cdmi_capabilities/ HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-capability
```

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: application/cdmi-capability

{
  "objectType": "application/cdmi-capability",
  "objectID": "00007E7F00104BE66AB53A9572F9F51E",
  "objectName": "cdmi_capabilities/",
  "parentURI": "/",
  "parentID": "00007E7F0010128E42D87EE34F5A6560",
  "capabilities": {
    "cdmi_domains": "true",
    "cdmi_export_nfs": "true",
    "cdmi_export_iscsi": "true",
    "cdmi_queues": "true",
    "cdmi_notification": "true",
    "cdmi_query": "true",
    "cdmi_metadata_maxsize": "4096",
    "cdmi_metadata_maxitems": "1024"
  },
  "childrenrange": "0-3",
  "children": [
    "domain/",
    "container/",
    "dataobject/",
    "queue/"
  ]
}
```

EXAMPLE 2: GET to the root container capabilities URI to read the capabilities and children of the container:

```
GET /cdmi_capabilities/?capabilities;children HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-capability
```

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: application/cdmi-capability

{
  "capabilities": {
    "cdmi_domains": "true",
    "cdmi_export_nfs": "true",
    "cdmi_export_iscsi": "true",
    "cdmi_queues": "true",
    "cdmi_notification": "true",
    "cdmi_query": "true",
    "cdmi_metadata_maxsize": "4096",
    "cdmi_metadata_maxitems": "1024"
  },
  "children": [
    "domain/",
    "container/",
    "dataobject/",
    "queue/"
  ]
}
```

EXAMPLE 3: GET to the root container capabilities URI to read the first two children contained within a domain:

```
GET /cdmi_capabilities/?childrenrange;children:0-1 HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-capability
```

2738

The following shows the response.

```
HTTP/1.1 200 OK
Content-Type: application/cdmi-capability

{
  "childrenrange" : "0-1",
  "children" : [
    "domain/",
    "container/"
  ]
}
```

## Clause 13

# Exported protocols

### 13.1 Overview

Container objects can be exported via multiple storage protocols. This is specified by adding an export field to the container object. Each container object export field has a set of well-defined items that include:

- zero or more export types, which specifies the export protocol;
- for each export type, the identity of the container as standardized by the export protocol;
- for each export type, the domain of the protocol name server for the clients being served;
- for each export type, the list of who may mount that container via that protocol, identified as standardized by that protocol or optionally by leveraging the name mapping protocol (see 13.1.3) and specifying CDMI user or groupnames;
- for each export type, required protocol-specific export parameters;
- for each export type, optional protocol-specific export parameters; and
- for each export type, export control parameters.

All cloud storage systems shall support the storage of the export field. The ability to export containers via a specified protocol is determined by the presence or absence of a `cdmi_export_<protocol>` system wide capabilities, which are listed in 12.1.7. The ability to export a specific container via a specified protocol is indicated by the `cdmi_export_<protocol>` capability in the specified container.

Container object exports are represented as a JSON object, containing zero or more named protocol exports.

EXAMPLE 1: CDMI container exports

```
{
  "exports" : {
    "nfs" : {<BR>
    "hosts" : { "*.mycollege.edu", "derf.cs.myuni.edu" },
    "domain" : "lab.mycollege.edu",
    "usermap" : {
      { <cdminame>, <map>, <nfsname> },
      { "jimsmith", "<-->", "jims" },
      { [ordered list of CDMIname/operator/NFSname triples] },
      { "*", "<-->", "*" }
    }
    "groupmap" : {
      { "admins", "<-", "wheel" },
      { "everyone", "<-", "*" }
    }
  }
  "cifs" : {
    "hosts" : "*",
    "domain" : "lab.mycollege.edu",
    "usermap" : {
      { "jimsmith", "<-->", "james.smith" }
```

(continues on next page)



(continued from previous page)

```

        { [ordered list of CDMIname/operator/NFSname triples] },
        { "*", "<-->", "*" }
    }
    "groupmap" : {
        { "admins", "<--", "Administrators" },
        { "everyone", "<--", "*" }
    }
}
}
}

```

The meaning, use, and permitted values for the fields associated with each export type are described later in this clause.

### 13.1.1 Container object export addressing

Container object exports are addressed in CDMI in two ways:

- by name (e.g. <https://cloud.example.com/container/?exports>); and
- by ID (e.g. [https://cloud.example.com/cdmi\\_objectid/00007ED900104E1D14771DC67C27BF8B/?exports](https://cloud.example.com/cdmi_objectid/00007ED900104E1D14771DC67C27BF8B/?exports)).

See 9.1 for more details on container object addressing.

### 13.1.2 Container object export fields

The export of a container, via data path protocols other than CDMI, is accomplished by creating or updating a container and supplying one or more export protocol structures, one for each such protocol. In this International standard, all such protocols are referred to as foreign protocols.

This International standard defines JSON export structures for several well known foreign protocols. All depend on the following user and groupname mapping feature in the case that multi-protocol access to the container is desired. However, name mapping is not required if an external domain is used, or if CDMI is used only to provision containers to be used exclusively by foreign protocols.

Implementations that support authenticated and authorized access to CDMI objects via both CDMI and foreign protocols need a way to support the setting of security on a per-object basis. The numerous methods of doing this include:

- Defining or adopting a security scheme and mapping all requests into that scheme. CDMI implementations that adopt this scheme shall use a name mapping technique to accomplish it, as (a) this mapping is easier for administrators to manage than straight id-to-id mapping, and (b) it is desired that interoperable CDMI implementations behave similarly in this respect. This means that the name of the principal in an incoming request is mapped to the name of a principal in the security domain, and that principal's id is acquired and used in the authorization procedure.
- Allowing each protocol to set its own security, which implies that an object might be accessible to a given user via one protocol but not another.
- Using the security scheme of the last protocol that was used to set permissions on the object. This method also requires mapping the principal in the incoming request to a principal in the security domain of the object. As in the first case, the server shall use a name mapping procedure to obtain the id that is used to authorize the user against the desired object's ACL.

CDMI does not mandate which method shall be used. It does, however, specify how users and groups shall be mapped between protocols.

### 13.1.3 Mapping names from CDMI to another protocol

Clients wishing to restrict exports via foreign protocols to mounting only by certain users and groups may be required to provide user and groupname mapping information to the server. This mapping information is also required if access to the container is desired by multiple protocols, e.g., both CDMI and NFS. The mapping is done as follows.

1. When a network share on a CDMI container is created, the server should use the appropriate mechanism, e.g., Powershell `WmiClass.Create( )` on the Windows platform or `/etc/exports` on Unix, to limit permitted mounts of the share from other servers, as specified in the “hosts” line of the “exports” property. The syntax of the hosts line follows the syntax of `/etc/exports` in the Linux operating system, as encoded in a JSON string. If the CDMI server is unable to limit mounts as specified by the hosts line, an error shall result, but the success or failure of the operation depends on the implementation.
  2. When any request requiring the use of a CDMI principal name comes in via a foreign protocol, the foreign domain controller to which the foreign server belongs shall be queried for the principal name corresponding to the user id given in the request. Failure to procure the principal name shall cause the original request to fail.
  3. The usermap list for that protocol shall be searched, in order, for an entry matching the username gotten from the foreign domain controller (see 13.1.8 for details on the search). If no match is found, the request shall be denied. The search results may be kept in the same cache entry as the information from the preceding step.
  4. The CDMI principal name gotten from the first matching usermap entry during this search is then used to authorize the user request via the security mechanism of the protocol whose security governs access to the object.
- Groupname mapping for each foreign protocol shall be specified in a groupname field of the foreign protocol export specification. Its syntax is identical to the syntax for the username field.

### 13.1.4 Administrative users

By default, the following users shall be considered “root”, or administrative users, and equivalent to each other:

- root (Unix/NFS/LDAP),
- Administrator (Windows/AD/CIFS), and
- the domain owner (CDMI).

Servers shall automatically map these users to the root user of the target protocol unless otherwise instructed by the usermaps.

As an automatic mapping does not meet strict security standards, servers shall override these built-in entries with any usermap entries that apply to one or more root users.

In the following example, root gets mapped to nobody, and everyone else is mapped to a user of the same name in the NFS domain and the CDMI domain.

#### EXAMPLE 2: NFS export user mapping

```
PUT /MyContainer HTTP/1.1
Host: cloud.example.com
Accept: application/vnd.org.snia.cdm.container+json
Content-Type: application/vnd.org.snia.cdm.container+json

{
  "exports": {
    "nfs": {
      "usermap": [
        [
          "nobody",
          "<-",
          "root"
        ],
        [
          "*",
          "<-->",
          "*"
        ]
      ]
    }
  }
}
```

### 13.1.5 Mapping domains from CDMI to another protocol

The internet domain name corresponding to each export shall be given as a JSON-formatted string in the “domain” child of the protocol export specification. If it is not present, it shall be assumed that the domain is the same as that of the server hosting the CDMI implementation.

### 13.1.6 User and domain mapping caching

The lookup to a foreign domain controller can be quite expensive, especially for stateless protocols such as NFS v3, in which it can be theoretically required for nearly every operation. It shall be permissible to cache the results of this lookup. The recommended lifetime of a username cache entry is 30 minutes. Implementations should use this value or less when possible. Servers shall flush this cache whenever a change is made to the exports metadata concerning the protocol being cached. A client may request that the cache be flushed by reading in the usermap data for one or more protocols and writing them back without change. Servers shall flush their username mapping caches, as part of the rewrite operation, for any protocol for which the usermap information has been changed or reset.

For authorization by group to operate via a foreign protocol, a similar mapping exercise must be performed. Multiple lookups to the foreign domain controller may be required to get all the groupnames for a given user (e.g., it is common for an NFS user to be a member of several groups). A groupname cache may be used to mitigate the cost of these lookups. The recommended lifetime of a groupname cache entry is 12 hours. Implementations should use this value or less when possible. Clients may force a flush of the cache by reading in and resetting the group map information. Servers shall immediately flush their groupname mapping cache, as part of the rewrite operation, for any protocol for which the group map information has been changed or reset.

### 13.1.7 Permissions mapping

The permissions sets of file-serving protocols, unfortunately, do not map on a one-to-one basis to each other. NFSv4 ACLs, Windows ACLs, POSIX ACLs, NFSv3 perms and object-based capabilities all are capable of representing security conditions that the others are not, except NFSv3, which is the least expressive. The primary area of concern is in representing the possibly rich set of permissions in a CDMI ACL in a more restricted perms-based system, such as NFSv3, for display to users.

As there are a number of possible ways to coordinate the permissions/ACLs and CDMI ACLs, this International specification does not mandate a particular method. However, all mappings of user and groupnames between domains shall use the name mapping mechanism specified in 13.1.8.

### 13.1.8 User and groupname mapping syntax and evaluation rules

A BNF-style grammar for name mapping is as follows:

```
name_mapping_list = protocol protocol mapping_list
protocol = "cdmi" | "nfs" | "cifs" | "ldap"
mapping_list = name mapping_operator name
name = pattern | utf8_name | quoted_utf8_name
quoted_utf8_name = " utf8_name "
utf8_name = <any legal utf8 character sequence not including the characters ",',\,/,,:*,?>
pattern = <utf8_name> * | *
mapping_operator = "<--" | "<-->" | "-->"
```

To restate this in English, a mapping entry consists of two names separated by a directional indicator. As most environments use the same usernames and groupnames across administrative domains, the most common mapping is “\* <--> \*”, which maps any name to the same name in the foreign protocol domain, and vice versa. It is highly recommended that this be both the default map and the last entry on all more complex maps.

CDMI specifies pattern matching on names in the name map, but only prefix matching is required. The symbol “\*” at the end of a character string shall match zero or more occurrences of any non-whitespace character.

Evaluation of the name mapping list shall proceed in order; once a match is made, evaluation shall cease and the result of the match shall be returned.

If no matches are found on the match list, the result is system dependent. However, it is recommended that servers either deny access altogether or map the user in question to the equivalent of “anonymous” on the destination protocol. It is also recommended that an entry be devoted to the special user “EVERYONE”.

## 13.2 NFS exported protocol

To export a container via NFS, the information required is exactly what the server implementation will use to do the export. Normally, this information is contained in the `/etc/exports` file on a server or the equivalent. Administrators should be aware that lines may be automatically added to that file for each CDMI container that is exported.

Required members of the protocol structure for NFS are described in [Table 136](#).

Table 136: Required members of the NFS protocol structure

Member	Description
protocol	The protocol being requested. This value shall be “NFSv3”, “NFSv4”, “NFSv4.1”, or any subsequent NFS version enshrined in a major IETF RFC. Version 2 of NFS is not supported by CDMI.
exportpath	The pathname to which the export should be surfaced. This value shall be a UTF8 string of the form [ <code>&lt;server&gt;</code> ]:/ <code>&lt;path&gt;</code> , where the <code>&lt;server&gt;</code> component is optional, (e.g., “ <code>eeserver:/lessons/number1</code> ”). The <code>&lt;server&gt;</code> component of the path must be obtained from an administrator of the service running the CDMI implementation.
exportdomain	The internet domain of the protocol name server for the clients being served. This value is normally the name of the LDAP domain for the organization, e.g., “ <code>iti.edu</code> ”. A value of “.” shall be interpreted to be the DNS name of the domain occupied by the CDMI server.
mode	This value shall be “ <code>ro</code> ”, “ <code>rw</code> ”, “ <code>root</code> ” or “ <code>rpc_gsssec</code> ” and becomes the default export mode. Hosts requiring different access shall be specified in the optional “ <code>rw_mode</code> ”, “ <code>ro_mode</code> ”, and “ <code>root_mode</code> ” structure members. However, the “ <code>rpc_gsssec</code> ” mode overrides all other modes, and all other mode members and their contents shall be ignored if it is specified.
control	Export control for the container. This value shall be “ <code>immediate</code> ”, “ <code>off</code> ”, “ <code>on</code> ”, or <code>&lt;n&gt;</code> (a number). Servers may set the value to <code>on</code> , but clients shall not. A numeric value ( <code>&lt;n&gt;</code> ) indicates that the export should be shut down in <code>&lt;n&gt;</code> seconds, possibly after a message has been sent to clients mounting the export. If a client specifies a value for <code>&lt;n&gt;</code> but the server does not support delayed shutdown of exports, then <code>&lt;n&gt;</code> shall be interpreted to mean <code>off</code> .

Optional export parameters for NFS are described in [Table 137](#).

Table 137: Optional members of the NFS protocol structure

Parameter	Description
domain_servers	A list of server names or IP addresses that function as name servers for the domain given in “ <code>domain</code> ”. If given, this list shall override the names obtainable by the CDMI server via other programmatic means.
mount_name	The name the client should use to surface the export. This name replaces the last name in the path string, (e.g., mounting “ <code>eeserver:/lessons/number1</code> ” with a mountname of “ <code>1</code> ” over the directory <code>/somepath/lessons/num1</code> should result in a <code>/somepath/lessons/1</code> directory on the client).
hosts	A list of hosts that can access the container in the mode given in “ <code>mode</code> ”. The default shall be “ <code>*</code> ”; other values restrict the possibilities.
root_hosts	A list of hosts that can access the container in superuser mode. The default shall be an empty list.
rw_hosts	A list of hosts that can access the container in <code>r/w</code> mode. The default shall be an empty list.
ro_hosts	A list of hosts that can access the container in <code>r/o</code> mode only. The default shall be an empty list.
mount_type	One of the two strings “ <code>hard</code> ” or “ <code>soft</code> ”. Clients hang when a server serving a hard mount becomes unresponsive. Clients with soft mounts generate error messages. The default is implementation dependent.
recurse	This value shall be either “ <code>true</code> ” or “ <code>false</code> ”. The default shall be “ <code>true</code> ”. When true, recurse indicates that mounts within the CDMI directory structure (presumably put there by other NFS operations) shall be followed and the mounted directory exposed as though it were part of the CDMI container actually being exported. This parameter is equivalent to the Linux “ <code>crossmnt</code> ” parameter.

Other export parameters for NFS are not specified by the CDMI protocol but may be included in the export structure.

These parameters include Linuxisms, such as “sync”, “no\_wdelay”, “insecure\_locks”, and “no\_acl”, as well as any other parameters used by a given server operating system. In all such cases, the parameter shall be specified as a JSON tuple in which “true” and “false” are explicitly called out for binary flags, and a JSON-formatted string or list is used for other parameters.

EXAMPLE 3: NFS export options

```
"exports":
{ "nfs":
{
...
{"no_wdelay": "true" },
{"refer": "otherserver://path/leaf"},
...
}
}
```

### 13.2.1 Export control

Export control is accomplished with the use of a single member, named “control”:

- The value “immediate” shall indicate to the server that the export shall be made successfully before the PUT operation returns. Servers shall reset the value to “on” and place that in the reply.
- The value “off” shall indicate to the server that the export, if new, shall not be enabled, and if existing, shall be shut down and all client connections forcibly broken.
- A numeric value <n> shall indicate that the server shall wait <n> seconds before forcibly shutting down the export and breaking client connections. Whether the server sends a warning message to clients, giving them a chance to exit from the connection gracefully, is recommended but implementation dependent. Once the export has been shut down, the server shall also change the value of “control” to “off” in the export structure.

Servers shall support wildcard matching on the “\\*” and “?” characters in the hosts lists (this is standard practice), so that “\\*\\*.cs.ucsc.edu” matches all servers in the cs.ucsc.edu department.

Servers may support netgroup names in the various hosts lists. When this functionality is supported, these names shall resolve to ordinary lists of hostnames via queries to the domain nameserver.

Servers may also support IP address ranges in the various lists of hosts. These IP addresses shall be augmented by the same wildcard matching as is used for ordinary host names (e.g., “192.168.1.\*” exports to all the machines on a default home network). Client-side developers should note that “exporting to” only means making a container available for export. The client must still mount the exported container before there is a connection with the server.

Users wishing to use optional and vendor-specific settings are responsible for determining from the CDMI product vendor the legal settings and their format. Servers shall return an HTTP status code of 400 Bad Request when an export setting does not conform to an allowable setting on the server.

## 13.3 CIFS exported protocol

To export a container via CIFS, the information required is exactly what the server implementation will use to do the export. Where this information is contained on a server is implementation dependent. The server may add or delete lines automatically to and from that file for each CDMI container that is exported or unexported.

Required members of the protocol structure for CIFS are described in [Table 138](#)

Table 138: Required members of the CIFS protocol structure

Member	Description
share_name	The name that CIFS shall use to discover the share.
exportdomain	The domain of the protocol name server for the clients being served. This value is normally the name of the Active Directory LDAP domain for the organization, e.g. "iti.edu". A value of "." shall be interpreted to be the domain occupied by the CDMI server.
mode	This value shall be either "ro" or "rw".
control	Export control for the container. This value shall be "immediate", "off", or <n> (a number). Servers may set the value to on, but clients shall not. The semantics and normative requirements are exactly the same as for NFS, as documented in the paragraph "13.2.1" in the subclause on NFS Exports (see <a href="#">13.2</a> ).

There is no protocol specification; CDMI assumes that normal SMB protocol negotiation will take place.

An optional export parameter is "comment," which is often used as a user-friendly share name on the client.

Other export parameters for CIFS are not specified by the CDMI protocol but may be included in the export structure. These parameters include vendor settings such as "forcegroup", "umask", "caching", and "oplocks", as well as any other parameters used by a given server operating system. In all such cases, the parameter shall be specified as a JSON tuple in which "true" and "false" are explicitly called out for binary flags, and a JSON-formatted string or list is used for other parameters.

### EXAMPLE 4: CIFS export

```
"exports":
{ "cifs":
  {
    ...
    {"caching": { "manual", "document", "program" },
    {"oplocks": "true"},
    ...
  }
}
```

Users wishing to manipulate vendor-specific settings are responsible for determining from the CDMI product vendor the legal settings and their format. Servers shall return an HTTP status code of 400 Bad Request when an export setting does not conform to an allowable setting on the server.

For more detail on the use of the OCCI export protocol structure attributes, see [13.4](#). Because the actual networking and access control is under the control of a hidden, common infrastructure implementing both OCCI and CDMI, the normal permission structure shall not be provided.

## 13.4 OCCI exported protocol

Container objects can be exported via multiple protocols. This is especially useful when CDMI is being used as a storage interface in a cloud computing environment, as illustrated in Fig. 8 below.




Figure 8: CDMI and OCCI in an integrated cloud computing environment

In this example, CDMI containers may also be used as virtual disks by virtual machines in the cloud computing environment. The cloud computing infrastructure management is shown as implementing both an Open Cloud Computer Interface (OCCI) and CDMI interfaces. With the internal knowledge of the network and the virtual machine manager's mapping of drives, this infrastructure may associate the CDMI containers to the guests using the appropriate exported protocol.

To support exported protocols and improve their interoperability with CDMI, CDMI provides a type of exported protocol that contains information obtained via the OCCI interface. In addition, OCCI provides a type of storage that corresponds

to a CDMI container that is exported with a specific type of protocol used by OCCI. A client of both interfaces performs operations that align the architectures, including the following:

- The client creates a CDMI container through the CDMI interface and exports it as an OCCI export protocol type. The CDMI container object ID is returned as a result.
- The client creates a virtual machine through the OCCI interface and attaches a storage volume of type CDMI using the object ID and protocol type. The OCCI virtual machine ID is returned as a result.
- The client updates the export protocol structure of the CDMI container object with the OCCI virtual machine ID to allow the virtual machine access to the container.
- The client starts the virtual machine through the OCCI interface.

CDMI defines an export protocol structure for the Open Cloud Computing Interface (13.4) as follows:

- The protocol is “OCCI/<protocol standard>” (e.g., “OCCI/NFSv4”).
- The identifier is the CDMI object ID.
- A JSON array of URIs to OCCI compute resources shall have access (permissions) to the exported container.

### EXAMPLE 5: OCCI export

```
"OCCI/iSCSI":
{
  "identifier": "00007E7F00104BE66AB53A9572F9F51E",
  "permissions":
  [
    "https://example.com/compute/0/",
    "https://example.com/compute/1/"
  ]
}
```

For more detail on using the OCCI export protocol structure attributes, see 13.1. Because the actual networking and access control is under the control of a hidden, common infrastructure that implements both OCCI and CDMI, the normal permission structure shall not be provided.



## 13.5 iSCSI exported protocol

CDMI defines the export of a container using the iSCSI protocol (see RFC 7143 [4]). Each container is exported as a single SCSI Logical Unit as a Logical Unit Number (LUN). One or more iSCSI initiators import the LUN through an iSCSI target node and port using one or more iSCSI network portals (IP addresses).

The export is described by the presence of an export field structure on the container that specifies the

- export protocol ("Network/iSCSI");
- iSCSI target information (IP addresses or fully qualified domain names, target identifier, and LUN);
- logical unit world-wide name; and
- iSCSI initiators having access.

The target identifier may be in *iqn*, *naa*, or *eui* format and shall have the target portal group tag appended in hexadecimal.

### 13.5.1 Read container

All of the information in the export structure is returned:

EXAMPLE 6: iSCSI export

```
"exports" :
{
  "Network/iSCSI": {
    "portals": [
      "192.168.1.101",
      "192.168.1.102"
    ],
    "target_identifier": "iqn.2010-01.com.cloudprovider:acmeroot.container1,t,0x0001",
    "logical_unit_number": "3",
    "logical_unit_name": "0x60012340000000000000000000000001",
    "permissions": [
      "iqn.2010-01.com.acme:host1",
      "iqn.2010-01.com.acme:host2"
    ]
  }
}
```

### 13.5.2 Create and update containers

The following export field contents, when included in a container create or update, indicates that the container shall be exported via iSCSI. Support for either of these operations is indicated by the `cdmi_export_iscsi` capability on the parent container of the created container or of the existing container, respectively.

EXAMPLE 7: iSCSI export creation

```
"exports" :
{
  "Network/iSCSI": {
    "permissions": [
      "iqn.2010-01.com.acme:host1",
      "iqn.2010-01.com.acme:host2"
    ]
  }
}
```

For these export creation operations, the CDMI implementation selects the IP portals, iSCSI target, logical unit number, and logical unit name; these are not supplied. Only the list of initiator identifiers that are to have access to the container are specified.

### 13.5.3 Modify an export

The following code modifies an export on an existing container. Support for this operation is indicated by the `cdmi_export_iscsi` on the parent container of the existing container. For this operation, only the current list of initiator identifiers that are to have access to the container are specified.

EXAMPLE 8: iSCSI export modification

```
"exports" :
{
  "Network/iSCSI": {
    "permissions": [
      "iqn.2010-01.com.acme:host2"
    ]
  }
}
```

## 13.6 WebDAV exported protocol

CDMI defines an export protocol structure for the WebDAV standard as follows (see RFC 4918 [6]):

- The protocol is “Network/WebDAV”.
- The path of the WebDAV mount point is as presented to clients (including server host name).
- The list of who may access the share is determined by the standard CDMI ACLs for each resource as exported via WebDAV.

### EXAMPLE 9: WebDAV export

```
"Network/WebDAV" :  
{  
  "identifier": "/users",  
  "permissions": "domain"  
}
```

In this example, the value “domain” in the permissions field indicates that user credentials should be mapped through the domain membership in the domain of the CDMI container being exported.

WebDAV supports locking, but it is up to implementations to support any locking of access through CDMI as a result, and the interaction between the two protocols is purposely not described in this International standard.

2986 **Clause 14**

2987 **CDMI snapshots**

2988 **14.1 Overview**

2989 A snapshot is a point-in-time copy (image) of a container and all of its contents, including subcontainers and all data  
2990 objects and queue objects. The client names a snapshot of a container at the time the snapshot is requested. A snapshot  
2991 operation creates a new container to contain the point-in-time image. The first processing of a snapshot operation also  
2992 adds a `cdmi_snapshots` child container to the source container. Each new snapshot container is added as a child of  
2993 the `cdmi_snapshots` container. The snapshot does not include the `cdmi_snapshots` child container or its contents  
2994 (see Fig. 9).

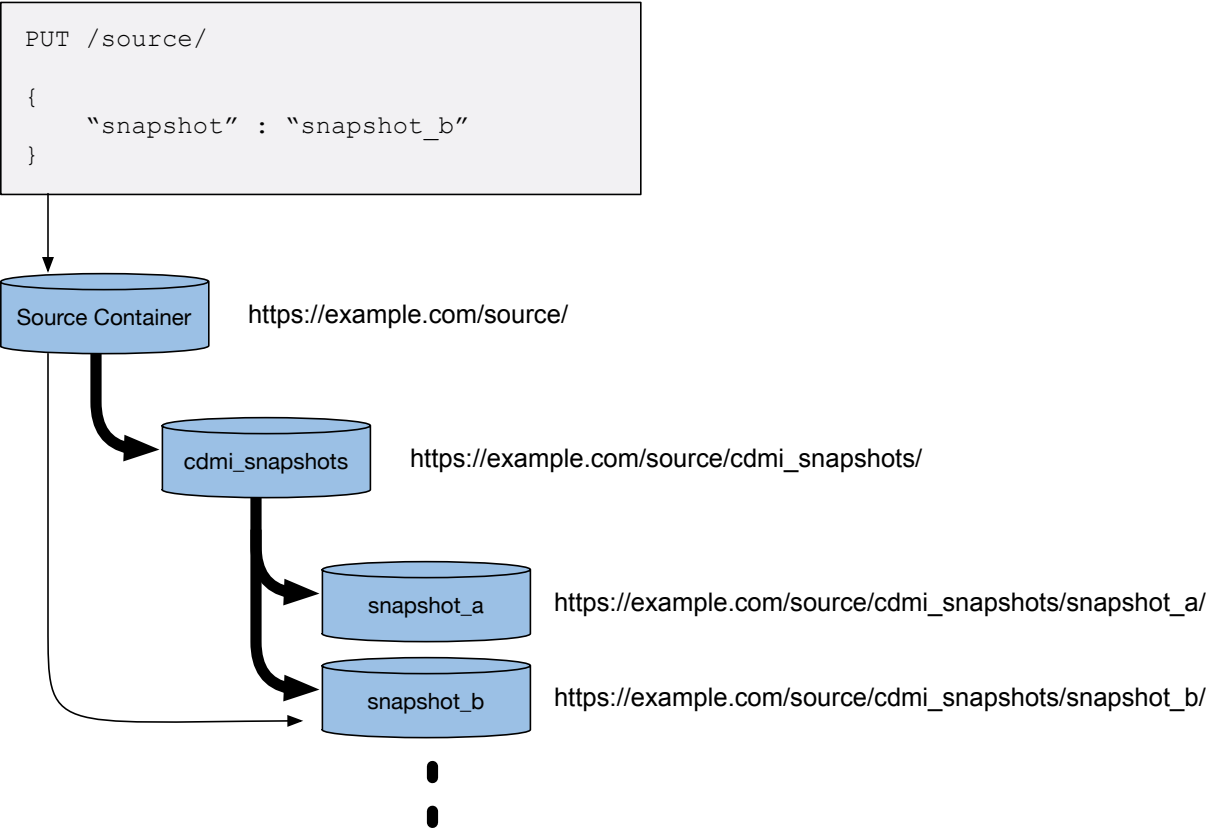


Figure 9: Snapshot container structure

## 14.2 Creating a snapshot

A snapshot operation is requested using the container update operation (see 9.4), in which the snapshot field specifies the requested name of the snapshot.

A snapshot may be accessed in the same way that any other CDMI™ object is accessed. An important use of a snapshot is to allow the contents of the source container to be restored to their values at a previous point in time using a CDMI copy operation.

### 14.2.1 Example

EXAMPLE 1: PUT to an existing container to create a snapshot:

```
PUT /MyContainer/ HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-container

{
  "snapshot" : "MySnapshot"
}
```

The following shows the response.

```
HTTP/1.1 201 Created
```

## 14.3 Deleting a snapshot

A snapshot can be deleted by performing a CDMI container delete operation on the corresponding child container in the `cdmi_snapshots` container, or by performing a CDMI container delete operation on the snapshot Object ID.

### 14.3.1 Example

EXAMPLE 1: DELETE to an existing snapshot:

```
DELETE /MyContainer/cdmi_snapshots/MySnapshot HTTP/1.1
Host: cloud.example.com
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

## Clause 15

# Serialization/deserialization

### 15.1 Overview

Bulk data movement is often needed between, into, or out of clouds. When moving bulk data, cloud serialization operations provide a means to normalize data to a canonical, self-describing format, which includes:

- data migration between clouds,
- data migration during upgrades (or replacements) of cloud implementations, and
- robust backup.

The canonical format of serialized data describes how the data is to be represented in a byte stream. As long as this byte stream is not changed during the transfer from source to destination, the data may be reconstituted on the destination system.

## 15.2 Canonical format

Support for CDMI serialization using JSON as the canonical format requires the presence of the `cdmi_serialization_json` capability.

The canonical format shall represent specified data objects and container objects as they exist within the storage system. Each object shall be represented by the metadata for the object, identifiers, and the data stream contents of the data object. Because data and storage system metadata is inherited from enclosing container objects, all parent metadata shall be represented in the top-level of the canonical format. To preserve the actual metadata values that apply to the data object that is being serialized, the non-overridden metadata is included from both the immediate parent container object of the specified object and from the parent of each higher-level container object.

The canonical format shall have the following characteristics:

- recursive JSON for the data object, consistent with the rest of CDMI;
- user and data system metadata for each data object/container object;
- data stream contents for each data object and queue object;
- binary data represented using escaped JSON strings; and
- typing of data values consistent with CDMI JSON representations.

### 15.2.1 Example JSON canonical serialized format

EXAMPLE 1: In this example, a data object and a queue object in a container object have been selected for serialization:

```
{
  "objectType": "application/cdmi-container",
  "objectID": "00007E7F00102E230ED82694DAA975D2",
  "objectName": "MyContainer/",
  "parentURI": "/",
  "parentID": "00007E7F0010128E42D87EE34F5A6560",
  "domainURI": "/cdmi_domains/MyDomain/",
  "capabilitiesURI": "/cdmi_capabilities/container/",
  "completionStatus": "Complete",
  "metadata": {
    ...
  },
  "exports": {
    "OCCE/iSCSI": {
      "identifier": "00007E7F00104BE66AB53A9572F9F51E",
      "permissions": [
        "https://example.com/compute/0/",
        "https://example.com/compute/1/"
      ]
    },
    "Network/NFSv4": {
      "identifier": "/users",
      "permissions": "domain"
    }
  },
  "childrenrange": "0-1",
  "children": [
    {
      "objectType": "application/cdmi-object",
      "objectID": "00007ED900104F67307652BAC9A37C93",
      "objectName": "MyDataObject.txt",
      "parentURI": "/MyContainer/",
      "parentID": "00007E7F00102E230ED82694DAA975D2",
      "domainURI": "/cdmi_domains/MyDomain/",
      "capabilitiesURI": "/cdmi_capabilities/dataobject/",
      "completionStatus": "Complete",
      "mimetype": "text/plain",
      "metadata": {
        ...
      },
      "valuerange": "0-36",
    }
  ]
}
```

(continues on next page)



(continued from previous page)

```

        "valuetransferencoding": "utf-8",
        "value": "This is the Value of this Data Object"
    },
    {
        "objectType": "application/cdmi-queue",
        "objectID": "00007E7F00104BE66AB53A9572F9F51E",
        "objectName": "MyQueue",
        "parentURI": "/MyContainer/",
        "parentID": "00007E7F00102E230ED82694DAA975D2",
        "domainURI": "/cdmi_domains/MyDomain/",
        "capabilitiesURI": "/cdmi_capabilities/queue/",
        "completionStatus": "Complete",
        "metadata": {
            ...
        },
        "queueValues": "0-1",
        "mimetype": [
            "text/plain",
            "text/plain"
        ],
        "valuetransferencoding": [
            "utf-8",
            "utf-8"
        ],
        "valuerange": [
            "0-2",
            "0-3"
        ],
        "value": [
            "red",
            "blue"
        ]
    }
]
}

```

3038 To allow efficient deserialization in stream mode when serializing container objects to JSON, data object `value` fields  
 3039 and container `children` arrays should be the last items in the canonical serialized JSON format.

## 15.3 Exporting serialized data

A canonical encoding of the data is obtained by creating a new data object and specifying that the source for the creation is to serialize a given CDMI™ data object, container object, or queue object. On a successful serialization, the result shall be a data object that is created with the serialized data as its value. If a container object has an exported block protocol, the serialized data may contain the block-by-block contents of that container object along with its metadata.

The resulting data object that is produced is the canonical representation of the selected data object, container object and children, or queue object.

- If the source specified is a data object, the canonical format shall contain all data object fields, including the `value`, `valuetransferencoding`, and `metadata` fields.
- If the source being specified is a queue object, the canonical format shall contain all queue object fields, including the `value` and `valuetransferencoding` fields of enqueued items, along with the metadata of the queue object itself.
- If the source being specified is a container object, the canonical format shall contain all container object fields, recursively, including all children of the container object. If a user attempts to serialize a container object that includes children that the user, who is performing the serialization operation, does not have permission to read, these objects shall not be included in the resulting serialized object.

When performing a serialization operation, objects shall only be included if the principal initiating the serialization has sufficient permissions to read those objects.

## 15.4 Importing serialized data

Canonical data may be deserialized back into the cloud by creating a new data object, container object, or queue object and by specifying that the source for the creation is to deserialize a given CDMI data object or by specifying the serialized data in base 64 encoding in the `deserializevalue` field.

The destination may or may not exist previously. If not, a create operation is performed. If a container object already exists, an update operation with serialized children shall update the container object and all children. If the serialized container object does not contain children, only the container object is updated. Data objects are recreated as specified in the canonical format, including all metadata and the data object ID.

Table 139: Serialization import behaviour

User has <code>cross_domain</code>	User specifies <code>domainURI</code>	Description
No	No	The <code>domainURI</code> of the parent object shall match the <code>domainURI</code> in each serialized object being deserialized.  If the <code>domainURI</code> in any serialize object does not match the <code>domainURI</code> of the parent object, the entire deserialize operation shall fail, and an HTTP status code of 400 <code>Bad Request</code> shall be returned.
No	Yes	The specified <code>domainURI</code> shall be used, overriding the original <code>domainURI</code> in each serialized object being deserialized.  If a <code>domainURI</code> other than the <code>domainURI</code> of the parent is specified, the entire deserialize operation shall fail, and an HTTP status code of 400 <code>Bad Request</code> shall be returned.
Yes	No	The original <code>domainURI</code> in each serialized object being deserialized shall be used.  If any of the original <code>domainURI</code> in each serialized object being deserialized is not valid in the context of the storage system on which the deserialization operation is being performed, the entire deserialize operation shall fail, and an HTTP status code of 400 <code>Bad Request</code> shall be returned.
Yes	Yes	The specified <code>domainURI</code> shall be used, overriding the original <code>domainURI</code> in each serialized object being deserialized.  If a <code>domainURI</code> that is specified is not valid in the context of the storage system on which the deserialization operation is being performed, the entire deserialize operation shall fail, and an HTTP status code of 400 <code>Bad Request</code> shall be returned.

Deserialization operations shall restore all metadata from the specified source. If the original provider of the serialized data-supported vendor extensions is through custom metadata keys and values, then these customized requirements shall be restored when deserialized. However, the custom metadata keys and values may be treated as user metadata (preserved, but not interpreted) by the destination provider. Preservation allows custom data requirements to move between clouds without losing this information.

## Clause 16

# Metadata

### 16.1 Overview

CDMI metadata allows for additional information to be associated with stored objects. JSON objects, strings and arrays are used to transfer metadata in CDMI operations, which allows for metadata to be hierarchical. CDMI servers may place a restriction on the number of metadata items, maximum size per metadata item, and total size of metadata items, as specified in the `cdmi_metadata_maxitems`, `cdmi_metadata_maxsize`, and `cdmi_metadata_maxtotalsize` capabilities. CDMI servers shall not place a restriction on the depth of the metadata hierarchy and number of array items, outside of the above restrictions.

When objects are created, object metadata is created according to the following process:

1. Metadata items specified in the create operation are added, overriding pre-existing metadata items
2. Storage System metadata items are added to the object, overriding pre-existing metadata items subject to the restrictions described in [Section 16.2](#)

When objects are updated, object metadata is updated according to the following process:

1. Existing metadata items are deleted, changed and/or added, as specified in the update operation
2. Storage System metadata items are updated for the object, overriding pre-existing metadata items subject to the restrictions described in [Section 16.2](#)

When objects are read, object metadata is returned according to the following process:

1. Data System Metadata items is inherited from the parent container
2. Metadata items stored with the object are returned, overriding any inherited Data System Metadata items

## 16.2 Support for storage system metadata

After an object has been created or updated, the storage system metadata, as described in Table 140, shall be generated or updated by the cloud storage system, and shall immediately be made available to a CDMI client in the metadata that is returned as a result of the create operation and any subsequent retrievals.

Which storage system metadata is supported by the CDMI server defined in 12.1.8. Storage system metadata that is not supported by the CDMI server shall be preserved.

Table 140: Storage system metadata

Metadata name	Type	Description	Requirement
cdmi_size	JSON string	The number of bytes consumed by the object.  This storage system metadata item is computed by the storage system, and any attempts to set or modify it will be ignored.	Optional
cdmi_ctime	JSON string	The time when the object was created, in ISO-8601 point-in-time format, as described in 5.6.  For a newly created object, this value shall be set to the creation time.  This metadata value can only be updated by a client if it has the “backup_operator” privilege. If a client does not have the backup operator privilege, updates of this metadata item shall be ignored.	Optional
cdmi_atime	JSON string	The time when the object was last accessed in ISO-8601 point-in-time format, as described in 5.6. The access or modification of a child is not considered an access of a parent container (access/modify times do not propagate up the tree).  For a newly created object, this value shall be set to the creation time.  This metadata value can only be updated by a client if it has the “backup_operator” privilege. If a client does not have the backup operator privilege, updates of this metadata item shall be ignored.	Optional
cdmi_mtime	JSON string	The time when the object was last modified, in ISO-8601 point-in-time format, as described in 5.6. The modification of a child is not considered a modification of a container object (modification times do not propagate up the tree).  For a newly created object, this value shall be set to the creation time.  This metadata value can only be updated by a client if it has the “backup_operator” privilege. If a client does not have the backup operator privilege, updates of this metadata item shall be ignored.	Optional
cdmi_acount	JSON string	The number of times that the object has been accessed since it was originally created. Accesses include all reads, writes, and lists.  For a newly created object, this value shall be set to the value “0”.  This metadata value can only be updated by a client if it has the “backup_operator” privilege. If a client does not have the backup operator privilege, updates of this metadata item shall be ignored.	Optional

Continued on next page

Table 140 – continued from previous page

Metadata name	Type	Description	Requirement
cdmi_mcount	JSON string	<p>The number of times that the object has been modified since it was originally created. Modifications include all value and metadata changes. Modifications to metadata resulting from reads (such as updates to <code>atime</code>) do not count as a modification.</p> <p>For a newly created object, this value shall be set to the value “0”.</p> <p>This metadata value can only be updated by a client if it has the “<code>backup_operator</code>” privilege. If a client does not have the backup operator privilege, updates of this metadata item shall be ignored.</p>	Optional
cdmi_hash	JSON string	<p>The hash of the value of the object, encoded using Base16 encoding rules described in RFC 4648 [19]. This metadata field shall be present when the “<code>cdmi_value_hash</code>” data system metadata for the object or a parent object indicates that the value of the object should be hashed.</p> <p>This storage system metadata item is computed by the storage system, and any attempts to set or modify it will be ignored.</p>	Optional
cdmi_owner	JSON string	<p>The name of the principal that has owner privileges for the object.</p> <p>If not specified when the object is created, this principal associated with the user creating the object shall be used.</p> <p>This metadata value can be updated by users with appropriate permissions.</p>	Optional
cdmi_acl	JSON array of JSON objects	<p>Standard ACL metadata as described in 17.1.</p> <p>If not specified when the object is created, the ACL metadata shall be generated in by the system.</p> <p>This metadata value can be updated by users with appropriate permissions.</p>	Optional
cdmi_dac_uri	JSON string	<p>Contains the URI used to submit a DAC request for the data object.</p> <p>URI schemes supported is defined in the <code>cdmi_dac_methods</code> capability.</p> <p>Both <code>cdmi_dac_certificate</code> and <code>cdmi_dac_uri</code> shall be included for delegated access control to be enabled for a given object.</p>	Optional
cdmi_dac_certificate	JSON object	<p>A JSON object, containing a JWE JWK which shall include a public key that is used to submit a DAC request for the data object, and should contains a X.509 certificate or certificate chain used to verify the identity of the DAC provider.</p> <p>Both <code>cdmi_dac_certificate</code> and <code>cdmi_dac_uri</code> shall be included for delegated access control to be enabled for a given object.</p>	Optional
cdmi_enc_signature	JSON object	<p>Contains JWS compact serialization of a signature for the entire object (value and metadata). See <a href="#">clause 23.7</a> for more details.</p>	Optional

## 16.3 Support for data system metadata

When specified, data system metadata, as described in Table 141, provides guidelines to the cloud storage system on how to provide storage data services for data managed through the CDMI interface.

Data system metadata is inherited from parent objects to any children objects. If a child object explicitly contains data system metadata, the metadata value of the child object data system metadata shall override any corresponding inherited metadata value of the parent object data system metadata.

Which data system metadata is supported by the CDMI server defined in 12.1.9. Data system metadata that is not supported by a CDMI server shall be preserved.

Table 141: Data system metadata

Metadata name	Type	Description	Requirement
cdmi_data_redundancy	JSON string	<p>If this data system metadata item is present and set to a positive numeric string, it indicates that the client is requesting a desired number of complete copies.</p> <p>Additional copies may be made to satisfy demand for the value. When this data system metadata item is absent, or is present and is not set to a positive numeric string, this data system metadata item shall not be used.</p>	Optional
cdmi_immediate_redundancy	JSON string	<p>If this data system metadata item is present and set to "true", it indicates that the client is requesting that at least the number of copies indicated in <code>cdmi_data_redundancy</code> contain the newly written value before the operation completes. This metadata is used to make sure that multiple copies of the data are written to permanent storage to prevent possible data loss. When this data system metadata item is absent, or is present and is not set to "true", this data system metadata item shall not be used.</p> <p>If the requested number of copies cannot be created within the HTTP timeout period, the transaction shall complete, but the <code>cdmi_immediate_redundancy_provided</code> data system metadata shall be set to "false".</p>	Optional
cdmi_assignedsize	JSON string	<p>If this data system metadata item is present and set to a positive numeric string, it indicates that the client is specifying the size in bytes that is desired to be reported for a container object exported via other protocols (see 9.1.3). The system is not required to reserve this space and may thin-provision the requested space. Thus, the requested value may be greater than the actual storage space consumed. When this data system metadata item is absent, or is present and is not set to a positive numeric string, this data system metadata item shall not be used.</p> <p>This data system metadata item is only applied against container objects and is not inherited by child objects.</p>	Optional

Continued on next page

Table 141 – continued from previous page

Metadata name	Type	Description	Requirement
cdmi_infrastructure_redundancy	JSON string	<p>If this data system metadata item is present and set to a positive numeric string, it indicates that the client is requesting a desired number of independent storage infrastructures supporting the multiple copies of data. This metadata is used to convey that, of the copies specified in <code>cdmi_data_redundancy</code>, these copies shall be stored on this many separate infrastructures.</p> <p>When this data system metadata item is absent, or is present and is not set to a positive numeric string, this data system metadata item shall not be used.</p>	Optional
cdmi_data_dispersion	JSON string	<p>If this data system metadata item is present and set to a positive numeric string, it indicates that the client is requesting a minimum desired distance (in km) between the infrastructures supporting the multiple copies of data. This metadata is used to separate the (<code>cdmi_infrastructure_redundancy</code> number of) infrastructures by a minimum geographic distance to prevent data loss due to site disasters.</p> <p>When this data system metadata item is absent, or is present and is not set to a positive numeric string, this data system metadata item shall not be used.</p>	Optional
cdmi_geographic_placement	JSON array of JSON strings	<p>If this data system metadata item is present and set to zero or more geopolitical identifiers, it indicates that the client is requesting restrictions on the geographic regions where the object is permitted to be stored. Each geopolitical identifier shall be in the form of either a string containing a valid ISO 3166 country/country-subdivision code, which indicates that storage is permitted within that geopolitical region, or in the form of a string starting with the “!” character in front of a valid ISO 3166 country/country-subdivision code, which excludes that country/country-subdivision from the previous list of geopolitical regions.</p> <p>The list is evaluated, in order, from left to right, with evaluation of each candidate storage location stopping when the candidate location is a permitted or prohibited region or is contained within a permitted or prohibited region. In addition to the ISO 3166 codes, “*” shall indicate all regions. If a candidate location does not match any of the entries in the list, the candidate location shall be considered to be prohibited.</p> <ul style="list-style-type: none"> <li>• When this data system metadata item is absent, this data system metadata item shall not be used.</li> <li>• When this data system metadata item is present and does not contain valid geopolitical identifiers, the create, update, or deserialize operation shall fail with an HTTP status code of 400 <code>Bad Request</code>.</li> <li>• When this data system metadata item is present and valid, but no available storage locations are permitted, the create, update, or deserialize operation shall fail with an HTTP status code of 403 <code>Forbidden</code>.</li> </ul>	Optional

Continued on next page



Table 141 – continued from previous page

Metadata name	Type	Description	Requirement
cdmi_retention_id	JSON string	<p>If this data system metadata item is present and not an empty string, it indicates that the client is requesting that the string be used to tag a given object as being managed by a specific retention policy. This data system metadata item is not required to place an object under retention, but is useful when needing to be able to perform a query to find all objects under a specific retention policy.</p> <p>When this data system metadata item is absent, or is present and an empty string, this data system metadata item shall not be used.</p>	Optional
cdmi_retention_period	JSON string	<p>If this data system metadata item is present and contains a valid ISO 8601:2004 time interval (as described in ), it indicates that the client is requesting that an object be placed under retention (see 18.3). When this data system metadata item is absent, this data system metadata item shall not be used. When this data system metadata item is present but does not contain a valid ISO 8601:2004 time interval, the create, update, or deserialize operation shall fail with an HTTP status code of 400 <i>Bad Request</i>.</p> <p>If this data system metadata item is updated and the new end date is before the current end date, the update operation shall fail with an HTTP status code of 403 <i>Forbidden</i>.</p>	Optional
cdmi_retention_autodelete	JSON string	<p>If this data system metadata item is present and set to “true”, it indicates that the client is requesting that an object under retention be automatically deleted when retention expires.</p> <p>When this data system metadata item is absent, or is present and is not set to “true”, this data system metadata item shall not be used.</p>	Optional
cdmi_hold_id	JSON array of JSON strings	<p>If this data system metadata item is present and not an empty array, it indicates that the client is requesting that an object be placed under hold (see 18.4). Each string in the array shall contain a unique user-specified hold identifier.</p> <p>When this data system metadata item is absent, or is present and is an empty JSON array, this data system metadata item shall not be used.</p> <p>If this data system metadata item is updated, and a previously existing hold string has been removed or changed in the update, the update operation shall fail with an HTTP status code of 403 <i>Forbidden</i>. (See 18.4 concerning releasing holds.)</p>	Optional

Continued on next page

Table 141 – continued from previous page

Metadata name	Type	Description	Requirement
cdmi_encryption	JSON string	<p>If this data system metadata item is present and not an empty string, it indicates that the client is requesting that the object be encrypted while at rest. If encrypted, all data and metadata related to the object shall be encrypted. Supported algorithm/mode/length values are provided by the <code>cdmi_encryption</code> capability.</p> <p>When this data system metadata item is absent, this data system metadata item shall not be used.</p> <p>If this data system metadata item is present but does not contain a valid encryption algorithm/mode/length string, the system is free to choose to ignore the data system metadata, to fail with an HTTP status code of 400 Bad Request, or to select an encryption algorithm/mode/length of the system's choice.</p> <p>Supported encryption algorithms are expressed as a string in the form of <code>ALGORITHM_MODE_KEYLENGTH</code>, where:</p> <ul style="list-style-type: none"> <li>• “ALGORITHM” is the encryption algorithm (e.g., “AES” or “3DES”).</li> <li>• “MODE” is the mode of operation (e.g., “XTS”, “CBC”, or “CTR”).</li> <li>• “KEYLENGTH” is the key size in bytes (e.g., “128”, “192”, “256”).</li> </ul> <p>To improve interoperability between CDMI implementations, the following designators should be used for the more common encryption combinations:</p> <ul style="list-style-type: none"> <li>• “3DES_ECB_168” for the three-key TripleDES algorithm, the Electronic Code Book (ECB) mode of operation, and a key size of 168 bits;</li> <li>• “3DES_CBC_168” for the three-key TripleDES algorithm, the Cipher Block Chaining (CBC) mode of operation, and a key size of 168 bits;</li> <li>• “AES_CBC_128” for the AES algorithm, the CBC mode of operation, and a key size of 128 bits;</li> <li>• “AES_CBC_256” for the AES algorithm, the CBC mode of operation, and a key size of 256 bits;</li> <li>• “AES_XTS_128” for the AES algorithm, the XTS mode of operation, and a key size of 128 bits; and</li> <li>• “AES_XTS_256” for the AES algorithm, the XTS mode of operation, and a key size of 256 bits.</li> </ul>	Optional

Continued on next page

Table 141 – continued from previous page

Metadata name	Type	Description	Requirement
cdmi_value_hash	JSON string	<p>If this data system metadata item is present and not an empty string, it indicates that the client is requesting that the system hash the object value using the hashing algorithm and length requested. The result of the hash shall be provided in the <code>cdmi_hash</code> storage system metadata item.</p> <p>Supported algorithm/length values are provided by the <code>cdmi_value_hash</code> storage system capability.</p> <p>When this data system metadata item is absent, this data system metadata item shall not be used.</p> <p>If this data system metadata item is present but does not contain a valid hash algorithm/length string, the system is free to choose to ignore the data system metadata, to fail with an HTTP status code of 400 <code>Bad Request</code>, or to select a hash algorithm/length of the system's choice.</p> <p>Supported hash algorithms are expressed as a string in the form of ALGORITHM LENGTH, where:</p> <ul style="list-style-type: none"> <li>• "ALGORITHM" is the hash algorithm (e.g., "SHA").</li> <li>• "LENGTH" is the hash size in bytes (e.g., "160", "256").</li> </ul> <p>To improve interoperability between CDMI implementations, the following designators should be used for the more common encryption combinations:</p> <ul style="list-style-type: none"> <li>• "SHA160" for SHA-1, and</li> <li>• "SHA256" for SHA-2.</li> </ul>	Optional
cdmi_latency	JSON string	<p>If this data system metadata item is present and set to a positive numeric string, it indicates that the client is requesting a desired maximum time to first byte, in milliseconds. This metadata is the desired latency (in milliseconds) to the first byte of data, as measured from the edge of the cloud and factoring out any propagation latency between the client and the cloud. For example, this metadata may be used to determine, in an interoperable way, from what type of storage medium the data may be served.</p> <p>When this data system metadata item is absent, or is present and is not set to a positive numeric string, this data system metadata item shall not be used.</p>	Optional
cdmi_throughput	JSON string	<p>If this data system metadata item is present and set to a positive numeric string, it indicates that the client is requesting a desired maximum data rate on retrieve, in bytes per second. This metadata is the desired bandwidth to the data, as measured from the edge of the cloud and factoring out any bandwidth capability between the client and the cloud. This metadata is used to stage the data in locations where there is sufficient bandwidth to accommodate a maximum usage.</p> <p>When this data system metadata item is absent, or is present and is not set to a positive numeric string, this data system metadata item shall not be used.</p>	Optional

Continued on next page

Table 141 – continued from previous page

Metadata name	Type	Description	Requirement
cdmi_sanitization_method	JSON string	<p>If this data system metadata item is present and not an empty string, it indicates that the client is requesting that the system use a specific sanitization method to delete data such that the data is unrecoverable after an update or delete operation. Supported sanitization method values are provided by the <code>cdmi_sanitization_method</code> capability.</p> <p>When this data system metadata item is absent, this data system metadata item shall not be used.</p> <p>If this data system metadata item is present but does not contain a valid sanitization method string, the system is free to choose to ignore the data system metadata, to fail with an HTTP status code of 400 <i>Bad Request</i>, or to select a sanitization method of the system's choice.</p> <p>Supported sanitization methods are defined as system-specific strings.</p>	Optional
cdmi_RPO	JSON string	<p>If this data system metadata item is present and set to a positive numeric string, it indicates that the client is requesting a largest acceptable duration in time between an update or create and when the object may be recovered, specified in seconds. This metadata is used to indicate the desired backup frequency from the primary copy or copies of the data to the secondary copy or copies. It is the maximum acceptable time period before a failure or disaster during which changes to data may be lost as a consequence of recovery.</p> <p>When this data system metadata item is absent, or is present and is not set to a positive numeric string, this data system metadata item shall not be used.</p>	Optional
cdmi_RTO	JSON string	<p>If this data system metadata item is present and set to a positive numeric string, it indicates that the client is requesting the largest acceptable duration in time to restore data, specified in seconds. This metadata is used to indicate the desired maximum acceptable duration to restore the primary copy or copies of the data from a secondary backup copy or copies.</p> <p>When this data system metadata item is absent, or is present and is not set to a positive numeric string, this data system metadata item shall not be used.</p>	Optional
cdmi_enc_key_id	JSON string	<p>If this data system metadata item is present and not an empty string, it indicates that the client is requesting that the system associate with the object a key identifier (e.g. KMIP Identifier) for the symmetric key used to encrypt and decrypt the object.</p>	Optional
cdmi_enc_value_sign_id	JSON string	<p>If this data system metadata item is present and not an empty string, it indicates that the client is requesting that the system associate with the object a key identifier (e.g. KMIP Identifier) for the private key used for signing the value of the object.</p>	Optional
cdmi_enc_value_verify_id	JSON string	<p>If this data system metadata item is present and not an empty string, it indicates that the client is requesting that the system associate with the object a key identifier (e.g. KMIP Identifier) for the public key or certificate chain used for verifying the signature of the value of the object.</p>	Optional

Continued on next page

Table 141 – continued from previous page

Metadata name	Type	Description	Requirement
cdmi_enc_object_sign_id	JSON string	If this data system metadata item is present and not an empty string, it indicates that the client is requesting that the system associate with the object a key identifier (e.g. KMIP Identifier) for the private key used for signing the entire object.	Optional
cdmi_enc_object_verify_id	JSON string	If this data system metadata item is present and not an empty string, it indicates that the client is requesting that the system associate with the object a key identifier (e.g. KMIP Identifier) for the public key or certificate chain used for verifying the signature of the entire object.	Optional
cdmi_versioning	JSON string	<p>If this data system metadata item is present and not an empty string, it indicates that the client is requesting that versioning be enabled for the data object, and what level of versioning is requested.</p> <ul style="list-style-type: none"> <li>• If set to the value “value”, versions shall be created when the value is updated.</li> <li>• If set to the value “user”, versions shall be created when the value and/or user metadata is updated.</li> <li>• If set to the value “all”, versions shall be created when any update is performed against the version-enabled data object.</li> </ul> <p>This data system metadata item shall not be present in data object versions.</p>	Optional
cdmi_versions_count	JSON string	<p>If this data system metadata item is present and not an empty string, it indicates that the client is requesting limits on the maximum number of historical versions to be retained.</p> <ul style="list-style-type: none"> <li>• If <code>cdmi_versions_count</code> is not present, no limits should be placed on the number of versions that are retained.</li> <li>• If <code>cdmi_versions_count</code> is present and has a value of zero, only the current version should be retained.</li> <li>• If <code>cdmi_versions_count</code> is present and has a value greater than zero, up to the specified number of historical versions should be retained.</li> <li>• If the number of historical versions exceeds the value specified, historical versions should be deleted from the oldest to the newest until the number of historical versions equals the value contained in <code>cdmi_versions_count</code>.</li> </ul>	Optional
cdmi_versions_age	JSON string	<p>If this data system metadata item is present and not an empty string, it indicates that the client is requesting limits on the maximum age of the oldest historical version requested to be retained.</p> <ul style="list-style-type: none"> <li>• If <code>cdmi_versions_age</code> is not present, no limit should be placed on the age of versions that are retained.</li> <li>• If <code>cdmi_versions_age</code> is present, historical versions should be retained until their age in seconds since creation is greater than the value contained in <code>cdmi_versions_age</code>.</li> <li>• If the age of a historical version exceeds the value specified, that historical version should be deleted.</li> </ul>	Optional

Continued on next page

Table 141 – continued from previous page

Metadata name	Type	Description	Requirement
cdmi_versions_size	JSON string	<p>If this data system metadata item is present and not an empty string, it indicates that the client is requesting limits on the maximum amount of space to be used to retain historical versions.</p> <ul style="list-style-type: none"> <li>• If <code>cdmi_versions_size</code> is not present, no limit should be placed on the size of versions that are retained.</li> <li>• If <code>cdmi_versions_size</code> is present, historic versions should be retained until the total storage consumption in bytes of the historical versions exceeds the value contained in <code>cdmi_versions_size</code>.</li> <li>• If the total size consumed by historical versions exceeds the value specified, historical versions should be deleted from the oldest to the newest until the total storage consumption of historical versions is equal or less than the value contained in <code>cdmi_versions_size</code>.</li> </ul>	Optional

## 16.4 Support for provided data system metadata

For each metadata item in a data system, there is an actual value that the cloud service is able to achieve at this time, as shown in Table 142. Data system-provided metadata items are read only. Updates of these metadata items shall be ignored.

Table 142: Provided values of data system metadata

Metadata name	Type	Description	Requirement
cdmi_data_redundancy_provided	JSON string	Contains the current number of complete copies of the data object at this time	Optional
cdmi_immediate_redundancy_provided	JSON string	If present and set to “true”, indicates if immediate redundancy is provided for the object	Optional
cdmi_infrastructure_redundancy_provided	JSON string	Contains the current number of independent storage infrastructures supporting the data currently operating	Optional
cdmi_data_dispersion_provided	JSON string	Contains the current lowest distance (km) between any two infrastructures hosting the data	Optional
cdmi_geographic_placement_provided	JSON array of JSON strings	Contains an ISO-3166 identifier that corresponds to a geopolitical region where the object is stored	Optional
cdmi_retention_period_provided	JSON string	Contains an ISO-8601 time interval (as described in 5.6) specifying the period the object is protected by retention	Optional
cdmi_retention_autodelete_provided	JSON string	Contains “true” if the object will automatically be deleted when retention expires	Optional
cdmi_hold_id_provided	JSON array of JSON strings	Contains the user-specified hold identifiers for active holds	Optional
cdmi_encryption_provided	JSON string	Contains the algorithm used for encryption, the mode of operation, and the key size. (See <code>cdmi_encryption</code> in 16.3 for the format.)	Optional
cdmi_value_hash_provided	JSON string	Contains the algorithm and length being used to hash the object value. See <code>cdmi_value_hash</code> in 16.3 for the format.	Optional
cdmi_latency_provided	JSON string	Contains the provided maximum time to first byte	Optional
cdmi_throughput_provided	JSON string	Contains the provided maximum data rate on retrieve	Optional
cdmi_sanitization_method_provided	JSON string	Contains the sanitization method used. See <code>cdmi_sanitization_method</code> in 16.3 for the format.	Optional
cdmi_RPO_provided	JSON string	Contains the provided duration, in seconds, between an update and when the update may be recovered	Optional
cdmi_RTO_provided	JSON string	Contains the provided duration, in seconds, to restore data	Optional

Continued on next page

Table 142 – continued from previous page

Metadata name	Type	Description	Requirement
cdmi_authentication_methods_provided	JSON array of JSON strings	Contains a list of authentication methods enabled for the domain. See <a href="#">cdmi_authentication_methods</a> in 16.3 for the format.	Optional
cdmi_versioning_provided	JSON string	Contains the value “value”, “user”, or “all” if versioning is enabled for the data object.	Optional
cdmi_versions_count_provided	JSON string	Contains the maximum number of historical versions that will be retained.	Optional
cdmi_versions_age_provided	JSON string	Contains the oldest age of a historical version that will be retained, in seconds before the current time.	Optional
cdmi_versions_size_provided	JSON string	Contains the maximum amount of space that can be used to retain historical versions, in bytes.	Optional



## 16.5 Support for user metadata

All CDMI objects that support metadata shall permit the inclusion of arbitrary user-defined metadata items, with the restriction that the name of a user-defined metadata item shall not start with the prefix "cdmi\\_".

- The maximum number of user-defined metadata items is specified by the capability `cdmi_metadata_maxitems`.
- The maximum size of each user-defined metadata item is specified by the capability `cdmi_metadata_maxsize`.
- The maximum total size of user-defined metadata items for an object is specified by the capability `cdmi_metadata_maxtotalsize`.

## 16.6 Metadata update operations

CDMI permits a client to replace all metadata items or to perform operations against one or more individual metadata items.

Replacing all metadata items is accomplished by including the metadata field in the update request body JSON and not specifying specific metadata items in the update URI.

Adding, updating, and removing specific metadata items is accomplished by specifying the specific metadata item names in the update URI:

- To add a new metadata item to an existing object, the metadata item name shall be included in the update request URI, and the metadata item shall be included in the metadata field in the update request body JSON.
- To update the value of an existing metadata item, the metadata item name shall be included in the update request URI, and the metadata item shall be included in the metadata field in the update request body JSON.
- To remove an existing metadata item, the metadata item name shall be included in the update request URI, and the metadata item shall not be included in the metadata field in the update request body JSON.

When individual metadata items are specified in the update URI, metadata items included in the metadata field in the request body JSON that are not referred to in the update URI shall be ignored.

## Clause 17

# Access control

### 17.1 Overview

Access control defines the mechanisms by which access to objects are permitted or denied. The CDMI International Standard supports the following options for access control:

- No access control
- Access Control List (ACL) based access control (See 17.2)
- Domain based access control (See 10.1.5)
- Delegated access control (See clause 24)
- Vendor-defined access control extensions
- Combinations of the above

#### 17.1.1 Access control flow

Fig. 10 (see next page) illustrates the control flow for access control in an example CDMI implementation. As every aspect of access control is optional within a CDMI server, each different implementations will typically implement appropriate subsets of the below illustrated access control flow, in a manner appropriate to the internal architecture of their implementation.

1. The CDMI client initiates a CDMI operation by sending a CDMI request to a CDMI server. As part of the request, the CDMI client includes information about its identity and information to prove this identity (credentials). The method by which these credentials are presented and formatted is not specified in this International Standard, however, some guidance is provided in 5.4.3.
2. If the CDMI server supports Domains (see clause 10), the CDMI server obtains the domain associated with the object the CDMI operation is being performed against. If the CDMI system does not support domains, steps 2 - 8 are skipped.
3. The CDMI server obtains required information about the domain associated with the object.
4. Domain Information is returned for further use.
5. Domain information is used to resolve CDMI client credentials.
6. If the Domain is configured to delegate identity resolution to an external system (such as Active Directory), credentials are sent to this external system for resolution.
7. If the Domain is configured to use local membership, credentials are compared against the configured domain members (see 10.3).
8. The resolved principle (user, group, indication of validity) is returned for further use.
9. If the CDMI server supports ACLs (see 17.2), the CDMI server evaluates the object ACL. If the CDMI system does not support ACLs, steps 9 - 15 are skipped.
10. The CDMI ACL processing subsystem obtains the ACL for the object.

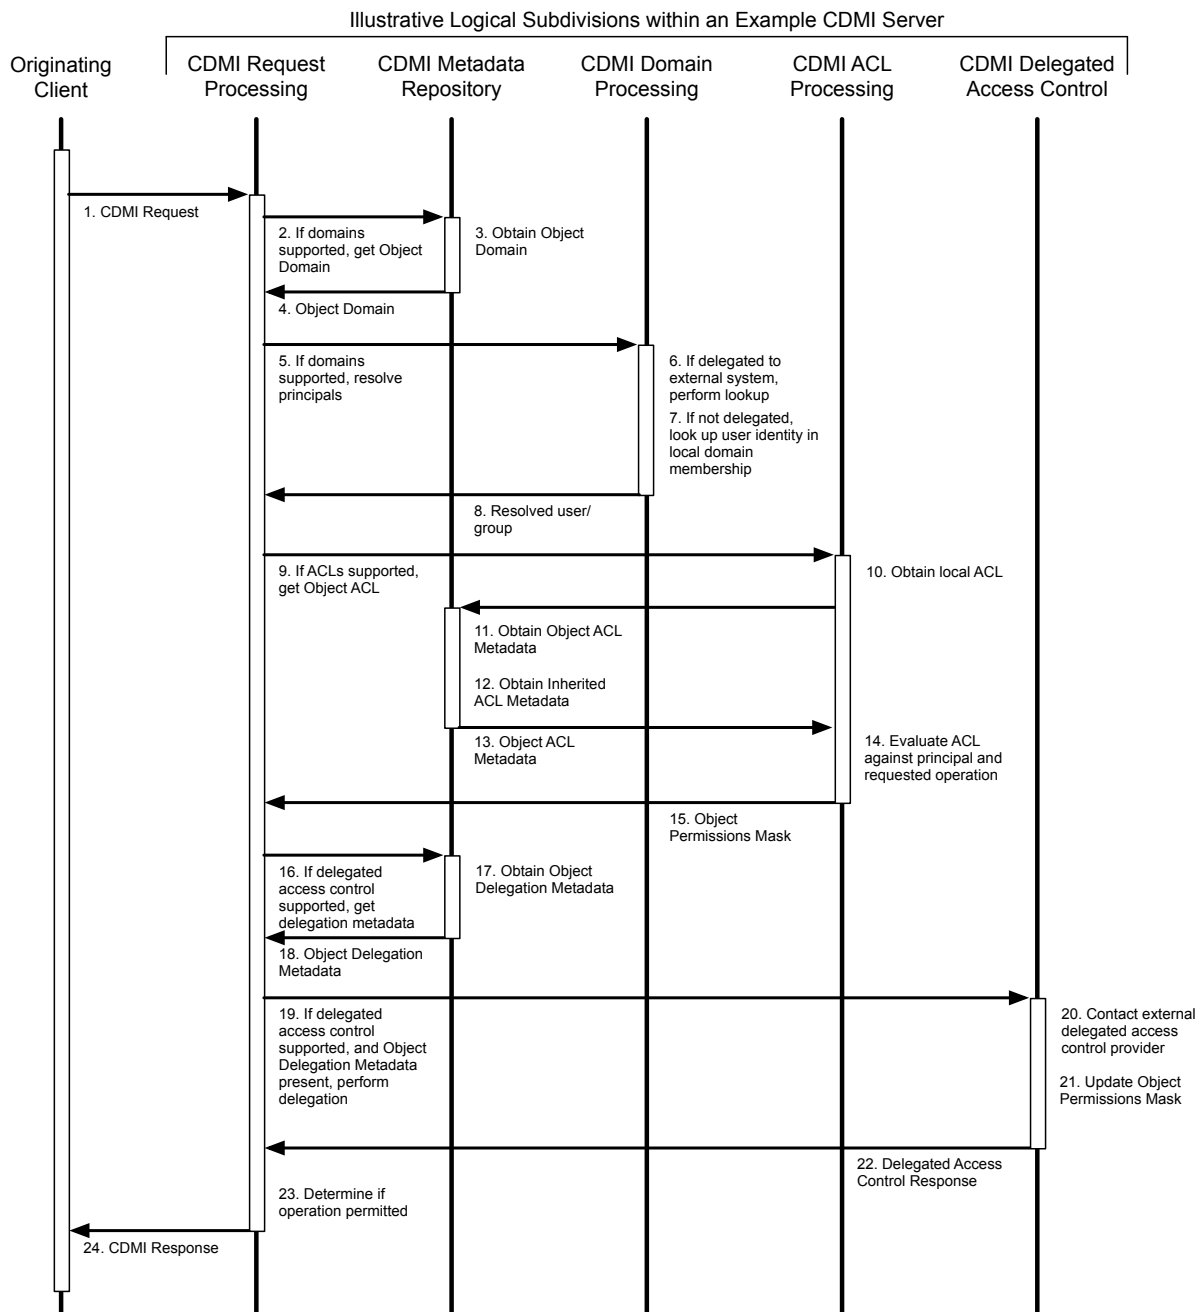


Figure 10: Access control flow

- 3165 11. The CDMI server obtains ACL metadata associated with the object.
- 3166 12. If the object is in a container, the CDMI server obtains ACL metadata for parent containers.
- 3167 13. The obtained ACL metadata is returned for further use.
- 3168 14. The CDMI ACL processing subsystem evaluates the resolved principals against the resolved ACL.
- 3169 15. The evaluated permission mask is returned for further use.
- 3170 16. If the CDMI server supports Delegated Access Control (DAC) (see clause 24), the CDMI server obtains DAC  
3171 metadata associated with the object the CDMI operation is being performed against. If the CDMI system does  
3172 not support DAC, steps 16 - 22 are skipped.
- 3173 17. The CDMI server obtains DAC metadata associated with the object.
- 3174 18. DAC metadata is returned for further use.
- 3175 19. If DAC metadata is present and indicates that DAC is to be used, the specified delegation is performed.
- 3176 20. The external DAC provider is contacted, including the evaluated Object permission mask.
- 3177 21. If a valid DAC response is received, the `dac_applied_mask` replaces the evaluated Object permission mask.
- 3178 22. The DAC results and Object permission mask is returned for further use.
- 3179 23. The Object permission mask is used to determine if the requested operation is permitted.
- 3180 24. The operation is permitted or denied, and the corresponding response returned to the CDMI Client.
- 3181 Steps 2 - 8, 9 - 15, and 16 - 18 may be performed in parallel.

## 17.2 Access control lists

CDMI uses the well-known mechanism of an Access Control List (ACL) as defined in the NFSv4 standard (see RFC 3530 [1]). ACLs are lists of permissions-granting or permissions-denying entries called Access Control Entries (ACEs).

### 17.2.1 ACL and ACE structure

An ACL is an ordered list of ACEs. The two types of ACEs in CDMI are `ALLOW` and `DENY`. An `ALLOW` ACE grants some form of access to a principal. Principals are either users or groups and are represented by identifiers. A `DENY` ACE denies access of some kind to a principal. For instance, a `DENY` ACE may deny the ability to write the metadata or ACL of an object but may remain silent on other forms of access. In that case, if another ACE `ALLOWs` write access to the object, the principal is allowed to write the object's data, but nothing else.

ACEs are composed of four fields: `type`, `who`, `flags` and `access_mask`, as per RFC 3530 [1]. The `type`, `flags`, and `access_mask` shall be specified as either unsigned integers in hex string representation or as a comma-delimited list of bit mask string form values taken from *ACE types*, *ACE flags*, and *ACE masks bits*.

### 17.2.2 ACE types

Table 143 defines the following ACE types, as specified in section 5.11.1 of RFC 3530 [1].

Table 143: ACE types

String form	Description	Constant	Bit mask
"ALLOW"	Allow access rights for a principal	"CDMI_ACE_ACCESS_ALLOW"	0x00000000
"DENY"	Deny access rights for a principal	"CDMI_ACE_ACCESS_DENY"	0x00000001
"AUDIT"	Generate an audit record when the principal attempts to exercise the specified access rights	"CDMI_ACE_SYSTEM_AUDIT"	0x00000002

The reason that the string forms may be safely abbreviated is that they are local to the ACE structure type, as opposed to constants, which are relatively global in scope.

The client is responsible for ordering the ACEs in an ACL. The server shall not enforce any ordering and shall store and evaluate the ACEs in the order given by the client.

### 17.2.3 ACE who

The special "who" identifiers need to be understood universally, rather than in the context of a particular external security domain (see *Who identifiers*). Some of these identifiers may not be understood when a CDMI client accesses the server, but they may have meaning when a local process accesses the file. The ability to display and modify these permissions is permitted over CDMI, even if none of the access methods on the server understands the identifiers.

Table 144: Who identifiers

Who	Description
"OWNER@"	The owner of the file
"GROUP@"	The group associated with the file
"EVERYONE@"	The world
"ANONYMOUS@"	Access without authentication
"AUTHENTICATED@"	Any authenticated user (opposite of "ANONYMOUS@")
"ADMINISTRATOR@"	A user with administrative status, e.g., "root"
"ADMINUSERS@"	A group whose members are given administrative status

To avoid name conflicts, these special identifiers are distinguished by an appended "@" (with no domain name).

## 17.2.4 ACE flags

CDMI allows for nested containers and mandates that objects and subcontainers be able to inherit access permissions from their parent containers. However, it is not enough to simply inherit all permissions from the parent; it might be desirable, for example, to have different default permissions on child objects and subcontainers of a given container. The flags in Table 145 govern this behavior.

Table 145: ACE flags

String form	Description	Constant	Bit mask
"NO_FLAGS"	No flags are set	"CDMI_ACE_FLAGS_NONE"	0x00000000
"OBJECT_INHERIT"	An ACE on which "OBJECT_INHERIT" is set is inherited by objects as an effective ACE: "OBJECT_INHERIT" is cleared on the child object. When the ACE is inherited by a container, "OBJECT_INHERIT" is retained for the purpose of inheritance, and additionally, "INHERIT_ONLY" is set.	"CDMI_ACE_FLAGS_OBJECT_INHERIT_ACE"	0x00000001
"CONTAINER_INHERIT"	An ACE on which "CONTAINER_INHERIT" is set is inherited by a subcontainer as an effective ACE. Both "INHERIT_ONLY" and "CONTAINER_INHERIT" are cleared on the child container.	"CDMI_ACE_FLAGS_CONTAINER_INHERIT_ACE"	0x00000002
"NO_PROPAGATE"	An ACE on which "NO_PROPAGATE" is set is not inherited by any objects or subcontainers. It applies only to the container on which it is set.	"CDMI_ACE_FLAGS_NO_PROPAGATE_ACE"	0x00000004
"INHERIT_ONLY"	An ACE on which "INHERIT_ONLY" is set is propagated to children during ACL inheritance as specified by "OBJECT_INHERIT" and "CONTAINER_INHERIT". The ACE is ignored when evaluating access to the container on which it is set and is always ignored when set on objects.	"CDMI_ACE_FLAGS_INHERIT_ONLY_ACE"	0x00000008
"IDENTIFIER_GROUP"	An ACE on which "IDENTIFIER_GROUP" is set indicates that the "who" refers to a group identifier.	"CDMI_ACE_FLAGS_IDENTIFIER_GROUP"	0x00000040
"INHERITED"	An ACE on which "INHERITED" is set indicates that this ACE is inherited from a parent directory. A server that supports automatic inheritance will place this flag on any ACEs inherited from the parent directory when creating a new object.	"CDMI_ACE_FLAGS_INHERITED_ACE"	0x00000080

## 17.2.5 ACE mask bits

The mask field of an ACE contains a 32 bit mask, as specified in section 5.11.2 of RFC 3530 [1].

Table 146: ACE masks bits

String form	Description	Constant	Bit mask
"READ_OBJECT"	<p>If true, indicates permission to read the value of an object.</p> <p>If false:</p> <ul style="list-style-type: none"> <li>• A CDMI GET that requests all fields shall return all permitted fields with the value field excluded.</li> <li>• A CDMI GET that requests specific fields shall return requested permitted fields with the value field excluded.</li> <li>• A CDMI GET for only the value field shall return an HTTP status code of 403 Forbidden.</li> <li>• A non-CDMI GET shall return an HTTP status code of 403 Forbidden.</li> </ul>	"CDMI_ACE_READ_OBJECT"	0x00000001
"LIST_CONTAINER"	<p>If true, indicates permission to list the children of an object.</p> <p>If false:</p> <ul style="list-style-type: none"> <li>• A CDMI GET that requests all fields shall return all permitted fields with the children field and childrenrange field excluded.</li> <li>• A CDMI GET that requests specific fields shall return the requested permitted fields with the children field and childrenrange field excluded.</li> <li>• A CDMI GET for only the children field and/or childrenrange field shall return an HTTP status code of 403 Forbidden.</li> </ul>	"CDMI_ACE_LIST_CONTAINER"	0x00000001
"WRITE_OBJECT"	<p>If true, indicates permission to modify the value of an object</p> <p>If false, a PUT that requests modification of the value of an object shall return an HTTP status code of 403 Forbidden.</p>	"CDMI_ACE_WRITE_OBJECT"	0x00000002
"ADD_OBJECT"	<p>If true, indicates permission to add a new child data object or queue object.</p> <p>If false, a PUT or POST that requests creation of a new child data object or new queue object shall return an HTTP status code of 403 Forbidden.</p>	"CDMI_ACE_ADD_OBJECT"	0x00000002

Continued on next page



Table 146 – continued from previous page

String form	Description	Constant	Bit mask
"APPEND_DATA"	<p>If true, indicates permission to append data to the value of a data object.</p> <p>If "APPEND_DATA" is true and "WRITE_OBJECT" is false, a PUT that requests modification of any existing part of the value of an object shall return an HTTP status code of 403 Forbidden.</p>	"CDMI_ACE_APPEND_DATA"	0x00000004
"ADD_SUBCONTAINER"	<p>If true, indicates permission to create a child container object or domain object.</p> <p>If false, a PUT that requests creation of a new child container object or new domain object shall return an HTTP status code of 403 Forbidden.</p>	"CDMI_ACE_ADD_SUBCONTAINER"	0x00000004
"READ_METADATA"	<p>If true, indicates permission to read the metadata of an object.</p> <p>If false:</p> <ul style="list-style-type: none"> <li>• A CDMI GET that requests all fields shall return all permitted fields with the metadata field excluded.</li> <li>• A CDMI GET that requests specific fields shall return the requested permitted fields with the metadata field excluded.</li> <li>• A CDMI GET for only the metadata field shall return an HTTP status code of 403 Forbidden.</li> </ul>	"CDMI_ACE_READ_METADATA"	0x00000008
"WRITE_METADATA"	<p>If true, indicates permission to modify the metadata of an object.</p> <p>If false, a CDMI PUT that requests modification of the metadata field of an object shall return an HTTP status code of 403 Forbidden.</p>	"CDMI_ACE_WRITE_METADATA"	0x00000010
"EXECUTE"	If true, indicates permission to execute an object.	"CDMI_ACE_EXECUTE"	0x00000020
"TRAVERSE_CONTAINER"	<p>If true, indicates permission to traverse a container object or domain object.</p> <p>If false, all operations against all children below the container shall return an HTTP status code of 403 Forbidden.</p>	"CDMI_ACE_TRAVERSE_CONTAINER"	0x00000020
"DELETE_OBJECT"	<p>If true, indicates permission to delete a child data object or child queue object from a container object.</p> <p>If false, all DELETE operations shall return an HTTP status code of 403 Forbidden.</p>	"CDMI_ACE_DELETE_OBJECT"	0x00000040

Continued on next page

Table 146 – continued from previous page

String form	Description	Constant	Bit mask
"DELETE_SUBCONTAINER"	<p>If true, indicates permission to delete a child container object from a container object or to delete a child domain object from a domain object.</p> <p>If false, all DELETE operations shall return an HTTP status code of 403 Forbidden.</p>	"CDMI_ACE_DELETE_SUBCONTAINER"	0x00000040
"READ_ATTRIBUTES"	<p>If true, indicates permission to read the attribute fields<sup>1</sup> of an object.</p> <p>If false:</p> <ul style="list-style-type: none"> <li>• A CDMI GET that requests all fields shall return all non-attribute fields and shall not return any attribute fields.</li> <li>• A CDMI GET that requests at least one non-attribute field shall only return the requested non-attribute fields.</li> <li>• A CDMI GET that requests only non-attribute fields shall return an HTTP status code of 403 Forbidden.</li> </ul>	"CDMI_ACE_READ_ATTRIBUTES"	0x00000080
"WRITE_ATTRIBUTES"	<p>If true, indicates permission to change attribute fields[#a]_ of an object.</p> <p>If false, a CDMI PUT that requests modification of any non-attribute field shall return an HTTP status code of 403 Forbidden.</p>	"CDMI_ACE_WRITE_ATTRIBUTES"	0x00000100
"WRITE_RETENTION"	<p>If true, indicates permission to change retention attributes of an object.</p> <p>If false, a CDMI PUT that requests modification of any non-hold retention metadata items shall return an HTTP status code of 403 Forbidden.</p>	"CDMI_ACE_WRITE_RETENTION"	0x00000200
"WRITE_RETENTION_HOLD"	<p>If true, indicates permission to change retention hold attributes of an object.</p> <p>If false, a CDMI PUT that requests modification of any retention hold metadata items shall return an HTTP status code of 403 Forbidden.</p>	"CDMI_ACE_WRITE_RETENTION_HOLD"	0x00000400
"DELETE"	<p>If true, indicates permission to delete an object.</p> <p>If false, all DELETE operations shall return an HTTP status code of 403 Forbidden.</p>	"CDMI_ACE_DELETE"	0x00010000

Continued on next page

Table 146 – continued from previous page

String form	Description	Constant	Bit mask
"READ_ACL"	<p>If true, indicates permission to read the ACL of an object.</p> <p>If false:</p> <ul style="list-style-type: none"> <li>A CDMI GET that requests all metadata items shall return all permitted metadata items with the "cdmi_acl" metadata item excluded.</li> <li>A CDMI GET that requests specific metadata items shall return the requested permitted metadata items with the "cdmi_acl" metadata item excluded.</li> <li>A CDMI GET for only the cdmi_acl metadata item shall return an HTTP status code of 403 Forbidden.</li> </ul> <p>If "READ_ACL" is true and "READ_METADATA" is false, then to read the ACL, a client CDMI GET for only the "cdmi_acl" metadata item shall be permitted.</p>	"CDMI_ACE_READ_ACL"	0x00020000
"WRITE_ACL"	<p>If true, indicates permission to write the ACL of an object.</p> <p>If false:</p> <ul style="list-style-type: none"> <li>If "WRITE_ACL" is false, a CDMI PUT that requests modification of the "cdmi_acl" metadata item shall return an HTTP status code of 403 Forbidden.</li> <li>If "WRITE_ACL" is true and "WRITE_METADATA" is false, then to write the ACL, a client CDMI PUT for only the "cdmi_acl" metadata item shall be permitted.</li> </ul>	"CDMI_ACE_WRITE_ACL"	0x00040000
"WRITE_OWNER"	<p>If true, indicates permission to change the owner of an object.</p> <p>If false:</p> <ul style="list-style-type: none"> <li>If "WRITE_OWNER" is false, a CDMI PUT that requests modification of the "cdmi_owner" metadata item shall return an HTTP status code of 403 Forbidden.</li> <li>If "WRITE_OWNER" is true and "WRITE_METADATA" is false, then to write the owner, a client CDMI PUT for only the "cdmi_owner" metadata item shall be permitted.</li> </ul>	"CDMI_ACE_WRITE_OWNER"	0x00080000

Continued on next page

Table 146 – continued from previous page

String form	Description	Constant	Bit mask
"SYNCHRONIZE"	If true, indicates permission to access an object locally at the server with synchronous reads and writes.	"CDMI_ACE_SYNCHRONIZE"	0x00100000

Implementations shall use the correct string form to display permissions, if the object type is known. If the object type is unknown, the "object" version of the string shall be used.

## 17.2.6 ACL evaluation

When evaluating whether access to a particular object O by a principal P is to be granted, the server shall traverse the object's logical ACL (its ACL after processing inheritance from parent containers) in list order, using a temporary permissions bitmask m, initially empty (all zeroes).

- If the object still does not contain an ACL, the algorithm terminates and access is denied for all users and groups. This condition is not expected, as CDMI implementations should require an inheritable default ACL on all root containers.
- ACEs that do not refer to the principal P requesting the operation are ignored.
- If an ACE is encountered that denies access to P for any of the requested mask bits, access is denied and the algorithm terminates.
- If an ACE is encountered that allows access to P, the permissions mask m for the operation is XORed with the permissions mask from the ACE. If m is sufficient for the operation, access is granted and the algorithm terminates.
- If the end of the ACL list is reached and permission has neither been granted nor explicitly denied, access is denied and the algorithm terminates, unless the object is a container root. In this case, the server shall:
  - allow access to the container owner, "ADMINISTRATOR@", and any member of "ADMINUSERS@"; and
  - log an event indicating what has happened.

When permission for the desired access is not explicitly given, even "ADMINISTRATOR@" and equivalents are denied for objects that aren't container roots. When an admin needs to access an object in such an instance, the root container shall be accessed and its inheritable ACEs changed in a way as to allow access to the original object. The resulting log entry then provides an audit trail for the access.

When a root container is created and no ACL is supplied, the server shall place an ACL containing the following ACEs on the container:

```
"cdmi_acl":
[
  {
    "acetype": "ALLOW",
    "identifier": "OWNER@",
    "aceflags": "OBJECT_INHERIT, CONTAINER_INHERIT",
    "acemask": "ALL_PERMS"
  },
  {
    "acetype": "ALLOW",
    "identifier": "AUTHENTICATED@",
    "aceflags": "OBJECT_INHERIT, CONTAINER_INHERIT",
    "acemask": "READ"
  }
]
```

As ACLs are storage system metadata, they are stored and retrieved through the metadata field included in a PUT or GET request. The syntax is as follows, using the constant strings from [ACE types](#), [ACE flags](#), and [ACE masks bits](#):

```
ACL = { ACE [, ACE ...] }
ACE = { acetype , identifier , aceflags , acemask }
acetype = uint_t | acetypeitem
```

(continues on next page)

<sup>1</sup> The value fields, children fields, and metadata field are considered to be non-attribute fields. All other fields are considered to be attribute fields.

(continued from previous page)

```

identifier = utf8string_t
aceflags   = uint_t | aceflagsstring
acemask    = uint_t | acemaskstring

acetypeitem = aceallowedtype | acedeniedtype | aceaudittype
aceallowedtype = "CDMI_ACE_ACCESS_ALLOWED_TYPE" | 0x0
acedeniedtype  = "CDMI_ACE_ACCESS_DENIED_TYPE" | 0x01
aceaudittype   = "CDMI_ACE_SYSTEM_AUDIT_TYPE" | 0x02

aceflagsstring = aceflagsitem [| aceflagsitem ...]
aceflagsitem   = aceobinherititem | acecontinherititem | acenopropagateitem |
↪aceinheritonlyitem

aceobinherititem = "CDMI_ACE_OBJECT_INHERIT_ACE" | 0x01
acecontinherititem = "CDMI_ACE_CONTAINER_INHERIT_ACE" | 0x02
acenopropagateitem = "CDMI_ACE_NO_PROPAGATE_INHERIT_ACE" | 0x04
aceinheritonlyitem = "CDMI_ACE_INHERIT_ONLY_ACE" | 0x08

acemaskstring = acemaskitem [| acemaskitem ...]
acemaskitem   = acereaditem | acewriteitem | aceappenditem | acereadmetaitem |
↪acewritemetaitem | acedeleteitem | acedelselfitem | acereadaclitem | acewriteaclitem |
↪aceexecuteitem | acereadatritem | acewriteattritem | aceretentionitem

acereaditem      = "CDMI_ACE_READ_OBJECT" | "CDMI_ACE_LIST_CONTAINER" | 0x01
acewriteitem     = "CDMI_ACE_WRITE_OBJECT" | "CDMI_ACE_ADD_OBJECT" | 0x02
aceappenditem    = "CDMI_ACE_APPEND_DATA" | "CDMI_ACE_ADD_SUBCONTAINER" | 0x04
acereadmetaitem  = "CDMI_ACE_READ_METADATA" | 0x08
acewritemetaitem = "CDMI_ACE_WRITE_METADATA" | 0x10
acedeleteitem    = "CDMI_ACE_DELETE_OBJECT" | "CDMI_ACE_DELETE_SUBCONTAINER" | 0x40
acedelselfitem   = "CDMI_ACE_DELETE" | 0x10000
acereadaclitem   = "CDMI_ACE_READ_ACL" | 0x20000
acewriteaclitem  = "CDMI_ACE_WRITE_ACL" | 0x40000
aceexecuteitem   = "CDMI_ACE_EXECUTE" | 0x80000
acereadatritem   = "CDMI_ACE_READ_ATTRIBUTES" | 0x00080
acewriteattritem = "CDMI_ACE_WRITE_ATTRIBUTES" | 0x00100
aceretentionitem = "CDMI_ACE_SET_RETENTION" | 0x10000000

```

When ACE masks are presented in numeric format, they shall, at all times, be specified in hexadecimal notation with a leading "0x". This format allows both servers and clients to quickly determine which of the two forms of a given constant is being used. When masks are presented in string format, they shall be converted to numeric format and then evaluated using standard bitwise operators.

When an object is created, no ACL is supplied, and an ACL is not inherited from the parent container (or there is no parent container), the server shall place an ACL containing the following ACEs on the object:

```

"cdmi_acl":
[
  {
    "acetype": "ALLOW",
    "identifier": "OWNER@",
    "aceflags": "OBJECT_INHERIT, CONTAINER_INHERIT",
    "acemask": "ALL_PERMS"
  }
]

```

### 17.2.7 Example ACE mask expressions

Example 1:

```
"READ_ALL" | 0x02
```

evaluates to 0x09 | 0x02 == 0x0

Example 2:

```
0x001F07FF
```

evaluates to 0x001F07FF == "ALL\_PERMS"

Example 3:

"RW\_ALL" | DELETE

evaluates to 0x000601DF | 0x00100000 == 0x000701DF

## 17.2.8 Canonical format for ACE hexadecimal quantities

ACE mask expressions may be evaluated and converted to a string hexadecimal value before transmission in a CDMI JSON body. Applications or utilities that display them to users should convert them into a text expression before display and accept user input in text format as well.

The following technique should be used to decompose masks into strings. A table of masks and string equivalents should be maintained and ordered from greatest to least:

Table 147: ACE bit mask/string

Hex form	Object string form	Container string form
0x001F07FF	"ALL_PERMS"	"ALL_PERMS"
0x0006006F	"RW_ALL"	"RW_ALL"
0x0000001F	"RW"	"RW"
	...	...
0x00000002	"WRITE_OBJECT"	"ADD_OBJECT"
0x00000001	"READ_OBJECT"	"LIST_CONTAINER"

Given an access mask M, the following is repeated until M == 0:

1. Select the highest mask m from the table such that M & m == m.
2. If the object is a container, select the string from the 3rd column; otherwise, select the string from the 2nd column.
3. Bitwise subtract m from M, i.e., set M = M xor m.
4. The complete textual representation is then all the selected strings concatenated with ", " between them, e.g., "``ALL\_PERMS, WRITE\_OWNER". The strings should appear in the order they are selected.

A similar technique should be used for all other sets of hex/string equivalents.

This algorithm, properly coded, requires only one (often partial) pass through the corresponding string equivalents table.

## 17.2.9 JSON format for ACLs

ACE flags and masks are members of a 32-bit quantity that is widely understood in its hexadecimal representations. The JSON data format does not support hexadecimal integers, however. For this reason, all hexadecimal integers in CDMI ACLs shall be represented as quoted strings containing a leading "0x".

ACLs containing one or more ACEs shall be represented in JSON as follows:

```
{
  "cdmi_acl" : [
    {
      "acetype" : "0xnn",
      "identifier" : "<user-or-group-name>",
      "aceflags" : "0xnn",
      "acemask" : "0xnn"
    },
    {
      "acetype" : "0xnn",
      "identifier" : "<user-or-group-name>",
      "aceflags" : "0xnn",
      "acemask" : "0xnn"
    }
  ]
}
```

(continues on next page)

(continued from previous page)

}

ACEs in such an ACL shall be evaluated in order as they appear.

3271

EXAMPLE 1: An example of an ACL embedded in a response to a GET request is as follows:

```
HTTP/1.1 200 OK
Content-Type: application/cdmi-object

{
  "objectType" : "/application/cdmi-object",
  "objectID" : "00007ED9001086A99CC6487FEE373D82",
  "objectName" : "MyDataItem.txt",
  "parentURI" : "/MyContainer/",
  "domainURI" : "/cdmi_domains/MyDomain/",
  "capabilitiesURI" : "/cdmi_capabilities/dataobject/",
  "completionStatus" : "Complete",
  "mimetype" : "text/plain",
  "metadata" : {
    "cdmi_size" : "17",
    "cdmi_acl" : [
      {
        "acetype" : "0x00",
        "identifier" : "EVERYONE@",
        "aceflags" : "0x00",
        "acemask" : "0x00020089"
      }
    ],
    ...
  },
  "valuerange" : "0-16",
  "value" : "Hello CDMI World!"
}
```

## Clause 18

# Retention and hold management

### 18.1 Overview

A cloud storage system may optionally implement retention management disciplines into the system management functionality of the cloud-based storage system. The implementation of retention and hold capabilities is indicated by the presence of the cloud storage system-wide capabilities for retention and hold capabilities.

Retention management includes implementing a retention policy, defining a hold policy to enable objects to be held for specific purposes (e.g., litigation), and defining how the rules for deleting objects are affected by placing either a retention policy and/or a hold on an object. CDMI™ object deletion is not a capability of retention management, per se, but rather is a general system capability. However, this clause describes what happens when placing either a retention policy and/or a hold on an object.

Retention management may be applied to the following object types:

- data objects,
- queue objects, and
- container objects.



## 18.2 Retention management disciplines

CDMI retention, deletion, and hold management affect any CDMI client that creates or deletes CDMI objects, as these disciplines mandate how a cloud storage system manages CDMI objects when they are created and until they are deleted.

CDMI retention management is comprised of three management disciplines: retention, hold, and deletion:

- CDMI retention uses retention time criteria to determine the time period during which object deletion from the CDMI-based system is prohibited. No changes to the object are allowed, even after the retention period has expired, except as specified below.
- CDMI hold prohibits object deletion and modification until all holds on the object have been released.
- A CDMI-based system shall not allow the deletion of a CDMI object before the CDMI retention time criteria are met or while holds exist. Any deletion attempts (e.g., by a CDMI application) shall return an error.
- After the CDMI retention time criteria have been met and all holds have been released, CDMI retention and holds shall no longer be a reason to prohibit object deletion.
- Once the retention period has started or if holds exist, changes to the object data and metadata shall not be allowed, with the exception of extensions to the retention and hold data system metadata. The retention data system metadata may be added or the retention period extended, and the hold data system metadata may be added or extended with additional holds. Any other attempt to modify the object shall return an error.

## 18.3 CDMI retention

### 18.3.1 Overview

CDMI retention only allows one retention policy to be applied to an object at a time.


Retention management uses time criteria to determine the time period during which CDMI object deletion from the CDMI-based system shall be prohibited. CDMI retention criteria shall be specified by the following data system metadata:

- a retention criteria identifier—a CDMI client-specified string that shall identify the retention records class (`cdmi_retention_id`); and
- a retention start time and retention period time—the start time, when used together with period, indicating when retention shall no longer be enforced (`cdmi_retention_period`).

When a CDMI client attempts to delete an object, the cloud storage system shall evaluate all such retention criteria and return an error, if any retention criteria have not been met.

When copying objects with a retention policy, retention properties shall not be transferred from the source CDMI object to the destination object, and the destination object shall not have a retention policy.

Fig. 11 shows how to establish time-based retention with a retention identifier. The value of the object data system metadata for the retention period shall not be reduced. Removing holds is outside the scope of the CDMI International Standard.



**Example:** Retention start date of 2010-04-28 with a duration of 730 days. No holds.

Figure 11: Object retention

A specific HTTP error code (403) shall be returned on operations to objects that are under retention period when the cloud storage system attempts to change or delete the object before the retention period criteria are met.

A cloud storage system shall not prevent metadata changes that increase the retention period, as there are valid business reasons to change a retention period for an object.

### 18.3.2 Examples

**EXAMPLE 1:** Place an existing object under retention:

```
PUT /MyContainer/MyDataObject.txt?metadata:cdmi_retention_id;metadata:cdmi_retention_
↪period HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdm-object

{
  "metadata" : {
    "cdmi_retention_id" : "1",
    "cdmi_retention_period" : "2010-04-28T00:00:00.000000Z/2012-04-27T00:00:00.
↪000000Z"
  }
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

3327 **EXAMPLE 2: Increase the duration of retention on an existing object under retention:**

```
PUT /MyContainer/MyDataObject.txt?metadata:cdmi_retention_period HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-object

{
  "metadata" : {
    "cdmi_retention_period" : "2011-04-28T00:00:00.000000Z/2013-04-27T00:00:00.
↪000000Z"
  }
}
```

3328 The following shows the response.

```
HTTP/1.1 204 No Content
```

3329 **EXAMPLE 3: Decrease the duration of retention on an existing object under retention:**

```
PUT /MyContainer/MyDataObject.txt?metadata:cdmi_retention_period HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-object

{
  "metadata" : {
    "cdmi_retention_period" : "2011-04-28T00:00:00.000000Z/2012-01-27T00:00:00.
↪000000Z"
  }
}
```

3330 The following shows the response.

```
HTTP/1.1 403 Forbidden
```

## 18.4 CDMI hold

### 18.4.1 Overview


CDMI hold enforces read-only data object access and prohibition of object deletion. A cloud storage system shall allow multiple holds to be applied to a single object to satisfy multiple hold orders. While an object is on hold, a cloud storage system shall strictly enforce read-only access to the object and prohibit object deletion.

When copying objects that are on hold, hold properties shall not be transferred from the source CDMI object to the destination object, and the destination object shall not be on hold.

Hold management uses a hold indicator to determine the time periods during which CDMI object revision (data and metadata) and deletion from the CDMI-based system shall be prohibited. CDMI hold criteria shall be specified by data system metadata, specifically, a hold criteria identifier that is a client-specified string that shall identify the holds and their order.

A CDMI client may place an object on hold by adding a hold identifier to the `cdmi_hold_id` data system metadata item. When an object is on hold, CDMI clients shall be subject to failures or unexpected state changes on operations, which would otherwise be successful if the object was not on hold.

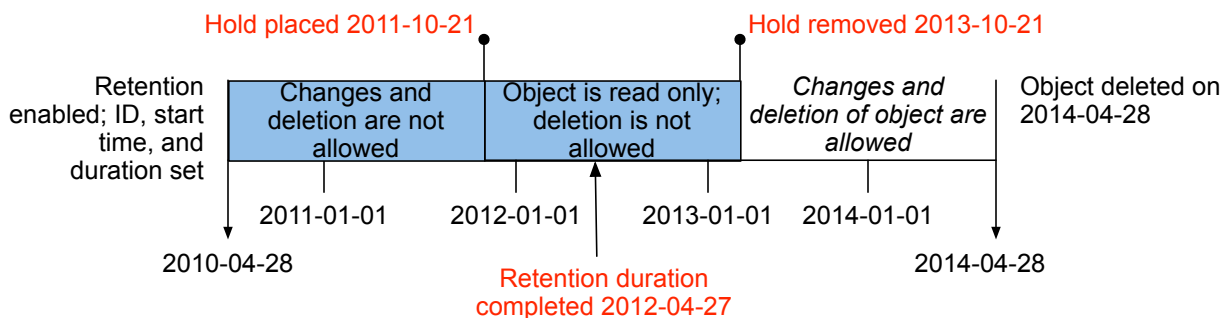
Fig. 12 shows how placing a hold on an object affects its read-only and deletion capability.



**Example:** Hold placed on the object on 2012-01-01 and removed on 2013-01-01

Figure 12: Object hold


Fig. 13 shows how to establish time-based retention with a retention identifier that has a hold placed on the object. The value of the object data system metadata for the retention period shall not be reduced, and the value of the object data system metadata for hold identifiers shall not permit holds to be removed. Removing holds is outside the scope of the CDMI International Standard.



**Example:** Start date of 2010-04-28 with a duration of 730 days; hold placed on the object

Figure 13: Object hold on object with retention

Fig. 14 shows how placing multiple holds on an object affects its read-only and deletion capability.



**Example:** Start date of 2010-04-28 with a duration of 730 days; holds placed on the object

Figure 14: Object with multiple holds

A cloud storage system shall maintain an on-hold object in read-only mode with respect to the application access to data and metadata and shall prohibit deletion, either automated or explicit.

- CDMI clients shall tolerate these object on-hold failures or state changes.
- Releases from hold are not part of this International Standard and are typically performed out of band using an additionally secured non-CDMI mechanism provided by the implementation.

A specific HTTP error code (403) shall be returned on operations to objects that are under a hold when the system attempts to change the object or attempts to delete the object before the hold is removed. This failure should be a an error to the application.

## 18.4.2 Examples

**EXAMPLE 1:** Place an existing object under hold:

```
PUT /MyContainer/MyDataObject.txt?metadata:cdmi_hold_id HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-object

{
  "metadata": {
    "cdmi_hold_id": {
      "case_7": ""
    }
  }
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

**EXAMPLE 2:** Attempt to remove a hold for an object under hold:

```
PUT /MyContainer/MyDataObject.txt?metadata:cdmi_hold_id HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-object

{
  "metadata": {
    "cdmi_hold_id": {}
  }
}
```

The following shows the response.

```
HTTP/1.1 403 Forbidden
```

3364 **EXAMPLE 3: Add a second hold to an object under hold:**

```
PUT /MyContainer/MyDataObject.txt?metadata:cdmi_hold_id HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdmi-object

{
  "metadata":{
    "cdmi_hold_id": {
      "case_7": "",
      "case_15": ""
    }
  }
}
```

3365 **The following shows the response.**

```
HTTP/1.1 204 No Content
```

## 18.5 CDMI auto-deletion

### 18.5.1 Overview

CDMI deletion controls cloud storage system actions with respect to object deletion. A cloud storage system may automatically delete a CDMI object after the retention time and hold criteria have been met. (See `cdmi_retention_autodelete` in *Data system metadata*.)

CDMI objects shall be automatically deleted by the system at the retention period expiration by setting the `cdmi_retention_autodelete` data system metadata item. The `cdmi_retention_autodelete` data system metadata item indicates to the system that the object shall be made unavailable for access after the retention criteria have been satisfied. The system shall ensure that the object is no longer available through the CDMI interface. If the system has satisfied the retention requirement and a hold is established for the object, the object shall not be made unavailable or deleted. When a hold and retention have been applied to an object, both need to be satisfied (retention period expired and no holds existing) for objects to be automatically deleted from the system.

EXAMPLE 1: Place an object under retention with autodelete:

```
PUT /MyContainer/MyDataObject.txt?metadata:cdmi_retention_period;metadata:cdmi_
↪retention_autodelete HTTP/1.1
Host: cloud.example.com
Content-Type: application/cdm-object

{
  "metadata":{
    "cdmi_retention_period": "2011-04-28T00:00:00.000000Z/2013-04-27T00:00:00.
↪000000Z",
    "cdmi_retention_autodelete": "true"
  }
}
```

The following shows the response.

```
HTTP/1.1 204 No Content
```

## 18.6 Retention security considerations

The accuracy and integrity of the retention start and elapsed times depend on the accuracy and integrity of the clock that is used to set their values. Equally important is the relative accuracy and security of the clock that determines if the retention period has elapsed when compared to the clock that sets the start time property. Relative time differences between these two clocks may lead to undesirable retention and deletion management behavior.

It is important to have a reliable source from which the system clock is set. A stratum 1 time is directly connected to a reference clock and is at the top of the time server hierarchy. Relative time differences between the system clock and the reference clock may lead to undesirable retention timestamps and difficulties with time action events.

EXAMPLE 1: An object is created in a cloud storage system at time 0 with a period of 8 years and autodelete of TRUE. At time 1 year, the system clock is adjusted forward to 9 years. Now, because the system time is 9 years, the retention time criterion is satisfied, even though only 1 year has actually elapsed. And, since autodelete is TRUE, the system automatically deletes the object.

The specification for accuracy and integrity of timekeeping is not within the scope of CDMI. However, to prevent undesirable retention and deletion management consequences, systems should maintain accurate clock time, with zero or minimal deviation to clock integrity.



## Clause 19

# Scope specification

### 19.1 Overview

CDMI™ provides a standardized mechanism to define sets of objects that match certain characteristics. This mechanism is known as a CDMI scope specification. Scope specifications are typically used to provide a CDMI client with a way to indicate in what set of CDMI objects it is interested.

Each JSON object within the scope specification represents a set of conditions that shall all be true in order for an object to be considered to match against the scope (a logical AND relationship). For queries, a matching object would be returned in the query results. An empty scope specification is considered to evaluate to true. Multiple JSON objects are used to express logical OR relationships, where if any JSON object in the scope evaluates to true, then the object shall be considered to have matched against the scope.

Each JSON object is constructed using the same structure that CDMI objects use. To show this structure, assume the following result from a CDMI GET for a data object:

```
HTTP/1.1 200 OK
Content-Type: application/cdmi-object

{
  "objectType" : "application/cdmi-object",
  "objectID" : "00007E7F0010EB9092B29F6CD6AD6824",
  "objectName" : "MyDataObject.txt",
  "parentURI" : "/MyContainer/",
  "parentID" : "00007E7F00102E230ED82694DAA975D2",
  "domainURI" : "/cdmi_domains/MyDomain/",
  "capabilitiesURI" : "/cdmi_capabilities/dataobject/",
  "completionStatus" : "Complete",
  "mimetype" : "text/plain",
  "metadata" : {
    "cdmi_size" : "108263",
    ...
  },
  "valuerange" : "0-108262",
  "value" : "..."
```

## 19.2 Examples

Each field inside a scope specification JSON object represents a condition that shall be met for a field.

EXAMPLE 1: A query to find all objects belonging to the domain “/cdmi\_domains/MyDomain/” is structured as follows:

```
[
  {
    "domainURI" : "== /cdmi_domains/MyDomain/"
  }
]
```

EXAMPLE 2: To query for all objects belonging to the domain “/cdmi\_domains/MyDomain/” AND are also located within the container “MyContainer”, the scope specification is structured as follows:

```
[
  {
    "parentURI" : "== /MyContainer/",
    "domainURI" : "== /cdmi_domains/MyDomain/"
  }
]
```

EXAMPLE 3: To query for all objects created within a certain time range, the scope specification is structured as follows:

```
[
  {
    "metadata": {
      "cdmi_ctime": [
        ">=2012-01-01T00:00:00",
        "<=2013-01-01T00:00:00"
      ]
    }
  }
]
```

When multiple matching expressions are specified for a given field or metadata item, all matching expression must evaluate true for an object to be considered a query result.

EXAMPLE 4: To query for all objects that belong to the domain “MyDomain” OR are located within the container “MyContainer”, the query is structured as follows:

```
[
  {
    "parentURI" : "== /MyContainer/",
  },
  {
    "domainURI" : "== /cdmi_domains/MyDomain/"
  }
]
```

Queries may match on any field within an object that a cloud storage system is capable of returning as a result of an object GET.

EXAMPLE 5: To query metadata items, the metadata object is included as an object within the query request. This query is shown as follows:

```
[
  {
    "metadata" : {
      "colour" : "== blue"
    }
  }
]
```

This approach allows matching against arbitrarily nested metadata structures. When a JSON object is included in the scope specification, matches are performed within that object, and when a JSON array is included in the scope specification, matches are performed within that array. Matching against the contents of arrays of objects is indicated by having an object within the array, as illustrated in Example 5.

3427 EXAMPLE 6: To query all objects with an ACE associated with the user “jdoe”:

```
[
  {
    "metadata" : {
      "cdmi_acl" : [
        {
          "identifier" : "== jdoe"
        }
      ]
    }
  }
]
```

3428 EXAMPLE 7: To query the value of objects, the value field is included within the query request. Values are always  
3429 represented using base 64 encoding in queries.

```
{
  [
    {
      "value": "== Ymx1ZQ=="
    }
  ]
}
```

3430 Query against the value of objects is optional and is indicated by the presence of the “cdmi\_query\_value” capability.

## 19.3 Query matching expressions

Query matching expressions are structured as “<operator>” or “<operator><sp><constant>”, and are defined in Table 148.

Table 148: Query matching expressions

Matching Expression	Description
“field”: “*”	The exists matching expression tests for the existence of the field. If the field is present, even if empty, the condition shall be considered to be met.
“field”: “!*”	The not exists matching expression tests for the non-existence of the field. If the field is absent, the condition shall be considered to be met.
“field”: “== constant”	The equals matching expression tests for the equality of the value of the field and a specified constant value. The equality test is case sensitive. If the constant value matches the value of the field, the condition shall be considered to be met.
“field”: “#== constant”	The numeric equals matching expression tests for the numeric equality of the value of the field and a specified constant value.
“field”: “!= constant”	The not equals matching expression tests for the non-equality of the value of the field and a specified constant value. The not-equals test is case sensitive. If the constant value does not match the value of the field, the condition shall be considered to be met.
“field”: “#!= constant”	The numeric equals matching expression tests for non-equality of the numeric equality of the value of the field and a specified constant value.
“field”: “> constant”	The greater than matching expression tests if the value of the field is lexicographically greater than a specified constant value. The greater than test is case sensitive. If the constant value is greater than the value of the field, the condition shall be considered to be met.
“field”: “#> constant”	The numeric greater than matching expression tests if the numeric value of the field is greater than a specified constant value.
“field”: “>= constant”	The greater than or equals to matching expression tests if the value of the field is lexicographically greater than or equal to a specified constant value. The greater than or equals to test is case sensitive. If the constant value is greater than or equal to the value of the field, the condition shall be considered to be met.
“field”: “#>= constant”	The numeric greater than or equals to matching expression tests if the numeric value of the field is greater than or equal to a specified constant value.
“field”: “< constant”	The less than operator tests if the value of the field is lexicographically less than a specified constant value. The less than test is case sensitive. If the constant value is less than the value of the field, the condition shall be considered to be met.
“field”: “#< constant”	The numeric less than operator tests if the numeric value of the field is less than a specified constant value.
“field”: “<= constant”	The less than or equals to matching expression tests if the value of the field is lexicographically less than or equal to a specified constant value. The less than or equal test is case sensitive. If the constant value is less than or equal to the value of the field, the condition shall be considered to be met.
“field”: “#<= constant”	The numeric less than or equals to matching expression tests if the numeric value of the field is less than or equal to a specified constant value.
“field”: “starts constant”	The starts with matching expression tests if the field value starts with a specified constant value. If the constant value is equal to the start of the value of the field, the condition shall be considered to be met.
“field”: “!starts constant”	The not starts with matching expression tests if the field value does not start with a specified constant value. If the constant value is not equal to the start of the value of the field, the condition shall be considered to be met.
“field”: “ends constant”	The ends with matching expression tests if the field value ends with a specified constant value. If the constant value is equal to the end of the value of the field, the condition shall be considered to be met.

Continued on next page

Table 148 – continued from previous page

Matching Expression	Description
"field": "!ends constant"	The not ends with matching expression tests if the field value does not end with a specified constant value. If the constant value is not equal to the end of the value of the field, the condition shall be considered to be met.
"field": "contains constant"	The contains matching expression tests if the field value contains a specified constant value. If the constant value is found as a substring within the value of the field, the condition shall be considered to be met. The contains operator is only supported if the "cdmi_query_contains" capability is present.
"field": "!contains constant"	The not contains matching expression tests if the field value does not contain a specified constant value. If the constant value is not found as a substring within the value of the field, the condition shall be considered to be met. The not contains operator is only supported if the "cdmi_query_contains" capability is present.
"field": "tag constant"	<p>The tag matching expression tests if the field value contains a specified constant tag value.</p> <p>The leading space character after the "tag" and before the constant value is not included in the comparison. The tag test is not case sensitive.</p> <p>If the constant value is found as a tag substring within the value of the field, the condition shall be considered to be met. Tag substrings start at the beginning of the value or a ",", and end at the next ",", or the end of the string. Whitespace before and after ",", characters shall be stripped for the purpose of comparisons. Tag matching expressions are only supported if the "cdmi_query_tags" capability is present.</p>
"field": "!tag constant"	<p>The not tag matching expression tests if the field value does not contain a specified constant tag value.</p> <p>The leading space character after the "!tag" and before the constant value is not included in the comparison. The not tag test is not case sensitive.</p> <p>If the constant value is not found as a tag substring within the value of the field, the condition shall be considered to be met. Tag substrings start at the beginning of the value or a ",", and end at the next ",", or the end of the string. Whitespace before and after ",", characters shall be stripped for the purpose of comparisons. Tag matching expressions are only supported if the "cdmi_query_tags" capability is present.</p>
"field": "=~ constant"	<p>The regular expression matching expression tests if the field value matches a specified constant regular expression value. If the regular expression evaluates to true against the value, the condition shall be considered to be met.</p> <p>Regular expression strings shall be processed according to the POSIX Extended Regular Expression (ERE) standard, as specified in IEEE 1003.1-2017 [41].</p> <p>Regex matching expressions are only supported if the "cdmi_query_regex" capability is present.</p>
"field": "!~ constant"	<p>The not regular expression matching expression tests if the field value does not match a specified constant regular expression value. If the regular expression evaluates to false against the value, the condition shall be considered to be met.</p> <p>Regular expression strings shall be processed according to the POSIX Extended Regular Expression (ERE) standard, as specified in IEEE 1003.1-2017 [41].</p> <p>Regex matching expressions are only supported if the "cdmi_query_regex" capability is present.</p>

3434 Numeric constant strings shall be processed according to the JSON number representation described in RFC 4627 [5].

3435 A numeric matching expression shall be considered to be non-matching against a non-numeric field value.

3436 All fields in objects that are not included in the scope specification shall be ignored for the purpose of matching objects.

3437 When a URI is used as the constant for the equals and not equals operators against the `parentURI`, `domainURI`, and  
 3438 `capabilitiesURI`, either a URI by path or URI by object ID can be specified and are considered interchangeable.

EXAMPLE 1: In a query to find all objects belonging to a specific domain, the following two query scopes are considered identical:

```
[
  {
    "domainURI" : "==" /cdmi_domains/MyDomain/"
  }
]
```

and

```
[
  {
    "domainURI" : "==" /cdmi_objectid/00007E7F001074C86AD256DA5C67180D/"
  }
]
```

EXAMPLE 2: Likewise, a query to find all objects with a given parent container would have two equivalent forms:

```
[
  {
    "parentURI" : "==" /MyContainer/"
  }
]
```

and

```
[
  {
    "parentURI" : "==" /cdmi_objectid/00007ED900100E358C3B312DB652C201/"
  }
]
```

If an object ID is used in a query scope in the `objectID` field or the `parentID` field, all object IDs shall be processed such that they are case insensitive.

## Clause 20

# Results specification

### 20.1 Overview

CDMI™ provides a standardized mechanism to define subsets of object contents. This mechanism is known as a CDMI results specification. Results specifications are typically used to provide a CDMI client with a way to indicate on what subset of the contents of CDMI objects it intends to retrieve or operate.

Each JSON object within the results specification represents a set of fields that are returned for each matching object.

The results JSON object shall be constructed using the same structure as is used for CDMI objects. To show this, assume the following result from a CDMI GET for a data object:

```
HTTP/1.1 200 OK
Content-Type: application/cdm-object

{
  "objectType" : "application/cdm-object",
  "objectID" : "00007E7F0010EB9092B29F6CD6AD6824",
  "objectName" : "MyDataObject.txt",
  "parentURI" : "/MyContainer/",
  "parentID" : "00007E7F00102E230ED82694DAA975D2",
  "domainURI" : "/cdmi_domains/MyDomain/",
  "capabilitiesURI" : "/cdmi_capabilities/dataobject/",
  "completionStatus" : "Complete",
  "mimetype" : "text/plain",
  "metadata" : {
    "cdmi_size" : "108263",
    ...
  },
  "valuerange" : "0-108262",
  "value" : "..."
```

## 20.2 Examples

Each field inside a results specification JSON object indicates that the field shall be included in the results.

EXAMPLE 1: The following results specification requests that the `objectID` and `cdmi_size` metadata fields be returned in the results:

```
{
  "cdmi_results_specification" : {
    "objectID" : "",
    "metadata" : {
      "cdmi_size" : ""
    }
  }
}
```

EXAMPLE 2: If an object is matched, the result JSON is enqueued as follows:

```
{
  "objectID" : "00007E7F0010EB9092B29F6CD6AD6824",
  "metadata" : {
    "cdmi_size" : "108263"
  }
}
```

For most common use cases, clients request either the `objectID`, the `objectName` and `parentURI`, or all three fields in the `cdmi_results_specification`. If the `parentURI` or `objectName` is requested, the field shall only be returned for objects existing in a container object.

EXAMPLE 3: To request all metadata items be returned for each matching object, the following `cdmi_results_specification` shall be used:

```
{
  "cdmi_results_specification" : {
    "metadata" : ""
  }
}
```

EXAMPLE 4: To request all fields and all metadata items be returned for each matching object, the following `cdmi_results_specification` shall be used:

```
{
  "cdmi_results_specification" : ""
}
```

The `value` field is always returned in base 64 encoding when included in a query result, where the `valuetransferencoding` field indicates the encoding that should be expected if a GET to read the object is performed.



# Clause 21

## Notification queues

### 21.1 Overview

A cloud storage system may optionally implement notification functionality. The implementation of notification is indicated by the presence of the cloud storage system-wide capabilities for notification, and requires support for CDMI™ queues.

Notification queues allow CDMI clients to efficiently discover what changes have occurred to the system. As queue data is persistent, no session state needs to be retained by the client. If different notification queues are used for different clients, then each client operates independently from the others (e.g., a storage management application may use a notification queue to keep its database current without having to do full scans of a container to discover what data objects have been added, modified, or removed).

When a client wishes to receive notifications, it may first check if the system is capable of providing notifications by checking for the presence of the `cdmi_notification` capability in the root container capabilities. If this capability is not present, creating a notification queue shall be successful, but no notifications shall be enqueued into the notification queue.

To create a notification queue, the client creates a regular CDMI queue and adds metadata instructing the storage system to treat the queue as a notification queue. This added metadata also instructs the system about what types of notifications shall be generated and what information shall be included with each notification.

After the notification queue is created, all subsequent matching events after the queue creation time shall result in notification results being enqueued into the queue. CDMI does not mandate any specific ordering of events, and clients must be able to handle events that arrive out of order.

#### 21.1.1 Required metadata

When creating a notification queue, the metadata described in Table 149 shall be provided. Attempts to change metadata in this table shall result in an HTTP status code of 403 `Forbidden`. After a notification queue has been created, with the exception of `cdmi_queue_type`, the metadata items in this table cannot be changed. `cdmi_queue_type` can only be removed, indicating to the system that the notification queue shall no longer receive notifications and shall be treated as a regular CDMI queue object.

Table 149: Required metadata for a notification queue

Metadata name	Type	Description	Requirement
<code>cdmi_queue_type</code>	JSON string	Queue type indicates how the cloud storage system shall manage the queue object. The type of <code>cdmi_notification_queue</code> is defined for notification queues.	Mandatory

Continued on next page

Table 149 – continued from previous page

Metadata name	Type	Description	Requirement
cdmi_notification_events	JSON array of JSON strings	<p>The notification events metadata contains a JSON array that indicates which events generate notifications. Defined values are:</p> <ul style="list-style-type: none"> <li>• <code>cdmi_create_processing</code> - Notifications are generated when a new object is created but is still in the “Processing” completion status.</li> <li>• <code>cdmi_create_complete</code> - Notifications are generated when a new object is created immediately or when a new object in the process of being created transitions from the “Processing” completion status. When an object transitions from “Processing” completion status, the “<code>cdmi_event_result</code>” is the HTTP result code that would have been returned if the create operation was not delayed.</li> <li>• <code>cdmi_read</code> - Notifications are generated when an object is read.</li> <li>• <code>cdmi_modify_processing</code> - Notifications are generated when an existing object is modified but is still in the “Processing” completion status.</li> <li>• <code>cdmi_modify_complete</code> - Notifications are generated when an existing object is modified and is in the “Complete” completion status. This notification is also generated when an existing object being modified transitions from “Processing” to “Complete”. When an object transitions from “Processing” completion status, the “<code>cdmi_event_result</code>” is the HTTP result code that would have been returned if the modify operation was not delayed.</li> <li>• <code>cdmi_rename</code> - Notifications are generated when an object is renamed as part of a move operation.</li> <li>• <code>cdmi_copy</code> - Notifications are generated for the newly created copied object when the copy is completed.</li> <li>• <code>cdmi_reference</code> - Notifications are generated when a reference is created.</li> <li>• <code>cdmi_delete</code> - Notifications are generated when an object is deleted.</li> <li>• <code>cdmi_export</code> - Notifications are generated when a container is exported.</li> <li>• <code>cdmi_snapshot</code> - Notifications are generated when a container snapshot is created.</li> <li>• &lt;implementor-specific events&gt;</li> </ul> <p>Clients may include the desired notification event types in the <code>cdmi_notification_events</code> JSON array. If all notifications events are desired, an empty JSON array shall be used.</p>	Mandatory
cdmi_scope_specification	JSON array of JSON objects	<p>The scope specification determines the set of objects on which operations trigger the generation of notifications. If notifications are desired for all objects, include an empty JSON array.</p> <p>See <a href="#">clause 19</a> for how to construct a scope specification.</p>	Mandatory

Continued on next page

Table 149 – continued from previous page

Metadata name	Type	Description	Requirement
cdmi_results_specification	JSON object	<p>The results specification contains the JSON fields to be returned for each object that matches the notification scope specification. See <a href="#">clause 20</a> for how to construct a results specification.</p> <p>In addition to the fields defined in <a href="#">clause 20</a>, for notifications, four additional fields are defined:</p> <ul style="list-style-type: none"> <li>cdmi_event - Indicates the event as specified in the "cdmi_notification_events" field that triggered the notification;</li> <li>cdmi_event_result - Indicates the status result of the event that triggered the notification. The status is the same as the status that was returned over the HTTP request, i.e., 200 OK, 404 Not Found, etc.;</li> <li>cdmi_event_time - Indicates the time of the event that triggered the notification. The time will be formatted in ISO-8601 time (see <a href="#">5.6</a> and <a href="#">ISO 8601-1:2019 [33]</a>); and</li> <li>cdmi_event_user - Indicates the principal (ACL name) of the user that caused the event that triggered the notification. If the system triggered the event, the name will be left as an empty string.</li> </ul>	Mandatory

EXAMPLE 1: The metadata associated with a notification queue is as follows:

```
{
  "metadata" : {
    "cdmi_queue_type" : "cdmi_notification_queue",
    "cdmi_notification_events" : [
      "cdmi_create_complete",
      "cdmi_read",
      "cdmi_modify_complete",
      "cdmi_delete"
    ],
    "cdmi_scope_specification" : [
      {
        "domainURI" : "== /cdmi_domains/MyDomain/",
        "parentURI" : "starts /sandbox",
        "metadata" : {
          "cdmi_size" : ">+100000"
        }
      }
    ],
    "cdmi_results_specification" : {
      "cdmi_event" : "",
      "cdmi_event_result" : "",
      "cdmi_event_time" : "",
      "objectID" : "",
      "metadata" : {
        "cdmi_size" : ""
      }
    }
  }
}
```

When notification results are stored in a notification queue, each enqueued value shall consist of a JSON object of MIME type "application/json". This JSON object contains the specified values requested in the cdmi\_results\_specification of the notification queue metadata.

EXAMPLE 2: A notification result JSON object is as follows:

```

{
  "cdmi_event" : "cdmi_read",
  "cdmi_event_result" : "200 OK",
  "cdmi_event_time" : "2010-11-15T13:12:52.342324Z",
  "objectID" : "00007E7F0010EB9092B29F6CD6AD6824",
  "metadata" : {
    "cdmi_size" : "108263"
  }
}

```

Objects shall only be included in the notification results if the user who created the notification queue is able to read the matching object.

If the administrator created the notification queue, then all matching objects that the administrator is allowed to read are included in the results. If user “jdoe” created the notification queue, then only matching objects that “jdoe” is allowed to read are included in the results.

### 21.1.2 System-created metadata

Table 150 describes the system-created metadata that provides details on the status of the notification queue.

Table 150: Notification status metadata

Metadata name	Type	Description	Requirement
cdmi_notification_status	JSON string	<p>A string indicating the state of the notification queue. Defined values are:</p> <ul style="list-style-type: none"> <li>Processing - Indicates that the notification queue is scanning for results;</li> <li>Halted - Indicates that new notifications will no longer be enqueued;</li> <li>Current - Indicates that the notification queue contained all notifications that can be found at this time; and</li> <li>Error - Indicates that the notification queue metadata is not valid, or other errors were encountered that prevented notification messages from being enqueued. Arbitrary vendor-defined text may follow the string “Error”.</li> </ul> <p>If this metadata item does not exist, then notifications have not yet started being enqueued.</p>	Mandatory

## Clause 22

# Query queues

### 22.1 Overview

A cloud storage system may optionally implement metadata and/or full-text query functionality. The implementation of query is indicated by the presence of the cloud storage system-wide capabilities for query and requires support for CDMI™ queues.

Query queues allow CDMI clients to efficiently discover what content matches a given set of metadata query criteria or full-content search criteria. Clients create or update a query queue by specifying metadata that defines the matching criteria (known as the query scope), along with what results should be returned for matching objects (known as the query results). The cloud service shall then perform the query using the content existing at the time the query is being processed, storing the query results in the query queue. As query results are found, they are added to the queue, and when the query is complete, the `cdmi_query_status` metadata of the queue is changed to indicate that the query has completed. Any matching objects created or modified while the query is being performed may or may not be included in the query results (e.g., as a consequence of eventual consistency).

When a client wishes to perform queries, it may first check if the system is capable of providing query functionality by checking for the presence of the `cdmi_query` capability in the root container capabilities. If this capability is not present, creating a query queue shall be successful, but no query results shall be enqueued into the query queue.

When creating a query queue, the metadata described in Table 151 shall be provided. Attempts to change metadata in this table shall result in an HTTP status code of 403 Forbidden. After a query queue has been created, with the exception of `cdmi_queue_type`, the metadata items in this table cannot be changed. If the value of `cdmi_queue_type` is changed from "cdmi\_query\_queue", this change indicates to the system that an in-process query shall be stopped, the query queue shall no longer receive query results, and the query queue shall be treated as a regular CDMI queue object. To start a new query with an existing queue, the value of the `cdmi_queue_type` shall be changed back to "cdmi\_query\_queue". This international standard does not define a mechanism to pause a running query or resume a stopped query.

Table 151: Required metadata for a query queue

Metadata name	Type	Description	Requirement
<code>cdmi_queue_type</code>	JSON string	The queue type indicates how the cloud storage system shall manage the queue object. The type of "cdmi_query_queue" is defined for query queues.	Mandatory
<code>cdmi_scope_specification</code>	JSON array of JSON objects	The scope specification determines which objects are included in the query results. This scope specification is similar to a "WHERE" clause in SQL-like languages. To query all objects, specify an empty JSON array. See Clause 19 for how to construct a scope specification.	Mandatory
<code>cdmi_results_specification</code>	JSON object	The results specification contains the JSON fields to be returned for each object that matches the query. This results specification is similar to a "SELECT" clause in SQL-like languages. See Clause 20 for how to construct a results specification.	Mandatory

EXAMPLE 1: An example of the metadata associated with a query queue is as follows:

```
{
  "metadata" : {
    "cdmi_queue_type" : "cdmi_query_queue",
    "cdmi_scope_specification" : [
      {
        "domainURI" : "== /cdmi_domains/MyDomain/",
        "parentURI" : "starts /sandbox",
        "metadata" : {
          "cdmi_size" : "#> 100000"
        }
      }
    ],
    "cdmi_results_specification" : {
      "objectID" : "",
      "metadata" : {
        "cdmi_size" : ""
      }
    }
  }
}
```

When results are stored in a query queue, each enqueued value shall consist of a JSON object of MIME type “application/json”. This JSON object contains the specified values requested in the cdmi\_results\_specification of the query queue metadata.

EXAMPLE 2: An example of a query result JSON object is as follows:

```
{
  "objectID" : "00007E7F0010EB9092B29F6CD6AD6824",
  "metadata" : {
    "cdmi_size" : "108263"
  }
}
```

Table 152 describes the system-created metadata that provides details on the status of the query queue.

Table 152: Query status metadata

Metadata name	Type	Description	Requirement
cdmi_query_status	JSON string	When present, this metadata item indicates the state of the query queue. Defined values are: <ul style="list-style-type: none"><li>Processing - Indicates that the query queue is scanning for results;</li><li>Halted - Indicates that new query results will no longer be enqueued;</li><li>Current - Indicates that the query queue contained all query results that can be found at this time; and</li><li>Error - Indicates that the query queue metadata was not valid, or other errors were encountered that prevented all query results from being enqueued. Arbitrary vendor-defined text may follow the string “Error”.</li></ul>	Mandatory

Objects shall only be included in the query results if the user who created the query queue is able to read the matching objects or metadata.

NOTE: If the administrator created the query queue, then all matching objects that the administrator is allowed to read are included in the results. If user “jdoe” created the query queue, then only matching objects that “jdoe” is allowed to read are included in the results.

## 22.2 Extending CDMI query

An implementor of a CDMI server may extend CDMI query by adding vendor-specific matching expressions. When an implementor adds vendor-specific metadata fields, these fields shall be queried using the standard query queue functionality.

An implementor of a CDMI server may extend CDMI query by allowing the creation of vendor-specific query queues with a type other than "cdmi\_query\_queue".

## Clause 23

# Encrypted objects

### 23.1 Overview

A cloud storage system may optionally implement additional operations against encrypted objects. Support for these operations are indicated by the presence of the cloud storage system-wide capabilities for encrypted objects.

Encrypted object operations include the ability to encrypt, re-encrypt, and decrypt objects that are already stored in the cloud (in-place), to sign and verify the signature of encrypted objects, and to access and update the plaintext associated with encrypted objects.

The CDMI International Standard does not specify the method by which keys are managed. Key management services are provided by an external key management system (KMS), and the use of the KMIP standard is given as an example of how a CDMI server interacts with an external KMS.

CDMI objects can contain values that are encrypted. Operations against an encrypted CDMI object are only supported if the encrypted object value is a valid CMS or JWE JSON format. The CMS or JWE JSON object shall include an embedded mimetype of the encrypted object. For JWE, the “cty” header shall be used for this purpose.



## 23.2 Encryption operations

The state transition diagram for encrypted objects is shown in [Fig. 15](#):




Figure 15: Encrypted object state transitions

The following eight encryption operations are defined:

### 23.2.1 Create a new encrypted object

Client-encrypted objects shall be stored to a CDMI server using a standard HTTP or CDMI PUT operation, as described in clauses 7.2 and 8.2. The client shall indicate that an object is encrypted by specifying a mimetype of “application/cms” or “application/jose+json”.

A client may register an encryption key, signing keys and/or verification keys with a Key Management System (KMS), and may indicate the Key IDs in `cdmi_enc_key_id`, `cdmi_enc_value_sign_id`, `cdmi_enc_object_sign_id`, `cdmi_enc_value_verify_id`, and/or `cdmi_enc_object_verify_id` metadata items. This allows the CDMI server to access the keys from the KMS on behalf of a client, when needed.

Creating an encrypted objects on a CDMI server does not require any encryption-specific capabilities to be supported, and is backwards compatible with earlier versions of the CDMI standard. This permits encrypted objects to be stored and transferred by CDMI servers that do not support encryption-specific functionality.

### 23.2.2 Delete an encrypted object

Encrypted objects shall be deleted using a standard HTTP or CDMI DELETE operation, as described in [clause 7.5](#) and [clause 8.5](#). Any client with sufficient permissions shall be permitted to delete an encrypted object, regardless of if they can access the decryption keys.

### 23.2.3 Encrypt an unencrypted object

Existing unencrypted objects shall be encrypted in-place by performing a CDMI PUT operation, as described in [clause 8.4](#), that changes the object mimetype to “application/cms” or “application/jose+json” and specifies a `cdmi_enc_key_id` metadata item. The client may also specify a `cdmi_enc_value_sign_id` and/or `cdmi_enc_value_verify_id` metadata item to indicate that the object is to be signed, and to provide signature verification information.

The CDMI Server shall use the client’s credentials (which are included in HTTP headers, and are out of scope of this International Standard) to retrieve the encryption and signing keys, and encryption and signing algorithm information from the KMS, and shall use the keys to encrypt and sign the value of the object. The mimetype of the encrypted value is stored in the CMS wrapper, or in a “`cty`” field of the JWE JSON.

### 23.2.4 Decrypt an encrypted object

Existing encrypted objects shall be decrypted in-place by performing a CDMI PUT operation, as described in [clause 8.4](#), that changes the object mimetype from “application/cms” or “application/jose+json” to the original mimetype as specified in the CMS wrapper, or in the “`cty`” field of the JWE JSON. Specifying any other fields or metadata shall return a “400 Bad Request” result code.

The CDMI Server shall use the client’s credentials (which are included in HTTP headers, and are out of scope of this International Standard) to retrieve the encryption, signing and verification keys, and encryption, signing and verification algorithm information from the KMS, and shall use the keys to decrypt and verify the encrypted value and user metadata included in the object.

### 23.2.5 Re-encrypt an encrypted object

Existing encrypted objects shall be re-encrypted in-place by performing a CDMI PUT operation, as described in [clause 8.4](#), that retains the object mimetype of “application/cms” or “application/jose+json”, or changes the object mimetype from “application/cms” to “application/jose+json”, or vice-versa. The client shall also specify a new `cdmi_enc_key_id`, `cdmi_enc_value_sign_id` and/or `cdmi_enc_value_verify_id` metadata item to indicate the new key(s) to be used. Specifying any other fields or metadata shall return a “400 Bad Request” result code.

The CDMI Server shall use the client’s credentials (which are included in HTTP headers, and are out of scope of this International Standard) to retrieve both the original encryption and signing keys using the original metadata values, and the new encryption and signing keys using the new metadata values from the KMS, and shall use these keys to decrypt, verify, encrypt and sign the value of the object, as needed.

If an encrypted object does not have an existing `cdmi_enc_key_id` metadata item, does not have a “`kid`” header, and no keys are associated with the Object ID, the specified metadata shall be added to the object, and no re-encryption operation shall be performed.

### 23.2.6 Access ciphertext of an encrypted object

The ciphertext content of an encrypted object shall be read by performing an HTTP GET operation, as described in [clause 6.3](#), with an Accept header value of “application/cms” or “application/jose+json”, depending on the mimetype of the encrypted object.

The ciphertext content of an encrypted object shall also be read by performing a CDMI GET operation, as described in [clause 8.3](#).

### 23.2.7 Access plaintext of an encrypted object

The plaintext value of an encrypted object shall be read by performing an HTTP GET operation, as described in [clause 6.3](#), with an Accept header value other than “application/cms” or “application/jose+json”, typically “\*/\*”. Object plaintext cannot be transparently accessed using a CDMI GET.

The CDMI Server shall use the client’s credentials (which are included in HTTP headers, and are out of scope of this International Standard) to retrieve the encryption, signing and verification keys, and encryption, signing and verification algorithm information from the KMS, and shall use the keys to decrypt and verify the encrypted value included in the object.

When an encrypted object is decrypted for access, the plaintext shall not be retained or cached by the CDMI server.

### 23.2.8 Update plaintext of an encrypted object

The plaintext value of an encrypted object shall be modified by performing an HTTP PUT operation, as described in [clause 6.4](#), with an Content-Type header value other than “application/cms” or “application/jose+json”, typically “\*/\*”, depending on the mimetype of the encrypted object. Object plaintext cannot be transparently modified using a CDMI GET.

The CDMI Server shall use the client’s credentials (which are included in HTTP headers, and are out of scope of this International Standard) to retrieve the encryption, signing and verification keys, and encryption, signing and verification algorithm information from the KMS, and shall use the keys to decrypt and verify the encrypted value, update the value, and re-encrypt/re-sign the updated value.

When an encrypted object is decrypted for update, the plaintext shall not be retained or cached by the CDMI server.

### 23.2.9 Other CDMI operations

Other operations specified by this International Standard (such as copying, serializing, querying, etc) treat an encrypted value the same way as a non-encrypted value.

## 23.3 Example uses of encrypted objects

Encrypted objects can be used with CDMI systems in the following ways:

- **Passthrough** – A client may store an encrypted object in any format in a CDMI server, with the ciphertext being accessible to the server and to other authorized clients. No access to the plaintext is provided. Passthrough use is compatible with all CDMI systems and is useful when the clients manage all security-related operations and want to protect against potentially untrustworthy clouds.
- **Server-side encryption and signing** – A client may instruct a CDMI server that supports encrypted object operations to take an existing CDMI object and encrypt or encrypt and sign it in place into CMS or JWE JSON representation, where the value of the object is persistently stored from that point on in an encrypted format. Server-side encryption and signing is useful when clients trust the CDMI server and want to increase object security without having to re-upload the data.
- **Server-side decryption** – A client may instruct a CDMI server that supports encrypted object operations to take an existing CDMI object and decrypt it in place from a CMS or JWE JSON representation, where the value of the object is persistently stored from that point on in a decrypted format. Server-side decryption is useful when a client trusts the CDMI server and wants to decrease object security without having to re-upload the data.
- **Client access decryption** – A CDMI server may automatically attempt to decrypt an encrypted object when accessed via HTTP. Client access decryption is useful to provide transparent access to authorized HTTP clients without requiring modifications to the HTTP clients.
- **Cloud access decryption** – A CDMI server may automatically decrypt encrypted objects when it has access to the decryption keys. Cloud access decryption is useful for cloud-resident data processing performed by the CDMI server, such as virus scanning, query, and analytics.
- **Signature verification** – A CDMI server can automatically verify signatures that are attached to encrypted objects that include a signature. Signature verification is useful for detecting corruption or alteration before delivering data to a client.

## 23.4 KMS integration

The encryption key is obtained from the KMS using a unique identifier that is stored in the `cdmi_enc_key_id` metadata item associated with the encrypted object. If this metadata item is not present, the CDMI object ID shall be used to locate the key.

When a client requests that an operation be performed that requires accessing the key for the object, the CDMI server evaluates the credentials provided by the client to determine if the client is authorized to perform the requested operation. If the operation is permitted, the CDMI server retrieves the key from the KMS to complete the requested operation. To retrieve the key, the client may be required to provide additional information in the HTTP request that the CDMI server can then use to authenticate to the KMS.

The CDMI International Standard does not specify the mechanism by which the CDMI server communicates with the KMS. In this International Standard, the KMIP protocol is used as an example. CMS and JWE strings for algorithms, key lengths, etc., need to be mapped to the strings used by the KMS (see KMIP clause 9.1.3.2.7).

All keys are created and managed externally to the CDMI server, typically by the client or a system operating on behalf of the client. As a consequence, the CDMI server requires read-only access to the KMS. The CDMI server shall not cache keys.

## 23.5 CMS format

Any valid CMS-formatted data can be stored to a CDMI server. However, encrypted object operations are only defined for the following subset of valid CMS-formatted data.

For encryption operations, the CDMI server shall support the following:

- EnvelopedData
- EncryptedContentInfo
- contentEncryptionAlgorithm value listed in the `cdmi_cms_encryption` capability of that CDMI server

For signature operations, the CDMI server shall support the following:

- AuthenticatedData
- SignedData
- digestAlgorithms value listed in the `cdmi_cms_digest` capability of that CDMI server
- SignerInfo
- signatureAlgorithm value listed in the `cdmi_cms_signature` capability of that CDMI server

The following CMS-formatted data may be ignored: `recipientInfos`

## 23.6 JOSE format

Any valid JWE-formatted data can be stored to a CDMI server. However, encrypted object operations are only defined for the following subset of valid JWE-formatted data.

For encryption operations, the CDMI server shall support the following:

- JWE with Direct Encryption (Symmetric Key from KMS)
- JWE with Key Encryption (Public Key from KMS)

For signature operations, the CDMI server shall support the following:

- JWS RSA (Private Key from KMS)
- JWS ECDSA (Private Key from KMS)
- JWS HMAC-SHA2 (Symmetric Key from KMS)

The following JOSE-formatted data may be ignored: Multiple recipients and multiple signatures

## 23.7 Signature/digest verification

If a signature is present as part of the CMS or JWE JSON value, the CDMI server shall verify that the signature of the value is valid before allowing plaintext access or modification.

If a whole-object signature is present, the CDMI server shall verify that the signature contained in the `cdmi_enc_signature` metadata item is valid before allowing any read operations for the object. Write operations are permitted for an object with an invalid or unverifiable whole-object signature.

When present, a whole-object signature shall be attached as a `"cdmi_enc_signature"` metadata item in JWS compact format, with the second field (the JWS payload field) replaced with an empty string as described in Appendix F of RFC 7515.

For signature creation and verification, payload field shall be computed using the following process:

1. Create a serialized representation of the CDMI object, as described in [clause 15](#)
2. Remove the following metadata items, if present:
  - `cdmi_atime`
  - `cdmi_acount`
  - `cdmi_enc_signature`
  - Any `*_provided` metadata items
3. Sort all JSON objects in the serialized CDMI object according to the following rules:
  - Within each JSON object, name/value pair entries shall be sorted lexicographically by name
  - Within each JSON array, the initial order shall be preserved
4. Remove all JSON whitespace
5. Base64 URL encode, according to the JWS RFC 7515 [17]



## 23.8 Error handling

If a decryption or signature validation operation is requested against a CDMI object containing an invalid CMS or JWE JSON representation, an HTTP status code of 500 `Internal Error` shall be returned to the client.

If a decryption or signature validation operation is requested against a CDMI object containing a valid CMS or JWE JSON representation that uses an unsupported algorithm or feature, an HTTP status code of 501 `Not Implemented` shall be returned to the client.

If a decryption or signature validation operation is requested against a CDMI object containing a valid CMS or JWE JSON representation, but the required keys are temporarily unavailable given the credentials presented, an HTTP status code of 408 `Request Timeout` shall be returned to the client.

If a decryption or signature validation operation is requested against a CDMI object containing a valid CMS or JWE JSON representation, but the required keys are unavailable given the credentials presented, an HTTP status code of 401 `Unauthorized` shall be returned to the client.

If a decryption or signature validation operation is requested against a CDMI object containing a valid CMS or JWE JSON representation, valid keys are available, and signature verification fails, an HTTP status code of 403 `Forbidden` shall be returned to the client.

# Clause 24

## Delegated access control

### 24.1 Overview

CDMI access control is based around Access Control Lists (ACLs) that are stored as object metadata. When a client requests to perform an operation against a CDMI object, the CDMI server shall validate the client's identity and credentials against the object ACL to determine if the operation is allowed. This request assumes that the CDMI server is trusted and capable of making these access control decisions.

Fig. 16 illustrates an ACL-based access control request:

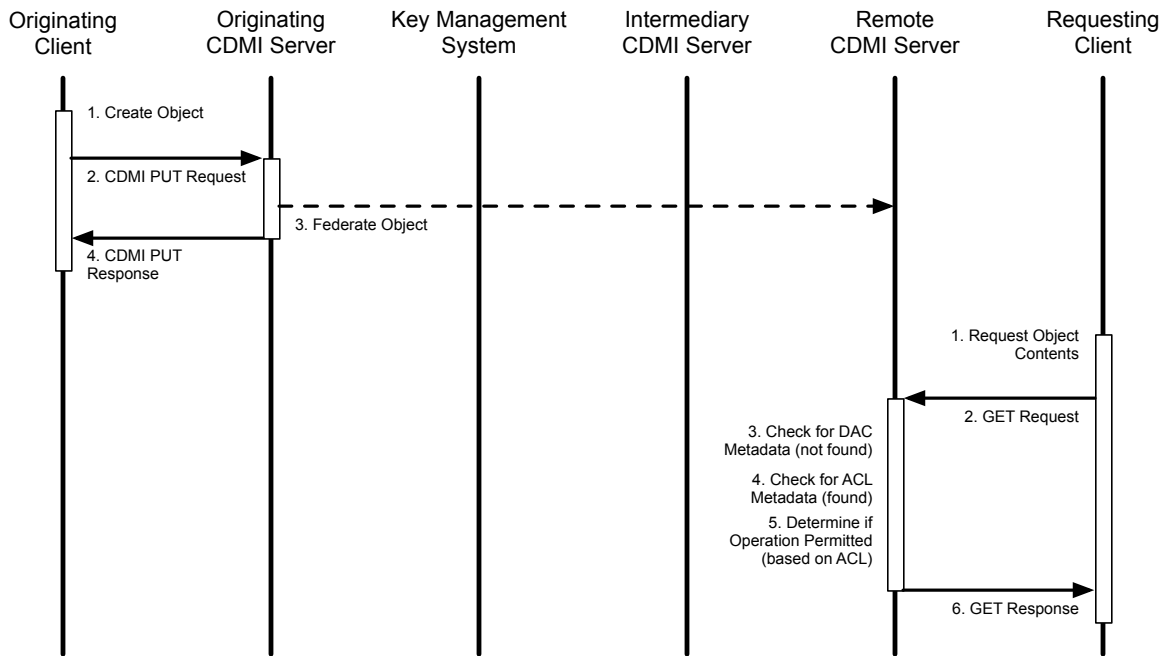


Figure 16: Non-delegated (ACL-based) access control data flow

When an access control decision needs to be made by a third party (such as by the originating CDMI server in Fig. 16), access control is delegated. When `cdmi_dac_uri` and `cdmi_dac_certificate` object metadata is present, as specified in clause 16.2, Delegated Access Control (DAC) shall be used.

An example of an object with DAC metadata is shown below:

```
{
  "objectType": "application/cdmi-object",
  "objectName": "MyObject.txt",
  "capabilitiesURI": "/cdmi_capabilities/dataobject/",
  "objectID": "0000000800182ADB37303732323136662D343564622D3462",
  "mimetype": "text/plain",
  "metadata": {
    "cdmi_size": "33",
    "cdmi_ctime": "2017-04-05T11:01:25",
    "cdmi_atime": "2017-04-05T11:44:28",
    "cdmi_dac_uri": "https://cloud.example.com/dac/",
    "cdmi_dac_certificate": {
      "kty": "EC",
      "x": "goqhRgM4hyEhlp-fD1oU15QAgdKXsBZTQ_0B-IgSz6M",
      "y": "cd8RTm8uLTGblIzioAzv8dzIkM85c08o23eksJrDt2Y",
      "crv": "P-256"
    }
  },
  "valueTransferEncoding": "utf-8",
  "valueRange": "33",
  "value": "This is an unencrypted text file."
}
```

The process by which objects are federated between systems is outside the scope of access control delegation and involves how objects are replicated, synchronized, mirrored, or migrated between CDMI servers. These processes are typically under the control of policies or external policy management systems. Federation is typically performed by third-party systems that use CDMI features including notification, serialization, and the preservation of globally unique object identifiers, which forms the basis for client-transparent interoperability.

## 24.2 Delegated access control (DAC)

A cloud storage system may implement support for DAC, which is indicated by the presence of the `cdmi_dac` system-wide capability.

DAC enables requests for operations against an object to be allowed or denied by a third-party DAC provider, in addition to ACL access control. When required by object metadata, DAC access control verification shall be performed after ACL evaluation, but before ACL enforcement, as the DAC provider may overrule local ACL evaluation results. When an encrypted object is accessed, the DAC provider may provide the decryption key. The decryption key enables access to encrypted objects, even if the CDMI server cannot access the keys directly.

Clients often have different degrees to which they trust the CDMI server with which they are interacting. Table 153 describes the four ways that DAC shall interact with stored objects.

Table 153: Access modes for DAC

Mode of access	Degree of trust
Client-side decryption	<p>CDMI server is not trusted with keys or to make delegated access control decisions.</p> <ol style="list-style-type: none"> <li>1. Client requests encrypted object from CDMI Server</li> <li>2. Client receives ciphertext from the CDMI Server</li> <li>3. Client is responsible for getting decryption keys out of band</li> <li>4. Client verifies signatures (if present)</li> <li>5. Client verifies correct object</li> <li>6. Client decrypts object</li> </ol> <p>This mode of access does not use any functionality indicated by the <code>cdmi_dac_capability</code> and is supported by all CDMI servers.</p>
Client-side decryption with DAC	<p>CDMI server is not trusted with keys and is used to establish an opaque channel of communication between the client and the DAC provider for key delivery.</p> <ol style="list-style-type: none"> <li>1. Client requests encrypted object from the CDMI Server, and includes custom DAC headers specifying information required for secure delivery of decryption key</li> <li>2. Client receives ciphertext from the CDMI Server, along with custom DAC header from the DAC provider for the decryption key</li> <li>3. Client extracts decryption key from DAC provider headers</li> <li>4. Client verifies signatures (if present)</li> <li>5. Client verifies correct object</li> <li>6. Client decrypts object</li> </ol> <p>This mode of access requires the <code>cdmi_dac</code> capability but does not require encrypted object support.</p> <p>In this mode, data is exchanged between the client and the DAC provider using one or more “CDMI-DAC-” headers, as described in <a href="#">clause 24.4</a>.</p>
Direct Client DAC	<p>CDMI server is not trusted with keys, and client establishes channel of communication between the client and the DAC provider for key delivery.</p> <ol style="list-style-type: none"> <li>1. Client requests encrypted object from CDMI Server</li> <li>2. Client receives ciphertext from CDMI Server</li> <li>3. Client sends DAC request directly to DAC Provider</li> <li>4. Client receive DAC response directly from DAC Provider</li> <li>5. Client verifies signatures (if present)</li> <li>6. Client verifies correct object</li> <li>7. Client decrypts object</li> </ol> <p>This mode of access requires the <code>cdmi_dac</code> capability but does not require encrypted object support.</p>

Continued on next page

Table 153 – continued from previous page

Mode of access	Degree of trust
Server-side decryption with DAC	<p>CDMI server is trusted with keys and to delegate access control decisions. DAC message exchange is used to get the decryption keys to decrypt the contents of the object, and keys are not revealed to the client.</p> <ol style="list-style-type: none"> <li>1. Client requests encrypted object from CDMI Server</li> <li>2. CDMI server contacts the DAC Provider to determine access control decision and gets decryption keys, where the keys are not revealed to the client.</li> <li>3. CDMI server verifies signatures (if present)</li> <li>4. CDMI server verifies correct object</li> <li>5. CDMI server decrypts object</li> <li>6. Client receives plaintext</li> </ol> <p>This mode of access requires DAC and encrypted object support.</p>
Plaintext objects with DAC	<p>CDMI server is trusted with plaintext and to not bypass delegated access control decisions.</p> <ol style="list-style-type: none"> <li>1. Client requests non-encrypted object from CDMI Server</li> <li>2. CDMI server contacts DAC provider to determine access control decision</li> <li>3. CDMI server verifies signatures (if present)</li> <li>4. CDMI server verifies correct object</li> <li>5. Client receives plaintext</li> </ol> <p>This mode of access requires DAC support.</p>

The `cdmi_dac_uri` metadata item indicates where delegated access control requests shall be submitted, and the `cdmi_dac_certificate` metadata item indicates how securely communication with the delegated access control provider shall be established. Both of these metadata items shall be present for DAC to be enabled for a given object.

DAC requests are submitted to a DAC provider using two typical methods:

- **Direct** - The DAC request shall be submitted directly to the absolute URI specified in the `cdmi_dac_uri` metadata item. This approach requires the host specified in the URI to be accessible from the CDMI server, and for the CDMI server making the request to have sufficient permissions to PUT the DAC request to that location.
- **Indirect** - The DAC request shall be sent to the DAC provider using an indirect route. Indirect routing is useful when the `cdmi_dac_uri` does not specify a host. An example of indirect routing is when the `cdmi_dac_uri` contains a mailto URI; the Internet mail system is then responsible for delivering the DAC request.

In other cases, the certificate included with the DAC request (taken from the `cdmi_dac_certificate` metadata) may be used by intermediary CDMI servers to determine the further routing of the DAC request. For example, DAC requests using a E.U.-issued certificate can be forwarded to a different intermediary CDMI server to those requests using a U.S.-issued certificate. How certificate fields are used to determine routing is not defined in this International Standard.

Both direct and indirect routing may be synchronous or asynchronous. If a DAC response is not received within the CDMI server or client timeout windows, the client request may time out; however a subsequent request may be processed locally if the DAC response allows response caching. When the CDMI server times out while waiting for a DAC response, it shall return an HTTP status code of 504 Gateway Timeout.

## 24.3 Delegated access control message exchange

When a client requests to access or modify an object containing DAC metadata on a CDMI server that supports DAC, the CDMI server shall create and send a DAC request as specified in [clause 24.5](#). Upon receiving a DAC response as specified in [clause 24.7](#), the CDMI server shall allow or deny the operation based on the contents of the response.

[Fig. 17](#) provides an example of access control delegation for a non-encrypted object. The black solid lines show indirect routing, and gray dashed lines show direct routing.

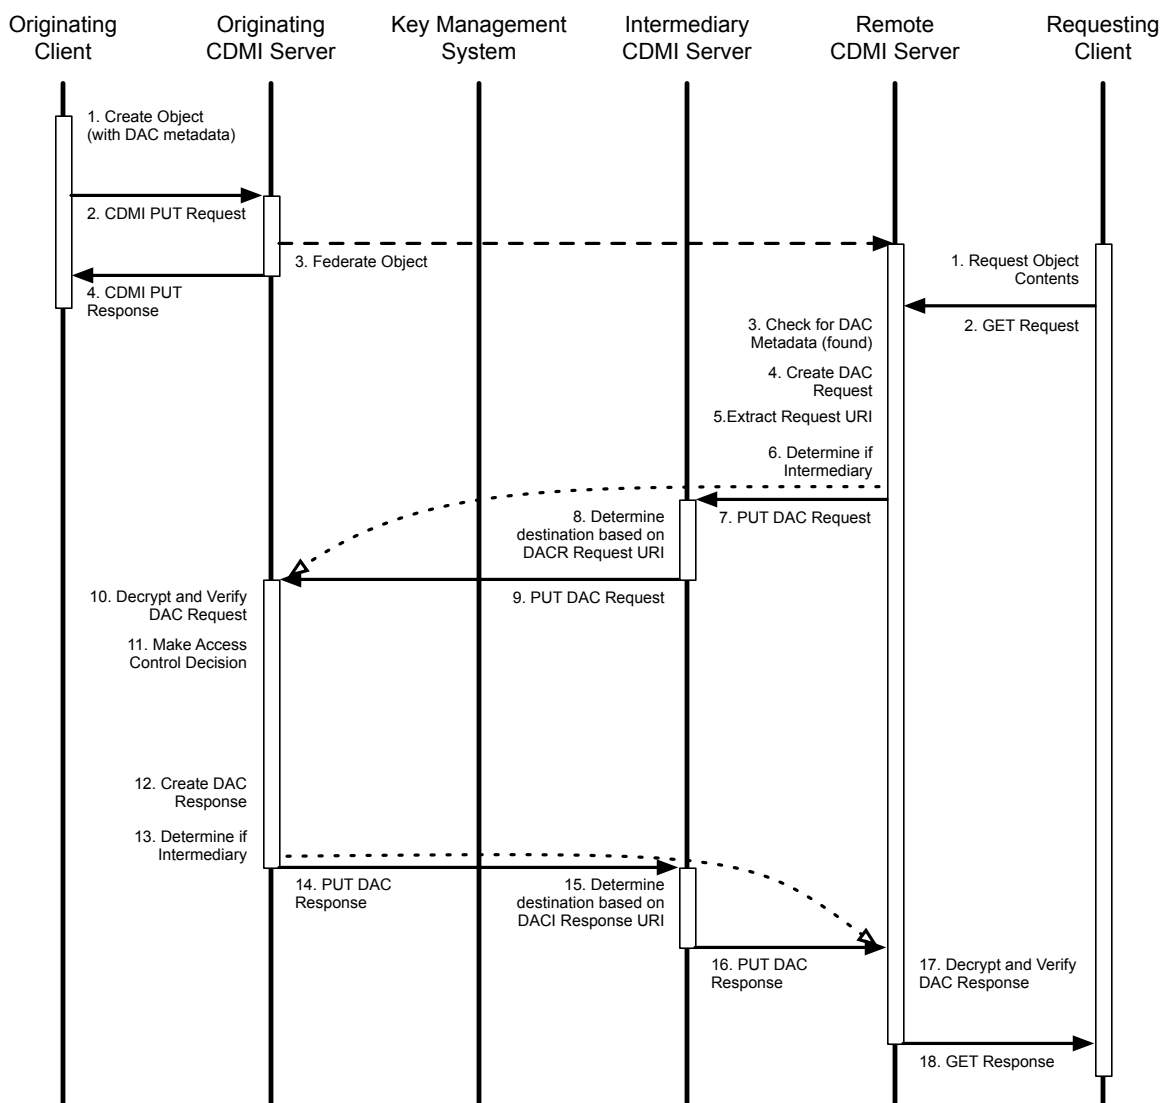


Figure 17: Delegated access control data flow example for non-encrypted object

For non-encrypted objects, an originating client indicates that DAC is requested by including the DAC metadata items. It is important to emphasize that for non-encrypted objects, DAC cannot be guaranteed to be enforced, as when an object with DAC metadata is accessed from a CDMI server that does not support DAC; only ACL-based access control shall be evaluated.

3798 Fig. 18 provides a second example of access control delegation for an encrypted object. The black solid lines show  
 3799 indirect routing, and gray dashed lines show direct routing.




Figure 18: Delegated access control data flow example for encrypted object

3800 For encrypted objects, as access to the decryption keys are provided in the DAC response, the plaintext is inaccessible  
 3801 unless the CDMI server supports DAC.

3802 When the DAC provider processes the DAC request, if the operation is allowed and the key is requested by the CDMI  
 3803 server, the object key, if present, shall be obtained and sent back as part of the DAC response. Upon receiving the DAC  
 3804 response, the CDMI server shall extract the key to perform the client operation.

## 24.4 Client header passthrough

The Delegated Access Control extension provides facilities to allow client-provided HTTP request headers to be passed through to the DAC provider, and for the DAC provider to pass HTTP response headers back to the client. These headers are identified by the “CDMI-DAC-” prefix.

The contents and full names of these headers are not defined in this International Standard. However, it is anticipated that these headers shall be used to allow the client to provide additional information that may be required for the access control decision-making process, for audit purposes, or for secure key exchange.

For example, when an operation is allowed by a DAC provider, the object key may be encrypted using the public key from a client-provided certificate (verified by the DAC provider), which is included in a “CDMI-DAC-” request header, with the encrypted object key being sent back to the client in a “CDMI-DAC-” response header. In this scenario, the CDMI server cannot decrypt the ciphertext but can securely pass on the encrypted object key to the client. The client can then use its private key to decrypt the response header to get the object key, which can then be used to decrypt the object.



## 24.5 DAC request

When a CDMI server that supports DAC needs to contact the DAC provider as specified in the DAC metadata, it shall construct a DAC request, as specified below:

Table 154: DAC request

Field name	Type	Description	Requirement
dac_request_version	JSON string	Indicates the version of the DAC request. This field shall be set to the value "1".	Mandatory
dac_request_id	JSON string	Contains a system-specified identifier that is used to match up the corresponding DAC response. This identifier shall be unique within the window that multiple DAC responses may be received.	Mandatory
server_identity	JSON object	A JSON object, containing a JWE JWK which shall include a public key that is used to submit a DAC response, and should contain a X.509 certificate or certificate chain used to verify the identity of the CDMI server that is generating the DAC request. This ensures that only the CDMI Server that generated the DAC request can read the DAC response.	Mandatory
client_identity	JSON object	A JSON object containing the following JSON entities: JSON String, "acl_name", containing the ACL name of the client requesting the operation. JSON Array, "acl_group", containing the ACL group(s) of the client requesting the operation.	Optional
acl_effective_mask	JSON string	A text or hexadecimal string representation of the ACE mask determined by ACL evaluation for the requested operation, as defined in 17.2.5.	Mandatory
client_headers	JSON object	A JSON object containing a JSON string for each HTTP header in the operation request that starts with "CDMI-DAC-", where the JSON string name is the header name, and the JSON string value is the header value. These headers can be used for tunneling information from the client to the DAC provider.	Mandatory
cdmi_objectID	JSON string	Contains the object ID of the object the operation is performed against.	Mandatory
cdmi_enc_key_id	JSON string	Contains the encryption key identifier (for example, a KMIP identifier) for the symmetric key that is used to encrypt and decrypt the object, which is used to indicate that the CDMI server is requesting the encryption key.	Optional
cdmi_operation	JSON string	Contains a string indicating which operation is being requested to be performed against the object. The following operations are defined: <ul style="list-style-type: none"> <li>"cdmi_read"</li> <li>"cdmi_modify"</li> <li>"cdmi_delete"</li> </ul>	Mandatory
dac_response_uri	JSON string	An optional URI that specifies where to send the DAC response. This URI is required for asynchronous DAC requests, such as when sent via email URIs. If this field is omitted, the DAC response shall be based on the context of the request, for example, as a message body returned for the request PUT when using HTTPS, or an email reply when using a mailto URI.	Optional

3821

An example of a DAC request is shown below:

```
{
  "dac_request_version": "1",
  "dac_request_id": "037130fa-da72-44f0-8a31-62073263ac95",
  "server_identity": {
    "kty": "EC",
    "x": "joyfi05KEI3hcOhJeOfny_TWsz9FFS1zUydFQhm3G78",
    "y": "Nsk3jXlph0FH8APR2k0XSu6pDZYyF7f_Okplf7hZ_8k",
    "crv": "P-256"
  },
  "client_identity": {
    "acl_name": "anonymous",
    "acl_group": ["users"]
  },
  "acl_effective_mask": "READ_ALL",
  "client_headers": {
    "cdmi-dac-header1": "This is a test header"
  },
  "cdmi_objectID": "0000000800182ADB37303732323136662D343564622D3462",
  "cdmi_operation": "cdmi_read"
}
```

## 24.6 Packaged DAC request

The above JSON (DAC request) shall be encrypted in JWE format, where the recipient is the public key of the DAC provider certificate (as specified in the DAC object `cdmi_dac_certificate` metadata), and is JWS signed using the private key of the CDMI server that corresponds to the server identity certificate included in the DAC request. The certificate of the DAC provider from the object is then attached:

Table 155: Packaged DAC request

Field name	Type	Description	Requirement
<code>dac_request</code>	JSON object	JOSE encrypted and signed request	Mandatory
<code>dac_request_dest_certificate</code>	JSON object	The <code>cdmi_dac_certificate</code> metadata value, which is used to indicate where the DAC request is being sent via indirect routing.	Mandatory
<code>dac_request_dest_uri</code>	JSON string	The <code>cdmi_dac_uri</code> metadata value, which is used to indicate where the DAC request is being sent via direct routing, or used to indicate the first location when being sent via indirect routing.	Mandatory

An example of a packaged DAC request is shown below<sup>1</sup>:

```
{
  "dac_request": {
    "protected": "
      eyJqd2siOiJ7XCJrdHlcIjpcIkVdXCIsXCJ4XCI6XCJqb3lmaTA1S0VJM2hjT2hK
      ZU9mbnlfVfDzWjlGRlMxelv5ZEZRaG0zRzc4XCIsXCJ5XCI6XCJOC2szalxcGgw
      Rkq4QVB3SMmswWFN1NnBEWl15RjdmX09rcGxmN2haXzhrXCIsXCJjcnZcIjpcIlAt
      MjU2XCJ9IiwiYWxnIjoiriVMYNTYifQ",
    "payload": "
      eyJwcm90ZWNOZWQioiOiJleUpoYkdjaU9pSkZRMfJFTFVWVElpd2laVzVqSWpvaVFU
      STFOa2REVFVNjC0ltVndheUk2ZXlKcmRlA2lPaUpGUXlJc0luZ2lPaUpuUkZOek1F
      cFRxBu5VVVRsWGVGWXRiRXhSVtJ4elFsY3lXazFvTmlkb1JrcDJTVmt4TWt4dlld
      TlJJaXdpZVZNJnklSkZkZVp1pXWkVkalZtdGtPVVZCvmpaVGYeE9NVZQyUkdsdlpV
      dHVZV3BLWmpsdsWVFOVlJrRlpoYmtFaUxDSmpjb1lpT2lKUUXUSTFOaUo5ZlEiLCJl
      bmNyeXB0ZWRfa2V5IjoIiIiwiaXYiOiJLRDlGRlBOcFh2cWNIYTdIIiwiY2lwaGVy
      dGw4dCI6Im42NlpmUzBXRmhjN3ZzT3Rnc1o5SXJtWU5paDI4RDVzT1psTk96dEdO
      tV5h4FRSGZTMGozMzUhrMUxPME9IbFBYMNvfyXVXtcWm2aDF2ZlIxSfNlOE13TmFq
      TFZfs29ZMndGXzlkadRtWfJlVXA4R1hpbm05MFE0ZWZmY1BLRml1cEo0dE94TTVS
      VjlLN2VvdWnXSkxzczJKbHclZUJhOVQ5WjFyS1pvQmIxvURSLSVVMRW9lQlNZRFA3
      NU1lSEFRSUW4UW5qOW04QjFhb18tNTFPNndKb2d6cHh5Ulhpd3g2SWdoYl1hSYmNX
      MWQ5bVVRtZkR3UFBoSe4zTUplUGUxbVBpelNLWnJ3NWNQM2lNZmhKWmNoT3gyZkZt
      Q3NMME5zSkphQWo3WEs0elFiMGVBd0RSS1BzeTJ6MnZCZzFQT1lhUHppOVphNjRK
      RHgyZ3hWRTA2Y0xERGx3TXY4dW9CbFUlTVdyZF9YRGdScUZsSFl1T19aZEtXqKRp
      MVQ1SW5HeDc2YzdCcmVobzFibnVqV200M0FsanpPRmIyTHBhdU5PQnlET19oVXFi
      WGRlSTUOWNZBUdU0MzVteHhZDRl1SYUpMZGxvFUEhNEhXfNdyloRGxOMeTfYUfF
      U0JtZUZtem5ITGfKzUNDYzI3cWNUO1ZVlZBMHhZMZYV2N2xzbmZMY3VyOHIOOflT
      SXRmZGNZbFVOLTh2c0xhSlZzbHhMSzc0VjdjdWnHbFNUbWJvYkTWTvV6TnZuU29K
      NHpldXBYZzItb192WnIwbkZlSUFWelIxZmJvUVA0clF4bXNSUWJNY2d4bpmSM21E
      eTJsQzY5dFN1TDJGYmlqUnZiYWM3XzFRA01CiwidGFniIjoiriVMYNTYifQ",
    "signature": "
      rGz9Cku3csTIj_p3qmHzUrPSLb1ZSD3ZlfaJDw0F-dNmJs6sgzizFC_jf5VgDVuo
      GT-wH2b2zVuP_O1HdCkPDQ"
  },
  "dac_request_dest_certificate": {
    "kty": "EC",
    "x": "goqhRgM4hyEhlp-fDloU15QAgdKXsBZTQ_0B-IgSz6M",
    "y": "cd8RTm8uLTGblIzioAzv8dzIkM85c08o23eksJrDt2Y",
    "crv": "P-256"
  },
  "dac_request_dest_uri": "https://cloud.example.com/dac/"
}
```

Once created, the packaged DAC request shall be submitted using the DAC request URI specified in the DAC ob-

<sup>1</sup> Decrypt with "d": "NnU0IEyV4JSyLoKwIzKNlFAXDvL6qqawAHlPkpwbMSY".

3829 ject metadata, for example, as an HTTP PUT operation of type “application/json”, or via an SMTP email. The  
3830 `dac_request_dest_certificate` and `dac_request_dest_uri` may be used to route the request through inter-  
3831 mediary hops, as needed.

## 24.7 DAC response

When a DAC provider receives a DAC request, it shall decrypt the request using its private key, verify the signature of the CDMI server, and shall evaluate the request. Based on the information provided, the DAC provider shall allow or deny operations by modifying or replacing the ACL mask that was initially determined by the CDMI server.

To indicate the result of the DAC request to the requesting CDMI server, the DAC provider shall construct a DAC response, as specified below:

Table 156: DAC response

Field name	Type	Description	Requirement
dac_response_version	JSON string	Indicates the version of the DAC response. This field shall be set to the value "1".	Mandatory
dac_response_id	JSON string	Contains the system-specified identifier specified in the corresponding <code>dac_request_id</code> .	Mandatory
dac_applied_mask	JSON string	A text or hexadecimal string representation of the ACE mask that shall be used, as defined in 17.2.5.	Mandatory
dac_object_key	JSON object	The key for the object in JWK format (See RFC 7517 [16]). This key is only disclosed when <code>cdmi_enc_key_id</code> is included in the DAC request and the DAC provider allows access.	Optional
dac_response_headers	JSON object	A series of headers that start with "CDMI-DAC-" to be returned to the client. These headers can be used to pass information from the DAC provider back to the client.	Optional
dac_key_cache_expiry	JSON string	The complete date/time when the object key is no longer to be cached, specified in ISO 8601 date/time format. If this field is not included, the key shall not be cached.	Optional
dac_response_cache_expiry	JSON string	The complete date/time when the DAC response is no longer to be cached, specified in ISO 8601 date/time format. If this field is not included, the response shall not be cached.	Optional
dac_redirect_objectID	JSON string	Indicates an alternate CDMI Object ID used to access the requested object. If present, the CDMI server shall send an HTTP Redirect to the client.	Optional
dac_audit_uri	JSON string	Indicates a URI to a CDMI queue where audit logging messages associated with the operations shall be submitted. When present, audit logging messages shall be generated for receiving the response, performing the operation, and determining when to purge the key. The format of these audit messages is not defined by this International Standard.	Optional

An example of a DAC response is shown below:

```
{
  "dac_response_version": "1",
  "dac_response_id": "037130fa-da72-44f0-8a31-62073263ac95",
  "dac_applied_mask": "ALL_PERMS",
  "dac_response_headers": {
    "CDMI-DAC-AuthInfo": "No key requested."
  },
  "dac_response_cache_expiry": "2017-04-06T15:06:01.554Z"
}
```

## 24.8 Packaged DAC response

The above JSON (DAC response) shall be encrypted in JWE format where the recipient is the public key of the CDMI server certificate (as specified in the DAC request), and is JWS-signed using the private key of the DAC provider that corresponds to the DAC provider identity certificate associated with the object (`cdmi_dac_certificate`), or with a different signing, included in a `jku/jwk/x5u` or `x5c` JOSE header to allow retrieval of the public signing verification key.

The certificate of the CDMI server is then attached:

Table 157: Packaged DAC response

Field name	Type	Description	Requirement
<code>dac_response</code>	JSON object	JOSE encrypted and signed response	Mandatory
<code>dac_response_dest_certificate</code>	JSON object	The contents of the DAC request <code>server_identity</code> field.	Mandatory
<code>dac_response_dest_uri</code>	JSON string	The contents of the DAC request <code>dac_response_uri</code> field, if present	Optional

An example of a packaged DAC response is shown below<sup>2</sup>:

```
{
  "dac_response": {
    "protected":
      "eyJqd2siOiJ7XCJrdHlcIjpcIkVdXCIsXCJ4XCi6XCJnb3FoUmdNNGh5RWgxcC1m
      RDEvVTE1UUFnZEtYc0JaVFFfMEItSWdTejZNXCIscXJ5XCi6XCJjZDhSVG04dUxU
      R2JsSXppb0F6djhkeklrTTglYzA4bzIzZWtzSnJEdDJZXCIsXCJjcnZcIjpcIlAt
      MjU2XCJ9IiwiaWxnIjoiaRVMyNTYifQ",
    "payload":
      "eyJwcm90ZWNOZWQioiOiJleUpoYkdjaU9pSkZRMFJJTFVWVlpd2laVzVqSWpvaVFU
      STFOa2REVFNjc0ltVndheUk2ZXlKcmRlIa2lPaUpGUX1Jc0luZ2lPaUpNVUVReWRX
      WmlkMUpmT0hoU2FWRlRNMWN3YUzSbU5tWn1XWEZDU0hWYU4xQTVUbEEzVFdaVFEy
      MDRJaXdpZVNJNklqWmhiMWgxUzJFeVYqZHNMTW93YlU5U1JUQmF1V0pQU2pKWlYy
      bzNOM1l3Wm5GWU1ESnBiRE5EVUVVaUxDSmpjb1lpT2lKUUXUSTFOaUo5ZlEiLCJl
      bmNyeXB0ZWRfa2V5IjoiaWwiaXJmMEh1wjlxbk5aOHY0d1JUaXBGS0RakpvlUhXOG82bzlm
      dGV4dCI6Ikd5dQXE1dnBCeUVaVERJcH1Wem5Gemxtbm1JU09sVks5uNGpSUUtKWjB5
      c0s0dzZPcDYtNE94cGtvVYy5WUFvbDhmdUVOeFFMdjFBQUpDWXB3M0ZFeRrMEpG
      VmU1NWEOU1NlVkhNslJmMEhiwjlxbk5aOHY0d1JUaXBGS0RakpvlUhXOG82bzlm
      czV2YmRVtGJPRk9Db3RTTGZuekdSQ3lMV3Z2TUZaSa3BHxzM1b21PeFpNcW1oN2Ro
      c3IxMmF6cHdkSnJKX084TTFkVHdDaWZxeURLWWFpNGM4M3U4TUhieDdETldRWkhH
      QnIzTlJ0bDhaWGJTQW90Q09fVWRpdU8zWXZmWmNiWU51TTY2UXBZbDFobENSaDJO
      eEZtLW12VUR0a1VoaxR5cTdyZ3BSbWZoYndKNklCaGdpdyIsInRhZyI6Ijh3YWx6
      T0Q4U3hWTC1STXY3OXlTZGcifQ",
    "signature":
      "8-09XlWUUDsXXqoEh5EKIAYEOTR-vtAYqauW1aNfdv2Io9B4RCuAL13zi7i27vbo
      TYvHxnFa7K6HJPYgsAvN5g "
  },
  "dac_response_dest_certificate": {
    "kty": "EC",
    "x": "joyfi05KEI3hcOhJeOfny_TWsz9FFS1zUydFQhm3G78",
    "y": "Nsk3jXlph0FH8APR2k0XSu6pDZYyF7f_Okplf7hZ_8k",
    "crv": "P-256"
  }
}
```

Once created, the packaged DAC response shall be returned as the response to the HTTPS/HTTP request, or submitted using the DAC response URI specified in the DAC request, for example, as an HTTP PUT operation or via an SMTP email. The `dac_response_dest_certificate` and `dac_response_dest_uri` may also be used to route the request through intermediary hops if needed, as determined by the routing system, which is out of scope of this standard.

When the CDMI server receives a packaged DAC response message, it shall decrypt it using its private key and shall verify the signature. If the decryption and signature verification are successful, the CDMI server shall use the provided `dac_applied_mask` in place of the ACL computed mask.

If the CDMI server supports key or DAC response caching, cache expiry values shall be honored. Cached responses and keys can only be used for identical client operations, where the client identity, objectID, operation, and "CDMI-DAC-"

<sup>2</sup> Decrypt with "d": "huCoV1iC24rZ3uF5q-1HHIGb2Uc6Ue9oNezEQNZUB8"

3855 request headers are identical. Otherwise, the cached response shall be expired. If an audit URI is present in the cached  
3856 response, audit messages shall also be generated for all operations allowed using the cached response.

3857 The CDMI server shall also implement audit logging when specified in the DAC response. If the CDMI server does not  
3858 support audit logging and it is required by a DAC response, the operation shall be denied.

3859 If a `dac_redirect_objectID` field is returned in the DAC response, the CDMI server shall return an HTTP redirect  
3860 to the specified Object ID. This redirect allows a DAC provider to create a client-operation-specific instance of the object  
3861 that is encrypted with a single-use key.

## 24.9 Error handling

In the following scenarios, the following HTTP response codes shall be returned to a client:

- When a DAC response denies the requested operation, an HTTP status code of 403 `Forbidden` shall be returned to the client along with any `dac_response_headers` included in the response.
- When a DAC response includes a `dac_redirect_objectID`, an HTTP status code of 302 `Found` shall be returned to the client along with any `dac_response_headers` included in the response.
- When a DAC request to access or modify an encrypted object is allowed, but the key is not included in the DAC response, an HTTP status code of 401 `Unauthorized` shall be returned to the client along with any `dac_response_headers` included in the response.
- When a DAC request to access or modify an encrypted object is allowed, but cannot be performed due to lack of support for an encryption algorithm, signing algorithm, or key type, an HTTP status code of 501 `Not Implemented` shall be returned along with any `dac_response_headers` included in the response.
- When a DAC request times out, an HTTP status code of 500 `Internal Server Error` shall be returned to the client.
- When a DAC request cannot be sent or routed because the DAC metadata is not supported or valid, an HTTP status code of 501 `Not Implemented` shall be returned to the client.
- When a DAC request cannot be sent or routed because an upstream system is unavailable, an HTTP status code of 500 `Internal Server Error` shall be returned to the client.



## 24.10 Examples

The following examples illustrate the primary ways that DAC requests are performed.

### EXAMPLE 1: GET ciphertext of encrypted object with delegated access control

The following CDMI operation is performed against an encrypted CDMI object with delegated access control metadata:

```
GET /MyContainer/MyEncryptedObject.txt HTTP/1.1
Host: cloud.example.com
Accept: application/cms, application/jose+json
```

The CDMI server verifies local access controls and determines that the request can proceed. The following DAC request is generated:

```
{
  "dac_request_version": "1",
  "dac_request_id": "5b801b19-479e-446d-882a-8483f7c4905c",
  "server_identity": {
    "kty": "EC",
    "x": "joyfi05KEI3hcOhJeOfny_TWsz9FFSlzUydFQhm3G78",
    "y": "Nsk3jXlph0FH8APR2k0XSu6pDZYyF7f_Okplf7hZ_8k",
    "crv": "P-256"
  },
  "client_identity": {
    "acl_name": "anonymous",
    "acl_group": ["guest"]
  },
  "acl_effective_mask": "READ_ALL",
  "client_headers": {},
  "cdmi_objectID": "0000000800182F9E64313363323731622D363536662D3465",
  "cdmi_operation": "cdmi_read"
}
```

This request is first JWE encrypted with the key in `cdmi_dac_certificate`. The result is JWS signed, using either the key in `server_identity`, or a different key embedded in the JWS header.

The DAC provider verifies, decrypts and processes the request and returns the following DAC response:

```
{
  "dac_response_version": "1",
  "dac_response_id": "5b801b19-479e-446d-882a-8483f7c4905c",
  "dac_applied_mask": "ALL_PERMS",
  "dac_response_headers": {
    "CDMI-DAC-AuthInfo": "No key requested."
  }
}
```

The `CDMI-DAC-AuthInfo` indicates a custom header.

Since the operation is allowed by the DAC provider, the following response is sent:

```
HTTP/1.1 200 OK
Content-Type: application/jose+json
Content-Length: 290
CDMI-DAC-AuthInfo: No key requested.

<JOSE+JSON Encrypted Object>
```

### EXAMPLE 2: GET ciphertext of encrypted object with passthrough key access

The following CDMI operation is performed against an encrypted CDMI object with delegated access control metadata:

```
GET /MyContainer/MyEncryptedObject.txt HTTP/1.1
Host: cloud.example.com
Accept: application/cms, application/jose+json
Authorization: Basic am9lOnBhc3N3b3Jk
CDMI-DAC-N: <vendor-specific header that indicates key passthrough>
```

The CDMI server verifies local access controls and determines that the request can proceed. The following DAC request is generated. The `CDMI-DAC-N` is a custom header that indicates that the client wants to obtain the object decryption key via header pass-through.

To demonstrate the power of such custom headers: the `CDMI-DAC-N` request header could contain a cell phone number. The matching response header would then contain a password-based encryption of the object key, while the password will be delivered via a message to the cell phone. It is up to the vendor to come up with and implement such mechanisms.

```
{
  "dac_request_version": "1",
  "dac_request_id": "77b54650-183f-4053-8512-be08f7c6c50e",
  "server_identity": {
    "kty": "EC",
    "x": "joyfi05KEI3hcOhJeOfny_TWsZ9FFS1zUydFQhm3G78",
    "y": "Nsk3jX1ph0FH8APR2k0XSu6pDZYyF7f_Okplf7hZ_8k",
    "crv": "P-256"
  },
  "client_identity": {
    "acl_name": "joe",
    "acl_group": ["users"]
  },
  "acl_effective_mask": "READ_ALL",
  "client_headers": {
    "CDMI-DAC-N": "<copy from headers>"
  },
  "cdmi_objectID": "0000000800182F9E64313363323731622D363536662D3465",
  "cdmi_operation": "cdmi_read"
}
```

This request is first JWE encrypted with the key in `cdmi_dac_certificate`. The result is JWS signed, either using the key in `server_identity`, or a different key embedded in the JWS header. Replication of these encrypted messages is not useful and will be skipped.

The DAC provider processes the request, obtains the object decryption key and embeds it as a `dac_response_header`, then returns the following DAC response:

```
{
  "dac_response_version": "1",
  "dac_response_id": "5b801b19-479e-446d-882a-8483f7c4905c",
  "dac_applied_mask": "ALL_PERMS",
  "dac_response_headers": {
    "CDMI-DAC-AuthInfo": "Key successfully retrieved from keyserver.",
    "CDMI-DAC-N": "<vendor-specific decryption key info>"
  }
}
```

Since the operation is allowed by the DAC provider, the following response is sent:

```
HTTP/1.1 200 OK
Content-Type: application/jose+json
Content-Length: 290
CDMI-DAC-AuthInfo: Key successfully retrieved from keyserver.
CDMI-DAC-N: <vendor-specific decryption key info>

<JOSE+JSON Encrypted Object>
```

The client can now parse the key in the `CDMI-DAC-N` header and use it to decrypt the ciphertext.

#### EXAMPLE 3: GET plaintext of encrypted object with delegated access control

The following CDMI operation is performed against an encrypted CDMI object with delegated access control metadata:

```
GET /MyContainer/MyEncryptedObject.txt HTTP/1.1
Host: cloud.example.com
Accept: */*
Authorization: Basic am9lOnBhc3N3b3Jk
```

The CDMI server verifies local access controls and determines that the request can proceed. The following DAC request is generated:

```
{
  "dac_request_version": "1",
  "dac_request_id": "b79d7619-1bbd-45a1-b2d3-5753f7fc5155",
  "server_identity": {
    "kty": "EC",
    "x": "joyfi05KEI3hcOhJeOfny_TWsz9FFS1zUydFQhm3G78",
    "y": "Nsk3jX1ph0FH8APR2k0XSu6pDZYyF7f_Okplf7hZ_8k",
    "crv": "P-256"
  },
  "client_identity": {
    "acl_name": "joe",
    "acl_group": ["users"]
  },
  "acl_effective_mask": "READ_ALL",
  "client_headers": {},
  "cdmi_objectID": "0000000800182F9E64313363323731622D363536662D3465",
  "cdmi_operation": "cdmi_read",
  "cdmi_enc_key_id": "0000000800182F9E64313363323731622D363536662D3465"
}
```

The DAC provider processes the request, obtains the object decryption key and returns the following DAC response:

```
{
  "dac_response_version": "1",
  "dac_response_id": "b79d7619-1bbd-45a1-b2d3-5753f7fc5155",
  "dac_applied_mask": "ALL_PERMS",
  "dac_object_key": {
    "kty": "oct",
    "kid": "0000000800182F9E64313363323731622D363536662D3465",
    "use": "enc",
    "alg": "dir",
    "k": "vBX81leh8ydyI08by7L13kZKNmfRHTAMZa5vJqMCHQU"
  },
  "dac_response_headers": {
    "CDMI-DAC-AuthInfo": "Key successfully obtained from KMS."
  },
  "dac_key_cache_expiry": "2017-04-05T14:58:58Z",
  "dac_response_cache_expiry": "2017-04-05T14:58:58Z"
}
```

Since the operation is allowed by the DAC provider and the key is provided, the object is decrypted by the CDMI server and the following response is sent:

```
HTTP/1.1 200 OK
Content-Type: text/plain
Content-Length: 252

<Decrypted contents of Encrypted Value>
```

#### EXAMPLE 4: RSA Example

In this example, there are two hospitals (A and B), that both have CDMI servers, and federate objects between them. At some point, the following encrypted object has been made at hospital A. It contains a `cdmi_dac_certificate` and `cdmi_dac_uri` that indicate how access can be requested at hospital A. The certificate contains a 2048-bit RSA encryption key, with a matching X.509 certification chain that can be used to verify the certificate.

```
{
  "objectType": "application/cdmi-object",
  "objectName": "MyEncryptedObject.txt",
  "capabilitiesURI": "/cdmi_capabilities/dataobject/",
  "objectID": "000000080018F34436313131393061372D613735302D3438",
  "mimetype": "application/jose+json",
  "metadata": {
    "cdmi_size": "306",
    "cdmi_dac_uri": "https://cdmi.hos-a.fr:9001/dac/",
    "cdmi_atime": "2017-04-06T14:06:34",
    "cdmi_enc_key_id": "encryption_key_1",

```

(continues on next page)

(continued from previous page)

```

"cdmi_dac_certificate": {
  "kty": "RSA",
  "kid": "cdmi.hos-a.fr_encrypt_public",
  "key_ops": [
    "wrapKey",
    "unwrapKey",
    "encrypt",
    "decrypt"
  ],
  "n":
    "uL7ANgD80H5sNqo3nHzovPRxgncQLhz0oQvGMVvULCkrYXMaXZ5sNv7fT6UdMSZi
    T-e0sthapMEqrpeV9RKHSiF3COgl2YndUHixpEkHp8ylggcH6iTzoBsgXMZ70LW-
    m290RCAXDTE2MTAyNzEyNDUwMfoYDzk50TkkMjMxMjM1OTU5WjA1MQswCQYDVQ
    JLstzkP8-cKS0BkEquLQEMbZVRM6U5uG69cj1i9OWvuRzPoaATKyt6Cc4f6PUu9L
    OyCBUAs9dXsRrt3B8H1qe7io7FAAcOpcUDKdNLFXS1Thc37DK_zEyKZcMttjCvEl
    Ovt-cIaokdnxJeggV9AFGQ",
  "e": "AQAB",
  "x5c": [
    "MIDMDCCAhigAwIBAgIBBDANBgkqhkiG9w0BAQsFAADBCMQswCQYDVQ
    MA8GA1UEChMIbGllc2RvbmsxDALBgNVBAsTBGNkbWkxETAPBgNVBAMTCHJzYSly
    b290RCAXDTE2MTAyNzEyNDUwMfoYDzk50TkkMjMxMjM1OTU5WjA1MQswCQYDVQ
    EwJmcjEOMAwGA1UEChMFaG9zLWExFjAUBGNVBAMTDWNBkbnWkuaG9zLWUuZnIwggEi
    MA0GCSqGSIb3DQEBAQUAA4IBDwAwggEKAoIBAQC4vsA2APzQfmw2qjecfOi89HGC
    dxAuHPSHC8Yxw9QsKSthcXpdmw2/t9PpR0xJmJP57Sy2FqmYSqul5X1EoeYIXcI
    4aXZid1QeLgkSQenzLWCbwfgJPogGyBcxns4tb6YnZfBeuh1PuT61yJNhixNIz16
    Aw9BIAQC3Scyc+Z01GcTIDb80nxjuU2QRfKNrRAkuy3OQ/z5wpLQGQSQ4tAQxtlV
    EzpTm4br1yPWL05a+5HM+hoBMRK3oJzh/o9S70s7IIFQCz1lexGu3cHwFwP7uKjs
    UABw61xQmOcsVdLVMdzfsMr/MTIplwy22MK8SU6+35whqiR2fEl6CC/0AUZAgMB
    AAGjPDAMAwGA1UdEwEB/wQCMAAwHQYDVR0OBbYEFBAIGICMR5H6KLKML1ZAEHCCc
    KWE9MASGA1UdDwQEAWIEMDANBgkqhkiG9w0BAQsFAAOCAQEAAANYSSryUU6112pYM
    r83M3GWNjzul6B+4KgmZ8kbey94zNPdmmwQdSe0Xmg+1Otc6VUB40ouNnwK8efB
    aWBtXwCA7Nb715nTqo2+rn+X+A0mGrYaKkToPEe8ZYwDcOlOpNC9JFE+QgP9/CJa
    AaWrf95W+4kra2WnA4Bhqu2WWXnQkL47/nKcGVZgQAH+mVnxPaIogELYdonXU/S2
    8HqxoyjyGL/vmyc46zUbXysgx/jiE7J0fJVP6Yk/3dlNYCCpLtV8VmzFAQAEccn8
    AWowFcd09a4SY09rn1MUv/rrvXpzflfn9j7PtRRFj2e/KhitmOH1zKDuYzREpUOu
    TDlPIQ==",
    "MIIDQDCCAigAwIBAgIBATANBgkqhkiG9w0BAQsFAADBCMQswCQYDVQ
    MA8GA1UEChMIbGllc2RvbmsxDALBgNVBAsTBGNkbWkxETAPBgNVBAMTCHJzYSly
    b290RCAXDTE2MTAyNzEyNDUwMfoYDzk50TkkMjMxMjM1OTU5WjA1MQswCQYDVQ
    EwJubDERMA8GA1UEChMIbGllc2RvbmsxDALBgNVBAsTBGNkbWkxETAPBgNVBAMT
    CHJzYSlyb290MIBIIBIjANBgkqhkiG9w0BAQEFAAOCAQ8AMIIBCgKCAQEAsrUj46dx
    5ojlazzk7YtOL6e+Q6JoG7gVMAxkJn1Sz1x9ND/8w4Pe01SQ2skukd0HA1QRmxdf
    zhccNTM5hmbcn8TAFWSYqQF1R7s78bVjtm6AQPlvSgiyZ8Ak+iYZEq3c2zVyYQ
    HKKXWxmFZt1HT8/H/B3bXveXQcERKE+Tq66h8pqVcocQUtzRFsEYmv0bR1rghtoq
    H8nhB5xnebgVlXjApW+et2SE7r6Fjv1aAbGI89ouJlgsMPeX56P8AUjacFtNkC44
    Obu6HRXY/jm6f2m1EUM84EUsJ+9b5+S2x4qPttdfSCasWYyz4mFJ8MwmFiBGUw
    geT2bUm6t7qqbQIDABoz8wPTAPBgNVHRMBAf8EBTADAQH/MB0GA1UdDgQWBBR+
    tEB2udkEXxX0k15Gztf/4o103jALBgNVHQ8EBAMCAQYwDQYJKoZIhvcNAQELBQAD
    ggEBAIjx1f9rJ2B+mDSA3L2GRhjrPRjfi6Un3Z51CeW9gO9PMQ5ws5pDJyB79de/
    Q8Uf1e8pZyJchTsRa8GRdnKyndN2imayOVUvPoTd3/ZSmfkurcbj3I4VW8sjHP7C
    E8fmUS8Xprdp02SxV7oneJC0vt5eyh8mgfJ/qSbwVaiXuH1Wxi6duAvdxdMXAxQ
    KPG1KKVM7CYfCdpX/HagCOHzcto+374zFqgnQ1Kx5rbgvxNSgm/PDDOMwP03+bbT
    R63KSK1VbdtLBuS4jgaPabwyxQz/FciwTu/HLOQn8TNqDWyoIbs+eQX2Mds2Apul
    8XH2+CakjBLMLL3Tlj2x+6tKR9o="
  ]
},
"cdmi_ctime": "2017-04-06T14:06:31"
},
"valueTransferEncoding": "json",
"value": {
  "protected":
    "eyJraWQwQioiJlbnNyeXB0aW9uX2tleV8xIiwiaWwWxniJjoiQTIiNktXIiwiaWwY3R5Ijoj
    dGV4dC9wbGFpbiIsImVuYyI6IkEyNTZHZHQ0ifQ ",
  "encrypted_key":
    "329yyozEo3JPCpXGPKyI_fa5hhFH9dmfB7kulglQ6NhoVAvdMDMclg",
  "iv": "9Gr5Hxzcs9hxPmPm",
  "ciphertext": "-sJkChcdQUXChEBLzm7UZya1RR2_IcpRocC-BmQfAuA3",
  "tag": "VIFJdMdZngtpLWWDx8vFw"
}

```

(continues on next page)

(continued from previous page)

```
}
}
```

This encrypted object has been federated to the CDMI server at hospital B. Now, one of its clients wants to transparently access the plaintext of this object by performing the following operation:

```
GET /cdmi/MyContainer/MyEncryptedObject.txt HTTP/1.1
Host: cdm1.hos-b.us:9002
Accept: */*
```

The CDMI server at hospital B will look up the object and find out that it is an encrypted object with DAC information attached. As a result it will generate the following (plain) DAC request:

```
{
  "dac_request_version": "1",
  "dac_request_id": "73da04e1-2182-447e-8342-f4b9f06ec936",
  "server_identity": {
    "kty": "RSA",
    "kid": "cdmi.hos-b.us_encrypt_public",
    "key_ops": [
      "wrapKey",
      "unwrapKey",
      "encrypt",
      "decrypt"
    ],
    "n":
      "oQMqkY85Uzw07K6H0QQNfAiRMN3ZfhK0aXEkx7YwvrCU9IKOquZ10YZ9Cv8556_8
      E8yZm02JDWOBoaSSGHU835jvXf12f4MywKGWj5FtIGL-j9kXF6SWq3zuLVY1XpMI
      KsJngHMFca_-ZhZ2vLsrnDR1aCNEC48gR26ewp6WX1ptnSc1W4x3Mj-ONMVzxVE
      7XNlwYysTgDtonmTQD-YG6_KhhAPx0YowMbUWv_cMQvXsi7MMDyZn6fxfq42QmQ2
      V5RtUy5msd6K3beDzS4LmZhsJmjU7YnhOj0pZby4Zckm43npjXPAuwPhzK2OW7qb
      fkv0qm4rsFWUcuNh81BsDw",
    "e": "AQAB",
    "x5c": [
      "MIIDMCCAhiGAWIBAgIBBTANBgkqhkiG9w0BAQsFADBCMqswCQYDVQQGEwJubDER
      MA8GA1UEChMIbGllc2RvbmsxDTBALBgNVBAsTBGNkbWkxETAPBgNVBAMTCHJzYS1y
      b290MCAXDTE2MTAyNzEyNDUwMFOYDzK5OTkxMjMxMjM1OTU5WjAlMQswCQYDVQQG
      EwJlczEOMAwGA1UEChMFaG9zLWlXfjAUBGNVBAMTDWNBkBWkuaG9zLWlUdXMwgGEi
      MA0GCSqGSIb3DQEBAQUAA4IBDwAwggEKAoIBAQCChAxCRjz1TPA7srofrBA18CJEW
      3dl+ErRqcQrHtjC+sJT0go6q5nU5hn0K\znnr\wTzJmbTYkNY4E5pKwYdTzfmO
      9d\XZ\gzLaOzaPkW0gYv6P2RcXpJarfO4tViVekwgqgmeAcxUVxr\5mFna8uy
      ucNHVoI0QLjyBHbp7CnpZfWm2dJzVbjHcyP440xXPFUTtc2XBjKxOA02ieZNAp5g
      br8qGEA\HRijAxtRa\9wxC9eyLswPjmf\F+rjNCZDZX1G1TLmax3ordt4PN
      LguZmGwmanTtieE6PS1lvLhlySbjeemNc8C7A+HMrY5bupt+S\SqbiuwVZRY42H
      yUGwPaGMBAAgJPDA6MAwGA1UdEwEB\wQCMAAwHQYDVROBBYEFH7NjvMIftQtZn
      nyiIdLNkjcGwSIMAsGA1UdDwQEAwIEMDANBgkqhkiG9w0BAQsFAAOCAQEAdiADiv
      0v09SUDcPL+BKysvchn\SGx5KBu7n9KFwE31Dhx2zvT6ruL8kXdekPH9cfrDafW
      6I\vnbzAVj02i5pM2cHayj13fTOWSVwpcQuvkoIF9eVIWONkemMMf7M7jpTw07z
      7S2T5usaDmMNPqj8y5pRpQo3PnBVxpEZJ0XaSdfuiHtVLDq8gDZCq6Hc2tt7JM3W
      njnQgs+1lSGRuqWocpmVONioqvhioLNDZV35Z7puRwqck1N2f1qyHHGBWxfCN9U4
      ci6q1BnWBIFV+hURge8NSbpqawolaNueUbTcKjN3JsMC4ZxhMF9rN3uuPn+UAYka
      yQkcSmGSMm07wcAkMg==",
      "MIIDQCCAigAWIBAgIBATANBgkqhkiG9w0BAQsFADBCMqswCQYDVQQGEwJubDER
      MA8GA1UEChMIbGllc2RvbmsxDTBALBgNVBAsTBGNkbWkxETAPBgNVBAMTCHJzYS1y
      b290MCAXDTE2MTAyNzEyNDUwMFOYDzK5OTkxMjMxMjM1OTU5WjAlMQswCQYDVQQG
      EwJubDERMA8GA1UEChMIbGllc2RvbmsxDTBALBgNVBAsTBGNkbWkxETAPBgNVBAMT
      CHJzYS1yb290MIIBIjANBgkqhkiG9w0BAQEFAAOCAQ8AMIIBCgKCAQEAsrUj46dx
      5oj1aZk7YtOL6e+Q6JoG7gVMAxKJn1SzlX9ND\8w4Pe01SQ2skukdOHALQrmxdF
      tzhcNtM5hmbcn8TAFWSYQqF1R7s78bVjtmata6AQPlvSgiyZ8Ak+iYZEq3c2VYy
      QHKKXWxmFzt1HT8\H\B3bXveXQcERKE+Tq66h8pqVcocQUtzRFsEYmv0bR1rgh
      toqH8nhB5xnebgV1XjApW+et2SE7r6Fjv1aAbGI89ouJlgsMPeX56P8AUjacFtNK
      c440bu6HRXy\jm6f2m1EUM84EUSJ+9b5+S2x4qPttJDfSCasWYyZ4mfJ8MwmFiB
      GUwfgE2t2Um6t7qbbQIDAQABoz8wPTAPBgNVHRMBaf8EBTADAQH\MB0GA1UdDgQ
      WBBR+EB2udkEXX0k15GztF\4o103jALBgNVHQ8EBAMCAQYwDQYJKoZIhvcNAQ
      ELBQADggEBAIjx1f9rJ2B+mDSA3L2GRhjrPRjfi6Un3Z51CeW9gO9PMQ5ws5pDJy
      B79dE\Q8Uf1e8pZyJchTsRa8GRdnKyndN2imayOVUvPoTd3\ZSmfKurbj3i4V
      W8sjHP7CE8fmUS8Xprdp02SxV7oneJC0vt5eyh8mgfJ\qSbwVaiXuH1Wxi6duAv
      dxddMXAxQKPG1KKVM7CYfCdpX\HagCOHzcto+374zFqgnQ1Kx5rbgvxNSgm\PD
```

(continues on next page)

(continued from previous page)

```

        DOMwP03+bbTR63KSK1VbdtLBU54jgaPabwyxQz\FciwTu\HLOQn8TNqDWyoIbs
        +eQX2Mds2Apul8XH2+CakjBLMLL3Tlj2x+6tKR9o="
    ],
    "client_identity": {
        "acl_name": "anonymous",
        "acl_group": ["guests"]
    },
    "acl_effective_mask": "READ_ALL",
    "client_headers": {},
    "cdmi_objectID": "000000080018F34436313131393061372D613735302D3438",
    "cdmi_operation": "cdmi_read",
    "cdmi_enc_key_id": "encryption_key_1"
}

```

This plain DAC request will be JWE encrypted using the key found in the object's `cdmi_dac_certificate` (key id 'cdmi.hos-a.fr\_encrypt\_public'). Then it will be JWS signed using hospital B's private signing key. Since this signing key is not equal to the encryption key (in `server_identity`) it is embedded in the JOSE protected header of the JWS (note: Base64 decode of the protected header reveals the signing key; Base64 decode of the payload reveals the JWE.)

```

{
  "dac_request": {
    "protected":
      "eyJraWQiOiJjZGlpLmhmcyYiLnVzX3NpZ25fcHJpdM0ZSIzImp3ayI6IntcImt0
      eVwiOlwiU1NBXCIsXCJraWRcIjpcImNkbWkuaG9zLWludXNfc2lnb19wcm12YXRl
      XCIsXCJrZX1fb3BzXCi6WlwidmVyaWZ5XCIsXCJzaWduXCJdLWwib1wiOlwicEVR
      aFFUMVF6QmdrV2RiVW56eVkwbkZmWjRVYXJnbFVpcGFxeG1XYXk5cGhnQ0x6Tmtj
      RHZ4eVdIdHFRSWE0ZHpvMDVaXzBiOXhNTElrYU15MTJheV83M1l0ZHpMMWVfaUVX
      Mi1OdVB4MHVSaFV3SzQ4WUo2MFlwVtDpN2ZpQWNKeVJoU1dlWGtnQXQyRndUYnkt
      Sjl5NW9DVldZemRfc0U3a2NMSkc0QmkwSEtQbVhrUEVwbXpOamhsU0VsdnloDHFL
      djRERG1JRk1JTDNrUGJueGNfX0RwenAyaVVPdGhvUFhpY1pJQXMTUDlYbGRGMkRE
      X0tzZW9SU3RQR2NuTEVYbWpKcXhoRU13Qm5UZE14TjdQNNh6bk5iQVNTdXNnR21F
      XzJXdVUyS09yLVBTYm5wTnNlcm14SHRHT2trc2pzdjFyVGhzRmkxNUZmSVQyQ1dU
      MnVRXCI5XCJlXCi6XCJBUUFCCXCI5XCJ4NWNCIjpbXCJNSU1ETURDQ0FoaWdBd01C
      QWdJQkF6QU5CZ2txaGtpRz13MEJBUXNGQURCQ01Rc3dDUVlEVlFRR0V3SnViREV5
      TUE4R0ExVUVDaE1JYkdsbGMyUnZibXN4RFRBTEJnTlZCQXNUQkdOa2JXa3hFVEFQ
      QmdOVk1JBTVRDSEp6WVMxeWlYOTBNQ0FYRFRFMk1UQXl0ekV5TkRvD01Gb1lEems1
      T1RreE1qTXhNak0xT1RVNVdqQTFNUXN3Q1FzFRFzRUUdF0p1YkRFRUk1BOEdBMVVF
      Q2hNRmFHOXpMV014RmpBVUJnTlZCQU1URFdOa2JXa3VhRz16TFdJdWRyTXdnZ0Vp
      TUEwR0NTcUdTSWlZRFFFQkFRVUFBNE1CRHdBd2dnR0tBb01CQVFDa1JDRk1JQVkrN
      R0NSWjF0U2ZQSmpTY1Y5bmnScXVDVlE2bHfYR0packwybUdBSXZNMlJ3Ty9IS1l1
      MnBBaHJoM09qVGxUL1J2M0V3c21Sb0gzWfpyTc92WmcxM04vVjcrSVJiYjQyNC9I
      UzVHRlRBcm42Z25yUmlsVHVMdCtJQnduSkdGSlo1ZVNBQzNZWEJ0dkw0bjJYbWdK
      WlpqTjMrdlR1Undza2JnR0xRY28rWmVROFNTYk0yT0dWSVNXL0tHMm9xL2dNT1ln
      VXdnndmVROXVmRnovoE9ut25hS1NLMkdndOWVKeGtnQ3o0L2JhVjBYWU1QOHF5YWhG
      SzaA4WnljclJlYU1tckdFUXpBR2ROMHFM3MvckhPYzFzQkthNn1BYV1UL1phNVRZ
      bz22NctadWVRMndxdWJFZTFvNlNtE5pLld0T0d3V0xYa1Y4aFBZSlpQYTVBZ01C
      QUfHa1BEQT2NQXdhQTFVZEV3RUIvd1FDTUFBd0hrWURWUjBPQk1JRUZCeHdnVzB4
      TFV3Q1RsaU1TMVZ2di9KvMNPmlFNQXNHQTFVZER3UUVBd01IZ0RBTk1J3Foa21h
      OXcwQkFRc0ZBQU9DQVFFQV15bzMxQmR6N080d21GcFE0eEZja1FkSktSaFBiNndk
      RHdyOTM5OXh5OFOUxV0VFOEpeQ0FvbW9nakJlQ2RLUWVqWE1oVjRlVXk1Y1llpSitj
      WVRxZXh0NTJNb0pJRmUySnozNC9LbFVkyU5ENE5Jdm5teC9mWS83Qk9qQWfky2Ny
      L0NQVUxmczE5OHUybG9GNVUSYwtPM2lGajhwWfY1Mj1DSFFmekpsaGh0c0VoL3p2
      TXgydXpNTlpaWFlkVWxyQ0NBZGFGBwtFRDBHUHM4SGZDR1VXUyt1QVNiN1ZnQXBC
      N0NYb1ZJLzVKA2JzL0ZreEF3TWlxSmE2RUpyYkRkSF10N2prVktwcmVpMw1xRVBj
      Ri9MU0pNZXhGVjJVCdVIsk9OMG1QMgFDclYwbFE2Nys3Mjg2Mkw1VmFjM0tjQTK0
      dVBZVTUxVEJlITFNBVXRLTDI5Uk9oz09XCIsXCJNSU1EUURDQ0FpaWdBd01CQWdJ
      QkFQU5CZ2txaGtpRz13MEJBUXNGQURCQ01Rc3dDUVlEVlFRR0V3SnViREV5TUE4
      R0ExVUVDaE1JYkdsbGMyUnZibXN4RFRBTEJnTlZCQXNUQkdOa2JXa3hFVEFQmdO
      VkJBTVRDSEp6WVMxeWlYOTBNQ0FYRFRFMk1UQXl0ekV5TkRRR01Gb1lEems1T1Rr
      eE1qTXhNak0xT1RVNVdqQkNNUXN3Q1FzFRFzRUUdF0p1YkRFRUk1BOEdBMVVFQ2hN
      SWJhBgxjMlJ2YmlzeERUQUxkZ05WQkFzVEJHTmtiV2t4RVRBUEJnTlZCQU1UQ0hK
      e1l1TmX1liMjkwTU1JQk1qQU5CZ2txaGtpRz13MEJBUXNGQURCQ01Rc3dDUVlEVlF
      R0V3SnViREV5TUE4R0ExVUVDaE1JYkdsbGMyUnZibXN4RFRBTEJnTlZCQXNUQkdO
      a2JXa3hFVEFQmdOQ0FRFRRUzclVqNDZkeDVvamxhWms3WXRPTDZlK1E2Sm9HN2dWTFY
      a0puMVN6bHg5TkQvOHc0UGVPMVNRmNrdWtkt0hBbFFSbXhkZnR6aGNjT1RNNWhT
      YmNuOFRBZlDTWXFRFjFSN3M3OGJWanRtYXQ2QVFMXZT2Z15WjhBaytpWVpFctNjMnp
      WeVlRSEtLeFd4bUZadDFIVDgVSC9CM2JYdmVYUWNFUktFK1RxnNjZoOHBxVmNvY1FVdH
      pSRnNF
    }
  }
}

```

(continues on next page)

(continued from previous page)

```

WW12MGJSMXJnaHRvcUg4bmhCNXhuZWJnVmXyYakFwVytldDJTRTdyNkZqdjFhQWJH
STg5b3VKMWdzTVBLWDU2UDhBVWphY0Z0TktjNDRPYnU2SFJYWS9qbT2mMm0xRVVt
ODRFVXNKKz1iNstTMng0cVB0dEpEZlNDYXNXWV16NG1GSjhNd21GaUJHVXdmZ2VU
MmJvBT20N3Fxy1fJREFRQUJvej3h3UFRBUEJnt1ZiUk1CQWY4RUJUQURBUUgvtUIw
R0ExVWREZ1FXQkJSK3RFQjJ1ZGtFWHhYMGsxNUd6dEYvNG9sMDNqQUxCz05WSFE4
RUJBTUNBUV13RFFZSKtvWklodmNOQVFFTEJRQURnZ0VCQUlqeDfMOXJKMkIrbURT
QTNMMkdSaGpyUFJgZkk2VW4zWjUxQ2VXOWdPOVBNUtV3czVwREp5Qjc5ZEUVUthV
ZjF1OHBAEWpjaFRzUme4R1Jkbkt5bmROMmltYX1PV1V2UG9UZDMvWlNtZmt1cmNi
ajNjNFZxOHnqSFA3Q0U4Zm1VUzhYcHJkcG8yU3hWN29uZUpDMHZ0NWV5aDhtZ2ZK
L3FTYndWYw1YdUgxV3hpNmR1QXZkeGRkTVhBeFFLUecxS0tWTTdDWWZDZHBYL0hh
Z0NPSHpjdG8rMzc0ekZxcw5RMUt4NXJiZ3Z4T1NnbS9QRERPTXdQMDMrYmJUUYjYz
S1NLMVZiZHRMqNVTNGpnYVBhYnd5eFF6L0ZjaXduS9ITE9RbjhUtnFEV31vSWJz
K2VRWDJNZHMyQXB1bDhYSDIrQ2FrakJMTUxMM1RsaJ4KzZ0S1I5bz1c11191Iiwi
YWxnIjoUlMyNTYifQ",
"payload":
"eyJwcm90ZWNOZWQiOiJleUpyYVdRaU9pSmpaRzFwTGlodmN5MWhMbVp5WDJWdVxz
SjVjSFJmY0hWawJHbGpJaXdpWVd4bk1qb2lVbE5CTFU5Q1JWQWlMQOpsYm1NaU9p
SkJNaUyUyJBTklmMCIsImVuY3J5cHRlZF9rZXkiOiJ0WkQ4ZkpUT1hGZ3NTVi11
RXN1ZS1VYURnVpQX3FFZWVUUFFyQmpObUF3UlpMSX1ONk1Uc3hhc1RqbDR3YjdH
UXltN0prOWw4QX1ld0Fkak1tby1yOUNULUHsBmR5N1c2Nkd4bW11TGhwOV9ncj12
WETjexH3QUhHeGiyZGo3dm1Iqj1IRElNazzEQZb1Ri1naWZ1M3VaeEhjdvVYfYkRG
VGZBU0M2TEwyLTlIQ29OMz1SM2VZcC1FeWc0clhwZ2t6SkpnLTdRSUpHV3pvRFF1
d2VfeEM3Vm5Q25vdDc2bW9RbEpSM1I5Snh6M2M4TE96ckp0YTY3Yjdhs3ZodzK5
eUNPX3REb3d0WDZpNUd1cmZVbE5GN2VucmUwOWd3cUlLT3ZxQWRDVTdlSU5YwNvO
Wj1ZX3Y5RjVvWlF1M21Bcedib0dVdVRvRWxUMFFKdl1dwaUR6bUEiLCjpd1i61nVJ
ODQyQ3ZsRUpBYTd4Y1YiLCJjaXB0ZXJ0ZXh0IjoivjBCCWtLTv93TmV6UVJuOGxv
b3V0dFhPem9tSjJLbElYaU1NT0VQSk1NTVZpd0lKY05IWHdBV01wUnBoM1RrMn1S
dUoweHBZbmRPdVBtUHDpVGljMWR4YXZvaGJEdHl0TEFQS09TNWIzSk9lQTgxvWU3
enRaeFdxOFRDVVRVfVkhBVTM5dERzZUQyZ0tUME0yOFr1dXpEZHhJcEtUMXJQRXZr
WThIQVBjakhZQ2pTeVpmcEhJSlpkUUhEclo0YUZBN0xETkKxWkxPOEhXMmtJbFB6
UzQOWFdKQzYwS243THJvSXRYZ0hCWnJfZEFpNWl1QTdmU0VCSFVSrFpPQ0pfdzlr
TldLMDZfvTN3S2xVNU9DeFFYNzFMRGxQC1lENlpaa2lwaExoR191RUxvZFdkREo1
Zmg4VnhDRUJ6ZVdqckR1M3UzMXNldhJEbmRiAFNst2VyeEswNFZuQXNET1c2RXNT
a2twLXp5cFh4d2ZBRVQ0VWxiR3V0UV9FX2xsX2M4VVZjSnJhRXJoZFGzX2VNaJdS
YThuD2VEDGxZwjlDLWVlQkdJYzE1OTZLaVZzQlJ0c1dPZm5SLWQzdVZFREmwZnVv
alp1U3M1Mm1mWjh5dU1yeE9QWDJodzFVVG1CeHRuODdyQjRnc1ppSkwkc2V1e1Iz
T3d4eEFMUfJRSZfZQVNweXJRaWljTVJBSUtGVzJUUDfjSjRmYt1IdmNTbzN0eTUx
cENSvGlmUV9CSFlLM01LaGdWMWVnZ0I2Mmw4SEs2NkE0dFlpcjJyUnhkNEZPb1U0
bXN0bVQ3aFpmOVhBRzFBMzNwSHNxSEFINng4LVJdC0FwOH1GeHd1SHg3X0s2ZzNo
ZzJNSTNHZUpBTGxMRXhyeUJ5NEF4OXQ5ZEg0SWkzZWZjZ1h2LVROcmFIS0tHd3VM
QUItby1jVG1WUWZfU3dGSBtHZGNjSWgwazFYt3VXdKfmeKducjJKMnhBVjBUQdzd
YWNyZHBmUzgyWkM2T2dSSVliTEs0RnhUYXpuMXy1dlZjRms4dU40NFgxV2VUM09f
RmR1U1lLZ0ZPbUFyTHd4WXBMYXJSS1QzX3Fpcjl0QnYtQ3FycEJPV2EwV1VtaXF5
TWxTb2VHR2p1aHcwLVNuMVBtA3dWb25jNTVYTWxYcEVZc2hJb3dHNOxnb0dBakx1
MU13cjKxNWxmWnZsajFpck1jbGk5MgtIWBKtjdZVTJEMVhaM2hEV3FLTTTLRWG1w
RnJaV5qdi0yUWN0ZGt1VklMaHqWn2JESU51aUpDZENS01lbWVWZ1hZMw1takxB
N1J6Tjg0SjN4UTIzVXN5eFdWLTfYRTJueEFmcmYzMG5KZG5HTi1KbEhqdfMTQW5E
cXl1jUGQ4VjhBcUw5UURrUm1tZmsxRnBwTUdEaW9XS11VM2JTtGhyTHh0RzMweUQ2
bzV1b0UtUfN0aW9XcGVsN11ZcVg0aGx1NGNoNWVhWkQyRF93Ni1VOUxLbmQ2NnZq
OVptNjUuYVZrYy1SSGhNWEpQczhMeElCM2hrQ1BSmi1KN21UaTBZyTh2d0160FJ5
LVQ2Z285b3VVU2F3MHJYbi0wNm9HNC00aXZhMnZ2UkZGaXFOs25KYtdrM1BrWDF2
cXVqQU15Ti1ESki2cmZNznZxT1ZmV0phNXBuNTZLY1g4TFBPeUFiSGo4WUNWT3Fo
QldtbGt1SjM4WUtpY1dWRkJfTlpsZEpseDRTUXdoQ2pLY1JtdmxENXJqeEk4NTRY
NTdzQkNQbmpaTzdvdDRpZ21Vbn1fZjNyRkZQdl1ieVfmdWxkM3hMRVdSWWZCWUpu
YWRvNzB0S1VXM11OcWFzUzR3SW9ibFhFSnR5a2RPTHdQLVkybGRVMXdxVbThsUHDp
VXpNZGtrY2xWkQ2Wkx6LU9GdDdLaE9EbUxPTUdPNWlhcz1zeG9DaG4zYzFMcnFQ
amd1U1i1RWFNoaUZrNmtYVU5CMn1IVet3dmZZT0ZjUjhEbDdleVJCQ3NVUUY4STdT
ZkFwNkRjRmM82ZU1jdEJGTxlUWfgydU9nWxEwOwtUM0ppckhzbG13dThZbZGM1bHpW
ZGViWn1qREIxbWl2LUNKR091ZjRxbEzXrZniYVBCQmhNa0ZFS21WU0FJZmdDaFlu
dUc3WkxvNmVMMW4zMnRGEDfZdEdUcnpuMk13a1NzektPMTJTY3VkcUZnTureQ28w
T19OREZ4Qmhl1RFB3aVNaDnBGHJWV201QktEbjBOQzNFcy1zoExGaDZJUFVtMW11
VFB2UmX0cERJZTljQTZqdHhCQUrAcVphdGtreXB5ZzZGRDZzUjhUQVRDRG0WWhT
Rzk4MVBmRFZ1S09FVjlqaTdZd31kLW9ST2Q1UXZCa11KUkxhckpEOGHedWZicHdy
M1V1NENBb1JKcnJWdGM4TmJDcm1LbmZGWVBQa18yN3p1RFZKZXBxQTdpOVVYTY0ZD
T25vUEX3NEFPvMxqQVRPU93MmNzR0FqM19tLthPR2cxM3BTUHNjwd1Sum5TV11Y
cjdfWNkRJaUVXNDZaNT6N1JTNHGOZ1dqZWFzNUpHdmRsQkkoQk42bFd4QTFLa0Vs
e190N1VOSVhGSWpZRTdyVF9oUzJsWmstWGNPe1ZuM0gxT11ZU3NpTVVLVDNvWm4t
Uml1VvhoMGZjbvNwTHFRMnVuX2dQV3MyT0VUeHB3U1Y2dk9leWRjOfcwQVdhSVph

```

(continues on next page)

(continued from previous page)

```

Vm43bW1MNy2aEctTVFLS192S3pfaTNxbkRrUVzVjhDQVE3NTN1cWpHUVBEMUs4
VGtZeTNucnhncEh1ODY3UDJNUGpVQWFxeldTVldNN2FQeVhCTkZNeFQ5V3Q5ZmVw
VEFheDdhLVolVGcwS1pxS01DQVvNYUo4OWJveG1ncG9rcGRFRVWtQs1pFMjR2NGph
ZE5pT21SSHJMeE11VzVrdG5kQk5Bb0dHbm9HRTNzYWZlc24wX2Npc3M2TTVxcEhR
Z1N6VWxacF9aREJ5Q09ldFo1TmEwMjVfb2t1T2drUklKVjN4MmZPbWM2U0wwU3Fy
R1lWSEhORmZKN21LeVNXTk5Sdk14UmhWTmJYVXFMR19OYUMxZDdNQmpkTHY3WlFp
Ym0yT3BmSWtmS1NNuZFKcFvSVNRakR4djK2TFhIOHBhQXBTVdZtWnlEa0VrTW1O
b2ZXSkTwm0pQNUFTTXZ2S3dsbkGxQU9rSWJYZjdRLXlyYkthY1BtWEhkZ2hjRmZa
Vi1ldU9ERzhfDTRWdGZwoEF0LVBRskJyX1V0cVZJbHpReDBnUE5MVVdMbzd1Rmlv
dEtPTEx3Vjh2MGpxcDdUu1BIVzBHC3pVWmFZb0RjYUZMaXZDaFJwVDktQks3d29w
RXhfN1FxUWUXSVhuMVNKQWRtZFPvZmdfSF9ZEU0MVJjSnZGeWp6SmZBOG1ranBs
STMxVW9SszVaT1lXSVExZ2lvUFh3U3JUVW1GeE45X3Z0aH15Z1pYTXQzVvZiUTNB
eFBBZj1lSSE5EN1pBOGFvAlk2R0pNb1BvX3FySkMyLW9hQzd0cTc0SURBekk3bpm
S09VMG9IZ3RtaGhGUVB6cloyWjZOYU1LaTJIWm9yLWc3LWctN1N3ZGRxd0RDSnlO
ZXDEb3pJN3NNRFM3XzM1RkoydnRnaGZLUHFTLWhFd3JlZGtaNE1ORUg1dnRXU3F3
NHRmZDRfQmNROExtR1ZmM2hfV1FWZjVqQ0lGcEzNOV9MYnV6eWVmeU9IMENSLXlJ
YlEOQXfLeS1ZclAOWYJTDXF2VXB1MEFyVEdZS3pMWk5FZHHFVzliemxsYUVZVNU
MTRBNXAZam9CUVNOQmtLRUDPMVh2RTdSVnJ3T1kyemItWWFIRkxjU0gzdDRtMF1Z
aVducFZ1RnBRZjExWVlqbTgzLXFEBHzuYUN5T19GQXBtMy1SRnhrVTZfVzRYcDZ3
eVlxNDF1amxGVHp4Q1RnMEXXek1qTGLyRm1VSGpKZU1hR3lVSzZ5TGpkYkhtR0Ns
LW15T3FZR3hwNERybhVITHBkaWZHenhmQnd6OUM4MXNNbkFfxYmNRWGMwcZ3IMVdO
al9uX0tMblJfDhRWFbleDV2VmZ6eWR1ZUZfNE93WEXsdxZ1M3Z1M3Z3VldWamRP
ZWN2a05nbkVFUDdocVhkbXpRbV1KZndxSGRwZnNSREXESDloOWpRcEI0LTNMRjg4
U2NSalRtef1NU291bFVqcHdia25hTWZQRVmtZkgTjVwQXNqVE0woVFTV0Y5NENE
bHDZ1ZhzY2UkdLcGdwU1ptMGf1dGNSG1KOGVRmjRuRzNkDEf6bGNkVC0zXNXC
TlFsZHRkY1hkRzFxb2ZlOUpUNE5iS1l1VXdwemlhnNjYNGQ5NkRqdm1VdWpqaN1
NjhCkXx2bzJQM1VYtkxUWDBjNmUtb3lHbzcwduRTN1lxb2ZlUuXhU05ORHRnZ3d1
V1IyZFRVUDdLRkk5N29Qb0Zix1Jjs3BQVvU4cDZstJjRYkVOYy00TWN1eWJYVS1U
ODR1ZnpqZkKdHdMdkYtQ3d6Nk9lS1V1NnBPeVn2MGVnbFNJUWlmb3JPU3FGV19y
QWdvX3NuNDVHVvHkU0dNc0FFT0NXN0lQX0E4Y1hvaFpGbi1NNW1UTU83b01fSWN1
Q0dVc1NuREV2X3NNdXFLdFk0cExLbHRMT2tCTXVHU292MmlDVHBMQjVnVnBvaGZf
UzBaUW9CW1V2UDkUw3ZLUVRuTG12QnZnVtd3OVVmy1dRTzBrNUZSjNfREZxNHVl
VWJpNXB5cFVCT0tXd0Z1a3lJLXZXVGRPTGdLOW40ZkdUcj1zOG9ZeicZdC1tVXZ5
a3pfdkRzWWFpWHZQbkd3My00SDhVW1lXZy1GSHJfSH1LTlFQnk1GSTVGdEFneW5s
QjF6SEQ2a0xXbGtGaWdfb0FjdXJ6enE3WXJJaU1OcNZFUzFUZUgWRXBwcZVjZHpF
WWh0M3ZHNFMtQTJqcE5PN1pJuk1zUmFUMHkxbFBOLVZEdWJxOHBfZmhyX1B0ZVM0
T21zNGNvLWxiOUJNR1BmeEp0cVRDUWJGSnJLTFBDbDJKNk5VVGtnTy0zVlgzUkIw
Rk9uUmdWR196aVQ0amF1UFFxZDRsQ3BOME12S2RCb29BWGtYNzKxU2JjbmJoMUhS
cFlhblNVZENiZmZ5MGg3NTJ1M1dmV3JjSUVDa29GWkRqelUtbG04bnF0OTZnS1ZH
OFpHejF1bGhIzZvFVvIXn1ZNElROVBjNHVCTk1yazRUVThNRjF6R1Rsalhnt25t
UVZ1d0tPWmtkcFB4eWhGdExtalUzeWtoejBKMm9oMzJpaVhpbHFFeVQyQjM2Nzhz
NHf3NHNka1NTdXdCWXXZzmtRN180SnIzNnNrUm5CQ1ZLa21mdlFtbld4S3NheWQ2
WnRPc1lwZ1hjTlRfOUpYXzE5RFZiekdybW1zTExONzRPQWg0QjZUVXciLCJOYwci
OiJnb1QzT0V0XzZ5eVJXDEtPdFpUTdBIn0",
"signature":
  "a9idq1HPLsfvg7IEKECF_XNEKKpU5jioicQTxvpQywpT8DM5M7sy0PfrOSEL4iKJ
  hGd32xQ-JwaIgtT5RR0QwAM3cy15y0IMh-KkHh1H3Amsy2RIsi2jkZVpIN24OF0T4
  Wf6fXfuuXjBydT-PjIjQako2Lsic-nKAJukm0f3uwHxGXN3JG1Izi4_QrYeRjI4V
  b9ymaZW6soF1VaDqXulfaL7rtN_B46-Yg_N8XK5XWuIW5PZ1ZEkyFaqGzPaoMi31
  ipAlvx-VlQdVxCAT0x2uQOT-um6NJA1vAigyWn45CRM5ETXja_V8WkpDPLHH4A5G
  SEzKUGz3tYC85Agd0-Aoig"
},
"dac_request_dest_certificate": {
  "kty": "RSA",
  "kid": "cdmi.hos-a.fr_encrypt_public",
  "key_ops": [
    "wrapKey",
    "unwrapKey",
    "encrypt",
    "decrypt"
  ],
  "n":
    "uL7ANgD80H5sNqo3nHzovPRxgncQLhz0oQvGMVvULCkrYXMaXZ5sNv7ft6UdMSzi
    T-e0sthapMEqrpeV9RKHSiF3COGL2YndUHixpEkHp8y1ggcH6i1TzoBsgXMZ70LW-
    mJ2RW3rodT7k-tcozYYsTSM5egMPQSAKgt0nMnPMdNRnEyA2_NJ8Y71NkEXyja0Q
    J1stzkP8-cKS0BkEquLQEMbZVRM6U5uG69cj1i90WvuRzPoaATKyT6Cc4f6PUu9L
    OyCBUAs9dXsRrt3B8H1qe7io7FAAcOpCUDKdNLFXS1THc37DK_zEyKZcMttjCvEl
    Ovt-cIaokdnxJeggv9AFGQ",

```

(continues on next page)



(continued from previous page)

```

    "e": "AQAB",
    "x5c": [
      "MIIDMCCAhigAwIBAgIBBDANBgkqhkiG9w0BAQsFAADBCMQswCQYDVQQGEwJubDER
      MA8GA1UEChMIbGllc2RvbmsxDTALBgNVBAsTBGNkbWkxETAPBgNVBAMTCHJzYS1y
      b290MCAXDTE2MTAyNzEyNDUwMfoYDzk5OTkxMjMxMjM1OTU5WjA1MQswCQYDVQQG
      EwJmcjEOMAwGA1UEChMFaG9zLWExFjAUBGNVBAMTDWNkbWkuaG9zLWExZnIwggEi
      MA0GCsGqS1b3DQEBAAQUAA4IBDwAwggEKAoIBAQC4vsA2APzQfmw2qjefcOi89HGC
      dxAuHPSHC8YxW9QsKSthcxdnmw2/t9PpR0xJmJP57Sy2FqmYSqul5X1EoeYIXcI
      4aXZi1d1QLGkSQenzLWCWfJqJP0GyBcxns4tb6YnZFbeuh1PuT61yJNhixNizl6
      Aw9BIAqC3Scyc+Z01GcTIDb80nxjuU2QRfKNrRAkuy3QQ/z5wpLQGQsq4tAQxtlV
      EzpTm4br1yPWL05a+5HM+hoBMrK3oJzh/o9S70s7IIFQCz11exGu3cHwfWp7uKjs
      UABw61xQMoOcsVdLVMdzfsMr/MTIplwy22MK8SU6+35whqir2fEl6CC/0AUZAgMB
      AAGjPDA6MAwGA1UdEwEB/wQCMAAwHQYDVRO0BBYEFBAIGICMR5H6KLM1ZAEHCCC
      KwE9MAsGA1UdDwQEAwIEMDANBgkqhkiG9w0BAQsFAAOCAQEAANYSSryUU6112pYM
      r83M3Gwnjz16B+4KqimZ8kbey94zNPdwmwQdSe0Xmg+1Otc6VUB40ouNnwKefB
      aWBtXwCA7Nb715nTqo2+rn+X+A0mGrYaKkToPEe8ZYwDcOlOpNC9JFE+QgP9/CJa
      AaWrf95W+4kra2WnA4Bhqu2WWXnQkL47/nKcGVZgQAH+mVnxPaI0gELYdonXU/S2
      8HqxoyjpGL/vmyc46zUbxYsgx/jiE7J0fJVP6Yk/3dlNYCCpLtV8VmzFAQAeCcn8
      AWowFcd09a4SY09rn1MUv/rvvXpzflfn9j7PtRRFj2e/KhitmOH1zKDuYzREpUOu
      TDlPIQ==",
      "MIIDQCCAiiGAwIBAgIBATANBgkqhkiG9w0BAQsFAADBCMQswCQYDVQQGEwJubDER
      MA8GA1UEChMIbGllc2RvbmsxDTALBgNVBAsTBGNkbWkxETAPBgNVBAMTCHJzYS1y
      b290MCAXDTE2MTAyNzEyNDUwMfoYDzk5OTkxMjMxMjM1OTU5WjBGMQswCQYDVQQG
      EwJubDERMA8GA1UEChMIbGllc2RvbmsxDTALBgNVBAsTBGNkbWkxETAPBgNVBAMT
      CHJzYS1yb290MIIBIjANBgkqhkiG9w0BAQEFAAOCAQ8AMIIBCgKCAQEAsrUj46dx
      5ojlaZk7YtOL6e+Q6JoG7gVMAxKJn1Sz1x9ND/8w4PeOlSQ2skukdOHALQRmxdft
      zhccNTM5hmbcn8TAFWSYqQF1R7s78bVjtmata6AQPlvSgiyZ8Ak+iYZEq3c2zVyYQ
      HKKXwXmFZt1HT8/H/B3bXveXQcERKE+Tq66h8pqVcocQUTzRFsEYmv0bR1rghToq
      H8nhB5xnebgVlXjApW+et2SE7r6Fjv1aAbGI89ouJlgsMPeX56P8AUjacFTNkC44
      Obu6HRXY/jm6f2m1EUm84EUSJ+9b5+S2x4qPttJdFSCasWYYz4mFJ8MwmFiBGUwf
      geT2bUm6t7qgbQIDAQABoz8wPTAPBgNVHRMBAf8EBTADAQH/MB0GA1UdDgQWBRR+
      tEB2udkEXx0k15GztF/4o103jALBgNVHQ8EBAMCAQYwDQYJKoZIhvcNAQELBQAD
      ggEBAAjx1f9rJ2B+mDSA3L2GRhjrPRjfI6Un3Z51CeW9g09PMQ5ws5pDJyB79de/
      Q8Uf1e8pZyjjchTsRa8GRdnKyndN2imayOVUvPoTd3/ZSmfkurcbj3I4VW8sjHP7C
      E8fmUS8Xprdp02SxV7oneJC0vt5eyh8mgfJ/qSbwVaiXuH1Wxi6duAvdxddMXAxQ
      KPG1KKVM7CYfCdpX/HagCOHczcto+374zFqqnQ1Kx5rbgvxNSgm/PDDOMwP03+bbT
      R63KSK1VbdtLBuS4jgaPabwyxQz/FciwTu/HLOqn8TNQDwyoIbs+eQX2Mds2Apul
      8XH2+CakjBLMLL3Tlj2x+6tKR9o="
    ]
  },
  "dac_request_dest_uri": "https://cdmi.hos-a.fr:9001/dac/"
}

```

The DAC provider at hospital A will retrieve the signing key from the JOSE protected header, validate it using the included X.509 certificates, and then verify/decrypt. It creates the following (plain) DAC response. Note that it included the object decryption key.

```

{
  "dac_response_version": "1",
  "dac_response_id": "73da04e1-2182-447e-8342-f4b9f06ec936",
  "dac_applied_mask": "ALL_PERMS",
  "dac_object_key": {
    "kty": "oct",
    "kid": "encryption_key_1",
    "use": "enc",
    "alg": "A256KW",
    "k": "lmk_8n9GZJTLDUEuxBYT-9GO8bc_fr2qqqt03rVSRFak"
  },
  "dac_response_headers": {
    "CDMI-DAC-AuthInfo": "Key successfully obtained from KMS."
  },
  "dac_key_cache_expiry": "2017-04-06T14:42:47.393Z",
  "dac_response_cache_expiry": "2017-04-06T14:42:47.393Z"
}

```

As before, DAC response will be JWE encrypted using the key in server\_identity. The result will be:

```
{
```

(continues on next page)

(continued from previous page)

```

"dac_response": {
  "protected":
    "eyJraWQioiJjZGlpLmhvcylhLmZyX3NpZ25fcHJpdmF0ZSIsImp3ayI6IntcImt0
    eVwio1wiUlNBXCIsXCJraWRcIjpcImNkbWkuaG9zLWUuZnJfc2lnb19wcm12YXRl
    XCIsXCJrZXlfb3BzXCI6WlwidmVyaWZ5XCIsXCJzaWduXCJdLWwib1wiOlwib1Q2
    RjclalVZGZkdQLXdramdneVhacFdvRWhxNmZTWkNfRENTYXNLLUdVdXdaUeTUJG
    cXVreUM4ZmsxajNBjY0JyTt1IZERIBnFrVzRfM2YzOVdZMlMtRfY5YjhkWWpRRHJl
    TDFfcHBnRVg2enJnN1hBWEJJa1ViT2hldXJOTVNBEN1QXlCY2xqWUUpQ3dvTWxv
    aUNgB2RsbFYzUDZwekVLNjduTHNfYWVfTHVaUmRzaFhXakotYm9qQzZiNGhJcWVx
    UFh0dXBPeKl5MDBKLvZydhNGUTBZaDN2dWRmbVvJTkZSTf16c1l1y1Y3TVR4TUd5
    S1hIMXE0ekdHSHZmSlpTfS0MxTETYYk9Dc3JhaHVRUFVhY0tBSFBqbnTNAkHpONTR6
    amVYTGxnaDZPT2x4X2EzalV4YlBFTl8zYjhHTDhOVUhuWf1jCUNjVDVXUUV2ZUZw
    OF1RXUCIsXCJlXCI6XCJBuUFCXCIsXCJ4NWNcIjpbXCJNSU1ETURDQ0FoawdBd01C
    QWdJQkFqQU5CZ2txaGtpRz13MEJBUXNGQURCQ01Rc3dDUVlEVlFRR0V3SnViREV5
    TUE4R0ExVUVDaE1JYkdsbGMyUnZibXN4RFRBTEJnTlZCQXNUQkdOa2JXa3hFVEFQ
    QmdOVkJBTVRDSEp6WVMxeWIyOTBNQ0YFRFRFMk1UQXl0ekV5TkRvd01Gb1lEems1
    T1RreE1qTxhNak0xT1RVNVdqQTFNUXN3Q1FZFRFZRUUdF0pY2pFT01Bd0BmVVF
    Q2hNRmFHOXpMV0V4RmpBVUJnTlZCQU1URFdOa2JXa3VhRz16TfDfVpuSXdnZ0Vp
    TUEwR0NTcUdTSW1zRFFfQkFRVUFBNELCRHdBd2dnRtBb01CQVFDZFBvWHZtS1VW
    OFkvN0NTT0NESmRtbGFnU0dycDlKa0w4TUUpKcXdyNfPnT0JXSWNNd0VXcT2ZUSUx4
    K1RXUGN0d0deZjBkME1lZXFSYmovZC9mMVpqWkw0Tl1gxdngxaU5BT3Q0dlgrbWt3
    UmZyT3VEDGNCY0VpU1JzNkc2NnMweE1EZ0o0RElGeVdOZ0dJTEENneVdpSUlXaDJX
    Vlhj1L3FuTVFycnVjdXo5cDc4dTVsRjJ5RmRhTW41dWlNTHB2aUVpcDZvOWUyNms3
    TWpMVFFuVnd1MndWRfJpSGURnTERw1FnMFZFdGpPeGl4eFhzeFBfd2JjCGNmV3Jq
    TVl1ZTtK4bGxJb0xvc3BkczRzZXhRZzROVJwd29BYytPYmNtSE0zbmpPTjVjdVd
    SG80N1hIOXJlU1RGczhRMy9kdndZdncxUWVvKzGh5b0p4UGxaQVM5NFdueGhBZ01C
    QUfHa1BEQT2NQXdhQTFVZEV3RUIvd1FDTUFBd0hRWURWUjBPQkZJRZUZNWThmOW9h
    aXhYTkfWSW1Nby9heE9kn205clFNQXNHQTFVZER3UUVBd01IZ0RBTkNa3Foa21h
    OXcWQkFRc0ZBQU9DQVFFQWJPejR5akdLektMOpVQTJ0WVVOeFhoQ1RkbnBVUGZB
    WDhRdGs5ZEJXTZzhVXNSzJlTkh1TzZjL2tKMg0vKzR4LzJiczJlU0dISmRldU9w
    dzd3QmQwTDB6L1hUeStFVHRQM0k1TEovYzhNdmN6NTQyZCtMQUpXdlZtU1htQU82
    QjZRZDJtSkJ5aGFMU3k2YWUxK3BwcmE5Yk1aYjFTVWVxc3ZtZ1R5Z2h2Y3NkVzky
    TUhBUH1xUlhxTmthaTNhNW5kbE14YjdgcWFhMzNCZkxTOHhVcE1FQ0wQURBj4L0dH
    WXcxZXIzSUFcvDBCeG9JQXp0c0tHQTFUOWh2cnaA4NctFNjJhcnZrU1lHUK8zQmpu
    SGRHcGdnRUhRMmhGNkR1UF1vRGVYv2d2R0MyREN0cFNxVzJTTklaUFprb2tsz3pE
    QnVXdVlHYj14d1Rwa0psWVRpSzF0dz09XCIsXCJNSU1EUURDQ0FpaWdBd01CQWdJ
    QkFqQU5CZ2txaGtpRz13MEJBUXNGQURCQ01Rc3dDUVlEVlFRR0V3SnViREV5TUE4
    R0ExVUVDaE1JYkdsbGMyUnZibXN4RFRBTEJnTlZCQXNUQkdOa2JXa3hFVEFQQmdO
    VkJBTVRDSEp6WVMxeWIyOTBNQ0YFRFRFMk1UQXl0ekV5TkRRd01Gb1lEems1T1Rr
    eE1qTxhNak0xT1RVNVdqQkNNUXN3Q1FZFRFZRUUdF0p1YkRFUk1BOEdBMVVFQ2hN
    sW1HbGxjmlJ2YmlzeERUQUxvCZ05WQkFzVEJHTmtiV2t4RVRBUEJnTlZCQU1UQ0hK
    ellTMXl1mjkwTUlJQklqQU5CZ2txaGtpRz13MEJBuUUGQURFQ0FROEFNSU1CQ2dL
    Q0FRUFzclVqNDZkeDVvamxhWms3WXRPTDZlK1E2Sm9HN2dWTFWfYa0puMVN6bHg5
    TkQvOHc0UGVPMVNRmNrdWtkt0hBbFFSbXhkZnR6aGNjTlRNNWhTymNuOFRBZ1dL
    WXFRRjFSN3M3OGJWanRtYXQ2QVFMXZT215WjhBaytpWVpFcTNjMnpWeVlRSEtL
    eFd14UzadF1VDGvSC9CM2JYdmVYUWNfUktFK1RxnJz0OHbXVmNvY1FvDhP5RnNF
    WW12MGJSMXJnaHRvcUg4bmhCNXhuZWJnVmxYakFwVytldDJTRtdyNkZqdjFhQWJH
    STg5b3VKMWdzTVB1WDU2UDhBVWphY0Z0TktjNDRPYnU2SFJYWS9qbTZmMm0xRVVt
    ODRFVXNKKZl1nStTMng0cVB0dEpEZ1NDYXNXWV16NG1GSjhNd21GaUJHVXdmZ2VU
    MmJVBtZ0N3Fxy1FJREFRQUJvejh3UFRBUEJnTlZlUk1CQWY4RUJUURBUUgVtU1w
    R0ExVWREZ1FXQkJSK3RFQjJ1ZGtFWHhYMGsxNUd6dEYvNG9sMDNqQUxvCZ05WSFE4
    RUJBTUNBUV13RFFZSkvtWklodmNOQVFFTEJRQURnZ0VCQU1qeDFmOXJKMkIrbURT
    QTNMMkdSaGpyUFJqZkk2VW4zWjUxQ2VXOWdPOVBNUtV3czVwREp5QjC5ZEUVtThV
    ZjF1OHBAeWpjaFRzUmE4R1Jkbkt5bmROMmltYXlPVlV2UG9UZDMvW1NtZmt1cmNi
    ajNjNFZSOHNqSFA3Q0U4Zm1VUzhYCHJkcG8yU3hWN29uZUpDMHZ0NWV5aDhtZ2ZK
    L3FTYndWYw1YdUgxV3hpNmR1QXZkeGRkTVhBeFFLUecxS0tWTTdDWWZDZHBYL0hh
    Z0NPShpjdg8rMzc0ekZxcw5RMUt4NXJiZ3Z4TlNnbS9QRERPTXdQMDMrYmJUJYz
    S1NLmVZiZHRMqNVTNGpnYVBhYnd5eFF6L0ZjaXds9S1TE9RbjhUTnFEV31vSWJz
    K2VRWDJNZHMYQXB1bDhYSDIrQ2FrakJMTUxMM1RsaJ4KzZ0S1I5bz1cI119Iiwi
    YWxnIjoiUlMyNTYifQ",
  "payload":
    "eyJwcm90ZWNOZWQioiJleUpyYVdRaU9pSmPaRzFwTGl0dmN5MWlMb1Z6WDJwVkdz
    SjVjSFJmY0hWawJHbGpJaXdpWVd4bk1qb21VbE5CTFU5Q1JWQWlMQ0psYmlNaU9p
    SkJNaUyUjBOTklUMCIImVuY3J5cHRlZF9rZXkiOiJlZ11XVDVDOVBiN0FEOVZH
    bTczSDRQSskFJUHNbnTN2ampNbmZJcXZBOEFLOEtNREowMTl1NqkVHeC1tUpIQTB0
    dGhyb1F0UXA3Nk1DR09GLXJaNnBxajNTYm1iNFNUc3FhUjQ5aTN2Y2x1dUhrS1ZM
    WVNOVKlWOHBjYzZLb1lDWfNGMzNldh1IZ0p0bFZjXzE2RTIyUGxeVE9iWjhPRzZz
    dklWalo1cktnLVNGS1ZzbVlWeC05UC04NU1ueThRMkYtR3VBcDNI0HVZVZzhFNXdx

```

(continues on next page)

(continued from previous page)

```

Z1AwY2NHVYVnjZF82cFUzdk83aUZGWLVLV8ycHlZQWdRMlo3Zkl0a0lSWnVPdFFf
RkpNaThLcFdKc3lMNFNDanZSNkxFSVBIEkdTMUtKtMNEVTvrU28yc3pkanQOWENH
Sk9WSKc4dFhfWmZRTjh6RzdpTktJcXhiTVYwRUZVOWRjUzBkeWciLCJpdii6Im9n
Z3hKM3g4YjBhb2xlSDQilCJjaXBoZXJ0ZXh0IjoilWWhkM1A3UGM5emtpQ0hhYnda
ZVhWU29rYmNXcmdqQlQyU0FzaFFGYzNMNFY4b2ZLYVNYR3UtaZf3T014OHhZY1FH
Slo4Mk12b2tEUXNGRnZaQjBNUG9hdEljQWtmTms0VURqREcxNi1zVGc3QURKeV9p
TkdmZlZkUyY1XU3QxWW1PZW5HaXBYNEZWM3ZyNU9URlE4X3hCRHdYSHRiU05KX3Jy
RlVWX2R3TnJQdUdiSlYzSXhOYi01RFJlRkl1Z0ptY0JXelJYRW1LeUEyZEdZTG1V
eEF3WFpWT18wYmlKVkxzZz1KalpuczU3YXNqR1Fna0c4ZzhCNDlnTEl0QWRfQXpJ
SFlrNTVpUkRvQk9WRnhGR284Vlc1ZFNrM05ldilmZ1JxYzFDZEdodeZuTnlzUjUy
akhtMFprMzRoeTJRLWZVNnNTVXBWmu5vVmw3Ytd1NWMxdEVZSGRNVU9fWWtHV0g0
dFIxewdmbG94dFY3ODRMXYdObGNCYlV6NFpvWXfKexG2amEzdGxHaUpxaWdLZ043
WENTT0R6RXF0QVRQeXBItNvfYXJHRm9LNmk3cGdKZXcwYUxMQWJQOFA0SFJHTVZ6
Q1V6dktDbDdKYS1XTE81U1UwS0ViLvHJZGdHX195a005WTJquVQZdi1MUjNNGR1q
WlE0ZDM4cG9BS11fSHJCWDg1Z1dMQXFzY3MyRHhwhRT1IZm9xaXk0T2tscV1aa3Js
MUPNT1dtQVZkWWNqckh3Y3YVjZGSUJGV2Vjc1Nwb185ZUsxOCisInRhZyI6Ikj6
NFZZeHdtNFkyU1Y4bDVUckMwUEEifQ",
"signature":
  "RjUCI3Q_zfBJyeHjYflsd6MppSDNUAIzC771sbM1MiKfLDi8oN0999gpByS7Sx6
  kCXqsNkV4T0z_qaqy4UY9JrdjMRTNFPXJMwhqbBem-s6dJT6VquF3GBQTU8wb4OK
  5E8rGvTcWw-Hd0SfpjGoJtgv5RmpzfVgdvANZcJfST-r0ra3EnPitOf8dJ95Db3t
  78mEbMfqdoobk1Dnc39DvpnzD61TioXWoZj3UGcBcvNpSl2XijS6yZlgAsrQbGnX
  xWx-PCwEACZoVekzt-YV5QUFH2JqblpGeUUCwoFv1ON90iXVusIdWWnJO51gSKwg
  i80ZxOSSBwF6b9WIeXHe5g"
},
"dac_response_dest_certificate": {
  "kty": "RSA",
  "kid": "cdmi.hos-b.us_encrypt_public",
  "key_ops": [
    "wrapKey",
    "unwrapKey",
    "encrypt",
    "decrypt"
  ],
  "n":
    "oQMqkY85Uzw07K6H0QQnfAiRMN3ZfhK0aXEkx7YwvrCU9IKOquZ10YZ9Cv8556_8
    E8yZm02JDWOB0aSaSGHU835jvXf12f4MywKGWj5FtIGL-j9kXF6SWq3zuLVYlXpMI
    KsJngHMVFca_-ZhZ2vLsrnDR1aCNEC48gR26ewp6WX1ptnSc1W4x3Mj-ONMVzxVE
    7XNlWYysTgDtonmTQD-YG6_KhhAPx0YowMbUWv_cMQvXsi7MMDyZn6fxfq4zQmQ2
    V5RtUy5msd6K3beDzS4LmZhsJmjU7YnhOj0pZby4Zckm43npjXPAuwPhzK20W7qb
    fkv0qm4rsFWUcuNh81BsDw",
  "e": "AQAB",
  "x5c": [
    "MIIDMCCAhiGAWIBAgIBBTANBgkqhkiG9w0BAQsFADBCMqswCQYDVQQGEWJubDER
    MA8GA1UEChMIbGllc2RvbmsxDALBgNVBAStBGknkbWkxETAPBgNVBAMTCHJzYS1y
    b290MCAxDTEmZTAyNzEyNDUwMmF0YDZk50TkxMjMxMjM1OTU5WjA1MQswCQYDVQQG
    EwJ1czEOMAwGA1UEChMFaG9zLWlIXFjaUBgNVBAMTDWNkbWkuaG9zLWludXMwggEi
    MA0GCsQGSib3DQEBAQUAA4IBDwAwggEKAoIBAQCChAxCRjz1TPA7srofrBA18CJEw
    3dl+ErRpcQrHtjC+sJT0go6q5nU5hn0K\znnr\wTzJmbTYkNY4E5pKwYdTzfmO
    9d\XZ\gzLAoZaPkW0gYv6P2RcXpJarfo4tViVekwqgmeAcxUVxr\5mFna8uy
    ucNHVoI0QLjyYBhp7CnpZfWm2dJzVbjHcyP440xXPfUTtc2XBjKxOAO2ieZNAp5g
    br8qGEA\HRijAxtRa\9wxC9eyLswWpJmfp\F+rjNCZDZXLG1TLmax3ordt4PN
    LguZmGwmanTtieE6PS1lvLhlySbjeemNc8C7A+HMrY5bupt+S\SqbiuwVZRY42H
    yUGwPAGMBAAGjPDA6MAwGA1UdEwEB\wQCMAAwHQYDVRO0BBYEFH7NJvMIftQtZn
    nyiIdLNkjCgWSiMAsGA1UdDwQEAwIEMDANBgkqhkiG9w0BAQsFAAOCAQEADiADiv
    0v09SUDcPL+BKysvchn\Sgx5KBu7n9KFwE31Dhx2zvT6ruL8kXdekPH9cfrDafW
    6I\vnbzAVj02i5pM2cHayj13fTOWSVwpcQuvkoIF9eVIWONkemMMf7M7jpTw07z
    7S2TSusaDmMnpqj8y5pRpQo3PnBVxpEZJ0XaSdfuiHtVLDq8gDZCq6Hc2tt7JM3W
    njnQgs+1lSGRuqWocpmVONIoqvhiolNDZV35Z7puRwqck1N2f1qyHHGBWXCfN9U4
    ci6q1BnWBIFV+hURge8NSbpqawolaNueUbTcKjN3JsMC4ZxhMF9rN3uuPn+UAYka
    yQkcSmGSMm07wcAkMg==",
    "MIIDQCCAiiGAWIBAgIBATANBgkqhkiG9w0BAQsFADBCMqswCQYDVQQGEWJubDER
    MA8GA1UEChMIbGllc2RvbmsxDALBgNVBAStBGknkbWkxETAPBgNVBAMTCHJzYS1y
    b290MCAxDTEmZTAyNzEyNDUwMmF0YDZk50TkxMjMxMjM1OTU5WjBCMqswCQYDVQQG
    EwJubDERMA8GA1UEChMIbGllc2RvbmsxDALBgNVBAStBGknkbWkxETAPBgNVBAMT
    CHJzYS1yb290MIBIjANBgkqhkiG9w0BAQEFAAOCAQ8AMIIBCgKCAQEAsrUj46dx
    5ojlaZk7YtOL6e+Q6JoG7gVMAxKJn1Szlx9ND\8w4Pe01SQ2skukdOHA1QRmxd
    fztzhcNtM5hmbcn8TafWSYqQF1R7s78bVjtmata6AQPlvSgiyZ8Ak+iYZEq3c2zVyY

```

(continues on next page)

(continued from previous page)

```

    QHKKxWxmFzt1HT8\H\B3bXveXQcERKE+Tq66h8pqVcocQUtzRFsEYmv0bRlrgH
    toqH8nhB5xnebgVlXjApW+et2SE7r6Fjv1aAbGI89ouJlgsMPeX56P8AUjacFtNK
    c44Obu6HRXY\jm6f2m1EUm84EUsJ+9b5+S2x4qPttJDfSCasWYYz4mFJ8MwmFiB
    GUwfgeT2bUm6t7qqbQIDAQABoz8wPTAPBgNVHRMBAf8EBTADAQH\MB0GA1UdDgQ
    WBBR+tEB2udkEXX0k15GztF\4o103jALBgNVHQ8EBAMCAQYwDQYJKoZIhvcNAQ
    ELBQADggEBAIjx1f9rJ2B+mDSA3L2GRhjrPRjfI6Un3Z51CeW9gO9PMQ5ws5pDJy
    B79dE\Q8Uf1e8pZyJchTsRa8GRdnKyndN2imayOVUvPoTd3\ZSmfkurcbj3I4V
    W8sjHP7CE8fmUS8Xprdp02Sxv7oneJC0vt5eyh8mgfJ\qSbwVaiXuH1Wxi6duAv
    dxddMXAxQKPG1KKVM7CYfCdpX\HagCOHzcto+374zFqqnQlKx5rbgvxNSgm\PD
    DOMwP03+bbTR63KSK1VbdtLBU54jgaPabwyxQz\FciwTu\HLOQn8TNqDWyoIbs
    +eQX2Mds2Apul8XH2+CakjBLMLL3Tl1j2x+6tKR9o="
  ]
}
}

```

3937

3938

The CDMI server at Hospital B can now decrypt this message, process the access control decision, and use the object key to decrypt the encrypted object:

```

HTTP/1.1 200 OK
Content-Type: text/plain
Content-Length: 33

This is an unencrypted text file.

```

3939 **Clause 25**

3940 **Data object versions**

3941 **25.1 Overview**

3942 Version-enabled data objects allow the previous state of a data object to be retained when an update is performed. In a  
3943 non-version-enabled data object, each update changes the state of the object, and the previous state is lost. This state  
3944 change is shown in Fig. 19.

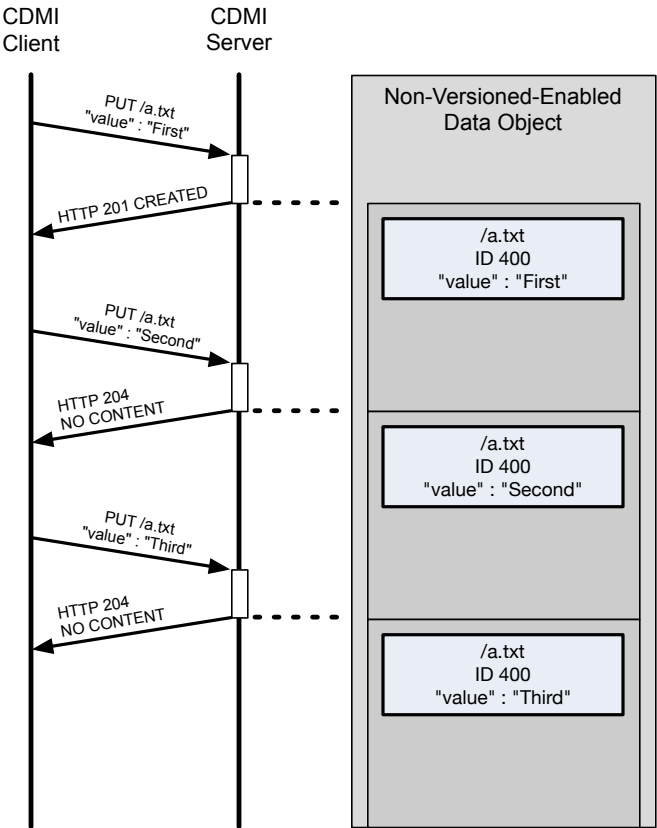


Figure 19: Updates to a non-version-enabled data object

3945 When a data object has versioning enabled, each update creates a new “current version” with the same contents of  
3946 the version-enabled data object, and the previous current version becomes a historical version. All versions can be  
3947 accessed via separate URIs and are immutable. The version-enabled data object continues to be mutable and has the  
3948 same behaviors to clients as a non-version-enabled data object. This behavior is shown in Fig. 20 from the perspective

3949 of a client.




Figure 20: Updates to a version-enabled data object

3950 Using this approach, CDMI clients that are not aware of versioning can continue to access version-enabled data objects  
 3951 the same way as non-version-enabled data objects, while CDMI clients that are aware of versioning can access and  
 3952 manage the immutable versions associated with the version-enabled data object.

3953 Versioning is enabled for a data object by adding a data system metadata item that indicates that versioning is desired.

3954 Version-enabled data objects and all associated versions contain additional storage system metadata items. These  
 3955 metadata items allow a client to discover the versions that are associated with a version-enabled data object and to  
 3956 iterate through these versions.

3957 The maximum number of versions to be retained, maximum age of versions to be retained, and the maximum space  
 3958 that can be consumed by versions is controlled by data system metadata.

3959 When a data object is version enabled, it always contains at least one version, the "current version". The current version  
 3960 has the same contents as the version-enabled data object but has a different identifier (URI and Object Identifier) and is  
 3961 immutable. When a version-enabled data object is changed, a new current version is created, and the previous current  
 3962 version becomes a historical version.

3963 Versioning has multiple client use cases:

- 3964 • Clients that need to preserve all data written to a data object over time can use versions to retain all updates  
 3965 made to a data object.
- 3966 • Clients can restore the contents of a historical version by copying it to the version-enabled data object.
- 3967 • Clients that retrieve a large data object across multiple parallel or sequential transactions or that need to be able  
 3968 to resume a retrieval at a later time can retrieve the URI for the current version of the data object. Clients can  
 3969 then use that URI to retrieve the data object itself. As the current version is immutable and retains its identifier,  
 3970 even if an update occurs (where the current version becomes a historical version), the client will always receive  
 3971 the same results and will not receive a mixture of the older and newer data object contents.
- 3972 • Clients can iterate through historical versions to detect where concurrent updates have occurred and can access  
 3973 any overwritten data.
- 3974 • Distributed CDMI implementations can also use versions to merge concurrent changes made on different, even-  
 3975 tually consistent nodes without resulting in data loss.

## 25.2 Traversing version-enabled data objects

Version-enabled data objects have multiple metadata items that allow a client to traverse through the data object versions.

When a client enables versioning for a data object, the following metadata items shall be added to the version-enabled data object:

- a `cdmi_version_object` metadata item that contains the URI to the corresponding version-enabled data object. This metadata item allows a client to detect that a given object is a version-enabled data object and not a data object version.
- a `cdmi_version_current` field that contains the URI to the current version of the version-enabled data object.
- a `cdmi_version_oldest` field that contains the URI of one or more of the oldest versions. More than one version can exist in this metadata item as explained in 25.3.

Each data object version shall contain the above three fields, with the same values as found in the version-enabled data object. Each data object version shall also contain the following two fields:

- a `cdmi_version_parent` field that contains the URI of the previous version. If the data object version does not have a parent, this field is omitted.
- `cdmi_version_children` field that contains the URIs of the versions created by modifying this version. If the data object version does not have any children, this metadata item shall be empty.

To visualize how these fields allow a client to traverse data object versions, the linkages between the version-enabled data object and data object versions in the final state of Fig. 20 is shown in Fig. 21.




Figure 21: Linkages between a version-enabled data object and data object versions

A client accessing the version-enabled data object (`/b.txt`) can traverse to the current version and to the oldest version.

A client accessing a data object version can traverse to the version-enabled data object, to the current version, to the parent version, to child versions, and to the oldest version.

## 25.3 Concurrent updates and version-enabled data objects

When multiple concurrent updates are performed against a version-enabled data object, each update is performed against the state of the object at the time the update starts. The change to the state resulting from the update to the object becomes visible to clients at the time the update completes.

Two different types of concurrent updates can occur: overlapping updates and nested updates.

Fig. 22 and Fig. 23 show the update sequence and resulting version linkages for overlapping updates:




Figure 22: Overlapping concurrent updates




Figure 23: Linkages for overlapping updates

In the sequence shown in Fig. 22, both the "Second" and "Third" updates are performed against the "First" state. As the "Third" update completes last, it becomes the current version. In this example, historical version 501 would have two children, versions 502 and 503. Both versions 502 and 503 would have the same parent 501.



4007 Fig. 24 and Fig. 25 show the update sequence and resulting version linkages for nested updates:




Figure 24: Nested concurrent updates

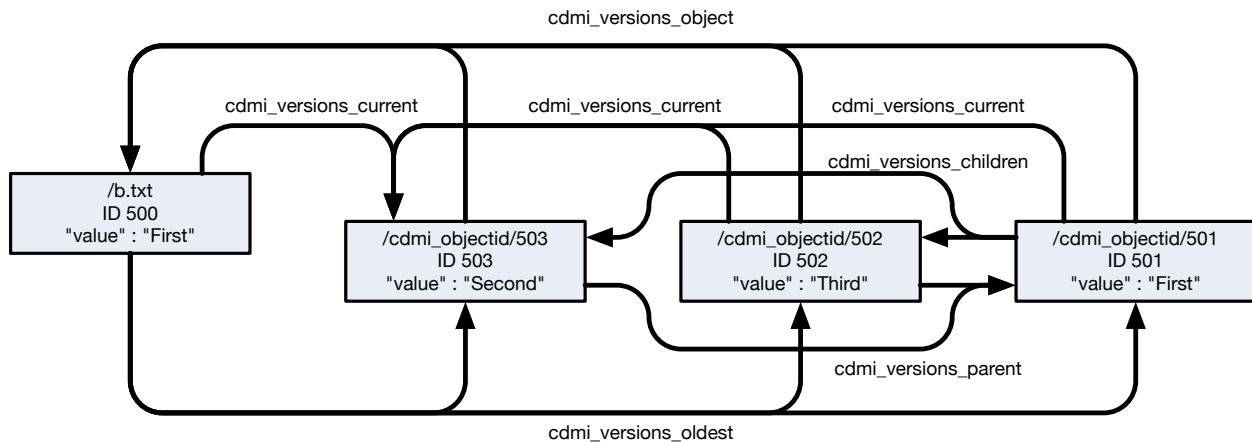


Figure 25: Linkages for nested updates

4008 In the sequence shown in these figures, both the "Second" and "Third" updates are performed against the "First" state.  
 4009 As the "Second" update completes last, it becomes the current version. In this example, historical version 501 would  
 4010 have two children, versions 502 and 503. Both versions 502 and 503 would have the same parent 501.

4011 Both of these data structures are equivalent, with the only difference being which update completed last.

## 25.4 Capabilities for version-enabled data objects

The relationship between version-enabled data objects, data object versions, and capabilities is shown in Fig. 26.




Figure 26: Version to `capabilityURI` relationships

Data object versions are immutable but may be deleted by a client or by the system, depending on the data system metadata specified.

## 25.5 Updates triggering version creation

If versioning is enabled by setting the value of the `cdmi_versions` metadata item in the version-enabled data object to “value”, the following updates will trigger the creation of a new version:

- changing the `mimetype`,
- changing the `value`, or
- changing the `valuetransferencoding`.

If versioning is enabled by setting the value of the `cdmi_versions` metadata item in the version-enabled data object to “user”, the following updates will trigger the creation of a new version:

- changing the `mimetype`,
- changing the `value`,
- changing the `valuetransferencoding`, or
- adding, modifying, or removing user metadata.

If versioning is enabled by setting the value of the `cdmi_versions` metadata item in the version-enabled data object to “all”, then all updates to the data object will trigger the creation of a new version.

While ACLs for historical versions are left unchanged, the effective ACL, owner, and domain of historical versions shall be the ACL, owner, and domain of the current version-enabled data object. This means that changing the ACL of a versioned data object also overrides the historical version ACL for all previous versions.

Modifications performed with the `X-CDMI-Partial` header shall not trigger the creation of a new version until the `completionStatus` is changed from “Processing” to “Complete”.

## 25.6 Operations on version-enabled data objects

Moving a version-enabled data object within a system is considered to be an update to the `name` and/or `parentURI` fields.

Moving a version-enabled data object between systems moves all data object versions associated with the version-enabled data object and preserves all identifiers. If the destination name and/or URI are different, the move is considered to be an update to the `name` and/or `parentURI` fields.

Copying a version-enabled data object shall only copy the current version of the version-enabled data object. Versions of the version-enabled data object are not copied.

Deleting a version-enabled data object shall also delete all versions associated with that version-enabled data object.

Disabling versioning for a version-enabled data object shall preserve all versions. Previously existing versioning meta-data shall remain present while versioning is disabled. Re-enabling versioning for a data object that previously was version-enabled shall result in the creation of a new current version.

If a version-enabled data object is placed under retention or hold, the retention behaviors of the version-enabled data object shall be applied to the data object versions.

No additional notifications are defined for version-enabled data objects. When a version-enabled data object is updated, an additional creation notification message shall be generated for the created data object version. Likewise, when a data object version is accessed or deleted, a notification message is generated. If a limited number, size, or age for versions is requested and a change to a version-enabled data object results in a version being automatically deleted, then the system shall generate a corresponding deletion notification message for the deleted data object version.

## 25.7 Operations on data object versions

A data object version is presented to the client as a standard CDMI data object.

Moving, copying over, deserializing over, and updating a data object version shall not be permitted and shall result in an HTTP status code of 403 Forbidden.

Copying a data object version is permitted. For example, to promote a version to become the current version of a version-enabled data object, the URI of the data object version is used in the copy field when performing an update to the URI of the version-enabled data object. Updates can also be performed as part of the copy operation.

Deleting a historical data object version shall be permitted if the client has ACL permissions to delete the version-enabled data object and the version-enabled data object.

Deleting the current version of a version-enabled data object shall revert the current version to the current version's parent. If there is no parent version, deleting the current version shall result in an HTTP status code of 403 Forbidden.

When an intermediate historical version is deleted, the parent and children metadata items of the parent and all child data object versions of the data object version being deleted must be updated.

In a version chain "C" -> "B" -> "A", where "C" is the newest and "A" is the oldest, deleting version "B" shall produce the following results:

- The `cdmi_version_parent` metadata item of "C" is set to the URI contained in the `cdmi_version_parent` metadata item of "B".
- The URI of "B" in the `cdmi_version_children` metadata item of "A" is replaced with the URIs contained in the `cdmi_version_children` metadata item of "B".

In pseudocode, the above translates to:

```
C->cdmi_version_parent = B->cdmi_version_parent
A->cdmi_version_children[B] = B->cdmi_version_children
Delete B
```

If the oldest version of a version-enabled data object is deleted and there are two or more children of that version, both of the children of the deleted oldest version will become the new oldest version.

When accessing a data object version, the `cdmi_acount` and `cdmi_atime` of the data object version shall be updated if present.

When accessing a historical version of a version-enabled data object, the ACL, owner, and domainURI of the version-enabled data object shall be in effect.

Standard notification messages are sent when data object versions are accessed and deleted.

## 25.8 Query of data object versions

As data object versions are regular CDMI objects, they will be included in query results unless explicitly excluded.

Querying for data object versions is performed by including the scope:

```
"metadata" : {  
  "cdmi_version_children" : "*"  
}
```

Querying for version-enabled data objects (but not their versions) is performed by including the scope:

```
"metadata" : {  
  "cdmi_versioning" : "*"  
}
```

Querying for non-versioned data objects with no versions is performed by including the scope:

```
"metadata" : {  
  "cdmi_version_current" : "!"  
}
```

Querying for non-versioned data objects with versions is performed by including the scope:

```
"metadata" : {  
  "cdmi_versioning" : "!",  
  "cdmi_version_current" : "*"  
}
```

## 25.9 Version-enabled data object serialization

Serializing a version-enabled data object shall serialize the data object, the versioning-related metadata, the current version, and all historical versions. The current version and all historical versions shall be serialized as data objects contained within a JSON array. These data objects shall replace the contents of the value field of the serialized representation of the version-enabled data object.

EXAMPLE 1: A version-enabled data object with three versions is serialized.

```
{
  "objectType" : "application/cdmi-object",
  "objectID" : "00007ED900100DA32EC94351F8970400",
  "objectName" : "MyVersionedDataObject.txt",
  "parentURI" : "/MyContainer/",
  "parentID" : "00007E7F00102E230ED82694DAA975D2",
  "domainURI" : "/cdmi_domains/MyDomain/",
  "capabilitiesURI" : "/cdmi_capabilities/dataobject/",
  "completionStatus" : "Complete",
  "mimetype" : "text/plain",
  "metadata" : {
    "cdmi_size" : "33",
    "cdmi_versioning" : "user",
    "cdmi_version_object" : "/cdmi_objectid/00007ED900100DA32EC94351F8970400",
    "cdmi_version_current" : "/cdmi_objectid/00007ED90010F077F4EB1C99C87524CC",
    "cdmi_version_oldest" : [
      "/cdmi_objectid/00007ED90010512EB55A9304EAC5D4AA"
    ],
    ...
  },
  "value" : [
    {
      "objectType" : "application/cdmi-object",
      "objectID" : "00007ED90010F077F4EB1C99C87524CC",
      "objectName" : "MyVersionedDataObject.txt",
      "parentURI" : "/MyContainer/",
      "parentID" : "00007E7F00102E230ED82694DAA975D2",
      "domainURI" : "/cdmi_domains/MyDomain/",
      "capabilitiesURI" : "/cdmi_capabilities/dataobject/dataobject_version/",
      "completionStatus" : "Complete",
      "mimetype" : "text/plain",
      "metadata" : {
        "cdmi_size" : "33",
        "cdmi_version_object" : "/cdmi_objectid/
↪00007ED900100DA32EC94351F8970400",
        "cdmi_version_current" : "/cdmi_objectid/
↪00007ED90010F077F4EB1C99C87524CC",
        "cdmi_version_oldest" : [
          "/cdmi_objectid/00007ED90010512EB55A9304EAC5D4AA"
        ],
        "cdmi_version_parent" : "/cdmi_objectid/
↪00007ED9001005192891EEBE599D94BB",
        "cdmi_version_children" : [ ],
        ...
      },
      "valuerange" : "0-32",
      "valuetransferencoding" : "utf-8",
      "value" : "Third version of this Data Object"
    },
    {
      "objectType" : "application/cdmi-object",
      "objectID" : "00007ED9001005192891EEBE599D94BB",
      "objectName" : "MyVersionedDataObject.txt",
      "parentURI" : "/MyContainer/",
      "parentID" : "00007E7F00102E230ED82694DAA975D2",
      "domainURI" : "/cdmi_domains/MyDomain/",
      "capabilitiesURI" : "/cdmi_capabilities/dataobject/dataobject_version/",
      "completionStatus" : "Complete",

```

(continues on next page)

(continued from previous page)

```

        "mimetype" : "text/plain",
        "metadata" : {
            "cdmi_size" : "34",
            "cdmi_version_object" : "/cdmi_objectid/
↪00007ED900100DA32EC94351F8970400",
            "cdmi_version_current" : "/cdmi_objectid/
↪00007ED90010F077F4EB1C99C87524CC",
            "cdmi_version_oldest" : [
                "/cdmi_objectid/00007ED90010512EB55A9304EAC5D4AA"
            ],
            "cdmi_version_parent" : "/cdmi_objectid/
↪00007ED90010512EB55A9304EAC5D4AA",
            "cdmi_version_children" : [
                "/cdmi_objectid/00007ED90010F077F4EB1C99C87524CC"
            ],
            ...
        },
        "valuerange" : "0-33",
        "valuetransferencoding" : "utf-8",
        "value" : "Second version of this Data Object"
    },
    {
        "objectType" : "application/cdmi-object",
        "objectID" : "00007ED90010512EB55A9304EAC5D4AA",
        "objectName" : "MyVersionedDataObject.txt",
        "parentURI" : "/MyContainer/",
        "parentID" : "00007E7F00102E230ED82694DAA975D2",
        "domainURI" : "/cdmi_domains/MyDomain/",
        "capabilitiesURI" : "/cdmi_capabilities/dataobject/dataobject_version/",
        "completionStatus" : "Complete",
        "mimetype" : "text/plain",
        "metadata" : {
            "cdmi_size" : "33",
            "cdmi_version_object" : "/cdmi_objectid/
↪00007ED900100DA32EC94351F8970400",
            "cdmi_version_current" : "/cdmi_objectid/
↪00007ED90010F077F4EB1C99C87524CC",
            "cdmi_version_oldest" : [
                "/cdmi_objectid/00007ED90010512EB55A9304EAC5D4AA"
            ],
            "cdmi_version_children" : [
                "/cdmi_objectid/00007ED9001005192891EEBE599D94BB"
            ],
            ...
        },
        "valuerange" : "0-32",
        "valuetransferencoding" : "utf-8",
        "value" : "First version of this Data Object"
    }
]
}

```

4093 Serializing a non-version-enabled data object that has versions shall serialize the data object, the versioning-related  
 4094 metadata, and all historical versions. The contents of the value field of the data object, the current version, and all  
 4095 historical versions serialized as data objects shall be contained within a JSON array. These data objects shall replace  
 4096 the contents of the value field of the serialized representation of the version-enabled data object.

4097 Deserializing either a version-enabled data object or a non-version-enabled data object with versions shall restore the  
 4098 data object and all serialized versions.

4099 Serializing and deserializing a data object version shall not be permitted.

4100 Attempting to deserialize a serialized version-enabled data object or non-version-enabled data object with versions onto  
 4101 a system that does not support versions shall result in an HTTP status code of 400 Bad Request. This error code  
 4102 results because a CDMI system that does not support versions expects a JSON string for the value field of a serialized  
 4103 data object, not a JSON array.



## Part V

# CDMI Annexes

## Clause 26

# Extensions

### 26.1 Overview

CDMI extensions describe non-normative additional functionality for extending the CDMI International Standard. Each extension is first written as a standalone document that describes the changes that are required to implement the functionality being added into this International Standard.

When one or more vendors have implemented a CDMI extension, it is eligible to be added to this annex. When multiple vendors have implemented a CDMI extension and demonstrated interoperability, the extension is eligible to be merged into the CDMI International Standard itself, at which point it becomes normative.

CDMI extensions shall not break or modify existing functionality, and thus do not result in compatibility problems with existing clients. Compatibility is typically accomplished by relaxing restrictions imposed in the current CDMI International Standard, adding new fields, or using reserved names for metadata. The clients that are using CDMI capabilities can identify the functionality that is associated with these CDMI extensions.

## 26.2 Summary metadata for bandwidth

### 26.2.1 Overview

Domain summaries provide summary measurement information about domain usage and billing. Some systems may track additional usage and billing information related to network bandwidth. This extension proposes a set of additional, optional contents for domain summary objects.

### 26.2.2 Changes to specification

Add new terms to [clause 3](#):

**private network segment** a single IP address or range of IP addresses that are considered internal (e.g., LAN)

**public network segment** a single IP address or range of IP addresses that are considered external (e.g., WAN)

Add table entries to end of [Table 79](#) in [10.2](#):

Metadata name	Type	Description	Requirement
cdmi_summary_network_bytes	JSON string	Total number of bytes read/written to/from public/private network segments	Optional
cdmi_summary_reads_private	JSON string	Total number of bytes read from private network segment	Optional
cdmi_summary_reads_private_min	JSON string	Minimum number of bytes read from private network segment for the given interval	Optional
cdmi_summary_reads_private_max	JSON string	Maximum number of bytes read from private network segment for the given interval	Optional
cdmi_summary_reads_private_avg	JSON string	Average number of bytes read from private network segment for the given interval	Optional
cdmi_summary_writes_private	JSON string	Total number of bytes written to private network segment	Optional
cdmi_summary_writes_private_min	JSON string	Minimum number of bytes written to private network segment for the given interval	Optional
cdmi_summary_writes_private_max	JSON string	Maximum number of bytes written to private network segment for the given interval	Optional
cdmi_summary_writes_private_avg	JSON string	Average number of bytes written to private network segment for the given interval	Optional
cdmi_summary_reads_public	JSON string	Total number of bytes read from public network segment	Optional
cdmi_summary_reads_public_min	JSON string	Minimum number of bytes read from public network segment for the given interval	Optional
cdmi_summary_reads_public_max	JSON string	Maximum number of bytes read from public network segment for the given interval	Optional
cdmi_summary_reads_public_avg	JSON string	Average number of bytes read from public network segment for the given interval	Optional
cdmi_summary_writes_public	JSON string	Total number of bytes written to public network segment	Optional

Continued on next page

Table 158 – continued from previous page

Metadata name	Type	Description	Requirement
cdmi_summary_writes_public_min	JSON string	Minimum number of bytes written to public network segment for the given interval	Optional
cdmi_summary_writes_public_max	JSON string	Maximum number of bytes written to public network segment for the given interval	Optional
cdmi_summary_writes_public_avg	JSON string	Average number of bytes written to public network segment for the given interval	Optional
cdmi_summary_reads_total	JSON string	Total number of bytes read from both public and private network segments	Optional
cdmi_summary_writes_total	JSON string	Total number of bytes written to both public and private network segments	Optional

## 26.3 Expiring access control entries (ACEs)

### 26.3.1 Overview

A common trait of cloud storage services is the ability to share an object with other clients for a limited time. This extension adds an attribute of ACEs used in ACLs that imposes a time limit (expiration) on the ACE. Once the ACE expires, the ACE is no longer valid or included in the authorization calculation for the object.

### 26.3.2 Changes to specification

Insert into 17.2.6:

After the bullet item:

- ACEs that do not refer to the principal P requesting the operation are ignored.

Insert bullet:

- ACEs that have an expiration value less than the current time are ignored.

Change 17.2.6:

Original text:

```
ACE = { acetype , identifier , aceflags , acemask , acetime }
```

Revised text:

```
ACE = { acetype , identifier , aceflags , acemask , acetime, expiration }
```

Insert into 17.2.6 after “acemask = uint\_t | acemaskstring”:

```
expiration = uint_t
```

Insert into 17.2.6 after “When ACE masks...”:

When ACE expiration is presented in string format, it shall be specified in ISO-8601 point-in-time format as described in 5.6.

Insert a new sub-clause after 17.2.9: “ACE expiration”

An ACE may have an optional expiration associated with it. The expiration is a point-in-time value, in ISO-8601 point-in-time format, as described in 5.6, which specifies that the ACE is no longer valid and shall be ignored after the time specified.

## 26.4 Group storage system metadata

### 26.4.1 Overview

ACLs in CDMI can refer to the owner of an object by specifying an ACE Who of "OWNER@". This reference corresponds to the contents of the `cdmi_owner` storage system metadata. However, no `cdmi_group` storage system metadata corresponds to an ACE Who of "GROUP@".

This extension defines a new storage system metadata item, `cdmi_group`, that allows an object to be associated with a group for ACL evaluation purposes.

### 26.4.2 Changes to specification

Add a new row at end of table [Table 126](#) in [12.1.9](#):

Capability name	Type	Definition
<code>cdmi_group</code>	JSON string	If present and "true", this capability indicates that the cloud storage system supports group storage system metadata to indicate a group associated with the object.

Add a new row below "`cdmi_owner`" in [Table 140](#) of [16.2](#):

Metadata name	Type	Description	Requirement
<code>cdmi_group</code>	JSON string	The name of the group that is associated with the object.	Optional

## 26.5 Header-based metadata

### 26.5.1 Overview

The CDMI protocol enables CDMI-aware clients to store and retrieve structured metadata using JSON bodies, but does not permit HTTP-based clients to access this metadata. This extension extends CDMI metadata to permit HTTP header metadata to be stored and retrieved as a subset of CDMI metadata.

Due to limitations associated with HTTP headers, certain restrictions must be placed on metadata that is accessible via headers.

### 26.5.2 Changes to specification

Add a new row at end of table [Table 9](#) in [6.2](#):

Header	Type	Description	Requirement
x-*-meta-*	Header string	<p>If the “cdmi_header_metadata” capability is present, for each request header matching the pattern “x-*-meta-*”, a new user metadata item shall be created, with the metadata name set to the header field-name, and the metadata value set to the header field-value.</p> <p>If the number of headers, the length of any of the headers, or the total size of the headers exceeds the limits specified in RFC 2616, or specified by the <code>cdmi_header_metadata_maxitems</code>, <code>cdmi_header_metadata_maxsize</code>, or the <code>cdmi_header_metadata_maxtotalsize</code> capabilities, a 400 Bad Request shall be returned to the client.</p>	Conditional

Add new example at end of [6.2](#):

EXAMPLE 4: PUT to the container URI the data object name, contents, and metadata:

```
PUT /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Content-Type: text/plain;charset=utf-8
X-CDMI-Meta-Colour: Yellow
X-Object-Meta-Shape: Square
Content-Length: 37

This is the Value of this Data Object
```

The following shows the response:

```
HTTP/1.1 201 Created
```

After [6.2](#), add a new clause “Inspect a Data Object using HTTP”:

### 26.5.3 Synopsis

To check for the presence of a data object, the following request shall be performed:

- HEAD <root URI>/<ContainerName>/<DataObjectName>

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate containers that already exist, with one slash (i.e., "/" ) between each pair of container names.
- <DataObjectName> is the name specified for the data object to be checked.

The object shall also be able to be checked at <root URI>/cdmi\_objectid/<objectID>.

### 26.5.4 Capabilities

The following capabilities describe the supported operations that may be performed when reading an existing data object:

- Support for the ability to read the metadata of an existing data object is indicated by the presence of the `cdmi_read_metadata` capability in the specified object.

### 26.5.5 Request headers

Request headers may be provided as per RFC 2616.

### 26.5.6 Request message body

A request message body shall not be provided.

### 26.5.7 Response headers

The HTTP response headers for checking for the presence of a data object using HTTP are shown in [Table 159](#).

Table 159: Response headers - Inspect a data object using HTTP

Header	Type	Description	Requirement
Content-Type	Header string	The content type returned shall be the <code>mimetype</code> field in the data object.	Mandatory
Location	Header string	The server shall respond with the URI that the reference redirects to if the object is a reference.	Conditional
x-*-meta-*	Header string	<p>If the <code>"cdmi_header_metadata"</code> capability is present, for each user metadata item in the <code>"metadata"</code> field with a metadata name that is a case-insensitive match to the pattern <code>"x-*-meta-"</code>, a corresponding response header shall be returned to the client where the header field-name shall be the metadata item name, and the header field-value shall be the metadata item value.</p> <p>If a header value to be return is not conformant with RFC 2616, the server may omit the field from the response headers.</p>	Conditional

### 26.5.8 Response message body

No response body shall be provided, as per RFC 2616.



## 26.5.9 Response status

The HTTP status codes that occur when checking the presence of a data object using HTTP are described in [Table 160](#).

Table 160: HTTP status codes - Inspect a data object using HTTP

HTTP status	Description
200 OK	The queue object content was returned in the response.
302 Found	The resource is a reference to another resource.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.

## 26.5.10 Example

EXAMPLE 1: HEAD to the data object URI to check for the presence of a data object with header metadata:

```
HEAD /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
```

The following shows the response:

```
HTTP/1.1 200 OK
Content-Type: text/plain
Content-Length: 37
X-CDMI-Meta-Colour: Yellow
X-Object-Meta-Shape: Square
```

Add a new row at end of table [Table 13](#) in [6.3](#):

Header	Type	Description	Requirement
x-*-meta-*	Header string	<p>If the "cdmi_header_metadata" capability is present, for each user metadata item in the "metadata" field with a metadata name that is a case-insensitive match to the pattern "x-*-meta-*", a corresponding response header shall be returned to the client where the header field-name shall be the metadata item name, and the header field-value shall be the metadata item value.</p> <p>If a header value to be return is not conformant with RFC 2616, the server may omit the field from the response headers.</p>	Conditional

Add new example at end of [6.3](#):

EXAMPLE 6: GET to the data object URI to read the value of the data object with header metadata:

```
GET /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
```

The following shows the response:

```
HTTP/1.1 200 OK
Content-Type: text/plain
Content-Length: 37
X-CDMI-Meta-Colour: Yellow
X-Object-Meta-Shape: Square

This is the Value of this Data Object
```

Add a new row at end of table [Table 16](#) in [6.4](#):

Header	Type	Description	Requirement
x*-meta-*	Header string	<p>If the "cdmi_header_metadata" capability is present, for each request header matching the pattern "x*-meta-*", a new user metadata item shall be created, or an existing metadata item shall be updated, with the metadata name set to the header field-name, and the metadata value set to the header field- value.</p> <p>If a metadata item already exists where the metadata name and the header field-name differ only in case, the existing metadata item value shall be updated.</p> <p>If an empty header field-value is specified, the corresponding metadata item shall be removed from the object.</p> <p>If the number of headers, the length of any of the headers, or the total size of the headers exceeds the limits specified in RFC 2616, or specified by the cdmi_header_metadata_maxitems, cdmi_header_metadata_maxsize, or the cdmi_header_metadata_maxtotalsize capabilities, a 400 Bad Request shall be returned to the client.</p>	Conditional

Add new example at end of [6.4](#):

EXAMPLE 3: PUT to the data object URI to update the value and metadata of the data object:

```
PUT /MyContainer/MyDataObject.txt HTTP/1.1
Host: cloud.example.com
Content-Type: text/plain;charset=utf-8
X-CDMI-Meta-Colour: Green
Content-Length: 41

This is the new Value of this Data Object
```

The following shows the response:

```
HTTP/1.1 204 No Content
```

Add a new table to Request Headers in [7.2](#):

Table 161: Request headers - Create a container object using HTTP

Header	Type	Description	Requirement
x-*-meta-*	Header string	<p>If the “cdmi_header_metadata” capability is present, for each request header matching the pattern “x-*-meta-*”, a new user metadata item shall be created, with the metadata name set to the header field-name, and the metadata value set to the header field-value.</p> <p>If the number of headers, the length of any of the headers, or the total size of the headers exceeds the limits specified in RFC 2616, or specified by the <code>cdmi_header_metadata_maxitems</code>, <code>cdmi_header_metadata_maxsize</code>, or the <code>cdmi_header_metadata_maxtotalsize</code> capabilities, a 400 Bad Request shall be returned to the client.</p>	Conditional

Add new example at end of 7.2:

EXAMPLE 2: PUT to the URI the container object name and metadata:

```
PUT /MyContainer/ HTTP/1.1
Host: cloud.example.com
X-CDMI-Meta-Colour: Yellow
```

The following shows the response:

```
HTTP/1.1 201 Created
```

After 7.2, add a new sub-clause “Inspect a container object using HTTP”:

### 26.5.11 Synopsis

To check for the presence of a container object, the following request shall be performed:

- HEAD <root URI>/<ContainerName>/<TheContainerName>/

Where:

- <root URI> is the path to the CDMI cloud.
- <ContainerName> is zero or more intermediate containers that already exist, with one slash (i.e., “/”) between each pair of container names.
- <TheContainerName> is the name specified for the container object to be checked.

The container object shall also be able to be checked at <root URI>/cdmi\_objectid/<objectID>.

### 26.5.12 Capabilities

The following capabilities describe the supported operations that may be performed when reading an existing container object:

- Support for the ability to read the metadata of an existing container object is indicated by the presence of the `cdmi_read_metadata` capability in the specified container object.

### 26.5.13 Request headers

Request headers may be provided as per RFC 2616.

### 26.5.14 Request message body

A request message body shall not be provided.

### 26.5.15 Response headers

The HTTP response headers for checking for the presence of a CDMI container object using HTTP are shown in [Table 162](#).

Table 162: Response Headers - Inspect a container object using HTTP

Header	Type	Description	Requirement
Content-Type	Header string	"application/cdm-container"	Mandatory
Location	Header string	The server shall respond with the URI that the reference redirects to if the object is a reference.	Conditional
x-*-meta-*	Header string	If the "cdmi_header_metadata" capability is present, for each user metadata item in the "metadata" field with a metadata name that is a case-insensitive match to the pattern "x-*-meta-*", a corresponding response header shall be returned to the client where the header field-name shall be the metadata item name, and the header field-value shall be the metadata item value.  If a header value to be return is not conformant with RFC 2616, the server may omit the field from the response headers.	Conditional

### 26.5.16 Response message body

No response body shall be provided, as per RFC 2616.

### 26.5.17 Response status

The HTTP status codes that occur when checking the presence of a container object using HTTP are described in [Table 163](#).

Table 163: HTTP status codes - Inspect a container object using HTTP

HTTP Status	Description
200 OK	The queue object content was returned in the response.
302 Found	The resource is a reference to another resource.
400 Bad Request	The request contains invalid parameters or field names.
401 Unauthorized	The authentication credentials are missing or invalid.
403 Forbidden	The client lacks the proper authorization to perform this request.
404 Not Found	The resource was not found at the specified URI.

### 26.5.18 Example

EXAMPLE 1: HEAD to the container object URI to check for the presence of a container object with header metadata:

```
HEAD /MyContainer/ HTTP/1.1
Host: cloud.example.com
```

The following shows the response:

```
HTTP/1.1 200 OK
Content-Type: application/cdm-container
X-CDMI-Meta-Colour: Yellow
```

Replace contents of 7.3 with:

### 26.5.19 Synopsis

Reading a container object using HTTP is not defined by this version of this international standard. 9.3 describes how to read a container object using CDMI.

A server implementation is free to respond to HTTP GETs for Container Objects in any way that conforms with RFC 2616.

If container object metadata items matching the pattern "x-\*-meta-\*" are present, these metadata items shall be returned as response headers as per 9.3.

Replace contents of 7.4 with:

### 26.5.20 Synopsis

Updating a container object using HTTP is not defined by this version of this international standard. clause 9.4 describes how to update a container object using CDMI.

A server implementation is free to respond to HTTP PUTs for existing Container Objects in any way that conforms with RFC 2616.

If container object metadata items matching the pattern "x-\*-meta-\*" are present, these metadata items shall be returned as response headers as per 9.4.

Add new rows to end of Table 124 in 12.1.7:

Capability name	Type	Definition
cdmi_header_metadata	JSON string	If present and "true", this capability indicates that the cloud storage system supports header-visible metadata.
cdmi_header_metadata_maxitems	JSON string	If present, this capability indicates the maximum number of user-defined header metadata items supported per object. If absent, there is no additional limit placed on the number of user-defined metadata items.
cdmi_header_metadata_maxsize	JSON string	If present, this capability indicates the maximum size, in bytes, of each user-defined header metadata item. If absent, there is no additional limit placed on the size of user-defined metadata items.

Continued on next page

Table 164 – continued from previous page

Capability name	Type	Definition
cdmi_header_metadata_maxtotalsize	JSON string	If present, this capability indicates the maximum size, in bytes, of all user-defined header metadata per object. If absent, there is no additional limit placed on the size of user-defined metadata.

---

Add to end of 16.5:

---

If metadata items with a name is a case-insensitive match to the pattern "x-\*meta-\*" are created or updated through a CDMI request, the following conditions shall be true, or else a 400 Bad Request result code shall be returned to the client:

- The metadata name shall be a valid HTTP header field-name
- The metadata value that is a valid HTTP header field-value
- The number of matching headers shall not exceed the limits specified by RFC 2616, and shall not exceed the number specified in the `cdmi_header_metadata_maxitems` capability.
- The size of each matching header shall not exceed the limits specified by RFC 2616, and shall not exceed the number specified in the `cdmi_header_metadata_maxsize` capability.
- The total size of all of the matching headers shall not exceed the limits specified by RFC 2616, and shall not exceed the number specified in the `cdmi_header_metadata_maxtotalsize` capability.

## 26.6 Immediate query

### 26.6.1 Overview

CDMI provides a query mechanism based around the concept of persistence. A query queue is created, metadata is specified that defines the query operation, the query is performed asynchronously, and results are populated in the queue and then read by the client as separate operations.

This architecture, while providing significant value, is complex for clients that do not need to persist the results of a query. Specifically, a client must: a) asynchronously poll the query queue to determine when results are present and when the query has completed, and b) delete the queue when results are no longer needed.

To provide a simpler interface for simple queries where a small number of results are expected and persistence is not required, the TWG has proposed the following approach to allow query queues to optionally not be persistent, with the results being returned immediately as the response to the initial query queue creation.

In addition, functionality that permits results to be returned immediately has been added to creating asynchronous query queues.

### 26.6.2 Changes to specification

Modify existing "cdmi\_query" entry in [Table 124](#) in [12.1.7](#):

Capability name	Type	Definition
cdmi_query	JSON string	If present and "true", the CDMI server supports persistent query queues.

Add a new row at end of table [Table 124](#) in [12.1.7](#):

Capability name	Type	Definition
cdmi_query_immediate	JSON string	If present and "true", the CDMI server supports immediate query queues.

Replace the first paragraph of Overview in [clause 22](#) with:

A cloud storage system may optionally implement metadata and/or full-text query functionality. The implementation of query is indicated by the presence of the cloud storage system-wide capabilities for query and requires support for CDMI queues when persisting query results.

Replace the third paragraph of Overview in [clause 22](#) with:

When a client wishes to perform queries, it shall first determine if the system is capable of providing query functionality by checking to see if the `cdmi_query` or `cdmi_query_immediate` capabilities are present in the root container capabilities. If these capabilities are not present and queues are supported, creating a query queue shall be successful, but no query results shall be enqueued into the query queue.

Modify existing "cdmi\_queue\_type" entry in [Table 151](#) in [22](#):

Table 165: Required metadata for a query queue

Metadata name	Type	Description	Requirement
cdmi_queue_type	JSON string	Queue type indicates how the cloud storage system shall manage the queue object. The defined values are: <ul style="list-style-type: none"> <li>"cdmi_query_queue" – Perform an asynchronous query, which may return none, some, or all results in the request response body. A new queue object shall be created.</li> <li>"cdmi_query_immediate" – Perform a synchronous query, returning all matching results in the request response body. The query queue object may not be accessible and shall be automatically deleted when the query completes.</li> </ul>	Mandatory

Add new clause "Immediate Queries" to end of [22](#):

If "cdmi\_query\_immediate" is specified in `cdmi_queue_type`, all query results shall be immediately returned in the response body as shown in the following example.

EXAMPLE 3: Perform an Immediate Query:

```
PUT /MyContainer/myQuery HTTP/1.1
Host: cloud.example.com
Accept: application/cdmi-queue
Content-Type: application/cdmi-queue

{
  "metadata" : {
    "cdmi_queue_type" : "cdmi_query_immediate",
    "cdmi_scope_specification" : [
      {
        "domainURI" : "==" /cdmi_domains/MyDomain/",
        "parentURI" : "starts /sandbox",
        "metadata" : {
          "cdmi_size" : "#> 100000"
        }
      }
    ],
    "cdmi_results_specification" : {
      "objectID" : "",
      "metadata" : {
        "cdmi_size" : ""
      }
    }
  }
}
```

The following shows the response.

```
HTTP/1.1 201 Created
Content-Type: application/cdmi-queue
```

(continues on next page)



(continued from previous page)

```

Location: https://cloud.example.com/MyContainer/myQuery

{
  "objectType" : "application/cdmi-queue",
  "objectID" : "00007E7F00104BE66AB53A9572F9F51E",
  "objectName" : "myQuery",
  "parentURI" : "/MyContainer/",
  "parentID" : "0000706D0010B84FAD185C425D8B537E",
  "domainURI" : "/cdmi_domains/MyDomain/",
  "capabilitiesURI" : "/cdmi_capabilities/queue/",
  "completionStatus" : "Complete",
  "metadata" : {
    "cdmi_queue_type" : "cdmi_query_immediate",
    "cdmi_scope_specification" : [
      {
        "domainURI" : "=/cdmi_domains/MyDomain/",
        "parentURI" : "starts /sandbox",
        "metadata" : {
          "cdmi_size" : "#> 100000"
        }
      }
    ],
    "cdmi_results_specification" : {
      "objectID" : "",
      "metadata" : {
        "cdmi_size" : ""
      }
    }
  },
  "queueValues" : "0-0",
  "mimetype": [ "application/json" ],
  "valuerange": [ "0-111" ],
  "valuetransferencoding": [ "base64" ],
  "value": "ew0KCQkJIm9iamVjdE1EIiA6IClwMDAwN0U3RjAwMTBFQjkwOTJCMj1lGNkNENkFENjgyNCIsDQoJCQkibWV0YWRhdGEiIDogew0KCQkJSjJkZG1pX3NpemUiIDogIjEwODI2MyINCgkJCX0NCgkJfQ0K"
}

```

4386 Where the value of the above base64 encoded value is:

4387 **EXAMPLE 1:** An example of the metadata associated with a query queue is as follows:

```

{
  "objectID" : "00007E7F0010EB9092B29F6CD6AD6824",
  "metadata" : {
    "cdmi_size" : "108263"
  }
}

```

## Part VI

# References

# Bibliography

- [1] Carl Beame, Robert Thurlow, Brent Callaghan, David Robinson, David Noveck, Mike Eisler, and Spencer Shepler. Network File System (NFS) version 4 Protocol. RFC 3530, April 2003. URL: <https://rfc-editor.org/rfc/rfc3530.txt>, doi:10.17487/RFC3530.
- [2] Tim Berners-Lee, Roy T. Fielding, and Larry M Masinter. Uniform Resource Identifier (URI): Generic Syntax. RFC 3986, January 2005. URL: <https://rfc-editor.org/rfc/rfc3986.txt>, doi:10.17487/RFC3986.
- [3] Scott O. Bradner. Key words for use in RFCs to Indicate Requirement Levels. RFC 2119, March 1997. URL: <https://rfc-editor.org/rfc/rfc2119.txt>, doi:10.17487/RFC2119.
- [4] Mallikarjun Chadalapaka, Julian Satran, Kalman Meth, and David L. Black. Internet Small Computer System Interface (iSCSI) Protocol (Consolidated). RFC 7143, April 2014. URL: <https://rfc-editor.org/rfc/rfc7143.txt>, doi:10.17487/RFC7143.
- [5] Douglas Crockford. The application/json Media Type for JavaScript Object Notation (JSON). RFC 4627, July 2006. URL: <https://rfc-editor.org/rfc/rfc4627.txt>, doi:10.17487/RFC4627.
- [6] Lisa M. Dusseault. HTTP Extensions for Web Distributed Authoring and Versioning (WebDAV). RFC 4918, June 2007. URL: <https://rfc-editor.org/rfc/rfc4918.txt>, doi:10.17487/RFC4918.
- [7] Roy Thomas Fielding. *REST: Architectural Styles and the Design of Network-based Software Architectures*. PhD thesis, University of California, Irvine, 2000. URL: <http://www.ics.uci.edu/~fielding/pubs/dissertation/top.htm>.
- [8] Professor John Franks, Phillip Hallam-Baker, Lawrence C. Stewart, Jeffery L. Hostetler, Scott Lawrence, Paul J. Leach, and Ari Luotonen. HTTP Authentication: Basic and Digest Access Authentication. RFC 2617, June 1999. URL: <https://rfc-editor.org/rfc/rfc2617.txt>, doi:10.17487/RFC2617.
- [9] Ned Freed and Dr. Nathaniel S. Borenstein. Multipurpose Internet Mail Extensions (MIME) Part One: Format of Internet Message Bodies. RFC 2045, November 1996. URL: <https://rfc-editor.org/rfc/rfc2045.txt>, doi:10.17487/RFC2045.
- [10] Ned Freed and Dr. Nathaniel S. Borenstein. Multipurpose Internet Mail Extensions (MIME) Part Two: Media Types. RFC 2046, November 1996. URL: <https://rfc-editor.org/rfc/rfc2046.txt>, doi:10.17487/RFC2046.
- [11] Tony Hansen and Alexey Melnikov. Additional Media Type Structured Syntax Suffixes. RFC 6839, January 2013. URL: <https://rfc-editor.org/rfc/rfc6839.txt>, doi:10.17487/RFC6839.
- [12] Russ Housley. Cryptographic Message Syntax (CMS). RFC 5652, September 2009. URL: <https://rfc-editor.org/rfc/rfc5652.txt>, doi:10.17487/RFC5652.
- [13] Russ Housley, Tim Polk, Dr. Warwick S. Ford, and David Solo. Internet X.509 Public Key Infrastructure Certificate and Certificate Revocation List (CRL) Profile. RFC 3280, May 2002. URL: <https://rfc-editor.org/rfc/rfc3280.txt>, doi:10.17487/RFC3280.
- [14] Karthik Jaganathan, Larry Zhu, and John Brezak. SPNEGO-based Kerberos and NTLM HTTP Authentication in Microsoft Windows. RFC 4559, June 2006. URL: <https://rfc-editor.org/rfc/rfc4559.txt>, doi:10.17487/RFC4559.
- [15] Michael Jones. JSON Web Algorithms (JWA). RFC 7518, May 2015. URL: <https://rfc-editor.org/rfc/rfc7518.txt>, doi:10.17487/RFC7518.
- [16] Michael Jones. JSON Web Key (JWK). RFC 7517, May 2015. URL: <https://rfc-editor.org/rfc/rfc7517.txt>, doi:10.17487/RFC7517.
- [17] Michael Jones, John Bradley, and Nat Sakimura. JSON Web Signature (JWS). RFC 7515, May 2015. URL: <https://rfc-editor.org/rfc/rfc7515.txt>, doi:10.17487/RFC7515.
- [18] Michael Jones and Joe Hildebrand. JSON Web Encryption (JWE). RFC 7516, May 2015. URL: <https://rfc-editor.org/rfc/rfc7516.txt>, doi:10.17487/RFC7516.

- [19] Simon Josefsson. The Base16, Base32, and Base64 Data Encodings. RFC 4648, October 2006. URL: <https://rfc-editor.org/rfc/rfc4648.txt>, doi:10.17487/RFC4648.
- [20] Larry M Masinter and Ernesto Nebel. Form-based File Upload in HTML. RFC 1867, November 1995. URL: <https://rfc-editor.org/rfc/rfc1867.txt>, doi:10.17487/RFC1867.
- [21] Keith McCloghrie, Jürgen Schönwälder, David T. Perkins, and Keith McCloghrie. Structure of Management Information Version 2 (SMIv2). RFC 2578, April 1999. URL: <https://rfc-editor.org/rfc/rfc2578.txt>, doi:10.17487/RFC2578.
- [22] Keith Moore. MIME (Multipurpose Internet Mail Extensions) Part Three: Message Header Extensions for Non-ASCII Text. RFC 2047, November 1996. URL: <https://rfc-editor.org/rfc/rfc2047.txt>, doi:10.17487/RFC2047.
- [23] Henrik Frystyk Nielsen, Jeffrey Mogul, Larry M Masinter, Roy T. Fielding, Jim Gettys, Paul J. Leach, and Tim Berners-Lee. Hypertext Transfer Protocol – HTTP/1.1. RFC 2616, June 1999. URL: <https://rfc-editor.org/rfc/rfc2616.txt>, doi:10.17487/RFC2616.
- [24] Eric Rescorla and Tim Dierks. The Transport Layer Security (TLS) Protocol Version 1.2. RFC 5246, August 2008. URL: <https://rfc-editor.org/rfc/rfc5246.txt>, doi:10.17487/RFC5246.
- [25] Krishna Sankar and Arnold Jones. Cloud Data Management Interface (CDMI) Media Types. RFC 6208, April 2011. URL: <https://rfc-editor.org/rfc/rfc6208.txt>, doi:10.17487/RFC6208.
- [26] Ross Williams. A painless guide to crc error detection algorithms. August 1993. URL: [http://www.ross.net/crc/download/crc\\_v3.txt](http://www.ross.net/crc/download/crc_v3.txt).
- [27] Jamie Zawinski, Larry M Masinter, and Martin J. Dürst. The 'mailto' URI Scheme. RFC 6068, October 2010. URL: <https://rfc-editor.org/rfc/rfc6068.txt>, doi:10.17487/RFC6068.
- [28] Efri Zeidner, Constantine Sapuntzakis, Mallikarjun Chadalapaka, Julian Satran, and Kalman Meth. Internet Small Computer Systems Interface (iSCSI). RFC 3720, April 2004. URL: <https://rfc-editor.org/rfc/rfc3720.txt>, doi:10.17487/RFC3720.
- [29] ISO/IEC Joint Directives Maintenance Team. Iso/iec directives, part 2 – principles and rules for the structure and drafting of iso and iec documents. ISO/IEC Directives, Part 2, 2018, 2018. URL: <https://www.iso.org/directives-and-policies.html>.
- [30] ISO/IEC JTC 1/SC 25 Interconnection of information technology equipment. Information technology – small computer system interface (scsi) – part 414: scsi architecture model-4 (sam-4). ISO/IEC 14776-414:2009, June 2009. URL: <https://www.iso.org/standard/53961.html>.
- [31] ISO/IEC JTC 1/SC 27 Information security, cybersecurity and privacy protection. Information technology – security techniques – storage security. ISO/IEC 27040:2015, January 2015. URL: <https://www.iso.org/standard/44404.html>.
- [32] ISO/IEC JTC 1/SC 38 Cloud Computing and Distributed Platforms. Information technology – cloud computing – overview and vocabulary. ISO/IEC 17788:2014, October 2014. URL: <https://www.iso.org/standard/60544.html>.
- [33] ISO/TC 154 Processes, data elements and documents in commerce, industry and administration. Date and time – representations for information interchange – part 1: basic rules. ISO 8601-1:2019, February 2019. URL: <https://www.iso.org/standard/70907.html>.
- [34] ISO/TC 154 Processes, data elements and documents in commerce, industry and administration. Date and time – representations for information interchange – part 2: extensions. ISO 8601-2:2019, February 2019. URL: <https://www.iso.org/standard/70907.html>.
- [35] ISO/TC 20/SC 13 Space data and information transfer systems. Space data and information transfer systems – open archival information system (oais) – reference model. ISO 14721:2012, August 2012. URL: <https://www.iso.org/standard/57284.html>.
- [36] ISO/TC 46 Information and documentation. Codes for the representation of names of countries and their subdivisions – part 1: country codes. ISO 3166-1:2013, November 2013. URL: <https://www.iso.org/standard/63545.html>.
- [37] ISO/TC 46 Information and documentation. Codes for the representation of names of countries and their subdivisions – part 2: country subdivision code. ISO 3166-2:2013, November 2013. URL: <https://www.iso.org/standard/63546.html>.
- [38] ISO/TC 46 Information and documentation. Codes for the representation of names of countries and their subdivisions – part 3: code for formerly used names of countries. ISO 3166-3:2013, November 2013. URL: <https://www.iso.org/standard/63547.html>.
- [39] ISO/TC 68/SC 8 Reference data for financial services. Codes for the representation of currencies. ISO 4217:2015, August 2015. URL: <https://www.iso.org/standard/64758.html>.
- [40] Open Grid Forum. Open cloud computing interface v1.1. June 2011. URL: <http://occi-wg.org/about/specification/>.

- 4484 [41] POSIX - Austin Joint Working Group. Ieee standard for information technology–portable operating system inter-  
4485 face (posix(r)) base specifications, issue 7. IEEE 1003.1-2017, December 2017. URL: [https://standards.ieee.org/  
4486 standard/1003\\_1-2017.html](https://standards.ieee.org/standard/1003_1-2017.html).
- 4487 [42] Storage Networking Industry Association. Tls specification for storage systems v1.0.1. November 2014. URL: [https:  
4488 //www.snia.org/tech\\_activities/standards/curr\\_standards/tls](https://www.snia.org/tech_activities/standards/curr_standards/tls).
- 4489 [43] Information technology – open systems interconnection – the directory – part 8: public-key and attribute certificate  
4490 frameworks. ISO/IEC 9594-8:2017, May 2017. URL: <https://www.iso.org/standard/72557.html>.